

Parallel Programming

Introduction to Threads and Synchronization

Multitasking/Multiprocessing

Multitasking

Concurrent execution of multiple tasks/processes

Time multiplexing of CPU

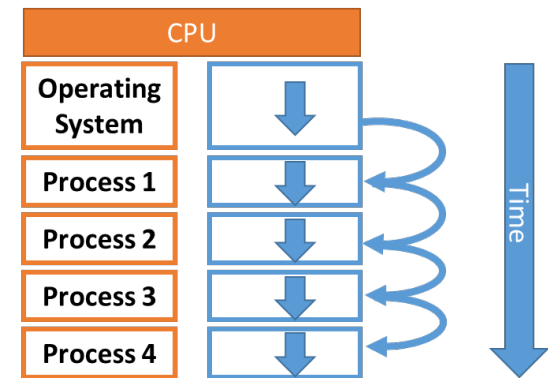
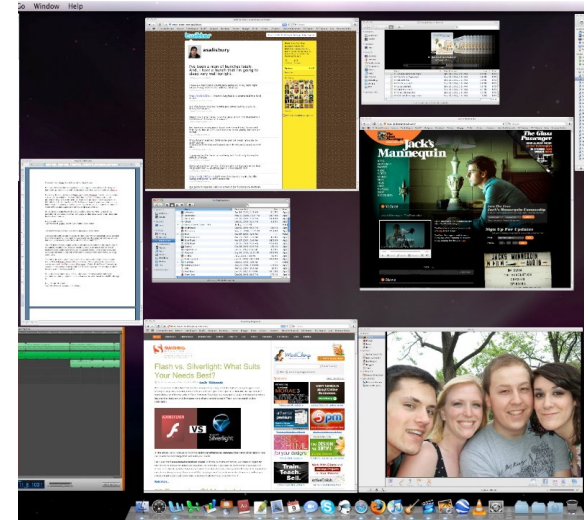
Creates impression of parallelism

Even on single core/CPU system

Allows for asynchronous I/O

I/O devices and CPU are truly parallel

10ms waiting for HDD allows other processes to execute $>10^{10}$ instructions



Process context

A **process** is (essentially) a program executing inside an OS

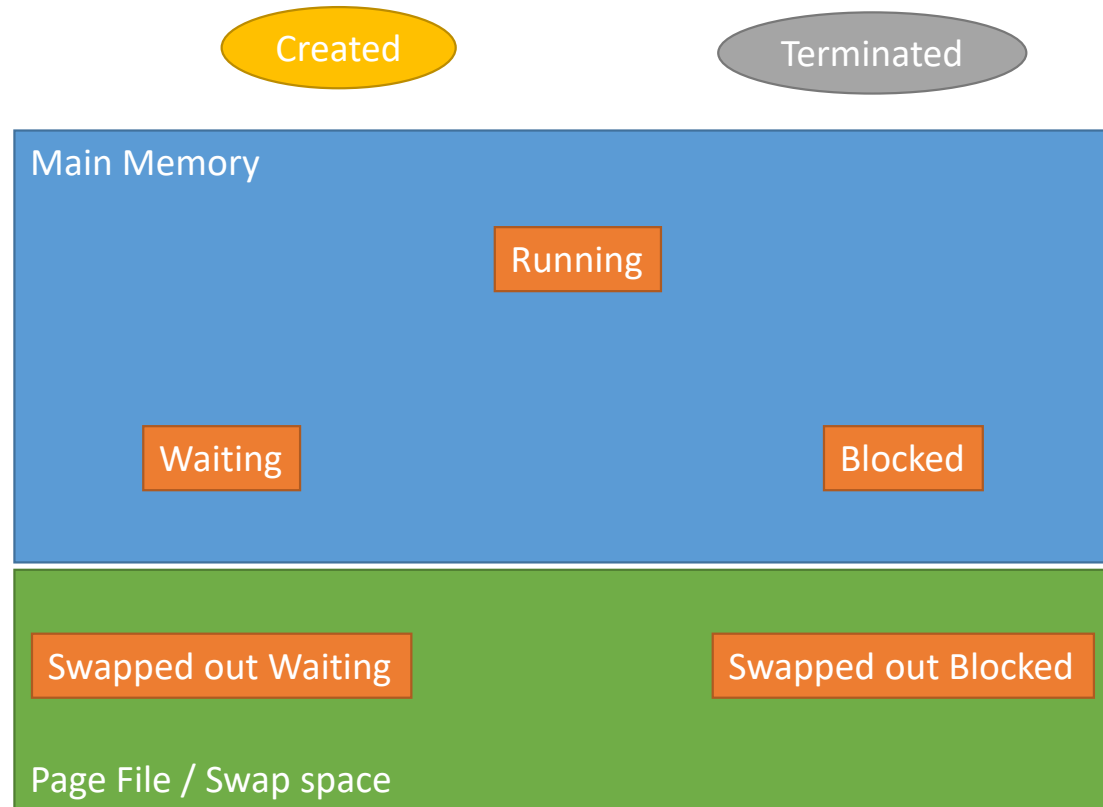
Each running instances of a **program** (e.g., multiple browser windows) is a separate process

Multiple applications (=processes) in parallel

Each process has a **context**:

- Instruction counter
- Values in registers, stack and heap
- Resource handles (device access, open files)
- ...

Process lifecycle states



Process management

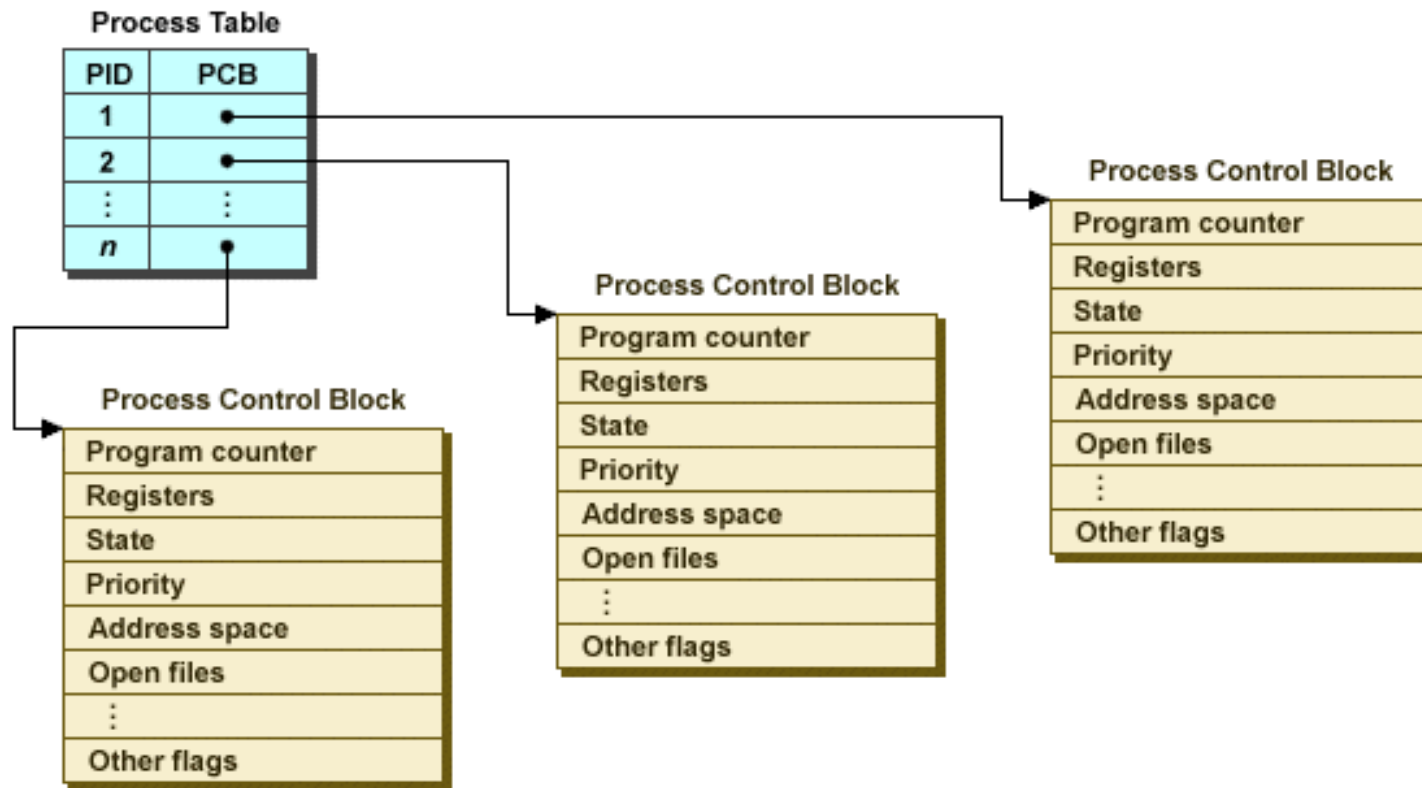
Processes need resources

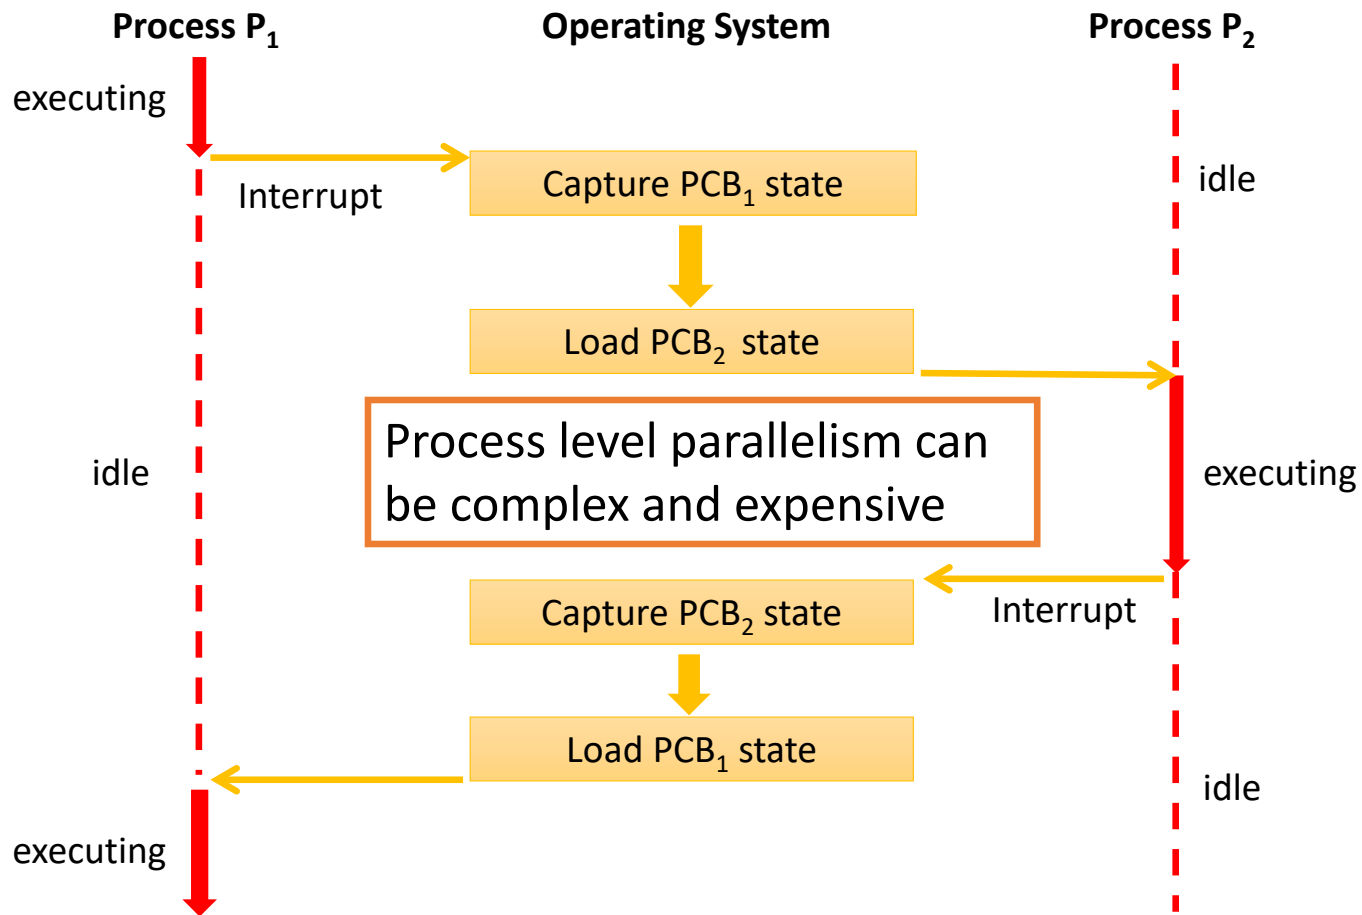
- CPU time, Memory, etc.

OS manages processes:

- Starts processes
- Terminates processes (frees resources)
- Controls resource usage (prevents monopolizing CPU time)
- Schedules CPU time
- Synchronizes processes if necessary
- Allows for inter process communication

Process control blocks (PCB)





Multithreading

Threads

Threads (of control) are

- independent sequences of execution
- running in the same OS process

Multiple threads **share the same address space**.

- Threads are not shielded from each other
- Threads share resources and can communicate more easily

More vulnerable for
programming mistakes

Context switching between threads is efficient

- No change of address space
- No automatic scheduling
- No saving / (re-)loading of PCB (OS process) state

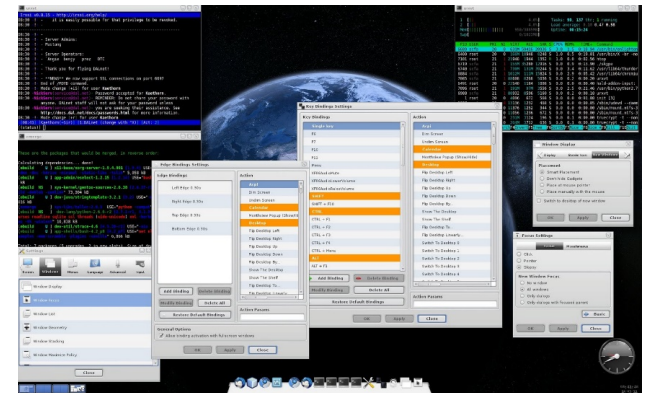

Usage of Multithreading

Reactive systems – constantly monitoring

More responsive to user input – GUI application can interrupt a time-consuming task


Server can handle multiple clients simultaneously

Take advantage of multiple CPUs/cores



Multithreading: 1 vs. many CPUs

Multiple
threads
sharing a
single CPU



Multiple
threads on
multiple
CPUs



Java Threads

Java Threads

Thread

- A set of instructions to be executed one at a time, in a specified order
- A special Thread class is part of the core language

(Some) methods of class `java.lang.Thread`

- `start()` : method called to spawn a new thread
 - Causes JVM to call `run()` method on object
- `interrupt()` : freeze and throw exception to thread

Create Java Threads: Option 1 (oldest)

Instantiate a subclass of `java.lang.Thread` class

- Override `run` method (must be overridden)
- `run()` is called when execution of that thread begins
- A thread terminates when `run()` returns
- `start()` method invokes `run()`
- Calling `run()` does not create a new thread

```
class ConcurrWriter extends Thread { ...  
    public void run() {  
        // code here executes concurrently with caller  
    }  
}  
ConcurrWriter writerThread = new ConcurrWriter();  
writerThread.start();    // calls ConcurrWriter.run()
```

Creating the Thread object does not start the thread!

Need to actually call `start()` to start it.

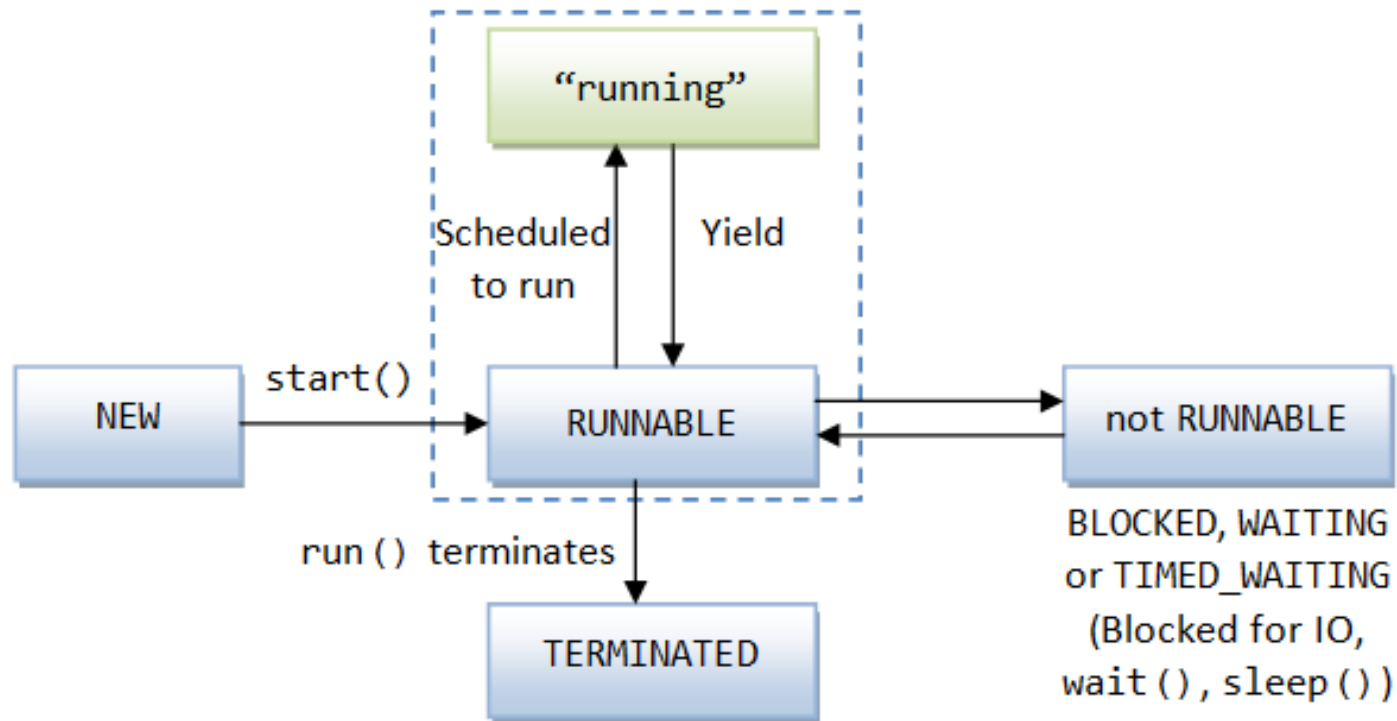
Create Java Threads: Option 2 (better)

Implement `java.lang.Runnable`

- Single method: `public void run()`
- Class implements `Runnable`

```
public class ConcurrWriter implements Runnable {  
    ...  
    public void run() { ...  
        // code here executes concurrently with caller  
    }  
}  
  
ConcurrReader readerThread = new ConcurrReader();  
Thread t = new Thread(readerThread);  
t.start();    // calls ConcurrWriter.run()
```


Thread state model in Java



http://pervasive2.morselli.unimo.it/~nicola/courses/IngegneriaDelSoftware/java/J5e_multithreading.html

java.lang.Thread (under the hood)

```
// Thread.java from OpenJDK:  
// https://hg.openjdk.java.net/jdk/jdk/file/tip/src/java.base/share/classes/java/lang/Thread.java
```

```
public class Thread implements Runnable {  
    static { registerNatives(); }
```

```
    private volatile String name;  
    private int priority;
```

```
    private boolean daemon = false;
```

```
    ...
```

```
    public static native void yield();  
    public static native void sleep(long millis) throws InterruptedException;
```

```
    private Thread(...) { ... }
```

```
    public synchronized void start() { ... }
```

```
    private native void start0();
```

```
    ...
```

A Thread is Runnable



Creates execution environment for the thread
(sets up a separate run-time stack, etc.)

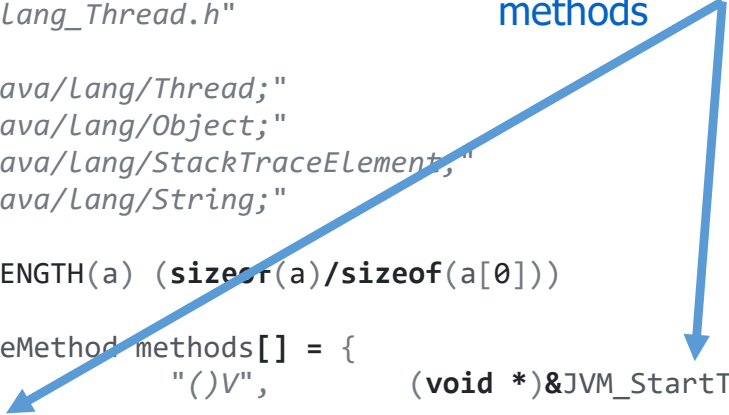


java.lang.Thread (under the hood)

```
// Thread.java from OpenJDK:  
// https://hg.openjdk.java.net/jdk/jdk/file/tip/src/java.base/share/native/libjava/Thread.c  
public class Thread implements Runnable {  
    static { registerNatives(); }  
  
    private volatile String name;  
    private int priority;  
  
    private boolean daemon = false;  
    ...  
  
    public static native void yield();  
    public static native void sleep(long ...)  
  
    private Thread(...) { ... }  
    public synchronized void start() { ..  
    private native void start0();  
    ...
```

```
// Thread.c from OpenJDK:  
// https://hg.openjdk.java.net/jdk/jdk/file/tip/src/java.base/share/native/libjava/Thread.c  
  
#include "jni.h"  
#include "jvm.h"  
  
#include "java_lang_Thread.h"  
  
#define THD "Ljava/lang/Thread;"  
#define OBJ "Ljava/lang/Object;"  
#define STE "Ljava/lang/StackTraceElement;"  
#define STR "Ljava/lang/String;"  
  
#define ARRAY_LENGTH(a) (sizeof(a)/sizeof(a[0]))  
  
static JNINativeMethod methods[] = {  
    {"start0", "()"V", (void *)&JVM_StartThread},  
    ...  
    {"yield", "()"V", (void *)&JVM_Yield},  
    {"sleep", "(J)V", (void *)&JVM_Sleep},  
    ...
```

Native C implementation
of Java's native thread
methods



Example: The parallel calculator

Create 10 threads: each calculates and prints multiplication tables between 1 -10

```
public class Calculator implements Runnable {  
    private int number;  
    public Calculator(int number) {  
        this.number = number;  
    }  
    public void run() { // Override run()  
        for (int i = 1; i <= 10; i++){  
            System.out.printf("%s: %d * %d = %d\n",  
                Thread.currentThread().getName(),  
                number, i, i*number);  
        }  
    }  
}
```

Example: The parallel calculator

Create 10 threads: each calculates and prints multiplication tables between 1 -10

```
public class Calculator implements Runnable {
    private int number;
    public Calculator(int number) {
        this.number = number;
    }
    public void run() { // Override run()
        for (int i = 1; i <= 10; i++){
            System.out.printf("%s: %d * %d = %d\n",
                Thread.currentThread().getName(),
                number, i, i*number);
        }
    }
}
```

```
public static void main(String[] args) {
    //Launch 10 threads that make the operation
    with a different number
    for (int i=1; i <= 10; i++){
        Calculator calculator = new Calculator(i);
        Thread thread = new Thread(calculator);
        thread.start();
    }
}
```

Example: The parallel calculator

Create 10 threads: each calculates and prints multiplication tables between 1 -10

```
public class Calculator implements Runnable {  
    private int number;  
    public Calculator(int number) {  
        this.number = number;  
    }  
    public void run() { // Override run()  
        for (int i = 1; i <= 10; i++){  
            System.out.printf("%s: %d * %d = %d\n",  
                Thread.currentThread().getName(),  
                number, i, i*number);  
        }  
    }  
}
```

```
public static void main(String[] args) {  
    //Launch 10 threads that make the operation  
    with a different number  
    for (int i=1; i <= 10; i++){  
        Calculator calculator = new Calculator(i);  
        Thread thread = new Thread(calculator);  
        thread.start();  
    }  
}
```

Sample output:

```
....  
Thread-9: 10 * 10 = 100  
Thread-4: 5 * 8 = 40  
Thread-4: 5 * 9 = 45  
Thread-4: 5 * 10 = 50  
Thread-5: 6 * 7 = 42  
Thread-2: 3 * 4 = 12  
Thread-5: 6 * 8 = 48  
Thread-0: 1 * 5 = 5  
....
```

Example: The parallel calculator

Create 10 threads: each calculates and prints multiplication tables between 1 -10


```
public class Calculator implements Runnable {  
    private int number;  
    public Calculator(int number) {  
        this.number = number;  
    }  
    public void run() { // Override run()  
        for (int i = 1; i <= 10; i++){  
            System.out.printf("%s: %d * %d = %d\n",  
                Thread.currentThread().getName(),  
                number, i, i*number);  
        }  
    }  
}
```

```
public static void main(String[] args) {  
    //Launch 10 threads that make the operation  
    with a different number  
    for (int i=1; i <= 10; i++){  
        Calculator calculator = new Calculator(i);  
        Thread thread = new Thread(calculator);  
        thread.start();  
    }  
}
```

Sample output:

```
....  
Thread-9: 10 * 10 = 100  
Thread-4: 5 * 8 = 40  
Thread-4: 5 * 9 = 45  
Thread-4: 5 * 10 = 50  
Thread-5: 6 * 7 = 42  
Thread-2: 3 * 4 = 12  
Thread-5: 6 * 8 = 48  
Thread-0: 1 * 5 = 5  
....
```

Note that threads do not appear
in the order they were created...



Java Threads: some key points

Every Java program has **at least one** execution thread

- First execution thread calls `main()`

Each call to **`start()`** method of a Thread object **creates an actual execution thread**

Program ends when all threads (non-daemon threads) finish.

Threads **can continue to run** even if `main()` returns

Creating a Thread object **does not start** a thread

Calling `run()` **doesn't start thread either** (need to call `start()`!)

(Some) Useful Thread attributes and methods

ID: this attribute denotes the unique identifier for each Thread.

```
Thread t = Thread.currentThread(); // get the current thread
System.out.println("Thread ID" + t.getId()); // prints the current ID.
```

Name: this attribute denotes the name of Thread.

```
t.setName("PP" + 2019); // can be modified like this
```

Priority: denotes the priority of the thread. Threads can have a priority between 1 and 10:

JVM uses the priority of threads to select the one that uses the CPU at each moment

```
t.setPriority(Thread.MAX_PRIORITY); // updates the thread's priority
```

Status: denotes the status the thread is in: one of new, runnable, blocked, waiting, time waiting, or terminated (we will discuss the different statuses in more detail later):

```
if (t.getState() == State.TERMINATED) //check if thread's status is terminated
```

Joining Threads

Results, please!

Common scenario:

- Main thread starts (*forks, spawns*) several *worker threads*...
- ... then needs to wait for the worker's results to be available

Previously:

- **Busy waiting** by *spinning* (looping) until each worker's state is TERMINATED
- Boilerplate code
- Inefficient! Main thread spinning uses up CPU time

```
...
finish = false;
while (!finish) {
    ...
    finish = true;
    for (int i=0; i<10; i++){
        finish = finish && (threads[i].getState() == State.TERMINATED);
    }
}
```

Wake me up when work is done

From main thread's perspective:

- Instead of busily waiting for the results (ready? now ready? now?) ...
- ... go to sleep and be woken up once the results are ready

```
...  
for (int i=0; i<10; i++) {  
    threads[i].join(); // May throw InterruptedException  
}
```

Performance trade-off:

- Join (sleep, wakeup) typically incurs context switch overhead
- If worker threads are short-lived, busy waiting may perform better
- Later in the course: SpinLock

Question: Is joining threads[0], ..., threads[9] optimal?

Shared Resources

synchronized

Shared memory interaction between threads

Two or more threads may read/write the same data (shared objects, global data). Programmer responsible for avoiding **bad interleaving** by **explicit synchronization**!

How do we synchronize? Via synchronization primitives.

In Java, all objects have an internal lock, called **intrinsic lock** or **monitor lock**

Synchronized operations (see next) lock the object: while locked, no other thread can successfully lock the object

Generally, if you access shared memory, make sure it is **done under a lock** (Java memory model is complicated!).

(can also use `volatile` keyword, more for experts writing concurrent collections)

Synchronized Methods

```
// synchronized method: locks on "this" object
public synchronized type name(parameters) { ... }

// synchronized static method: locks on the given class
public static synchronized type name(parameters) { ... }
```

A synchronized method grabs the object or class's lock at the start, runs to completion, then releases the lock

Useful for methods whose *entire* bodies are **critical sections** (recall Alice and Bob's farm), and thus should not be entered by multiple threads at the same time.

I.e. a synchronized method is a critical section with guaranteed **mutual exclusion**.

Synchronized Blocks

```
// synchronized block: uses the given object as a lock
synchronized (object) {
    statement(s); // critical sections
}
```

A synchronized method, e.g.

```
public synchronized void inc(long delta) {
    this.value += delta;
}
```

is syntactic sugar for

```
public void inc(long delta) {
    synchronized (this) {
        this.value += delta;
    }
}
```

Synchronized Blocks

```
// synchronized block: uses the given object as a lock
synchronized (object) {
    statement(s); // critical sections
}
```

Enforces **mutual exclusion w.r.t to some object**

Every Java object can *act* as a lock for concurrency:

A thread T_1 can ask to run a block of code, synchronized on a given object O .

- If no other thread has locked O , then T_1 locks the object and proceeds.
- If another thread T_2 has already locked O , then T_1 becomes blocked and must wait until T_2 is finished with O (that is, unlocks O). Then, T_1 is woken up, and can proceed.

Preview: Locks

In Java, all objects have an *internal* lock, called intrinsic lock or monitor lock, which are used to implement synchronized

Java also offers external locks (e.g. in package `java.util.concurrent.locks`)

- Less easy to use
- But support more sophisticated locking idioms, e.g. for reader-writer scenarios

Locks Are Recursive (Reentrant)

A thread can request to lock an object it has already locked

```
public class Foo {  
    public void synchronized f() { ... }  
    public void synchronized g() { ... f(); ... }  
}  
  
Foo foo = new Foo();  
synchronized(foo) { ... synchronized(foo) { ... } ... }
```

Examples: Synchronization granularity

```
public class SynchronizedCounter {  
    private int c = 0;  
    public synchronized void increment() { c++; }  
    public synchronized void decrement() { c--; }  
    public synchronized int value() { return c; }  
}
```

```
public void addName(String name) { synchronized(this) {  
    lastName = name;  
    nameCount++;  
}  
    nameList.add(name); // add synchronizes on nameList  
}
```

The advantage of not synchronizing the entire method is **efficiency** but need to be careful with correctness

Examples: Synchronization with different locks

```
public class TwoCounters {  
    private long c1 = 0, c2 = 0;  
    private Object lock1 = new Object();  
    private Object lock2 = new Object();  
    public void inc1() {  
        synchronized(lock1) {  
            c1++;  
        }  
    }  
    public void inc2() {  
        synchronized(lock2) {  
            c2++;  
        }  
    }  
}
```

The locks are disjoint – allows for more concurrency.

Examples: Synchronization with static methods

```
public class Screen {  
    private static Screen theScreen;  
  
    private Screen() {...}  
  
    public static synchronized getScreen() {  
        if (theScreen == null) {  
            theScreen = new Screen();  
        }  
  
        return theScreen;  
    }  
}
```

Which object does synchronized lock here?
What if Screen instances call getScreen()?

Interleavings: Examples

Suppose we have 2 threads, T1 and T2, both incrementing a shared counter. If we use **synchronized** (say on 'this' object), we will get the desired result of 2 by the time both threads have finished executing their code below.

```

                                c = 0

T1:                                T2:

synchronized (this)            synchronized (this)
{
  1:  t1 = c;
  2:  t1++
  3:  c = t1;
}
                                {
  4:  t2 = c;
  5:  t2++
  6:  c = t2;
}
```

For convenience, we use labels 1-6 to refer to the instructions. The possible interleavings / executions of this program are that either T1 runs before T2 or vice versa. So we will have:

Interleaving 1: 123456

Interleaving 2: 456123

Interleavings: Another example

Suppose the programmer **forgot to use synchronized** in thread T2. What is an example of an undesirable interleaving that we can see?

```

                                c = 0

T1:                                T2:

synchronized (this)
{
  1:  t1 = c;
  2:  t1++
  3:  c = t1;
}
                                4: t2 = c;
                                5: t2++
                                6: c = t2;
```

A possibly **bad interleaving** is: 4 1 2 3 5 6

This interleaving will result in the counter 'c' being set to 1 at the end of the interleaving.

Interleavings: Another example

Suppose the programmer now uses `synchronized` in thread T2 but not on 'this', but say another object 'p'. Does this prevent the bad interleaving we just saw?

```
                                c = 0

T1:                                T2:

synchronized (this)           synchronized (p)
{
  1:  t1 = c;
  2:  t1++
  3:  c = t1;
}
{
  4:  t2 = c;
  5:  t2++
  6:  c = t2;
}
```

No, the bad interleaving: 4 1 2 3 5 6 can still happen because 'p' and 'this' are different objects.