

# Concurrent programming


## Concurrent objects

Companion slides for

The Art of Multiprocessor Programming  
by Maurice Herlihy, Nir Shavit, Victor Luchangco,  
and Michael Spear

Modified by Piotr Witkowski

# Concurrent Computation




# Objectivism

- What is a concurrent object?
  - How do we **describe** one?
  - How do we **implement** one?
  - How do we **tell if we're right**?


# Objectivism

- What is a concurrent object?
  - How do we **describe** one?
  - How do we **tell if we're right**?


# FIFO Queue: Enqueue Method




# FIFO Queue: Dequeue Method



# Lock-Based Queue



# Lock-Based Queue






# A Lock-Based Queue

```
class LockBasedQueue<T> {  
    int head, tail;  
    T[] items;  
    Lock lock;  
    public LockBasedQueue(int capacity) {  
        head = 0; tail = 0;  
        lock = new ReentrantLock();  
        items = (T[]) new Object[capacity];  
    }  
}
```

# A Lock-Based Queue


```
class LockBasedQueue<T> {  
    int head, tail;  
    T[] items;  
    Lock lock;  
    public LockBasedQueue(int capacity) {  
        head = 0, tail = 0;  
        lock = new ReentrantLock();  
        items = (T[]) new Object[capacity];  
    }  
}
```



Fields protected by  
single shared lock

# Lock-Based Queue


Initially: **head = tail**



# A Lock-Based Queue


```
class LockBasedQueue<T> {  
    int head, tail;  
    T[] items;  
    Lock lock;
```

```
    public LockBasedQueue(int capacity) {  
        head = 0; tail = 0;  
        lock = new ReentrantLock();  
        items = (T[]) new Object[capacity];  
    }
```




Initially head = tail

# Lock-Based `deq()`




# Acquire Lock




# Implementation: `deq()`

```
public T deq() throws EmptyException {  
    lock.lock();  
    try {  
        if (tail == head)  
            throw new EmptyException();  
        T x = items[head % items.length];  
        head++;  
        return x;  
    } finally {  
        lock.unlock();  
    }  
}
```

Acquire lock at  
method start



# Check if Non-Empty







# Implementation: `deq()`

```
public T deq() throws EmptyException {  
    lock.lock();  
    try {  
        if (tail == head)  
            throw new EmptyException();  
        T x = items[head % items.length];  
        head++;  
        return x;  
    } finally {  
        lock.unlock();  
    }  
}
```

If queue empty  
throw exception




# Modify the Queue



# Implementation: `deq()`

```
public T deq() throws EmptyException {  
    lock.lock();  
    try {  
        if (tail == head)  
            throw new EmptyException();  
        T x = items[head % items.length];  
        head++;  
        return x;  
    } finally {  
        lock.unlock();  
    }  
}
```




Queue not empty?  
Remove item and update head

# Implementation: `deq()`


```
public T deq() throws EmptyException {  
    lock.lock();  
    try {  
        if (tail == head)  
            throw new EmptyException();  
        T x = items[head % items.length];  
        head++;  
        return x;  
    } finally {  
        lock.unlock();  
    }  
}
```

**return x;**


Return result



# Release the Lock




# Release the Lock



# Implementation: `deq()`

```
public T deq() throws EmptyException {  
    lock.lock();  
    try {  
        if (tail == head)  
            throw new EmptyException();  
        T x = items[head % items.length];  
        head++;  
        return x;  
    } finally {  
        lock.unlock();  
    }  
}
```



Release lock no  
matter what!

# Implementation: `deq()`

```
public T deq() throws EmptyException {  
    lock.lock();  
    try {  
        if (tail == head)  
            throw new EmptyException();  
        T x = items[head % items.length];  
        head++;  
        return x;  
    } finally {  
        lock.unlock();  
    }  
}
```


Should be correct because  
modifications are mutually exclusive...




# Now consider the following implementation

- The same thing without mutual exclusion
- For simplicity, only two threads
  - One thread **enq only**
  - The other **deq only**


# Wait-free 2-Thread Queue




# Wait-free 2-Thread Queue



# Wait-free 2-Thread Queue




# Wait-free 2-Thread Queue



# Wait-free 2-Thread Queue

```
public class WaitFreeQueue {  
  
    int head = 0, tail = 0;  
    items = (T[]) new Object[capacity];  
  
    public void enq(Item x) {  
        if (tail-head == capacity) throw  
            new FullException();  
        items[tail % capacity] = x; tail++;  
    }  
  
    public Item deq() {  
        if (tail == head) throw  
            new EmptyException();  
        Item item = items[head % capacity]; head++;  
        return item;  
    }  
}
```



**No lock needed**


# Wait-free 2-Thread Queue

```
public T deq() throws EmptyException {  
    lock.lock();  
    try {  
        if (tail == head)  
            throw new EmptyException();  
        T x = items[head];  
        head++;  
        return x;  
    } finally {  
        lock.unlock();  
    }  
}
```

How do we define “correct” when  
modifications are not mutually  
exclusive?

# What *is* a Concurrent Queue?

- Need a way to specify a concurrent queue object
- Need a way to prove that an algorithm implements the object's specification
- Lets talk about object specifications ...






# Correctness and Progress


- In a concurrent setting, we need to specify both the safety and the liveness properties of an object
- Need a way to define
  - when an implementation is correct
  - the conditions under which it guarantees progress

**Lets begin with correctness**




# Sequential Objects

- Each object has a ***state***
  - Usually given by a set of ***fields***
  - Queue example: sequence of items
- Each object has a set of ***methods***
  - Only way to manipulate state
  - Queue example: **enq** and **deq** methods




# Sequential Specifications

- If (precondition)
  - the object is in such-and-such a state
  - before you call the method,
- Then (postcondition)
  - the method will return a particular value
  - or throw a particular exception.
- and (postcondition, con't)
  - the object will be in some other state
  - when the method returns,




# Pre and PostConditions for Dequeue

- Precondition:
  - Queue is non-empty
- Postcondition:
  - Returns first item in queue
- Postcondition:
  - Removes first item in queue




# Pre and PostConditions for Dequeue

- Precondition:
  - Queue is empty
- Postcondition:
  - Throws Empty exception
- Postcondition:
  - Queue state unchanged




# Why Sequential Specifications Totally Rock

- Interactions among methods captured by side-effects on object state
  - State meaningful between method calls
- Documentation size linear in number of methods
  - Each method described in isolation
- Can add new methods
  - Without changing descriptions of old methods





# What About Concurrent Specifications ?

- Methods?
- Documentation?
- Adding new methods?





# Methods Take Time

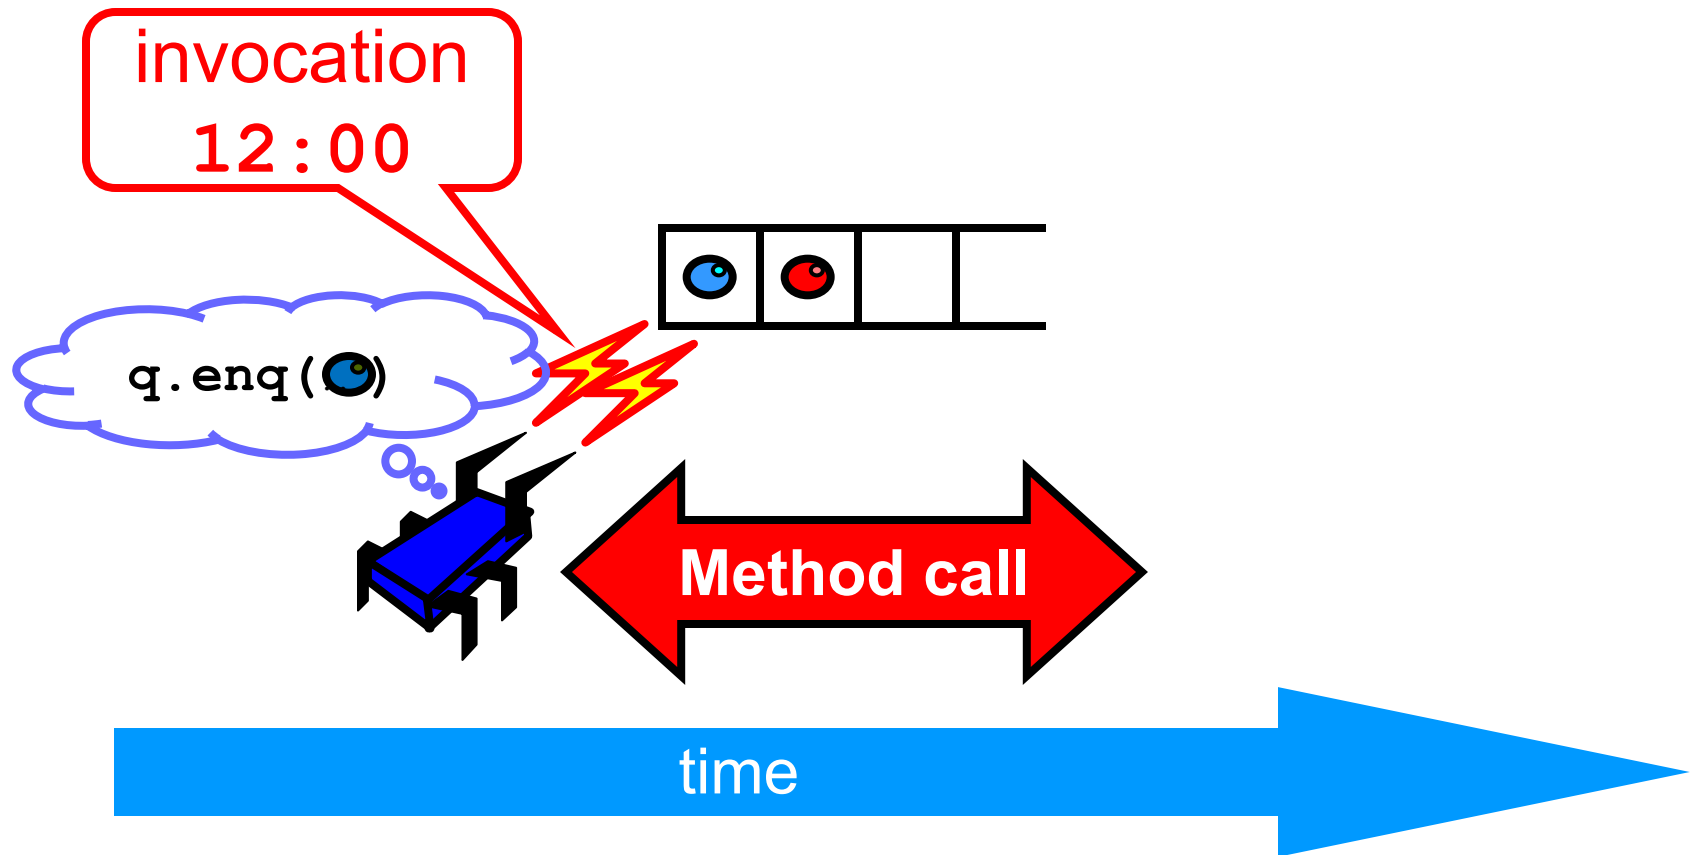





# Methods Take Time




# Methods Take Time



# Methods Take Time




# Methods Take Time





# Sequential vs Concurrent


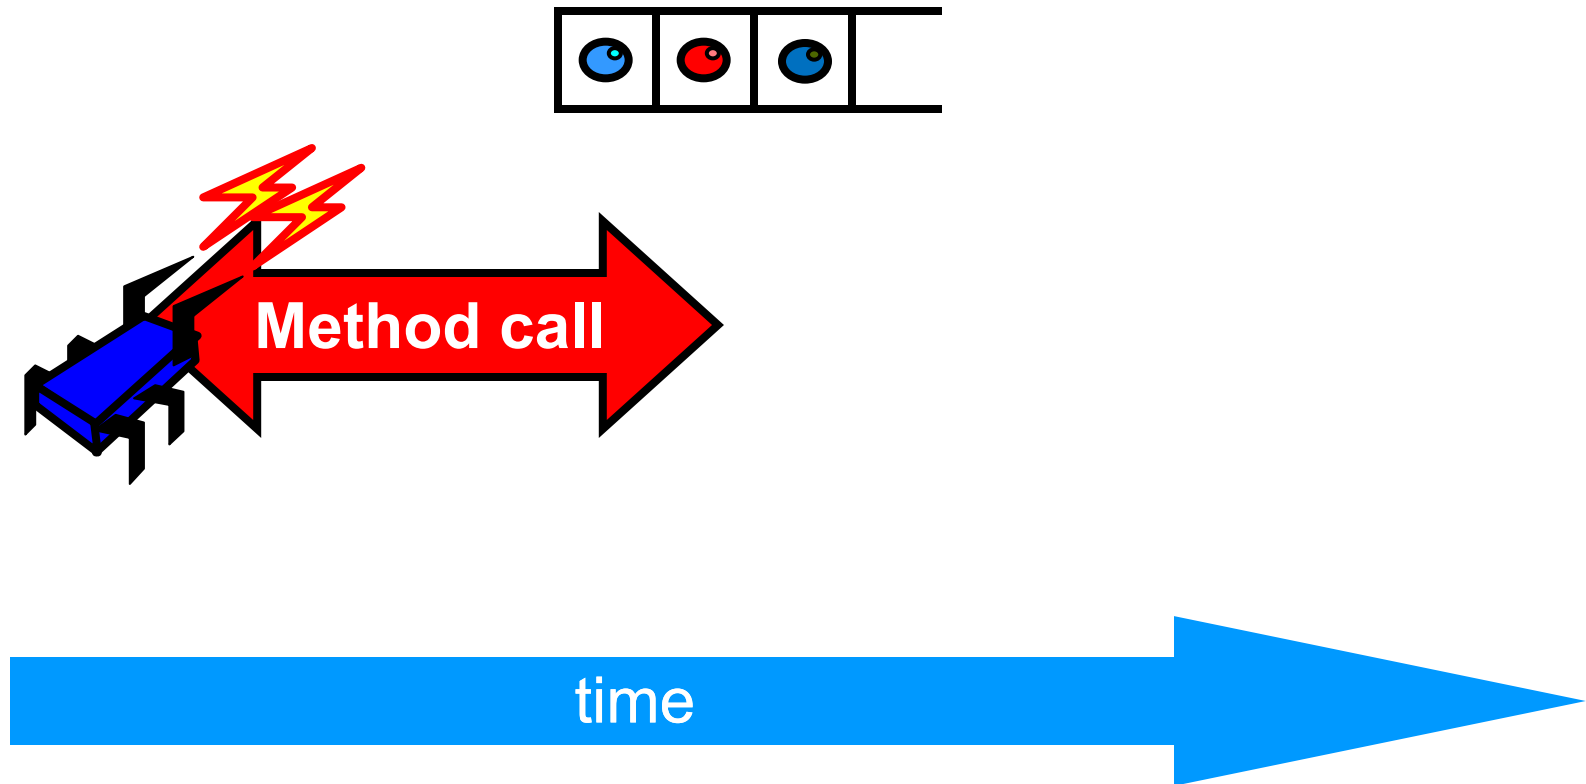
- Sequential
  - Methods take time? Who knew?
- Concurrent
  - Method call is not an event
  - Method call is an interval.




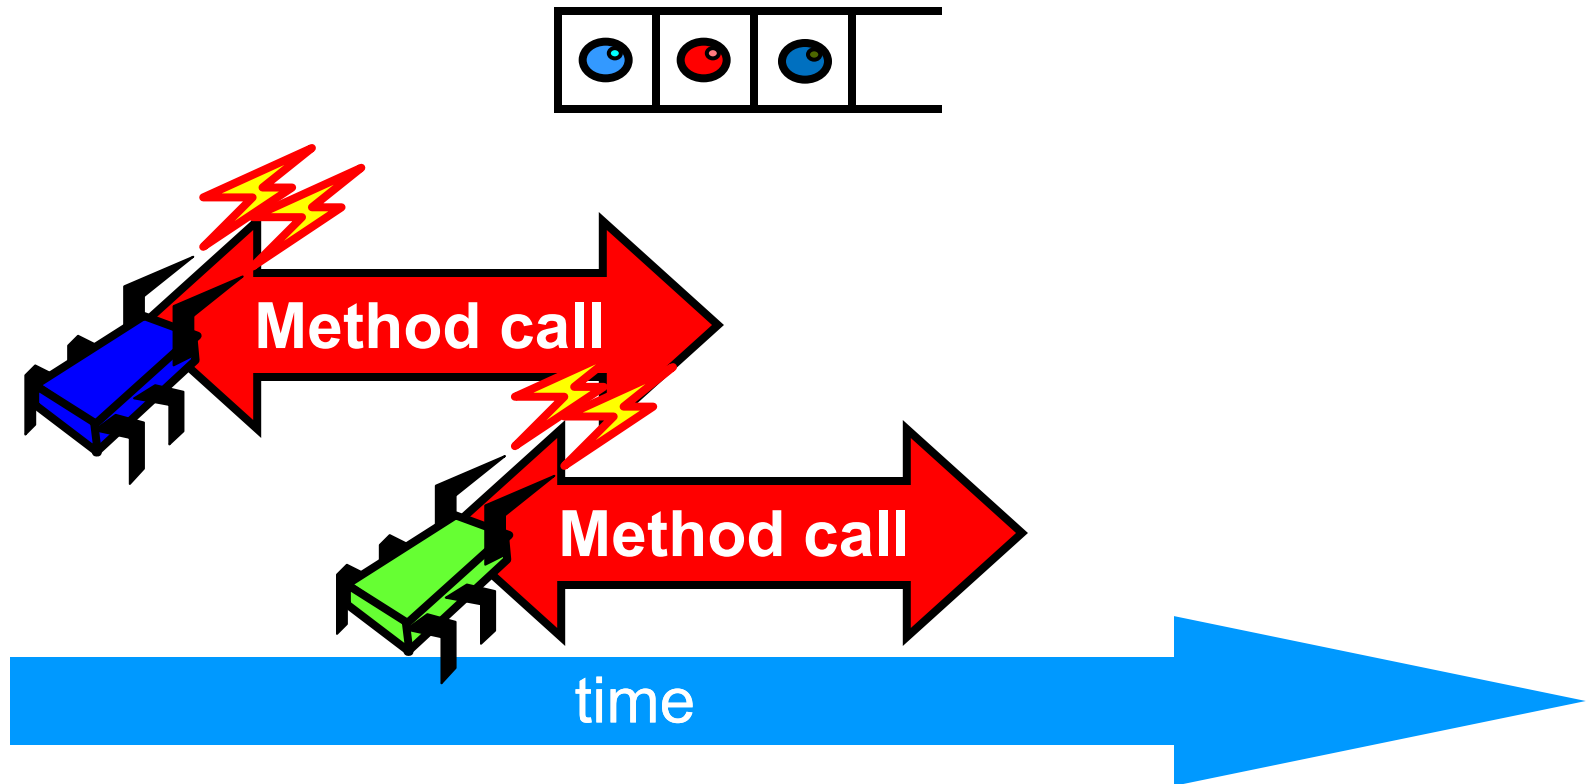
# Concurrent Methods Take Overlapping Time



# Concurrent Methods Take Overlapping Time


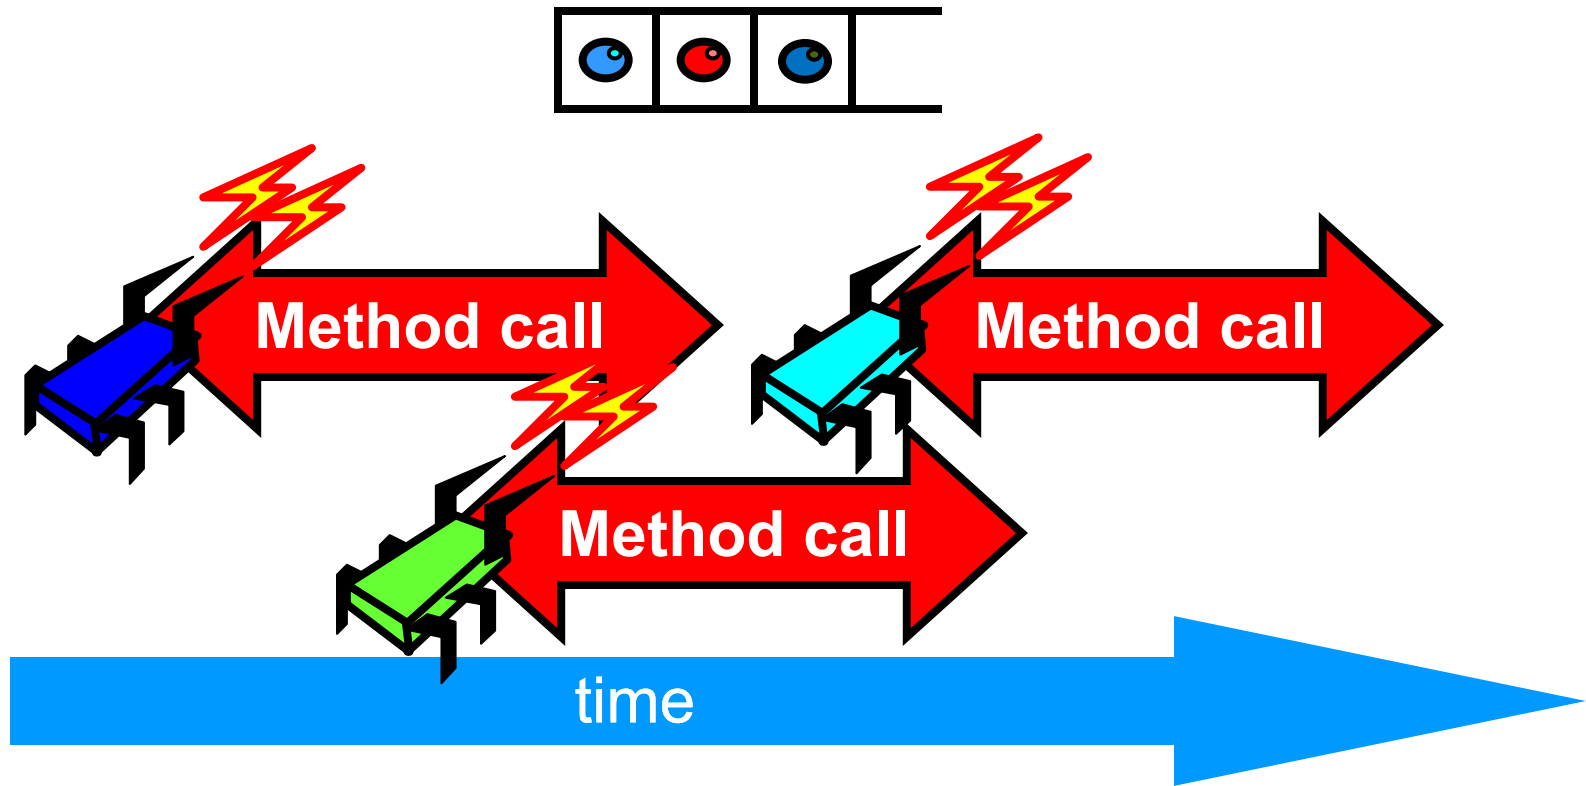


# Concurrent Methods Take Overlapping Time






# Concurrent Methods Take Overlapping Time




# Sequential vs Concurrent

- Sequential:
  - Object needs meaningful state only ***between*** method calls
- Concurrent
  - Because method calls overlap, object might ***never*** be between method calls




# Sequential vs Concurrent

- Sequential:
  - Each method described in isolation
- Concurrent
  - Must characterize **all** possible interactions with concurrent calls
    - What if two `enq()` calls overlap?
    - Two `deq()` calls? `enq()` and `deq()`? ...



# Sequential vs Concurrent


- Sequential:
  - Can add new methods without affecting older methods
- Concurrent:
  - Everything can potentially interact with everything else



# Sequential vs Concurrent


- Sequential:
  - Can add new methods without affecting older methods
- Concurrent:
  - Everything can potentially interact with everything else

**Panic!**




# The Big Question

- What does it **mean** for a *concurrent* **object** to be correct?
  - What *is* a concurrent FIFO queue?
  - FIFO means strict temporal order
  - Concurrent means ambiguous temporal order



# Intuitively...


```
public T deq() throws EmptyException {  
    lock.lock();  
    try {  
        if (tail == head)  
            throw new EmptyException();  
        T x = items[head % items.length];  
        head++;  
        return x;  
    } finally {  
        lock.unlock();  
    }  
}
```



# Intuitively...

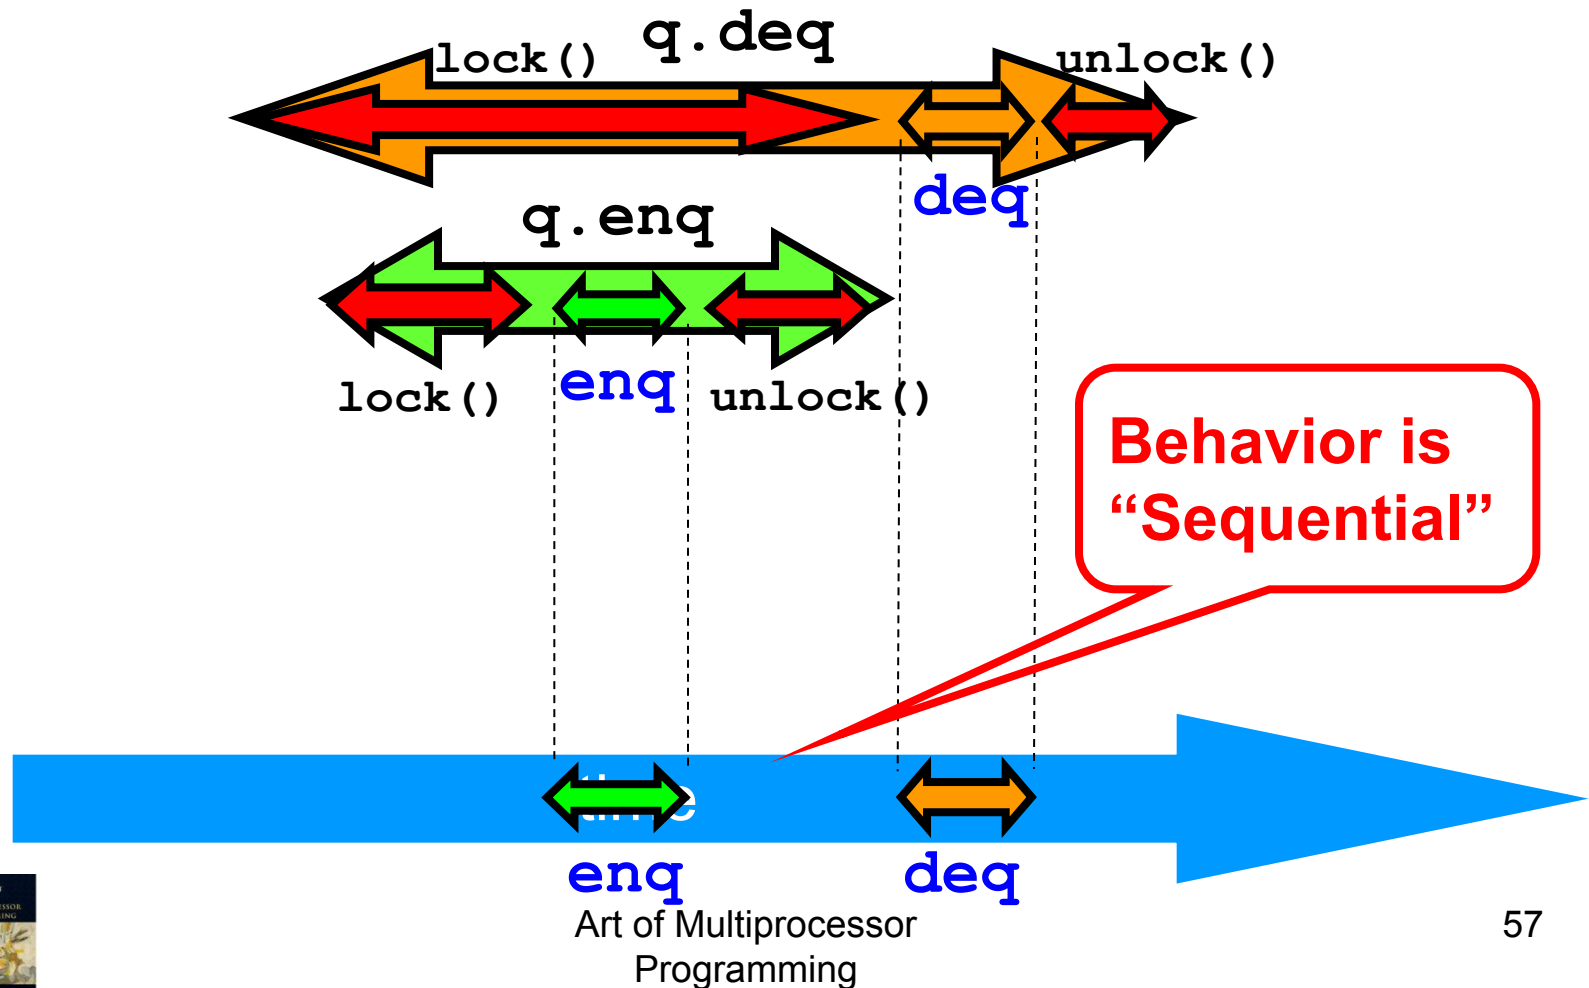
```
public T deg() throws EmptyException {  
    lock.lock();  
    try {  
        if (tail == head)  
            throw new EmptyException();  
        T x = items[head % items.length];  
        head++;  
        return x;  
    } finally {  
        lock.unlock();  
    }  
}
```

All queue modifications  
are mutually exclusive






Lets capture the idea of describing  
the concurrent via the sequential




# Linearizability

- Each method should
  - “take effect”
  - Instantaneously
  - Between invocation and response events
- Object is correct if this “sequential” behavior is correct
- Any such concurrent object is
  - **Linearizable<sup>TM</sup>**





# Is it really about the object?




- Each method should
  - “take effect”
  - Instantaneously
  - Between invocation and response events
- Sounds like a property of an execution...
- A linearizable object: one all of whose possible executions are linearizable






# Example




# Example



# Example


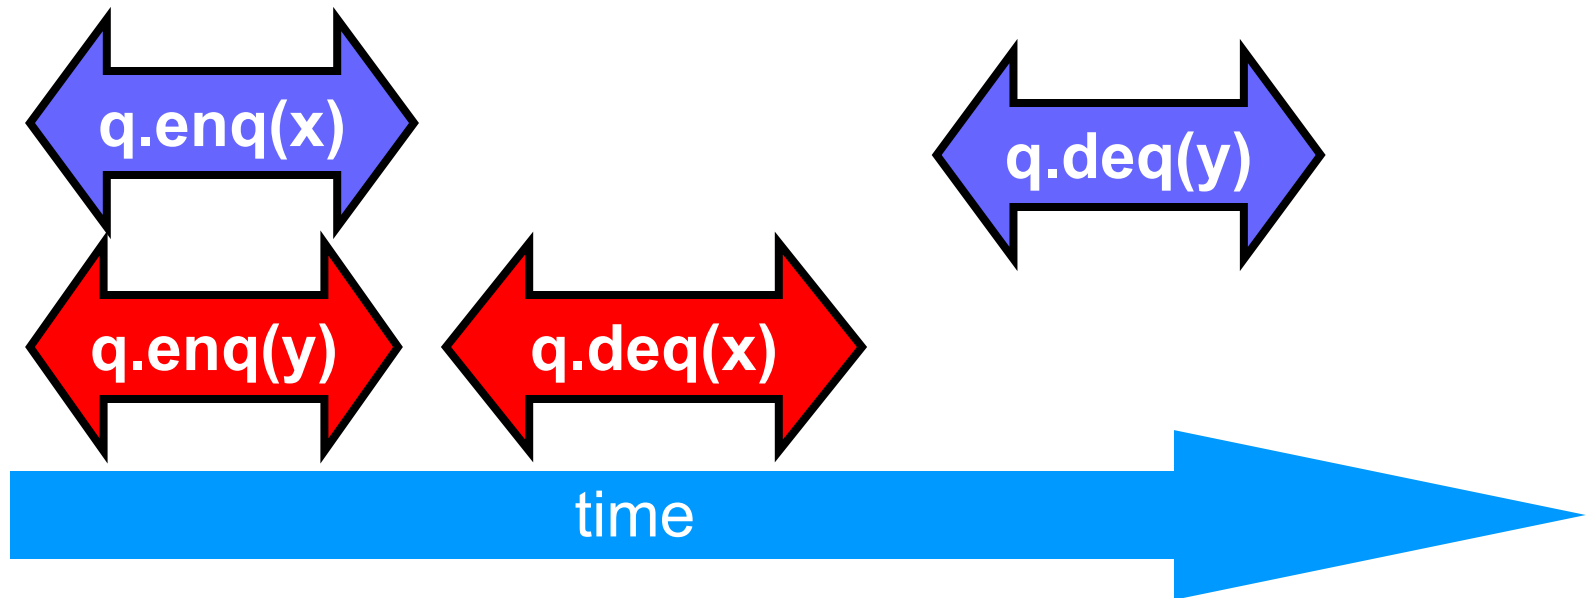



# Example



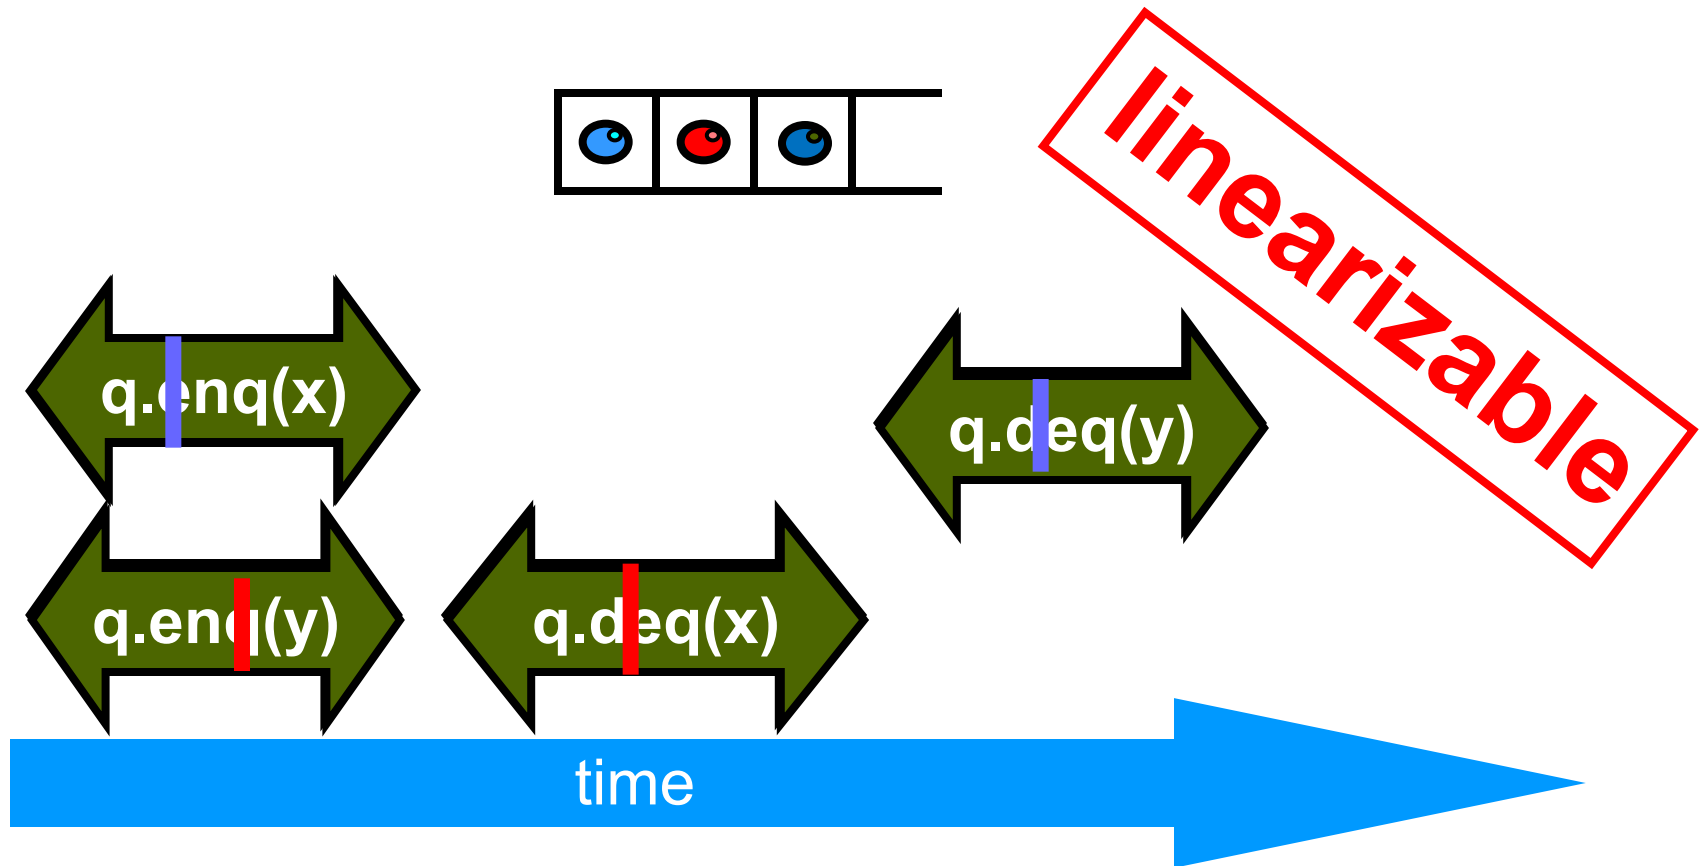


# Example

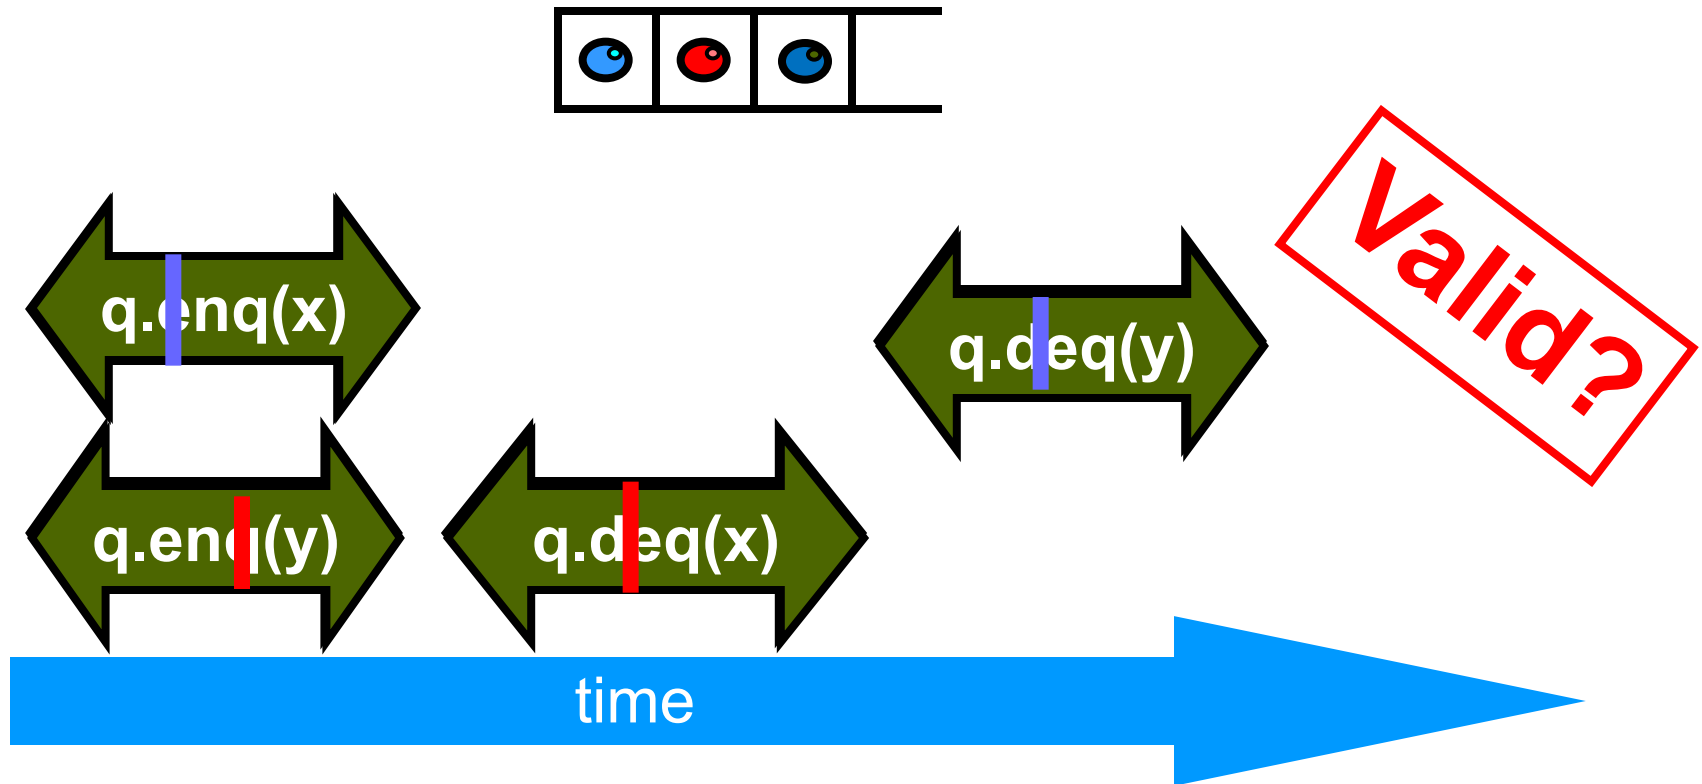






# Example






# Example







# Example



# Example









# Example










# Example






# Example









# Example





**not linearizable**








# Example


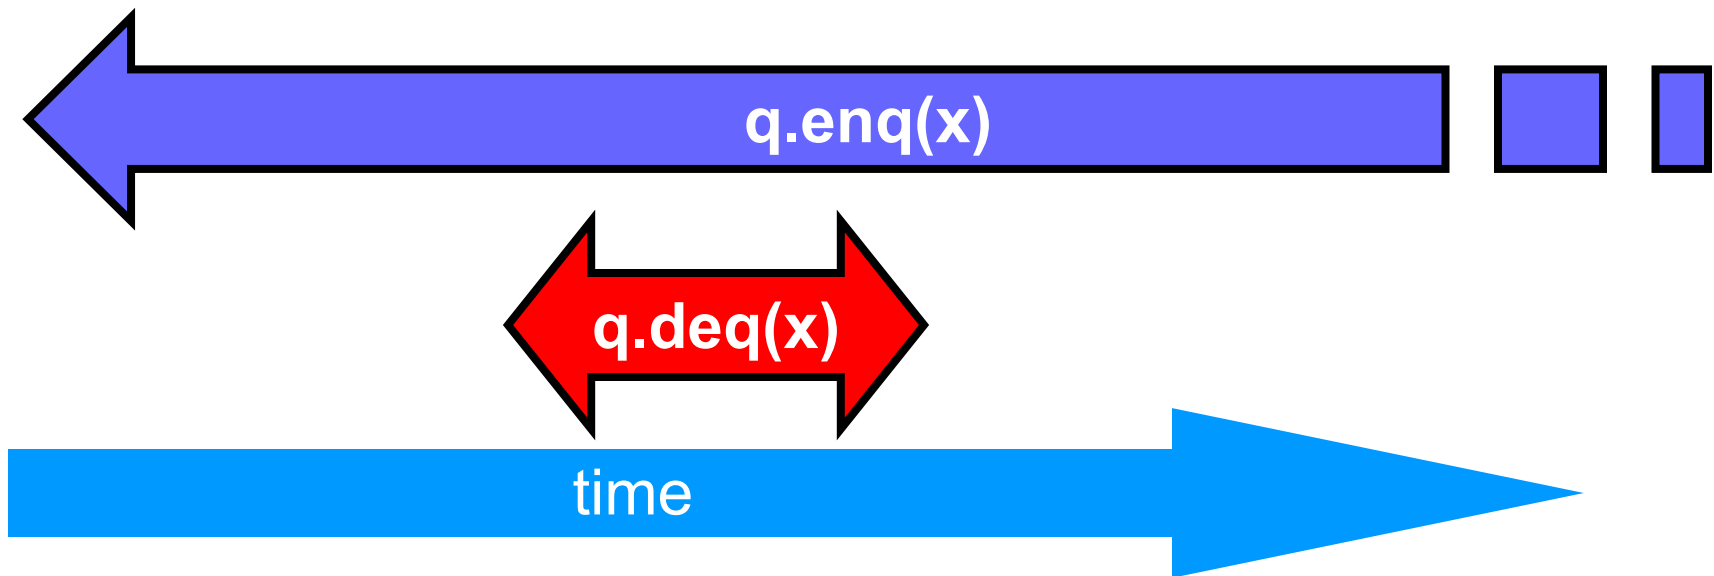



# Example




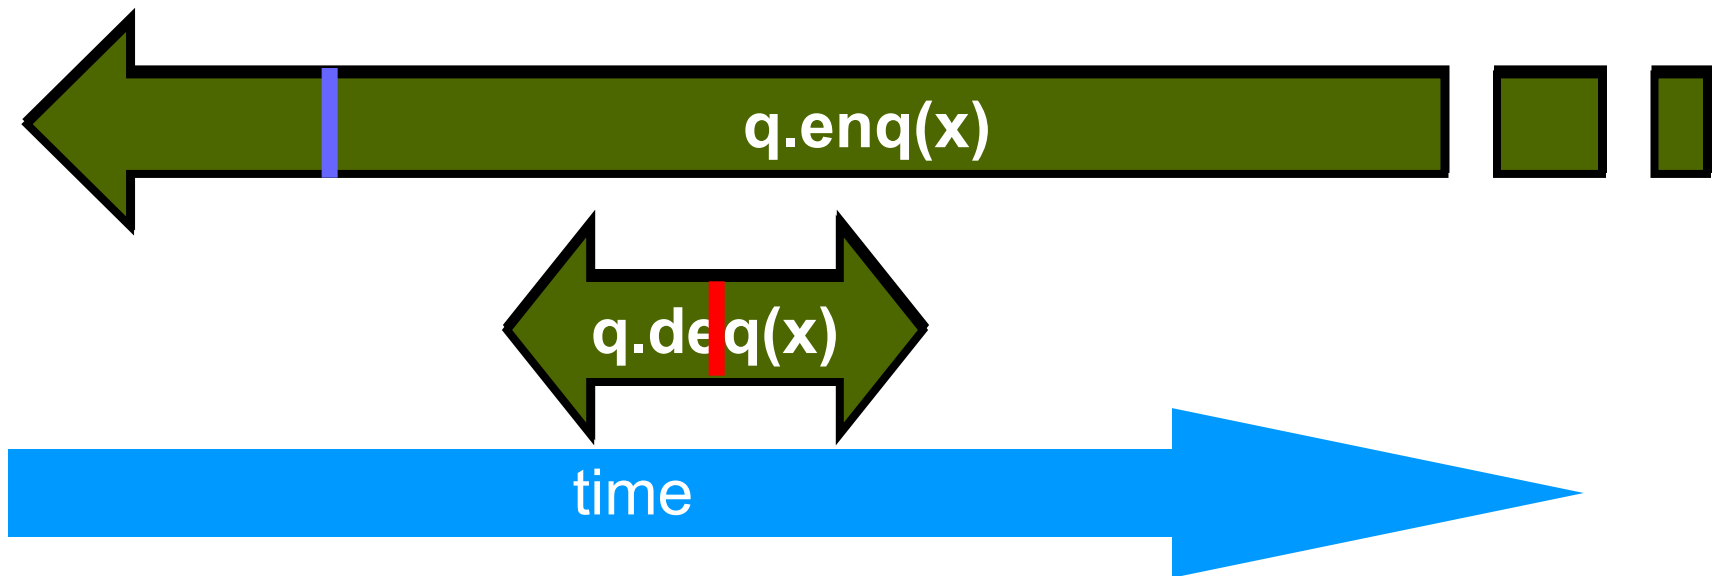



# Example






# Example


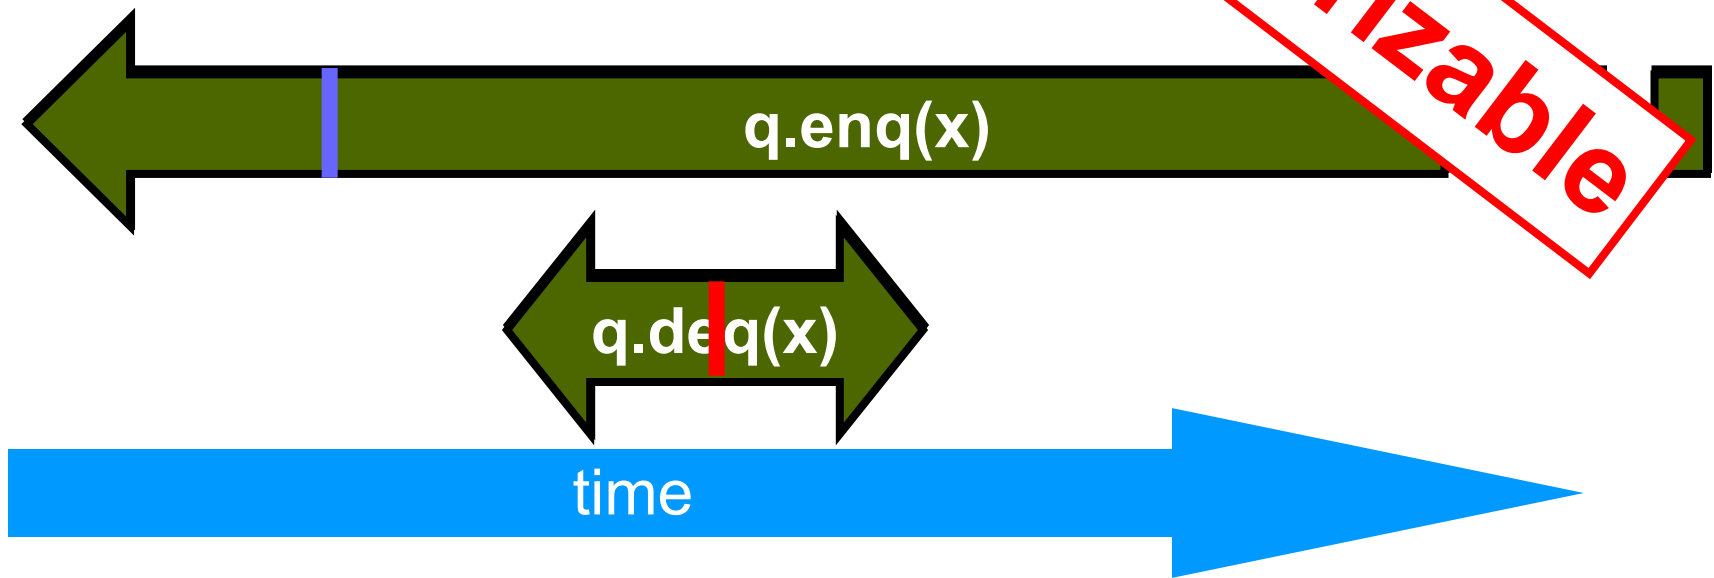







# Example






**linearizable**



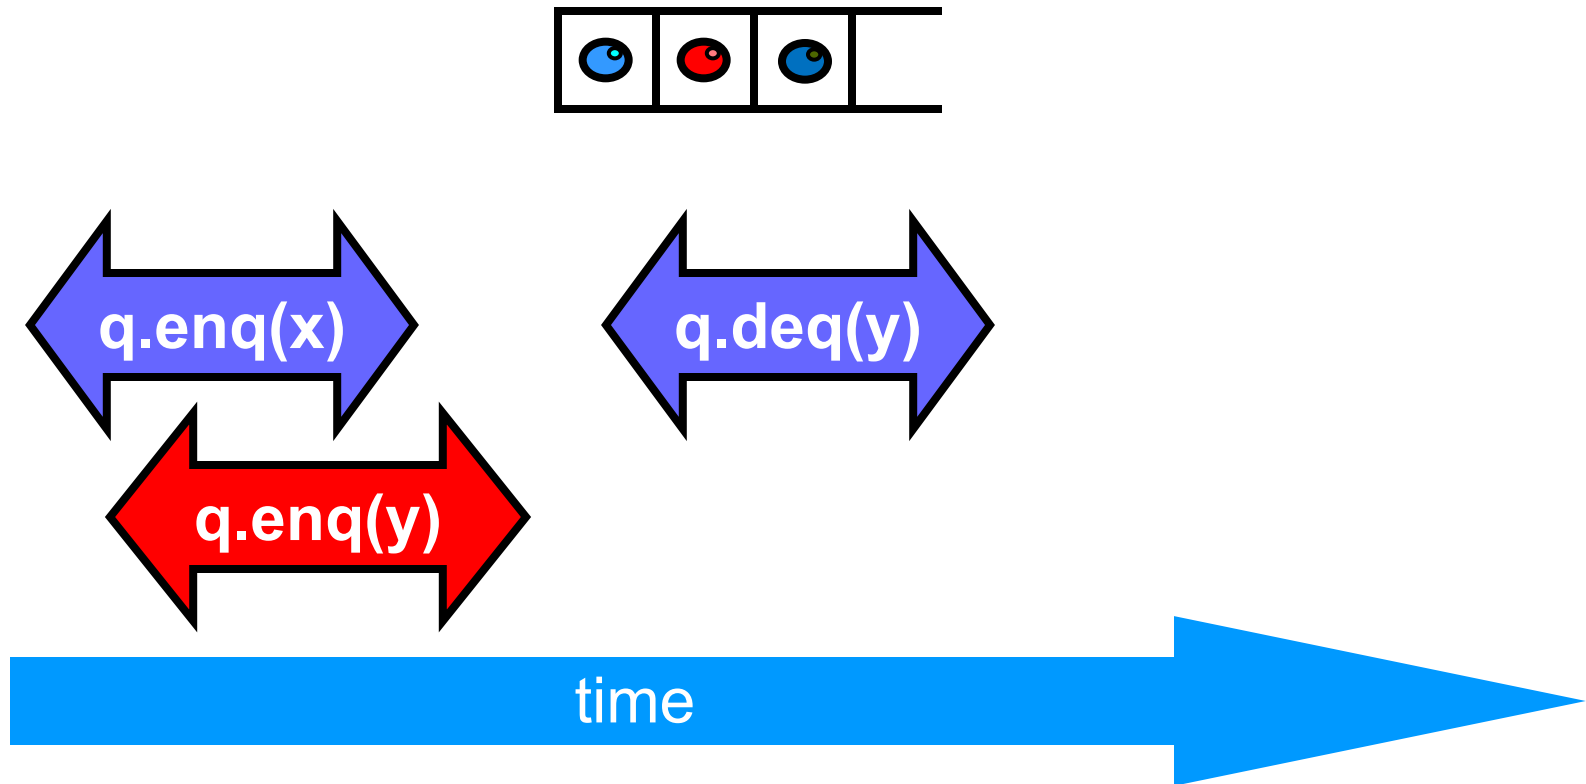
# Example



# Example




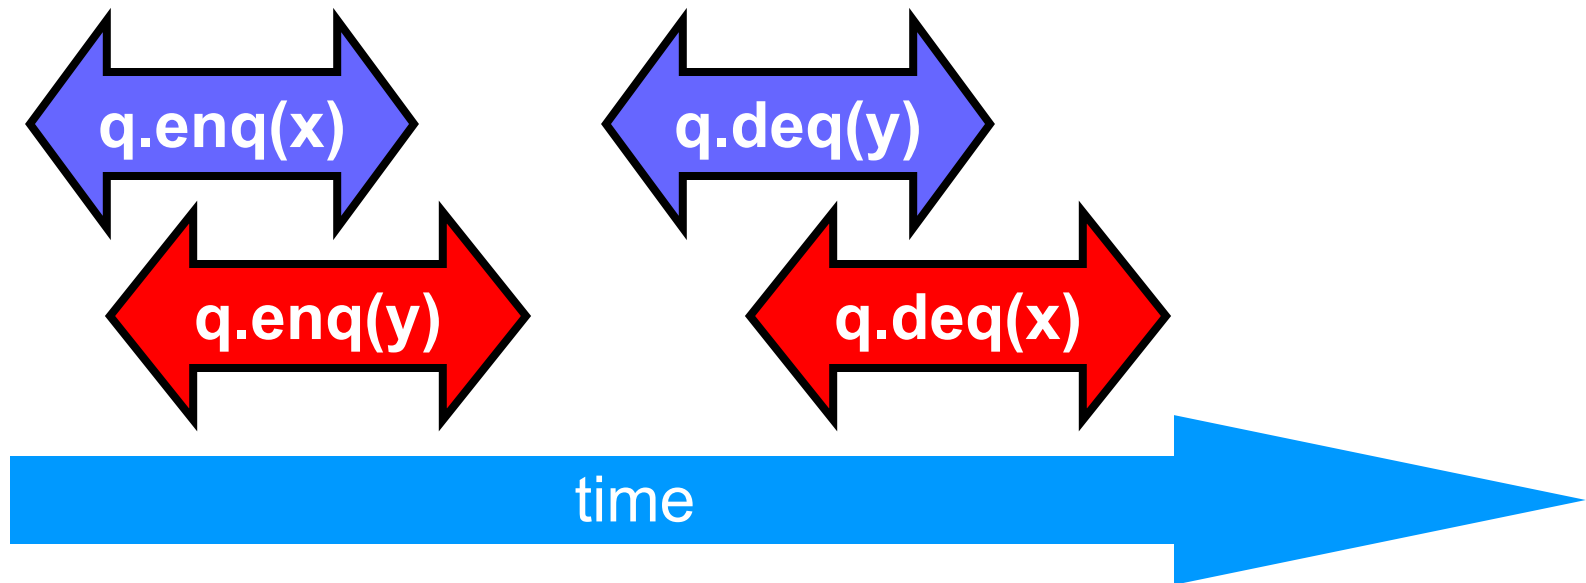

# Example





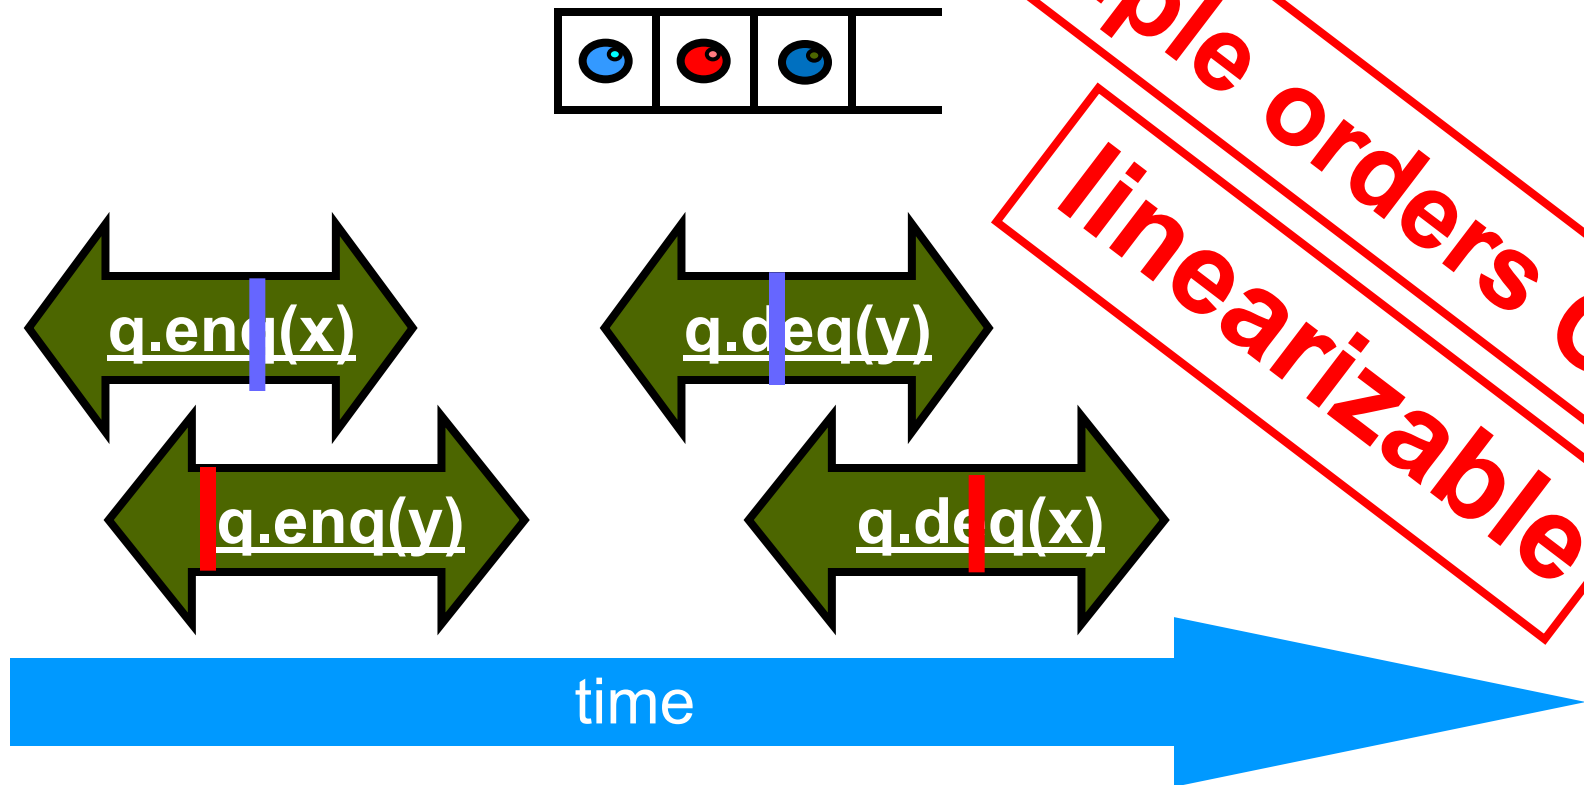


# Example




Comme ci  
Comme ça


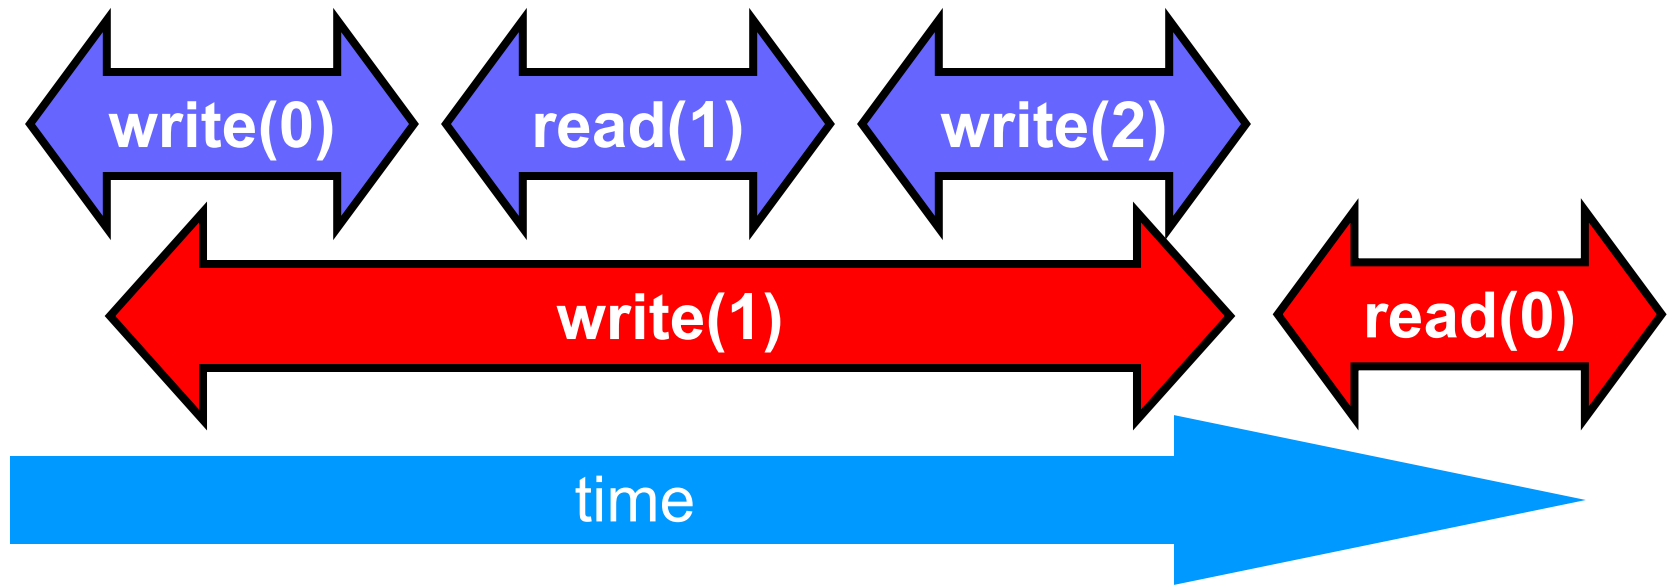
Example




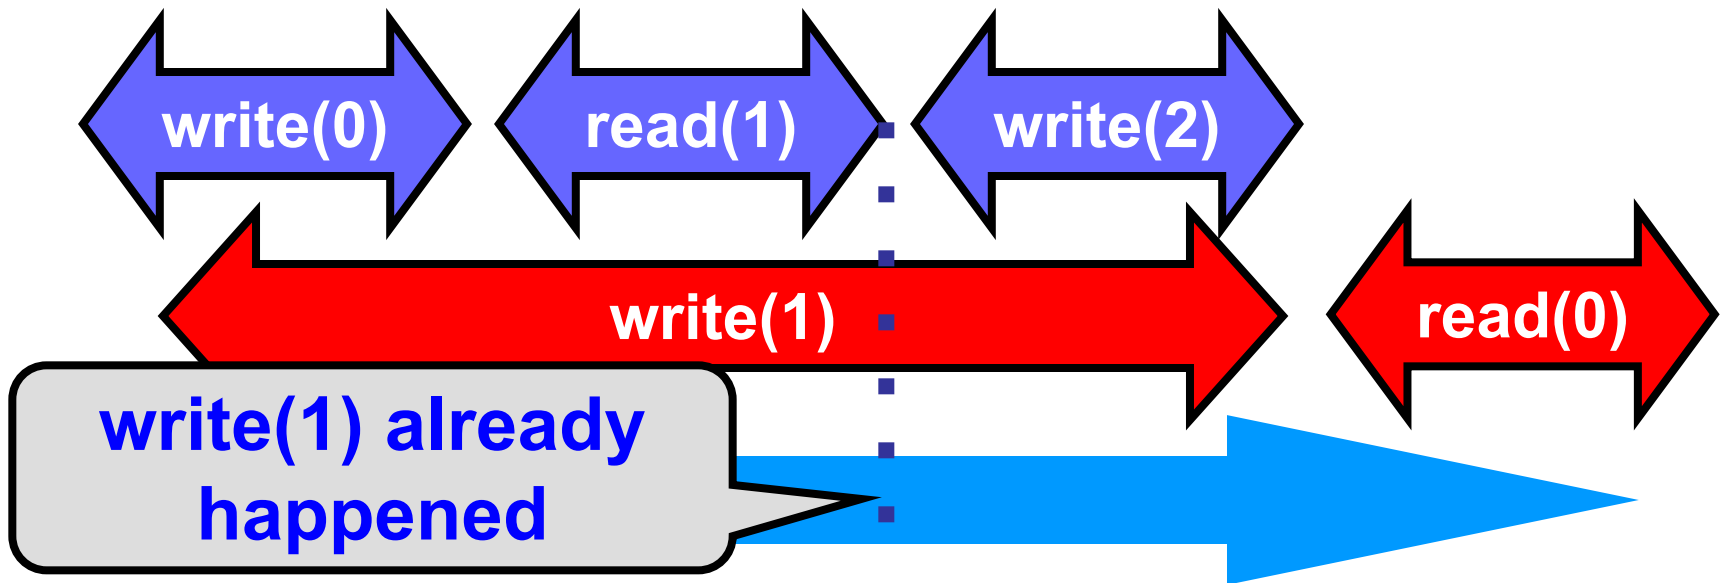
multiple orders OK  
linearizable




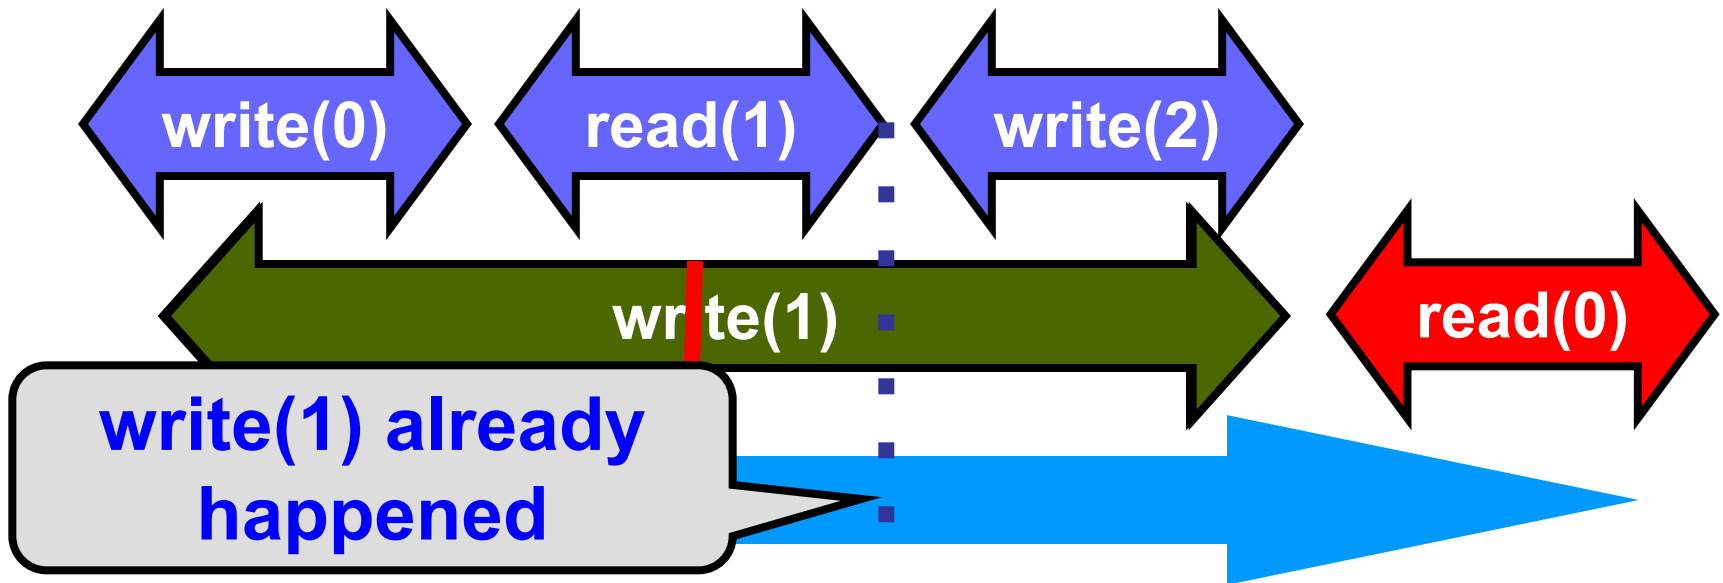
# Read/Write Register Example




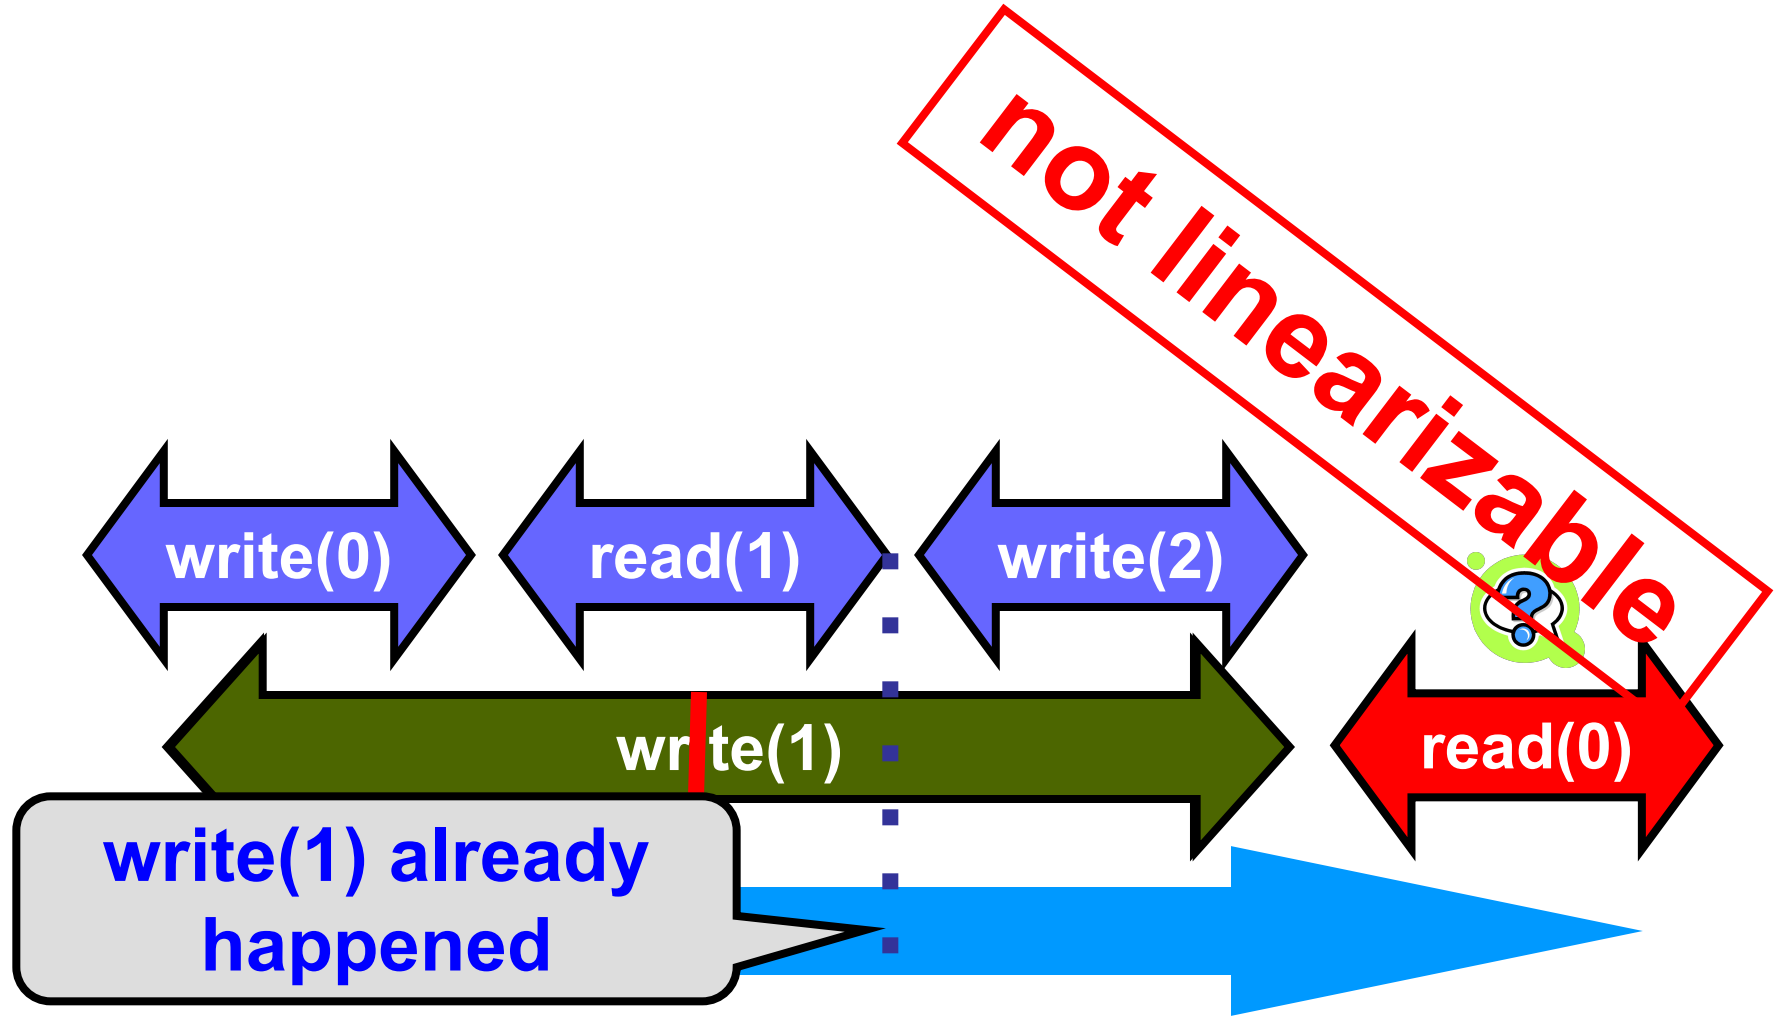
# Read/Write Register Example




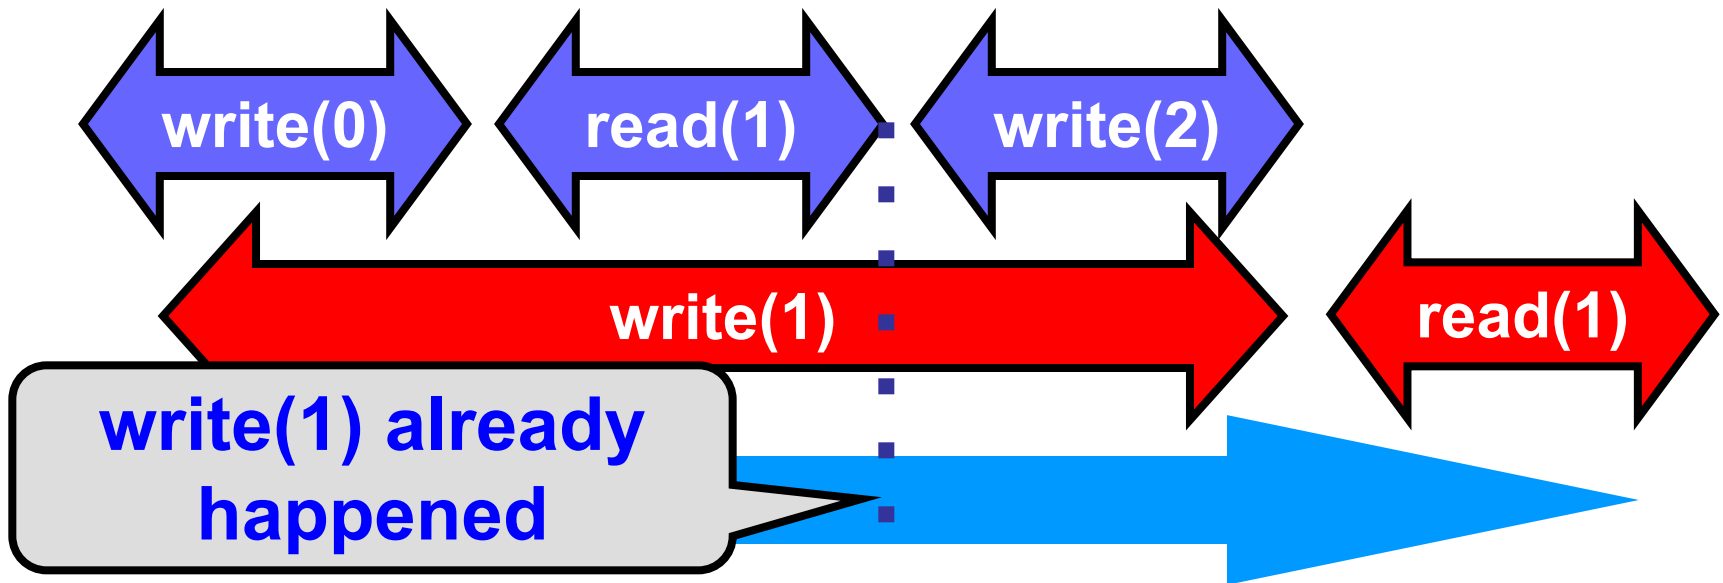
# Read/Write Register Example




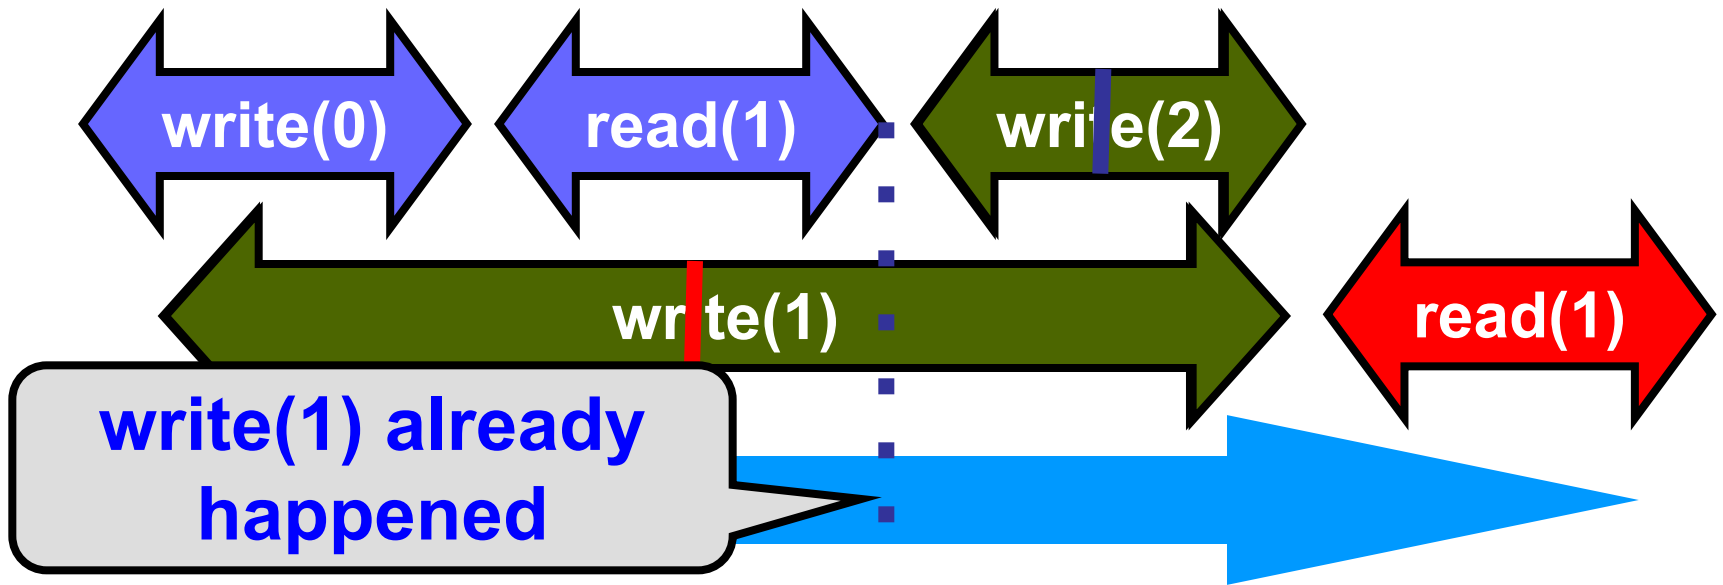
# Read/Write Register Example



# Read/Write Register Example


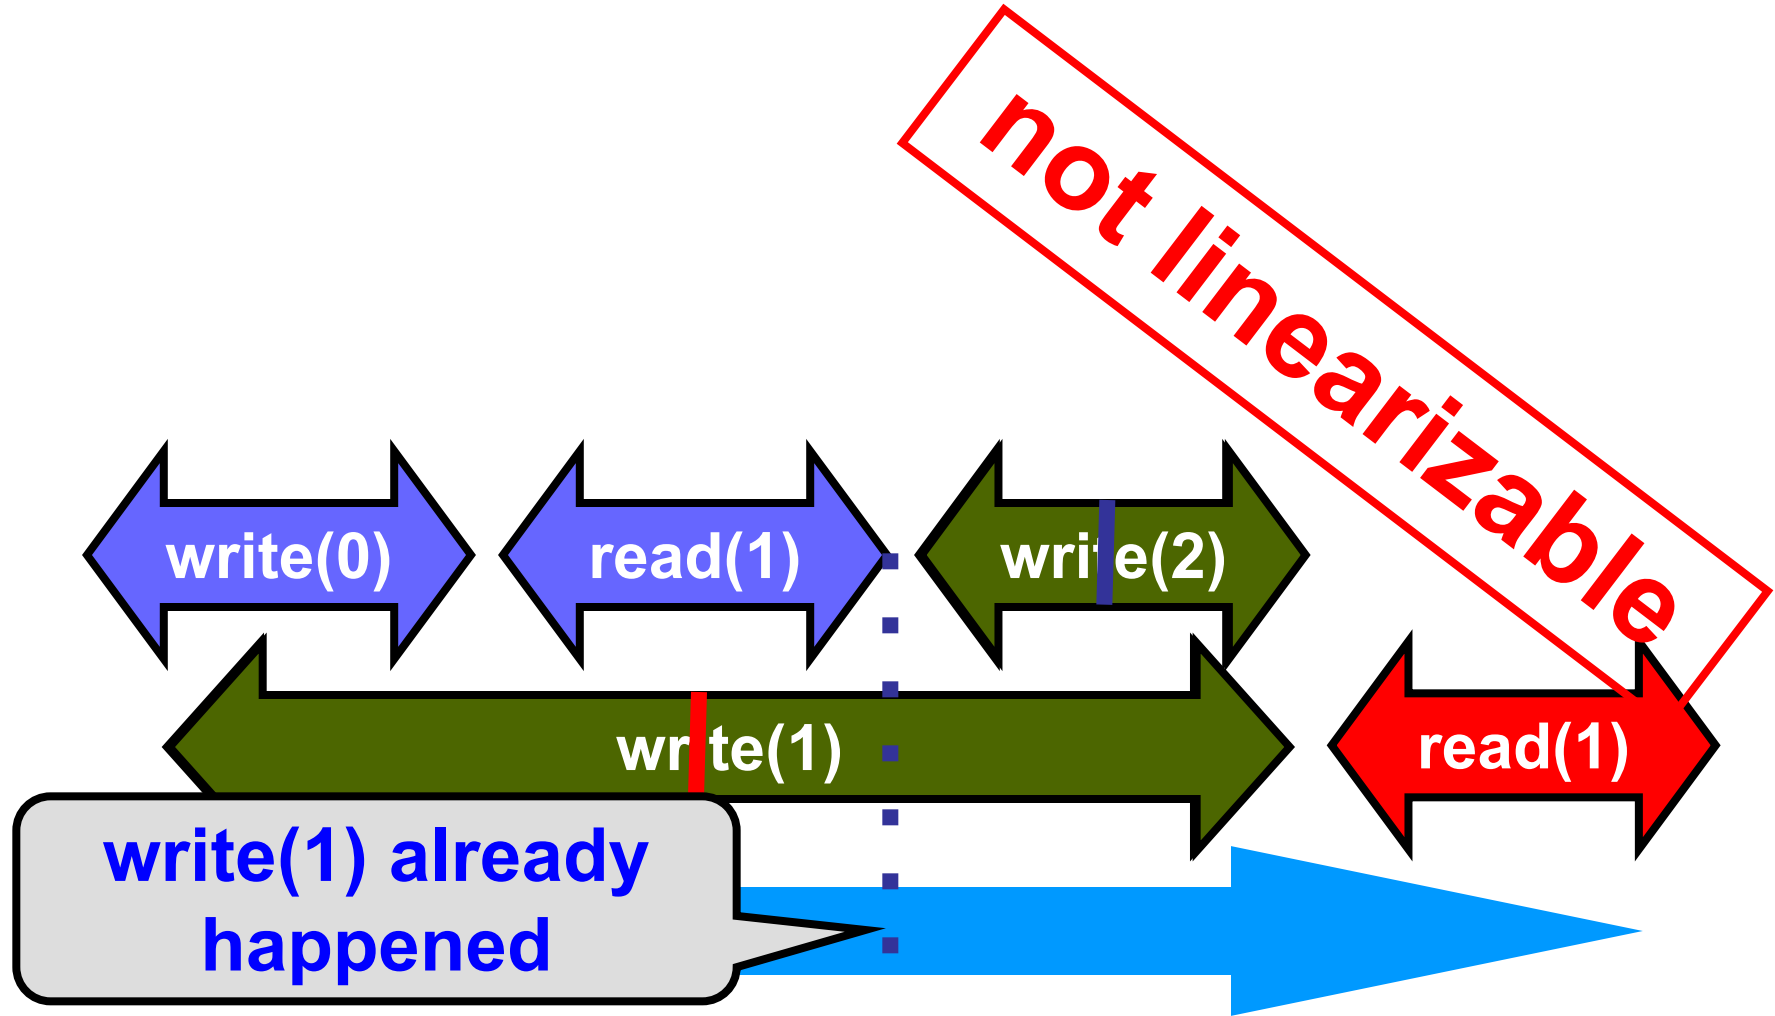


# Read/Write Register Example


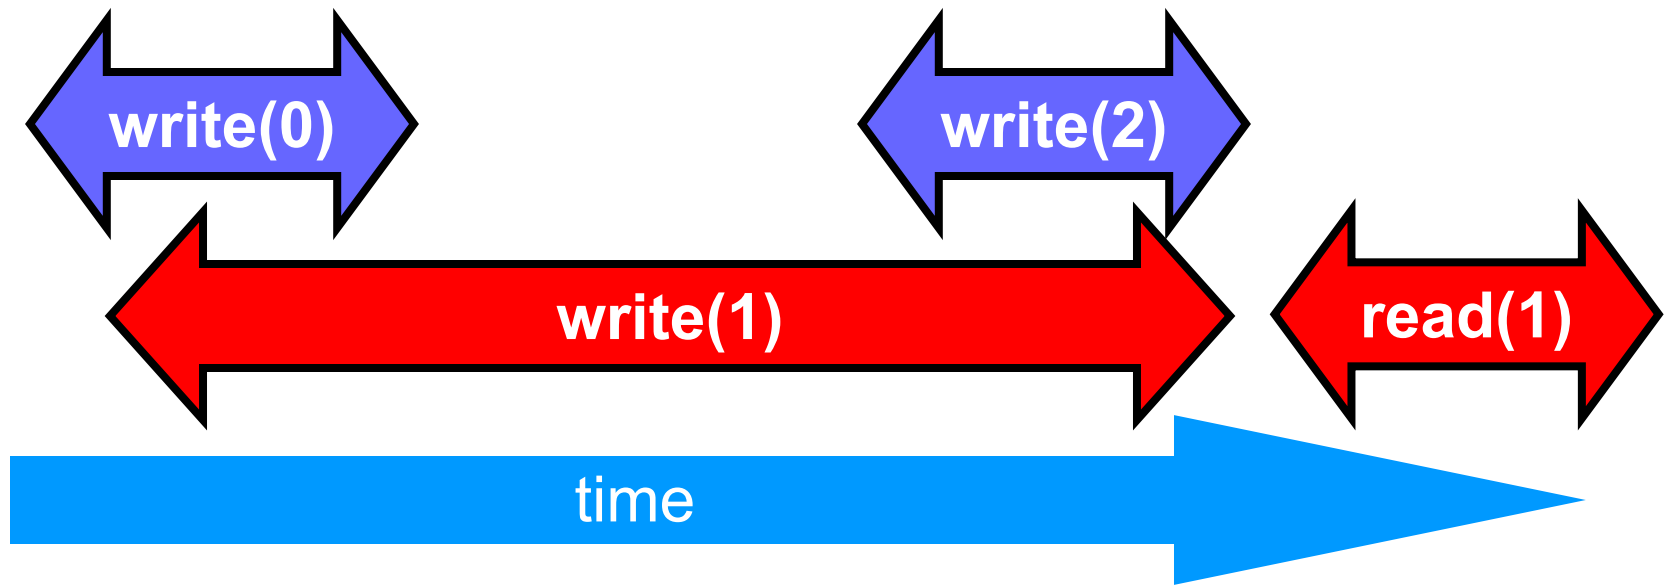





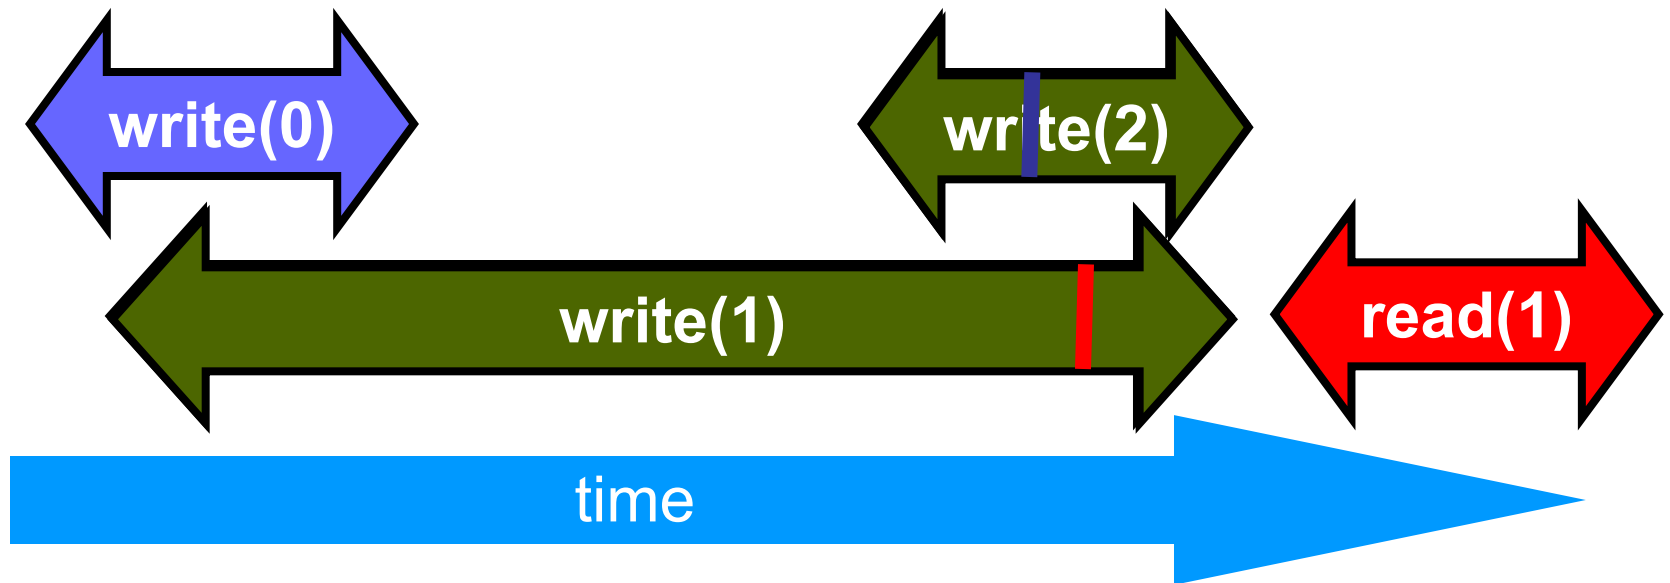
# Read/Write Register Example



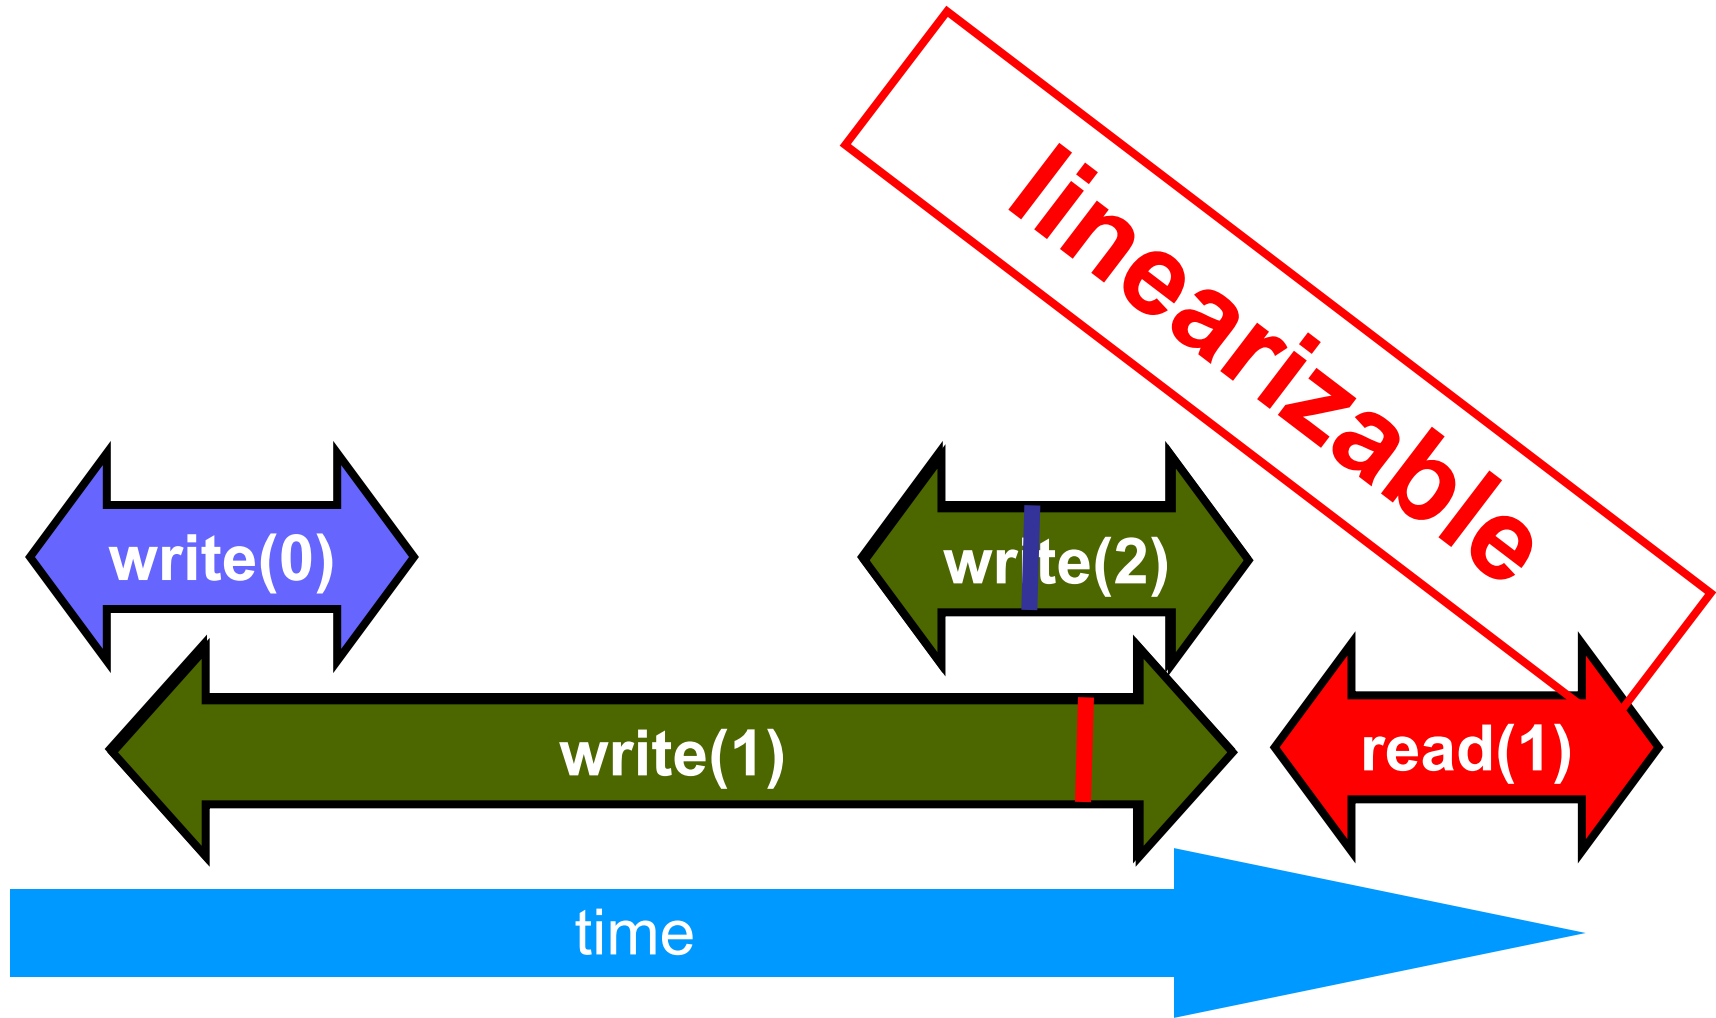
# Read/Write Register Example



# Read/Write Register Example




# Read/Write Register Example




# Talking About Executions

- Why?
  - Can't we specify the linearization point of each operation without describing an execution?
- Not Always
  - In some cases, linearization point ***depends on the execution***




# Formal Model of Executions

- Define precisely what we mean
  - Ambiguity is bad when intuition is weak
- Allow reasoning
  - Formal
  - But mostly informal
    - In the long run, actually more important
    - Ask me why!




# Split Method Calls into Two Events

- Invocation
  - method name & args
  - `q.enq(x)`
- Response
  - result or exception
  - `q.enq(x)` returns `void`
  - `q.deq()` returns `x`
  - `q.deq()` throws `empty`



# Invocation Notation


**A q.enq(x)**





# Invocation Notation

**thread** **A**q.enq(x)




# Invocation Notation

**Aq.enq(x)**


**thread**      **method**

The diagram illustrates the invocation notation **Aq.enq(x)**. The letter **A** is enclosed in a red rounded rectangle, with a red line pointing from it to the word **thread** in red text below. The text **q.enq(x)** is enclosed in another red rounded rectangle, with a red line pointing from it to the word **method** in red text below.

# Invocation Notation




# Invocation Notation



# Response Notation


**A q: void**




# Response Notation

**thread** **A** q: void

# Response Notation




# Response Notation






# Response Notation



# Response Notation




# History - Describing an Execution

**H =**


- A q.enq(3)
- A q:void
- A q.enq(5)
- B p.enq(4)
- B p:void
- B q.deq()
- B q:3

**Sequence of  
invocations and  
responses**



# Definition


- Invocation & response *match* if



# Object Projections

$H =$

A	q.enq(3)
A	q:void
B	p.enq(4)
B	p:void
B	q.deq()
B	q:3



# Object Projections


A `q.enq(3)`

A `q:void`

$H|q =$

B `q.deq()`


B `q:3`



# Thread Projections

$H =$


A	q.enq(3)
A	q:void
B	p.enq(4)
B	p:void
B	q.deq()
B	q:3



# Thread Projections

$H|B =$

- $B \text{ p.enq}(4)$
- $B \text{ p:void}$
- $B \text{ q.deq}()$
- $B \text{ q:3}$

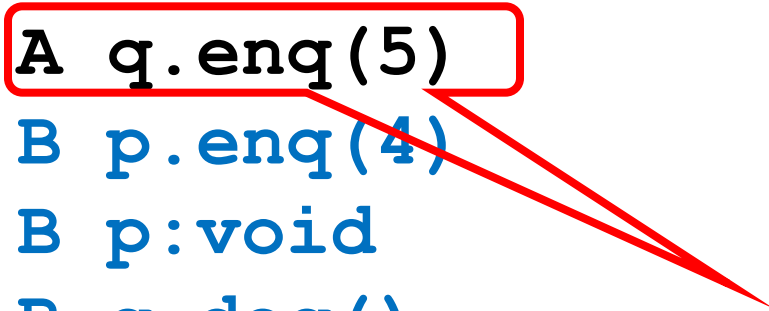




# Complete Subhistory

A q.enq(3)  
A q:void  
**A q.enq(5)**  
H = B p.enq(4)  
B p:void  
B q.deq()  
B q:3


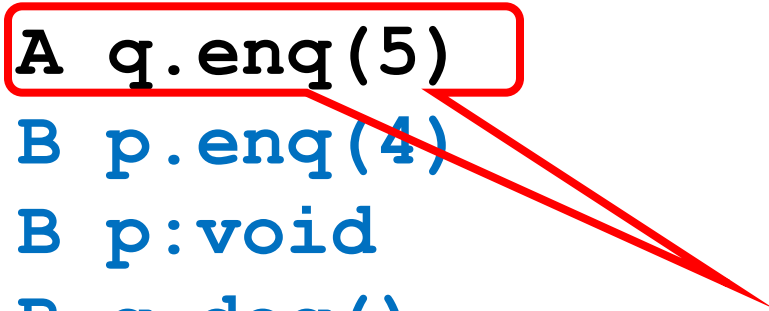
**An invocation is  
*pending* if it has no  
matching response**



# Complete Subhistory

A q.enq(3)  
A q:void  
**A q.enq(5)**  
H = B p.enq(4)  
B p:void  
B q.deq()  
B q:3

**May or may not  
have taken effect**


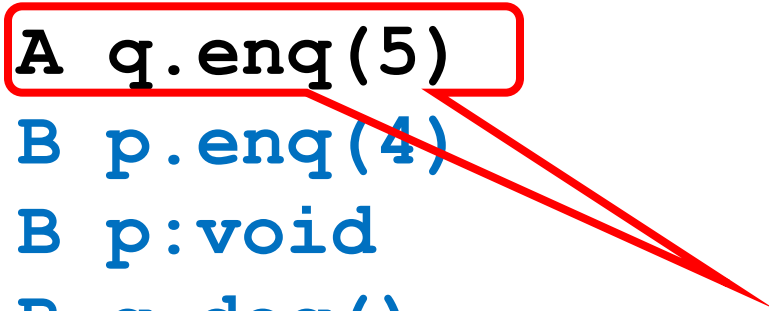


# Complete Subhistory

**H =**

A	q.enq(3)
A	q:void
A	q.enq(5)
B	p.enq(4)
B	p:void
B	q.deq()
B	q:3

**discard pending invocations**



# Complete Subhistory

A q.enq(3)


A q:void

**Complete(H) =** B p.enq(4)

B p:void

B q.deq()

B q:3



# Sequential Histories

A q.enq(3)

A q:void


B p.enq(4)

B p:void

B q.deq()

B q:3

A q:enq(5)



# Sequential Histories

A q.enq(3)

A q:void

B p.enq(4)


B p:void

B q.deq()

B q:3

A q:enq(5)

**match**



# Sequential Histories

A q.enq(3)

A q:void

**match**

B p.enq(4)


B p:void

**match**

B q.deq()

B q:3

A q:enq(5)



# Sequential Histories

A q.enq(3)

A q:void

**match**

B p.enq(4)

B p:void


**match**

B q.deq()

B q:3

**match**

A q:enq(5)





# Sequential Histories

A q.enq(3)

A q:void

**match**

B p.enq(4)

B p:void

**match**


B q.deq()

B q:3

**match**

A q:enq(5)

**Final pending  
invocation OK**



# Sequential Histories

A q.enq(3)

A q:void

B p.enq(4)

B p:void

B q.deq()

B q:3


A q:enq(5)

Method calls of different  
threads do not interleave

match


match

Final pending  
invocation OK



# Well-Formed Histories

H=  
A q.enq(3)  
B p.enq(4)  
B p:void  
B q.deq()  
A q:void  
B q:3



# Well-Formed Histories


## Per-thread projections sequential

**H=**

**A** `q.enq(3)`  
**B** `p.enq(4)`  
**B** `p:void`  
**B** `q.deq()`  
**A** `q:void`  
**B** `q:3`

**H | B=**

**B** `p.enq(4)`  
**B** `p:void`  
**B** `q.deq()`  
**B** `q:3`



# Well-Formed Histories

## Per-thread projections sequential

$H =$


A q.enq(3)  
B p.enq(4)  
B p:void  
B q.deq()  
A q:void  
B q:3

$H | B =$

B p.enq(4)  
B p:void  
B q.deq()  
B q:3

$H | A =$

A q.enq(3)  
A q:void



# Equivalent Histories

Threads see the same  
thing in both


$$\left\{ \begin{array}{l} H|A = G|A \\ H|B = G|B \end{array} \right.$$

H=

```
A q.enq(3)
B p.enq(4)
B p:void
B q.deq()
A q:void
B q:3
```


G=

```
A q.enq(3)
A q:void
B p.enq(4)
B p:void
B q.deq()
B q:3
```




# Sequential Specifications

- A sequential specification is some way of telling whether a
  - Single-thread, single-object history
  - Is legal
- For example:
  - Pre and post-conditions
  - But plenty of other techniques exist ...



# Legal Histories

- A sequential (multi-object) history  $H$  is legal if
  - For every object  $x$
  - $H|x$  is in the sequential spec for  $x$





# Precedence

A `q.enq(3)`

B `p.enq(4)`


B `p.void`

A `q:void`

B `q.deq()`

B `q:3`

A method call **precedes**  
another if response event  
precedes invocation  
event



# Non-Precedence

A q.enq(3)


B p.enq(4)

B p.void



B q.deq()

A q:void

B q:3




**Some method calls  
overlap one another**




# Notation

- Given
  - History  $H$
  - method executions  $m_0$  and  $m_1$  in  $H$
- We say  $m_0 \rightarrow_H m_1$ , if
  - $m_0$  precedes  $m_1$
- Relation  $m_0 \rightarrow_H m_1$  is a
  - Partial order
  - Total order if  $H$  is sequential




# Linearizability

- History  $H$  is *linearizable* if it can be extended to  $G$  by
  - Appending zero or more responses to pending invocations
  - Discarding other pending invocations
- So that  $G$  is equivalent to
  - Legal sequential history  $S$
  - where  $\rightarrow_G \subset \rightarrow_S$



# Remarks



- Some pending invocations
  - Took effect, so keep them
  - Discard the rest
- Condition  $\rightarrow_{\mathbf{G}} \subset \rightarrow_{\mathbf{S}}$ 
  - Means that **S** respects “real-time order” of **G**



# Ensuring $\rightarrow_G \subset \rightarrow_S$

$$\rightarrow_G = \{a \rightarrow c, b \rightarrow c\}$$

$$\rightarrow_S = \{a \rightarrow b, a \rightarrow c, b \rightarrow c\}$$



# Example

A q.enq(3)


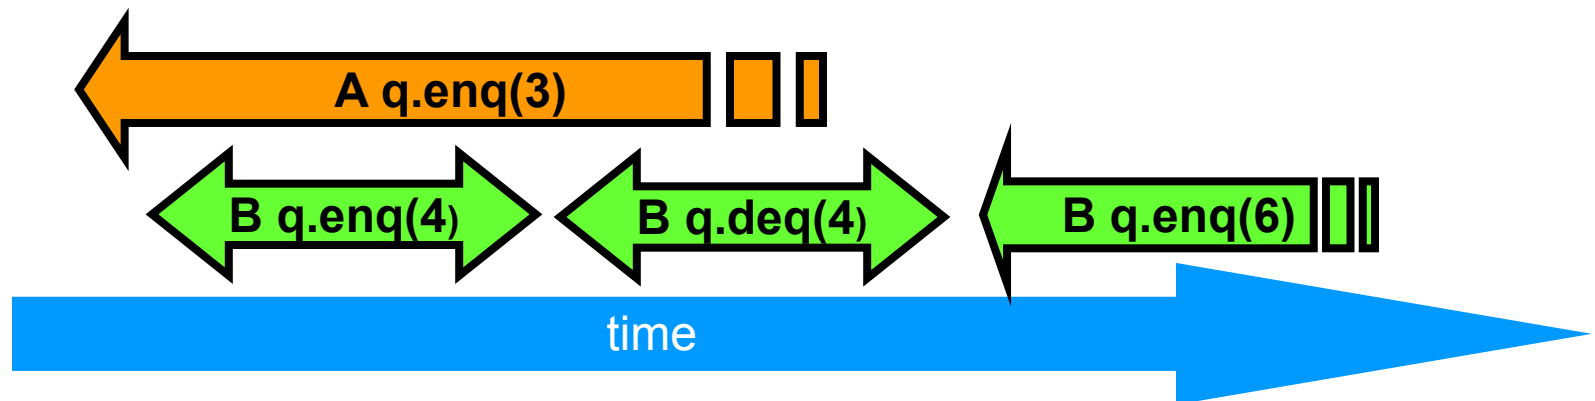
B q.enq(4)

B q:void

B q.deq()

B q:4

B q:enq(6)



# Example

A q.enq(3)

B q.enq(4)

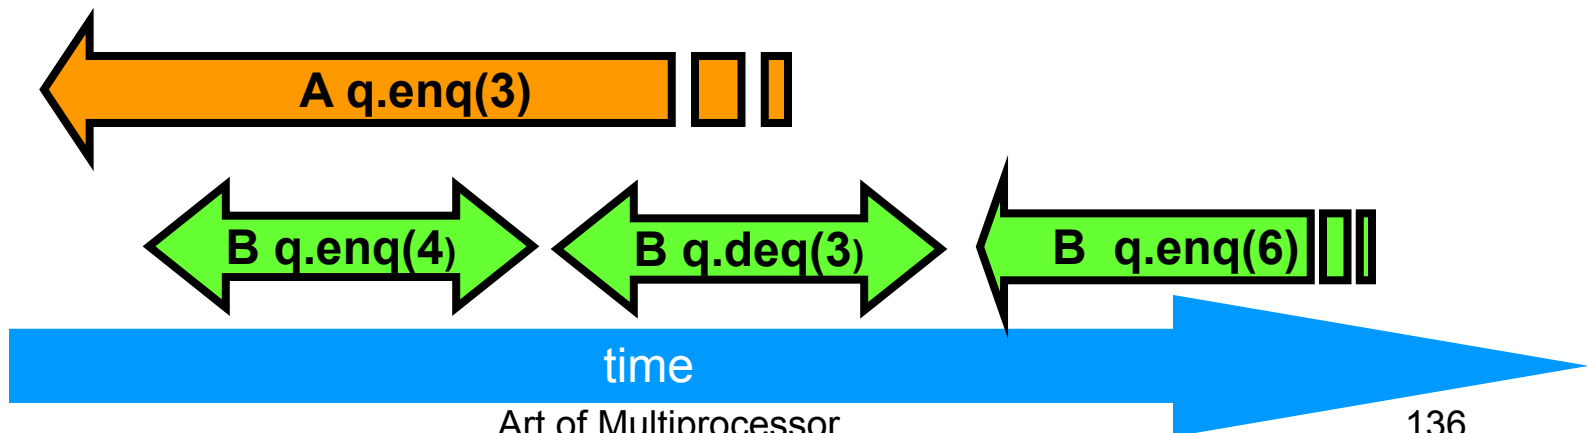
B q:void

B q.deq()

B q:4

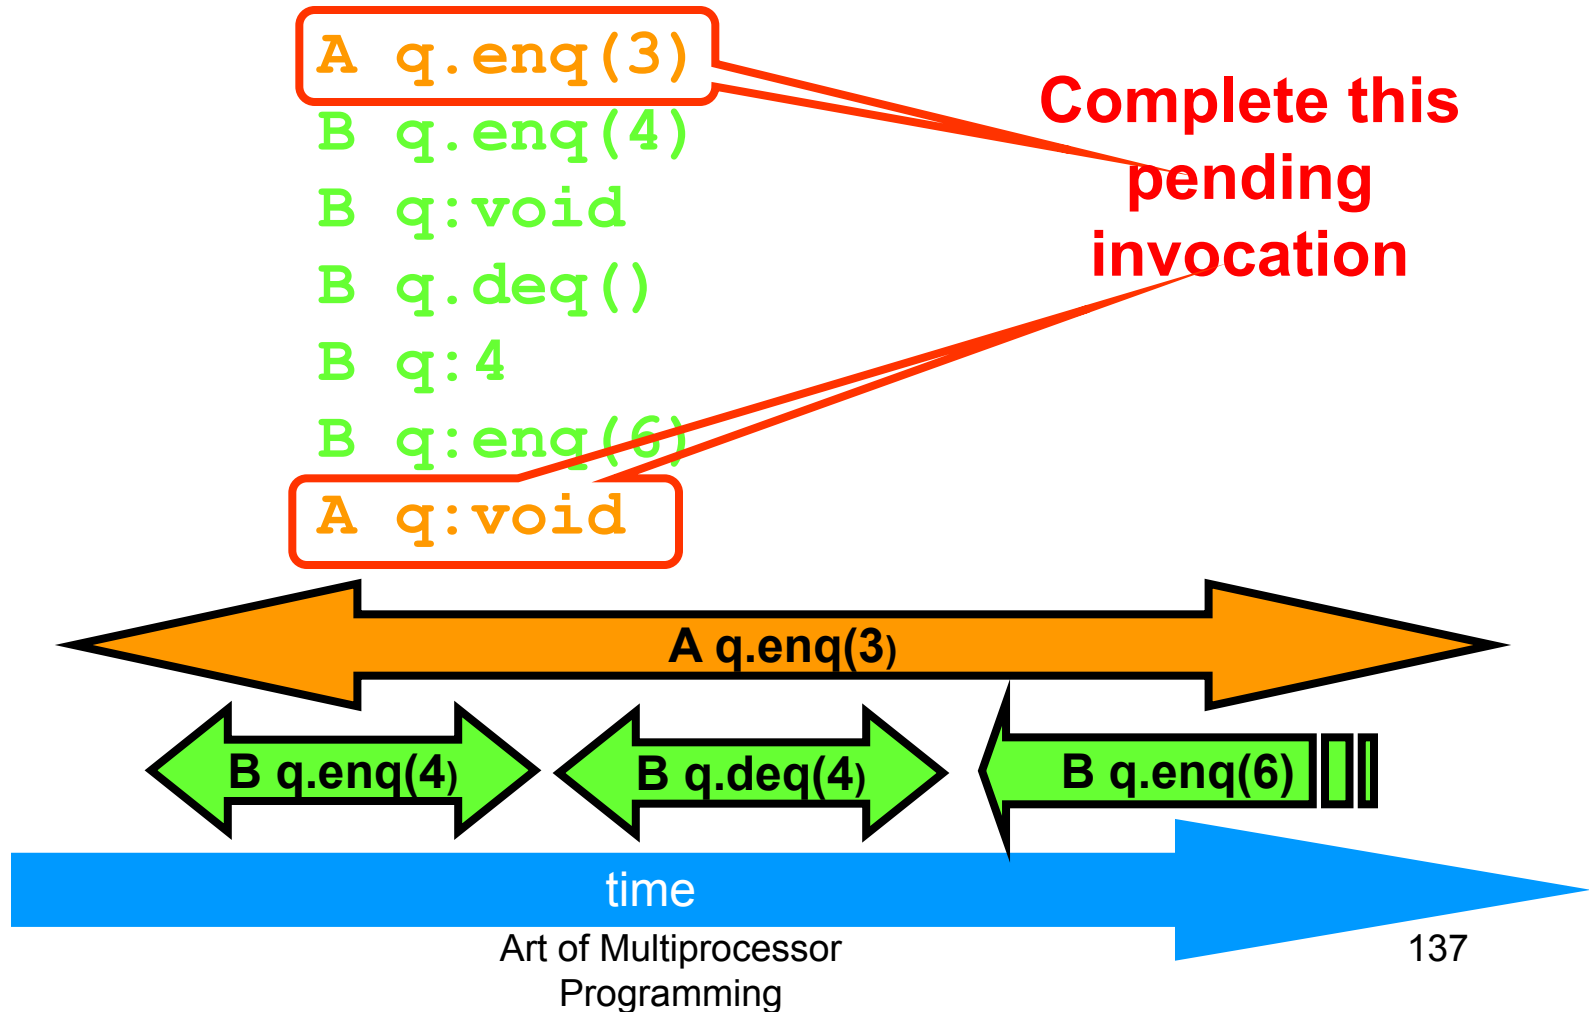
B q:enq(6)

Complete this  
pending  
invocation





# Example



# Example

discard this one

A q.enq(3)

B q.enq(4)


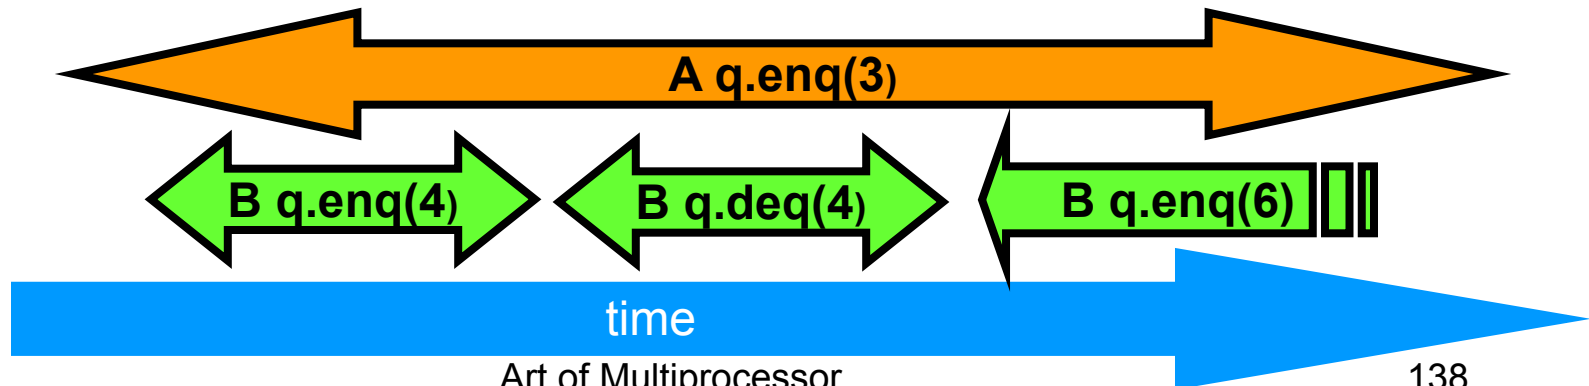
B q:void

B q.deq()

B q:4

B q:enq(6)

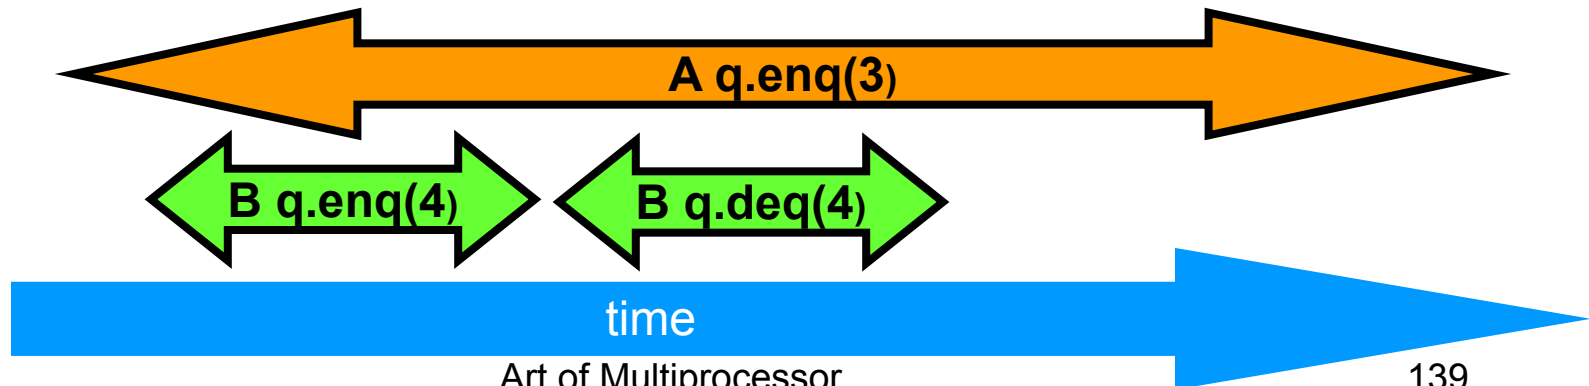

A q:void



# Example

discard this one

```
A q.enq(3)
B q.enq(4)
B q:void
B q.deq()
B q:4
A q:void
```



# Example

A q.enq(3)


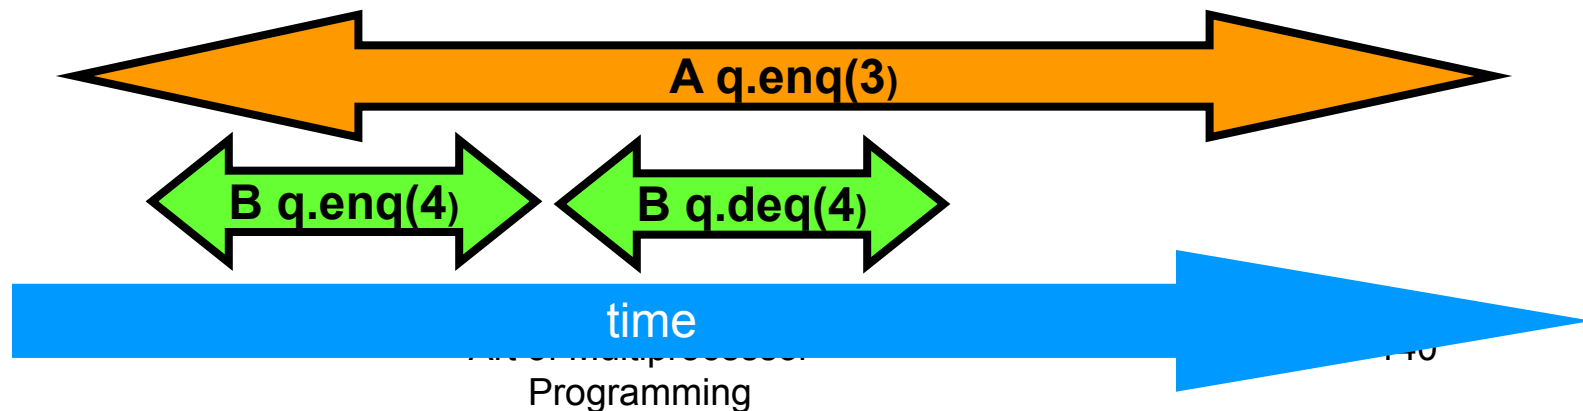
B q.enq(4)

B q:void

B q.deq()

B q:4

A q:void



# Example

A q.enq(3)

B q.enq(4)

B q:void

B q.deq()

B q:4

A q:void

B q.enq(4)


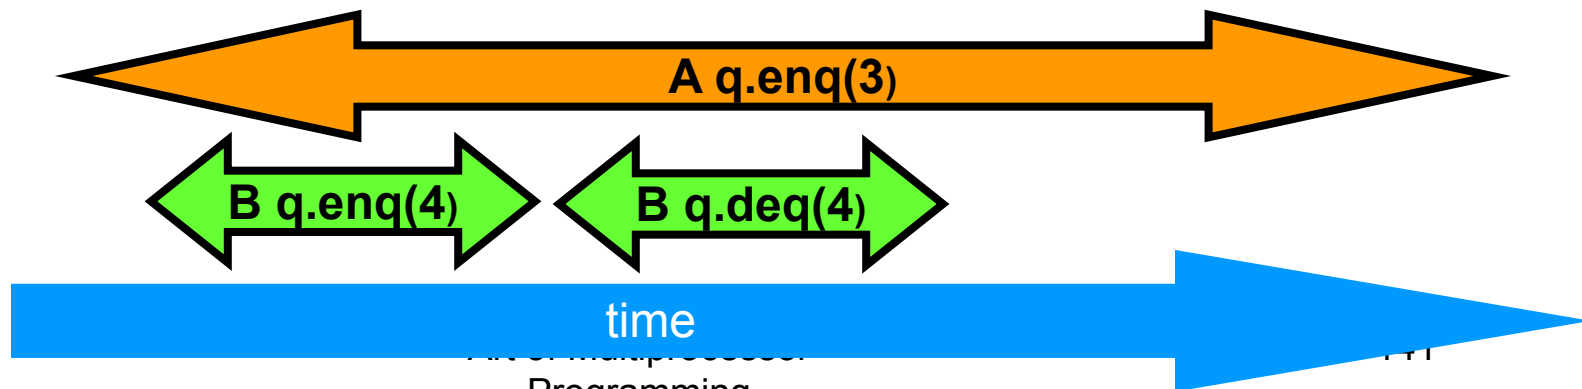
B q:void

A q.enq(3)

A q:void

B q.deq()

B q:4

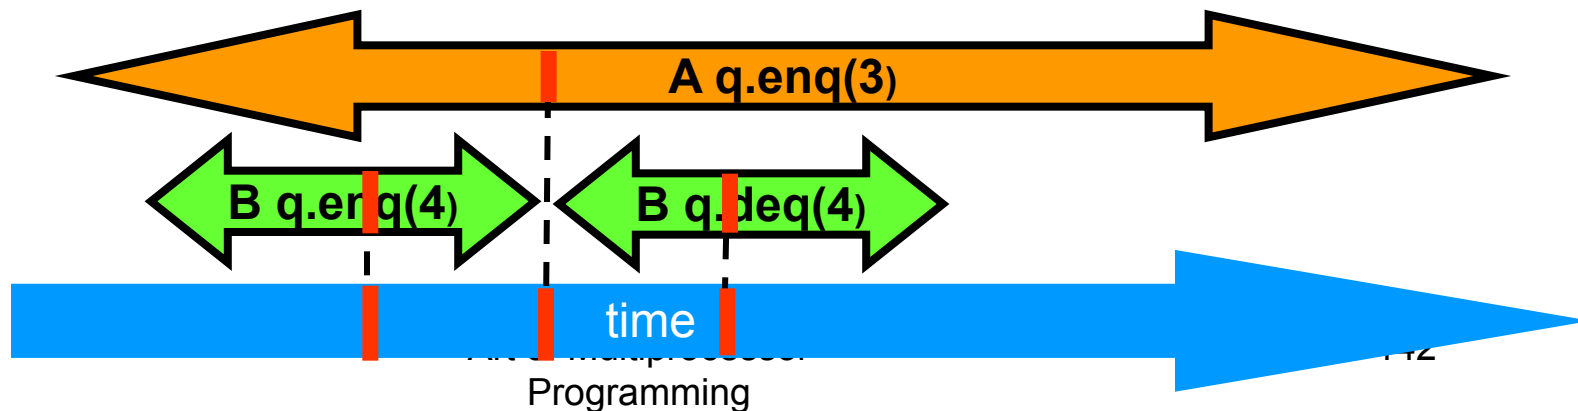


# Example

## Equivalent sequential history


A q.enq(3)  
B q.enq(4)  
B q:void  
B q.deq()  
B q:4  
A q:void

B q.enq(4)  
B q:void  
A q.enq(3)  
A q:void  
B q.deq()  
B q:4




# Concurrency

- How much concurrency does linearizability allow?
- When must a method invocation block?



# Concurrency


- Focus on ***total*** methods
  - Defined in every state
- Example:
  - `deq()` that throws **Empty** exception
  - Versus `deq()` that waits ...
- Why?
  - Otherwise, blocking unrelated to synchronization





# Concurrency

- **Question:** When does linearizability require a method invocation to block?
- **Answer:** never.
- Linearizability is *non-blocking*



# Non-Blocking Theorem

If method invocation

**A  $q.\text{inv}(\dots)$**


is pending in history  $H$ , then there exists a response

**A  $q:\text{res}(\dots)$**

such that


**$H + A\ q:\text{res}(\dots)$**

is linearizable




# Proof

- Pick linearization  $S$  of  $H$
- If  $S$  already contains
  - Invocation  $A \ q.\text{inv}(\dots)$  and response,
  - Then we are done.
- Otherwise, pick a response such that
  - $S + A \ q.\text{inv}(\dots) + A \ q:\text{res}(\dots)$
  - Possible because object is *total*.




# Composability Theorem

- History  $H$  is linearizable if and only if
  - For every object  $x$
  - $H|x$  is linearizable
- We care about objects only!
  - (Materialism?)




# Why Does Composability Matter?

- Modularity
- Can prove linearizability of objects in isolation
- Can compose independently-implemented objects




# Reasoning About Linearizability: Locking

```
public T deq() throws EmptyException {  
    lock.lock();  
    try {  
        if (tail == head)  
            throw new EmptyException();  
        T x = items[head % items.length];  
        head++;  
        return x;  
    } finally {  
        lock.unlock();  
    }  
}
```



# Reasoning About Linearizability: Locking


```
public T deq() throws EmptyException {  
    lock.lock();  
    try {  
        if (tail == head)  
            throw new EmptyException();  
        T x = items[head % items.length];  
        head++;  
        return x;  
    } finally {  
        lock.unlock();  
    }  
}
```



Linearization points  
are when locks are  
released

# More Reasoning: Wait-free

```
public class WaitFreeQueue {  
  
    int head = 0, tail = 0;  
    items = (T[]) new Object[capacity];  
  
    public void enq(Item x) {  
        if (tail-head == capacity) throw  
            new FullException();  
        items[tail % capacity] = x; tail++;  
    }  
  
    public Item deq() {  
        if (tail == head) throw  
            new EmptyException();  
        Item item = items[head % capacity]; head++;  
        return item;  
    }  
}
```





# More Reasoning: Wait-free

```
public class WaitFreeQueue {
```

```
    int capacity;
    int head = 0;
    int tail = 0;
    Item[] items = new Item[capacity];
```

Linearization order is  
order head and tail  
fields modified

Remember that there  
is only one enqueuer  
and only one dequeuer

```
    public void enq(Item x) {
        while (tail - head == capacity) throw
            new FullException();
        items[tail % capacity] = x;
    }
```

tail++;


```
    public Item deq() {
        if (tail == head) throw
            new EmptyException();
        Item item = items[head % capacity];
        return item;
```

head++;

```
}}
```


# Strategy

- Identify one atomic step where method “happens”
  - Critical section
  - Machine instruction
- Doesn't always work
  - Might need to define several different steps for a given method



# Linearizability: Summary

- Powerful specification tool for shared objects
- Allows us to capture the notion of objects being “atomic”
- Don’t leave home without it




# Alternative: Sequential Consistency

- History  $H$  is **Sequentially Consistent** if it can be extended to  $G$  by
  - Appending zero or more responses to pending invocations
  - Discarding other pending invocations
- So that  $G$  is equivalent to a
  - Legal sequential history  $S$


**Differs from  
linearizability**

~~Where  $G \subseteq S$~~





# Sequential Consistency




- No need to preserve real-time order
  - Cannot re-order operations done by the same thread
  - Can re-order non-overlapping operations done by different threads
- Often used to describe multiprocessor memory architectures







# Example



# Example









# Example











# Example






# Example









# Example








**not linearizable**






Exa

**Yet Sequentially  
Consistent**





# Theorem



Sequential Consistency is not  
composable




# FIFO Queue Example




# FIFO Queue Example



# FIFO Queue Example





**History H**






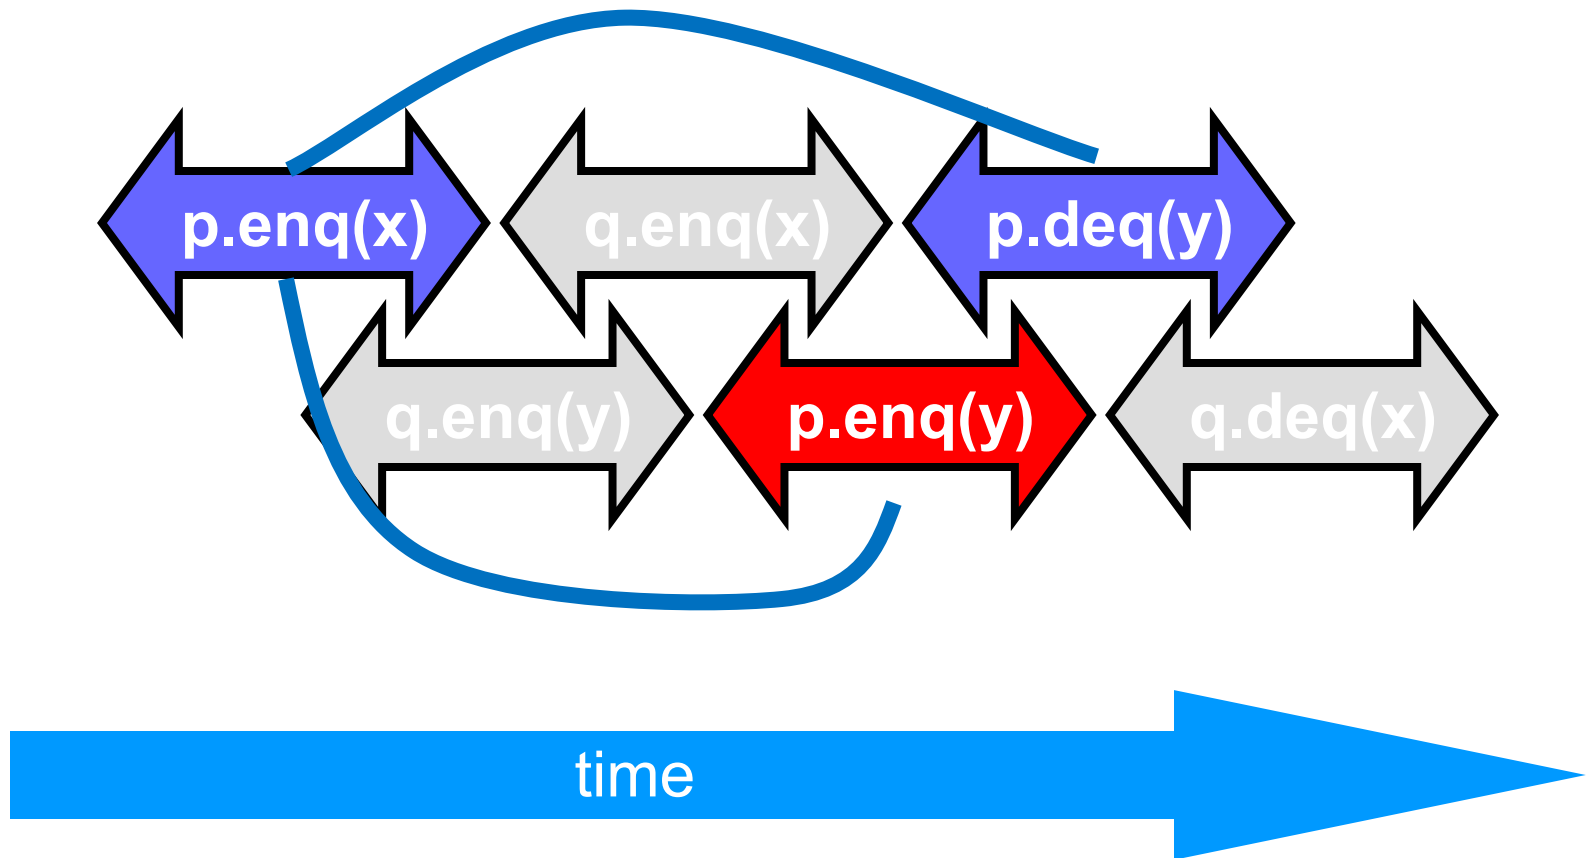
# H/p Sequentially Consistent



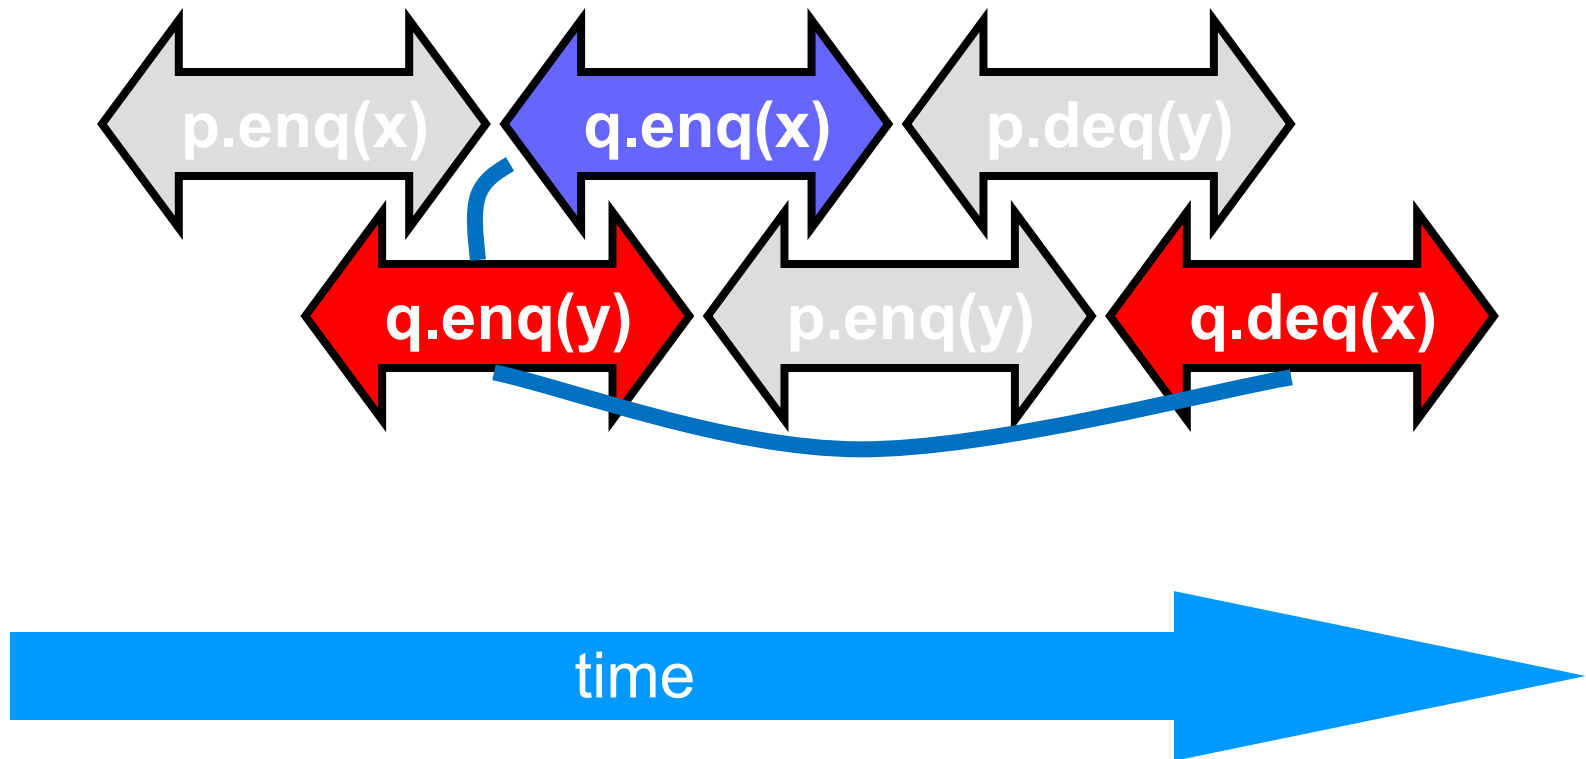
# H/q Sequentially Consistent



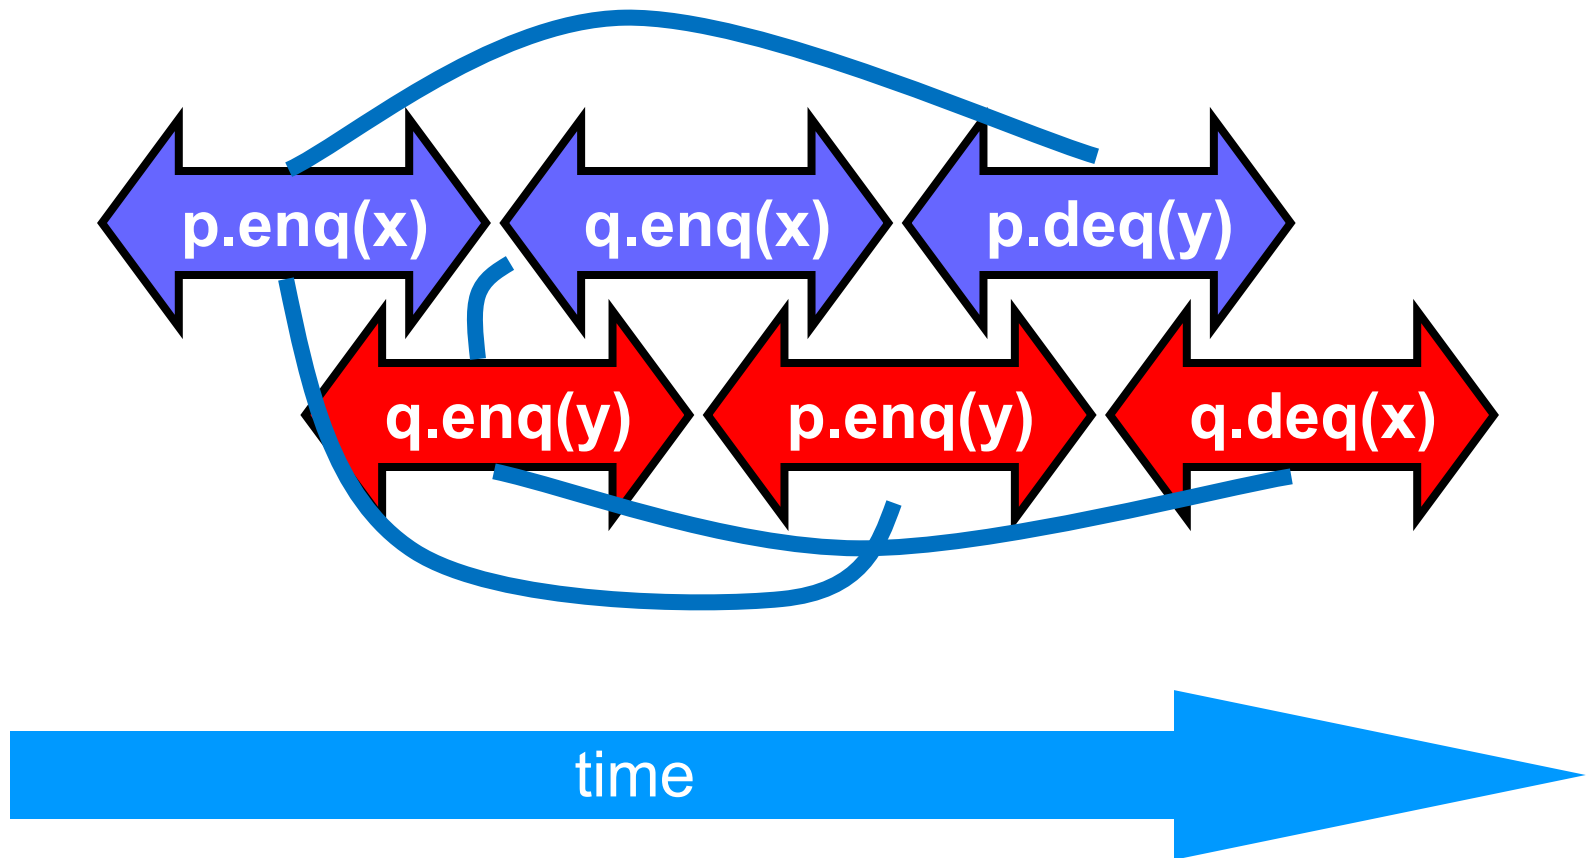
# Ordering imposed by p



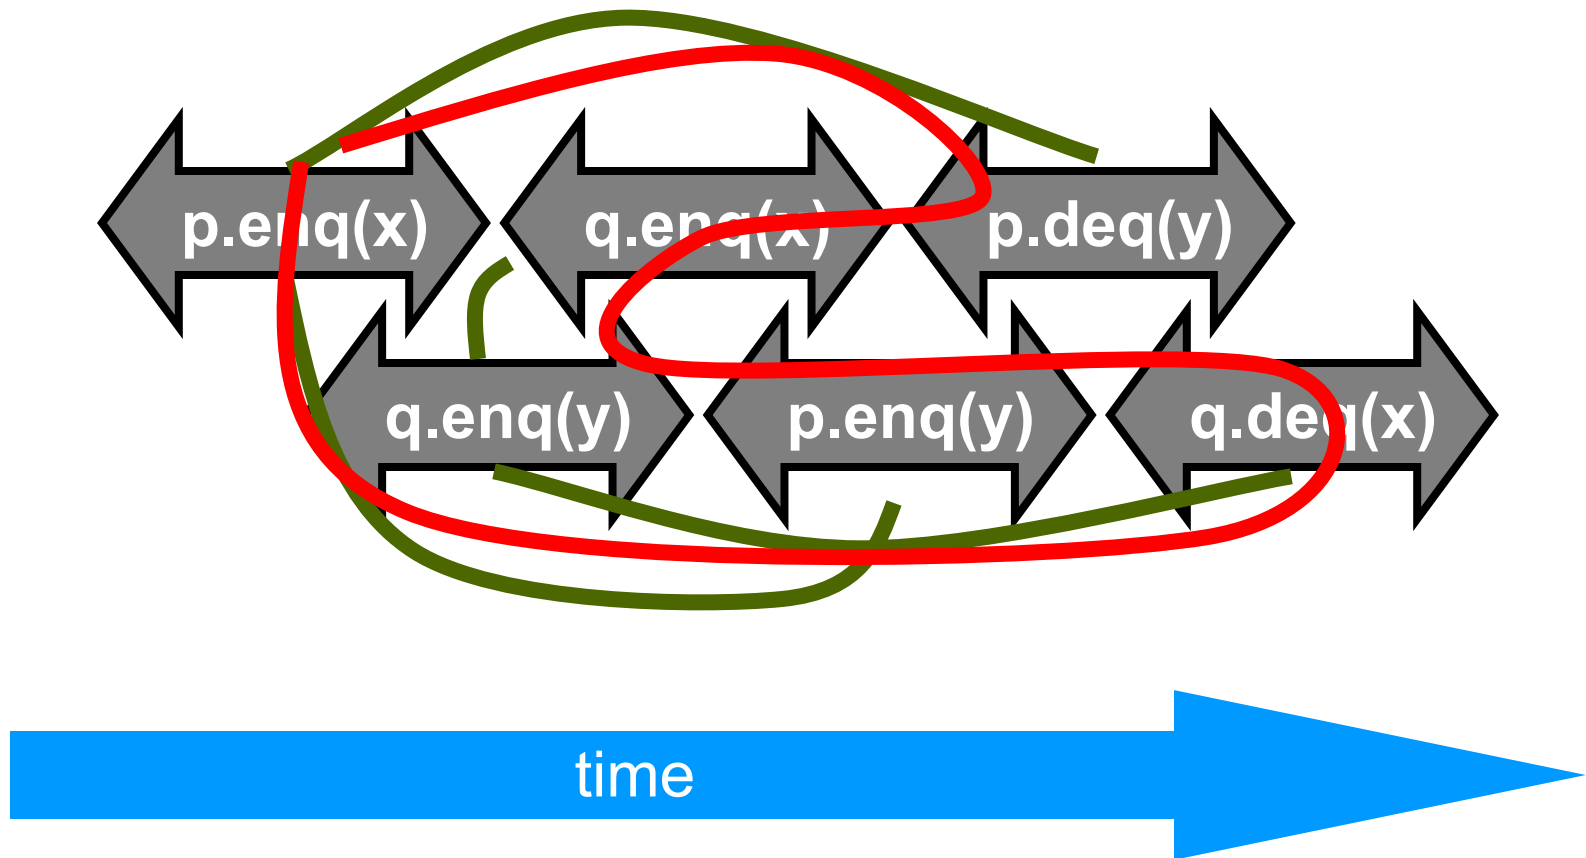
# Ordering imposed by q



# Ordering imposed by both




# Combining orders

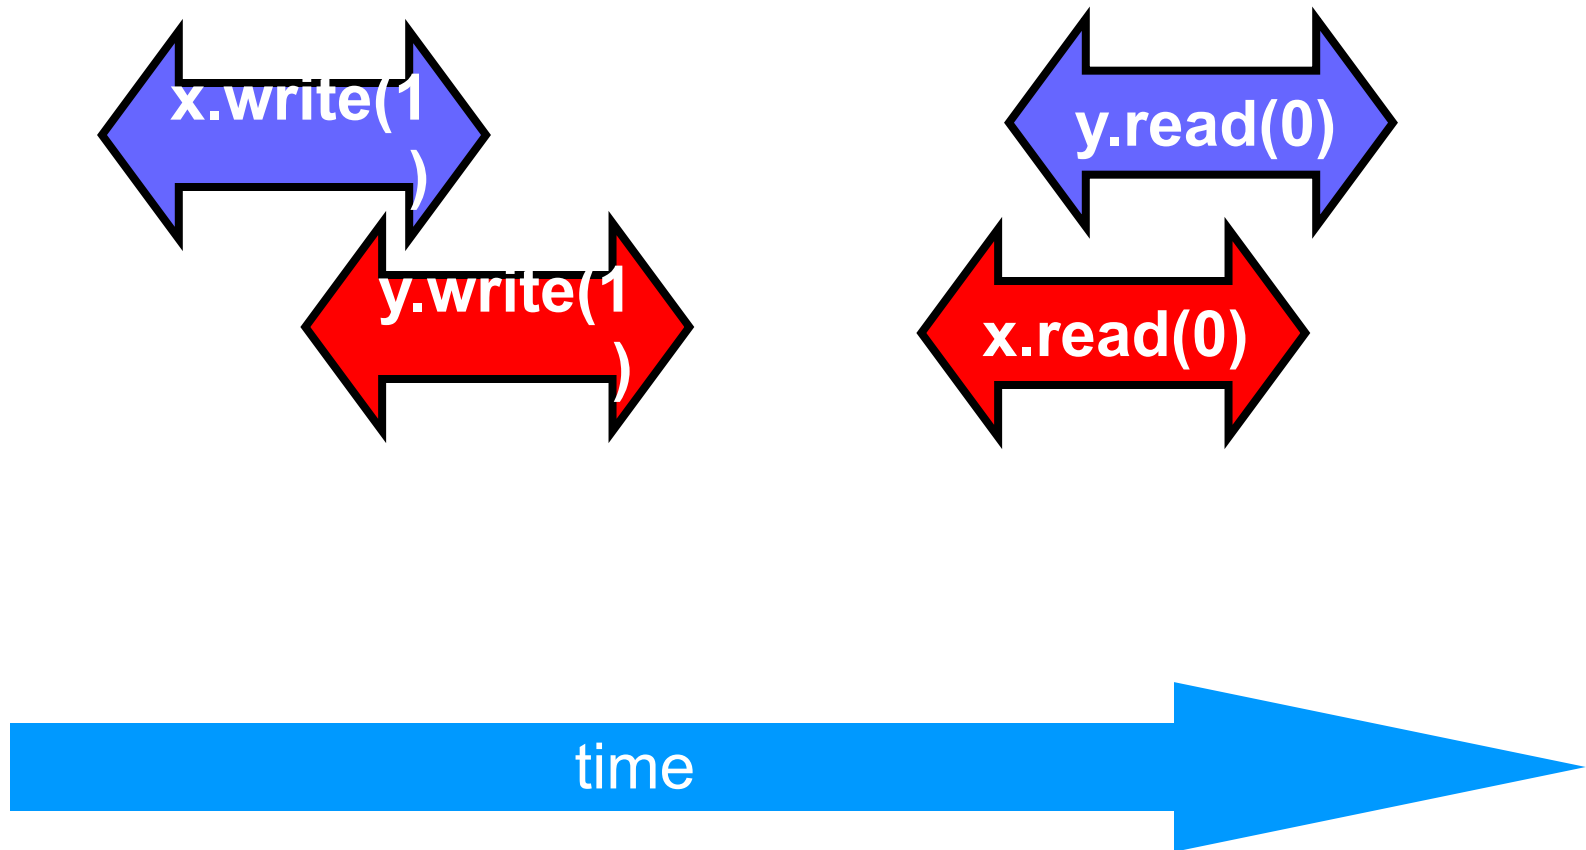


# Fact

- Most hardware architectures don't support sequential consistency
- Because they think it's too strong
- Here's another story ...




# The Flag Example







# The Flag Example




- Each thread's view is sequentially consistent
  - It went first




# The Flag Example



- Entire history isn't sequentially consistent
  - Can't both go first




# The Flag Example



- Is this behavior really so wrong?
  - We can argue either way ...

# Opinion: It's Wrong

- This pattern
  - Write mine, read yours
- Is exactly the flag principle
  - Beloved of Alice and Bob
  - Heart of mutual exclusion
    - Peterson
    - Bakery, etc.
- It's non-negotiable!



# Peterson's Algorithm

```
public void lock() {  
    flag[i] = true;  
    victim = i;  
    while (flag[j] && victim == i) {};  
}  
public void unlock() {  
    flag[i] = false;  
}
```

# Crux of Peterson Proof

(1)  $\text{write}_B(\text{flag}[B]=\text{true}) \rightarrow$

(3)  $\text{write}_B(\text{victim}=B) \rightarrow$

(2)  $\text{write}_A(\text{victim}=A) \rightarrow \text{read}_A(\text{flag}[B])$   
 $\rightarrow \text{read}_A(\text{victim})$


# Crux of Peterson Proof

- (1)  $\text{write}_B(\text{flag}[B]=\text{true}) \rightarrow$
- (3)  $\text{write}_B(\text{victim}=B) \rightarrow$
- (2)  $\text{write}_A(\text{victim}=A) \rightarrow \text{read}_A(\text{flag}[B])$   
 $\rightarrow \text{read}_A(\text{victim})$

Observation: proof relied on fact that if a location is stored, a later load by some thread will return this or a later stored value.

# Opinion: But It Feels So Right ...

- Many hardware architects think that sequential consistency is too strong
- Too expensive to implement in modern hardware
- OK if flag principle
  - violated by default
  - Honored by explicit request





# Hardware Consistency

**Initially,  $a = b = 0$ .**

Processor 0

```
mov 1, a    ;Store  
mov b, %ebx ;Load
```

Processor 1

```
mov 1, b    ;Store  
mov a, %eax ;Load
```


What are the final possible values of `%eax` and `%ebx` after both processors have executed?

Sequential consistency implies that no execution ends with  $\%eax = \%ebx = 0$

# Hardware Consistency

- No modern-day processor implements sequential consistency.
- Hardware actively reorders instructions.
- Compilers may reorder instructions, too.
- Why?
- Because most of performance is derived from a single thread's unsynchronized execution of code.

# Instruction Reordering




Program Order

Execution Order

Q. Why might the hardware or compiler decide to reorder these instructions?

A. To obtain higher performance by covering load latency — *instruction-level parallelism*.

# Instruction Reordering



Program Order


Execution Order

Q. When is it safe for the hardware or compiler to perform this reordering?

A. When  $a \neq b$ .

A'. And there's no concurrency.

# Hardware Reordering



- Processor can issue stores faster than the network can handle them  $\Rightarrow$  store buffer.
- Loads take priority, bypassing the store buffer.
- Except if a load address matches an address in the store buffer, the store buffer returns the result.

# X86: Memory Consistency

# Thread's Code

~~Storage 1~~  
~~Storage 2~~  
~~Storage 3~~  
~~Storage 4~~  
~~Storage 5~~

1. Loads **are *not*** reordered with loads.
2. Stores **are *not*** reordered with stores.
3. Stores **are *not*** reordered with prior loads.
4. A load ***may*** be reordered with a prior store to a different location **but *not*** with a prior store to the same location.
5. Stores to the same location **respect a global total order.**

# X86: Memory Consistency

Threa  
d's

Code

**Stor**

**Stor**

**Load**

**Load**

**Store**


**Store**

**Load**

**Load**

**Load**

**Load**



1. Loads **are not** reordered with loads.
2. Stores **are not** reordered with stores.

3. **Total Store Ordering (TSO)...weaker than sequential consistency**
4. **with a prior store to the same location.**

- OK!**
5. Stores to the same location **respect a global total order.**

# Memory Barriers (Fences)

- *A memory barrier (or memory fence) is a hardware action that enforces an ordering constraint between the instructions before and after the fence.*
- *A memory barrier can be issued explicitly as an instruction (x86: mfence)*
- *The typical cost of a memory fence is comparable to that of an L2-cache access.*



# X86: Memory Consistency

Threa  
d's

Code

**Stor**

**Stor**

**Load**

**Load**

**Stor**

**Stor**

**Barrier**

**Load**

**Load**

**Load**

**Load**

1. Loads **are not** reordered with loads.

2. Stor

3. Stor

load

4. A lo

store to

with a prior store to the same

location.

5. Stores to the same location

global total order.

**Total Store Ordering +  
properly placed memory  
barriers = sequential  
consistency**

store to the same location but not

with a prior store to the same


location.

Stores to the same location

global total order.


# Memory Barriers

- Explicit Synchronization
- Memory barrier will
  - Flush write buffer
  - Bring caches up to date
- Compilers often do this for you
  - Entering and leaving critical sections




# Volatile Variables

- In Java, can ask compiler to keep a variable up-to-date by declaring it `volatile`
- Adds a memory barrier after each store
- Inhibits reordering, removing from loops, & other “compiler optimizations”
- Will talk about it in detail in later lectures




# Summary: Real-World

- Hardware weaker than sequential consistency
- Can get sequential consistency at a price
- Linearizability better fit for high-level software




# Linearizability

- Linearizability
  - Operation takes effect instantaneously between invocation and response
  - Uses sequential specification, locality implies composability




# Summary: Correctness

- Sequential Consistency
  - Not composable
  - Harder to work with
  - Good way to think about hardware models
- We will use *linearizability* as our consistency condition in the remainder of this course unless stated otherwise




# Progress

- We saw an implementation whose methods were lock-based (deadlock-free)
- We saw an implementation whose methods did not use locks (lock-free)
- How do they relate?



# Progress Conditions


- *Deadlock-free*: some thread trying to acquire the lock eventually succeeds.
- *Starvation-free*: every thread trying to acquire the lock eventually succeeds.
- *Lock-free*: some thread calling a method eventually returns.
- *Wait-free*: every thread calling a method eventually returns.






# Progress Conditions

	Non-Blocking	Blocking
Everyone makes progress	Wait-free	Starvation-free
Someone makes progress	Lock-free	Deadlock-free



# Summary

- We will look at *linearizable blocking* and *non-blocking* implementations of objects.



This work is licensed under a [Creative Commons Attribution-ShareAlike 2.5 License](https://creativecommons.org/licenses/by-sa/3.0/).

- **You are free:**
  - **to Share** — to copy, distribute and transmit the work
  - **to Remix** — to adapt the work
- **Under the following conditions:**
  - **Attribution.** You must attribute the work to “The Art of Multiprocessor Programming” (but not in any way that suggests that the authors endorse you or your use of the work).
  - **Share Alike.** If you alter, transform, or build upon this work, you may distribute the resulting work only under the same, similar or a compatible license.
- For any reuse or distribution, you must make clear to others the license terms of this work. The best way to do this is with a link to
  - <http://creativecommons.org/licenses/by-sa/3.0/>.
- Any of the above conditions can be waived if you get permission from the copyright holder.
- Nothing in this license impairs or restricts the author's moral rights.