

# Concurrent programming

## Spin Locks and Contention

Companion slides for

The Art of Multiprocessor Programming

by Maurice Herlihy, Nir Shavit, Victor Luchangco,  
and Michael Spear

Modified by Piotr Witkowski

# Focus so far: Correctness and Progress

- Models
  - Accurate (we never lied to you)
  - But idealized (so we forgot to mention a few things)
- Protocols
  - Elegant
  - Important
  - But naïve

# New Focus: Performance

- Models
  - More complicated (not the same as complex!)
  - Still focus on principles (not soon obsolete)
- Protocols
  - Elegant (in their fashion)
  - Important (why else would we pay attention)
  - And realistic (your mileage may vary)

# Kinds of Architectures

- SISD (Uniprocessor)
  - Single instruction stream
  - Single data stream
- SIMD (Vector)
  - Single instruction
  - Multiple data
- MIMD (Multiprocessors)
  - Multiple instruction
  - Multiple data.

# Kinds of Architectures


- SISD (Uniprocessor)
  - Single instruction stream
  - Single data stream
- SIMD (Vector)
  - Single instruction
  - Multiple data

Our space




- MIMD (Multiprocessors)
  - Multiple instruction
  - Multiple data.

# MIMD Architectures



**Shared Bus**



**Distributed**

- Memory Contention
- Communication Contention
- Communication Latency

# Today: Revisit Mutual Exclusion

- Performance, not just correctness
- Proper use of multiprocessor architectures
- A collection of locking algorithms...

# What Should you do if you can't get a lock?

- Keep trying
  - “spin” or “busy-wait”
  - Good if delays are short
- Give up the processor
  - Good if delays are long
  - Always good on uniprocessor




# What Should you do if you can't get a lock?


- Keep trying
  - “spin” or “busy-wait”
  - Good if delays are short
- Give up the processor
  - Good if delays are long
  - Always good on uniprocessor

our focus

# Basic Spin-Lock




# Basic Spin-Lock




# Basic Spin-Lock

...lock suffers from contention



# Basic Spin-Lock


...lock suffers from contention



Notice: these are distinct phenomena

# Basic Spin-Lock


...lock suffers from contention



**Seq Bottleneck** □ **no parallelism**

# Basic Spin-Lock

...lock suffers from contention



# Review: Test-and-Set

- Boolean value
- Test-and-set (TAS)
  - Swap **true** with current value
  - Return value tells if prior value was **true** or **false**
- Can reset just by writing **false**
- TAS aka “getAndSet”



# Review: Test-and-Set

```
public class AtomicBoolean {  
    boolean value;  
  
    public synchronized boolean  
        getAndSet(boolean newValue) {  
        boolean prior = value;  
        value = newValue;  
        return prior;  
    }  
}
```

# Review: Test-and-Set

```
public class AtomicBoolean {
```

```
    boolean value;
```

```
    public synchronized boolean  
        getAndSet(boolean newValue) {  
        boolean prior = value;  
        value = newValue;  
        return prior;  
    }  
}
```

**Package**

**java.util.concurrent.atomic**

# Review: Test-and-Set

```
public class AtomicBoolean {  
    boolean value;
```

```
    public synchronized boolean  
        getAndSet(boolean newValue) {  
        boolean prior = value;  
        value = newValue;  
        return prior;  
    }
```

```
}
```

**Swap old and new  
values**

# Review: Test-and-Set

```
AtomicBoolean lock  
    = new AtomicBoolean(false)  
...  
boolean prior = lock.getAndSet(true)
```

# Review: Test-and-Set

```
AtomicBoolean lock  
= new AtomicBoolean(false)
```

```
boolean prior = lock.getAndSet(true)
```

**Swapping in `true` is called  
“test-and-set” or TAS**

# Test-and-Set Locks

- Locking
  - Lock is free: value is false
  - Lock is taken: value is true
- Acquire lock by calling TAS
  - If result is false, you win
  - If result is true, you lose
- Release lock by writing false

# Test-and-set Lock

```
class TASlock {  
    AtomicBoolean state =  
        new AtomicBoolean(false);  
  
    void lock() {  
        while (state.getAndSet(true)) {}  
    }  
  
    void unlock() {  
        state.set(false);  
    }  
}
```

# Test-and-set Lock

```
class TASlock {
```

```
    AtomicBoolean state =  
        new AtomicBoolean(false) ;
```

```
    void lock() {  
        while (state.getAndSet(true)) {}  
    }
```

```
    void unlock() {  
        state  
    }  
}
```

**Lock state is AtomicBoolean**



# Test-and-set Lock

```
class TASlock {  
    AtomicBoolean state =  
        new AtomicBoolean(false);
```

```
    void lock() {
```

```
        while (state.getAndSet(true)) {}
```

```
    }
```

```
    void unlock() {
```

```
        state.set(false);  
    }
```

**Keep trying until lock acquired**

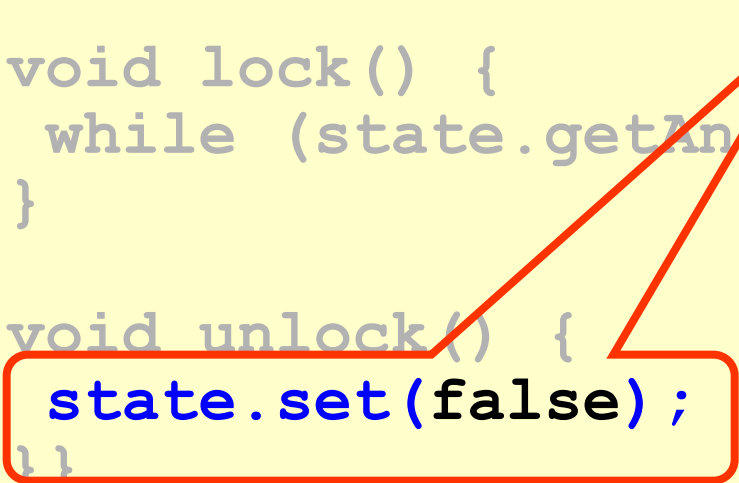
# Test-and-set Lock

```
class TA
{
    AtomicB state;
    new At

    void lock() {
        while (state.getAndSet(true)) {}
    }

    void unlock() {
        state.set(false);
    }
}
```

**Release lock by resetting state to false**




# Space Complexity

- TAS spin-lock has small “footprint”
- N thread spin-lock uses  $O(1)$  space
- As opposed to  $O(n)$  Peterson/Bakery
- How did we overcome the  $\Omega(n)$  lower bound?
- We used a RMW operation...


# Performance

- Experiment
  - $n$  threads
  - Increment shared counter 1 million times
- How long should it take?
- How long does it take?

# Graph



# Mystery #1



# Test-and-Test-and-Set Locks

- Lurking stage
  - Wait until lock “looks” free
  - Spin while read returns `true` (lock taken)
- Pouncing state
  - As soon as lock “looks” available
  - Read returns `false` (lock free)
  - Call TAS to acquire lock
  - If TAS loses, back to lurking


# Test-and-test-and-set Lock

```
class TTASlock {  
    AtomicBoolean state =  
        new AtomicBoolean(false);  
  
    void lock() {  
        while (true) {  
            while (state.get()) {}  
            if (!state.getAndSet(true))  
                return;  
        }  
    }  
}
```



# Test-and-test-and-set Lock

```
class TTASlock {  
    AtomicBoolean state =  
        new AtomicBoolean(false);  
  
    void lock() {  
        while (true) {  
            while (state.get()) {}  
            if (!state.getAndSet(true))  
                return;  
        }  
    }  
}
```



**Wait until lock looks free**

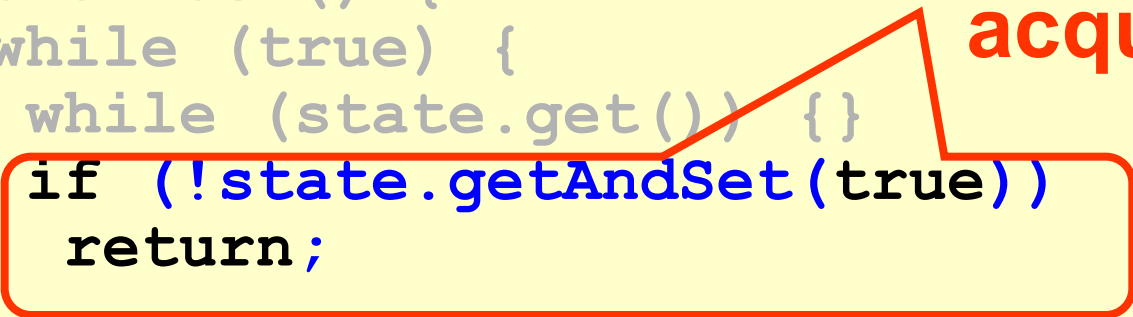
# Test-and-test-and-set Lock

```
class TTASlock {  
    AtomicBoolean state =  
        new AtomicBoolean(false);
```


```
    void lock() {  
        while (true) {  
            while (state.get()) {}
```

```
            if (!state.getAndSet(true))  
                return;  
        }  
    }
```

**Then try to  
acquire it**



# Mystery #2




# Mystery

- Both
  - TAS and TTAS
  - Do the same thing (in our model)
- Except that
  - TTAS performs much better than TAS
  - Neither approaches ideal


# Opinion

- Our memory abstraction is broken
- TAS & TTAS methods
  - Are provably the same (in our model)
  - Except they aren't (in field tests)
- Need a more detailed model ...

# Bus-Based Architectures




# Bus-Based Architectures



# Bus-Based Architectures

## Shared Bus

- Broadcast medium
- One broadcaster at a time
- Processors and memory all “snoop”






Bus-

## Per-Processor Caches

- Small
- Fast: 1 or 2 cycles
- Address & state information




# Granularity

- Caches operate at a larger granularity than a word
- Cache line: fixed-size block containing the address (today 64 or 128 bytes)


# Locality

- If you use an address now, you will probably use it again soon
  - Fetch from cache, not memory
- If you use an address now, you will probably use a nearby address soon
  - In the same cache line

# L1 and L2 Caches




# L1 and L2 Caches



# L1 and L2 Caches

**Larger and slower**  
**10s of cycles**  
**~128 byte line**



# Jargon Watch

- Cache hit
  - “I found what I wanted in my cache”
  - Good Thing™

# Jargon Watch

- Cache hit
  - “I found what I wanted in my cache”
  - Good Thing™
- Cache miss
  - “I had to shlep all the way to memory for that data”
  - Bad Thing™




# Cave Canem

- This model is **still** a simplification
  - But not in any essential way
  - Illustrates basic principles
- Will discuss complexities later


# Fully Associative Cache

- Any line can be anywhere in the cache
  - Advantage: can replace any line
  - Disadvantage: hard to find lines




# Direct Mapped Cache

- Every address has exactly 1 slot
  - Advantage: easy to find a line
  - Disadvantage: must replace fixed line




# K-way Set Associative Cache

- Each slot holds  $k$  lines
  - Advantage: pretty easy to find a line
  - Advantage: some choice in replacing line



# Multicore Set Associativity

- $k$  is 8 or even 16 and growing...
  - Why? Because cores share sets
  - Threads cut effective size if accessing different data



# Cache Coherence

- A and B both cache address  $x$
- A writes to  $x$ 
  - Updates cache
- How does B find out?
- Many cache coherence protocols in literature

# MESI

- Modified
  - Have modified cached data, must write back to memory

# MESI

- Modified
  - Have modified cached data, must write back to memory
- Exclusive
  - Not modified, I have only copy




# MESI

- Modified
  - Have modified cached data, must write back to memory
- Exclusive
  - Not modified, I have only copy
- Shared
  - Not modified, may be cached elsewhere


# MESI

- Modified
  - Have modified cached data, must write back to memory
- Exclusive
  - Not modified, I have only copy
- Shared
  - Not modified, may be cached elsewhere
- Invalid
  - Cache contents not meaningful


# Processor Issues Load Request




# Memory Responds




# Processor Issues Load Request




# Other Processor Responds



# Modify Cached Data



# Write-Through Cache





# Write-Through Caches

- Immediately broadcast changes
- Good
  - Memory, caches always agree
  - More read hits, maybe
- Bad
  - Bus traffic on all writes
  - Most writes to unshared data
  - For example, loop indexes ...


# Write-Through Caches

- Immediately broadcast changes
  - Good
    - Memory, caches always agree
    - More read hits, maybe
  - Bad
    - Bus traffic on all writes
    - Most writes to unshared data
    - For example, loop indexes ...
- “show stoppers”

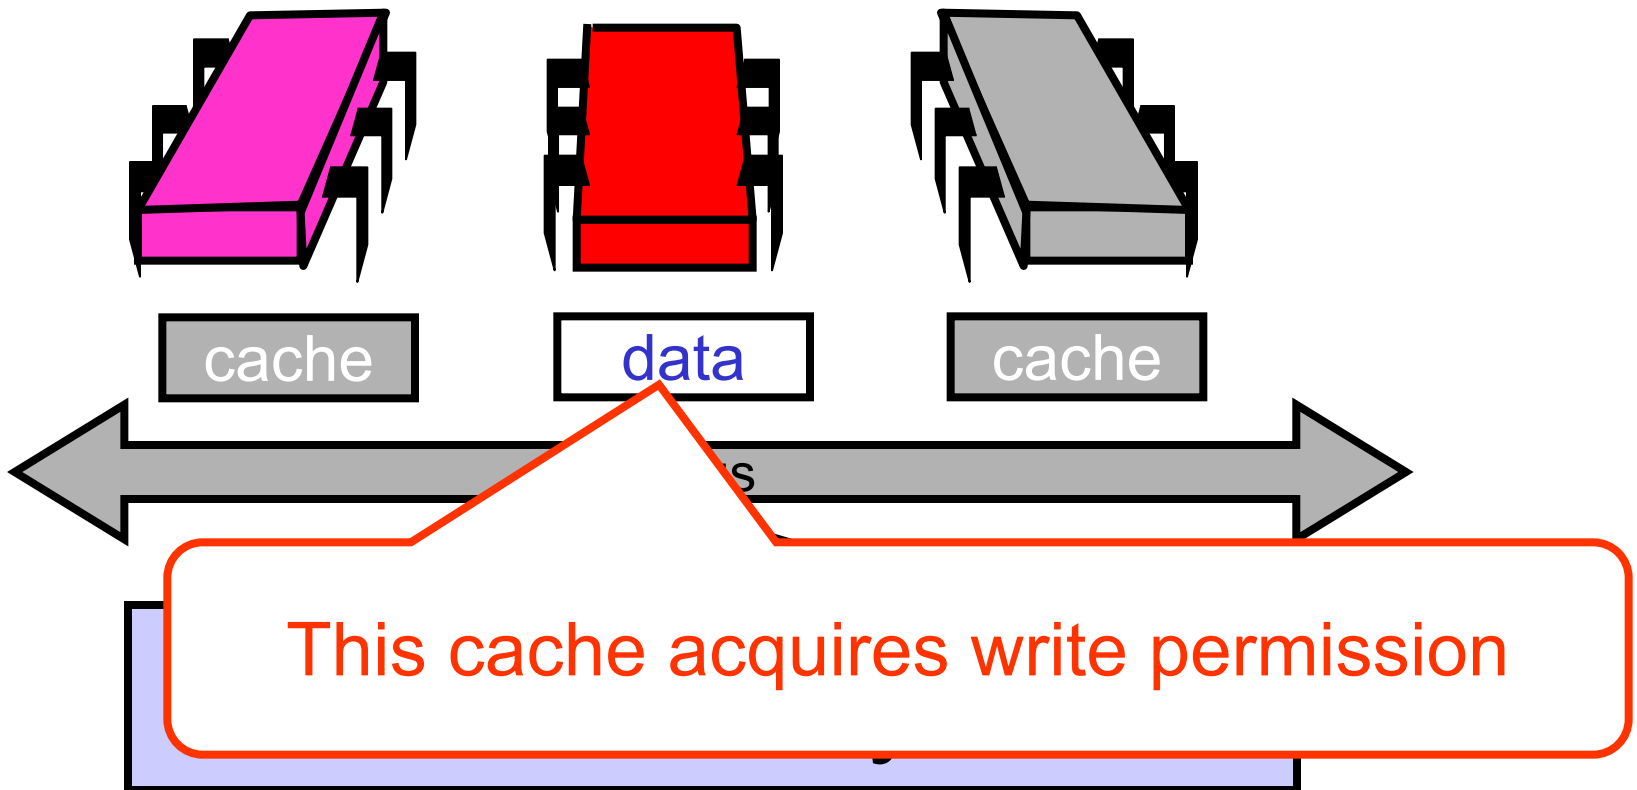
# Write-Back Caches

- Accumulate changes in cache
- Write back when line evicted
  - Need the cache for something else
  - Another processor wants it

# Invalidate




# Invalidate




# Invalidate

Other caches lose read permission



This cache acquires write permission

# Invalidate



# Mutual Exclusion

- What do we want to optimize?
  - Bus bandwidth used by spinning threads
  - Release/Acquire latency
  - Acquire latency for idle lock




# Simple TASLock

- TAS invalidates cache lines
- Spinners
  - Miss in cache
  - Go to bus
- Thread wants to release lock
  - delayed behind spinners


# Test-and-test-and-set

- Wait until lock “looks” free
  - Spin on local cache
  - No bus use while lock busy
- Problem: when lock is released
  - Invalidation storm ...

# Local Spinning while Lock is Busy




# On Release




# On Release

Everyone misses,  
rereads



# On Release

Everyone tries TAS




# Problems

- Everyone misses
  - Reads satisfied sequentially
- Everyone does TAS
  - Invalidates others' caches
- Eventually quiesces after lock acquired
  - How long does this take?

# Measuring Quiescence Time

- Acquire lock
- Pause without using bus
- Use bus heavily




If  $\text{pause} > \text{quiescence time}$ ,  
critical section duration independent of number of threads


If  $\text{pause} < \text{quiescence time}$ ,  
critical section duration slower with more threads



# Quiescence Time




# Mystery Explained




# Solution: Introduce Delay

- If the lock looks free
  - But I fail to get it
- There must be contention
  - Better to back off than to collide again



# Dynamic Example: Exponential Backoff



If I fail to get lock

- Wait random duration before retry
- Each subsequent failure doubles expected wait


# Exponential Backoff Lock

```
public class Backoff implements lock {  
    public void lock() {  
        int delay = MIN_DELAY;  
        while (true) {  
            while (state.get()) {}  
            if (!lock.getAndSet(true))  
                return;  
            sleep(random() % delay);  
            if (delay < MAX_DELAY)  
                delay = 2 * delay;  
        }  
    }  
}
```

# Exponential Backoff Lock

```
public class Backoff implements lock {  
    public void lock() {  
        int delay = MIN_DELAY;  
        while (true) {  
            while (state.get()) {}  
            if (!lock.getAndSet(true))  
                return;  
            sleep(random() % delay);  
            if (delay < MAX_DELAY)  
                delay = 2 * delay;  
        }  
    }  
}
```

**Fix minimum delay**



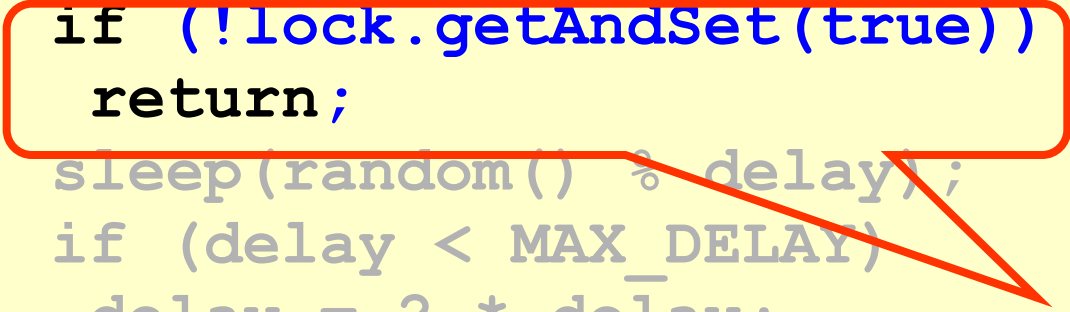
# Exponential Backoff Lock

```
public class Backoff implements lock {  
    public void lock() {  
        int delay = MIN_DELAY;  
        while (true) {  
            while (state.get()) {}  
            if (!lock.getAndSet(true))  
                return;  
            sleep(random() % delay);  
            if (delay < MAX_DELAY)  
                delay = 2 * delay;  
        }  
    }  
}
```

**Wait until lock looks free**

# Exponential Backoff Lock

```
public class Backoff implements lock {  
    public void lock() {  
        int delay = MIN_DELAY;  
        while (true) {  
            while (state.get()) {}  
            if (!lock.getAndSet(true))  
                return;  
            sleep(random() % delay);  
            if (delay < MAX_DELAY)  
                delay = 2 * delay;  
        }  
    }  
}
```



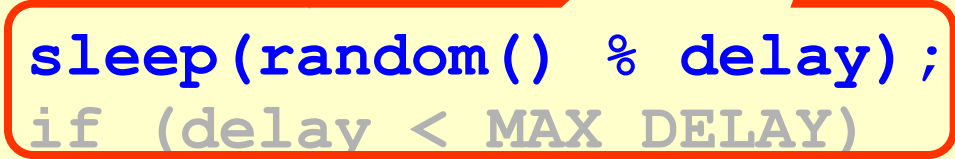
**If we win, return**



# Exponential Backoff Lock

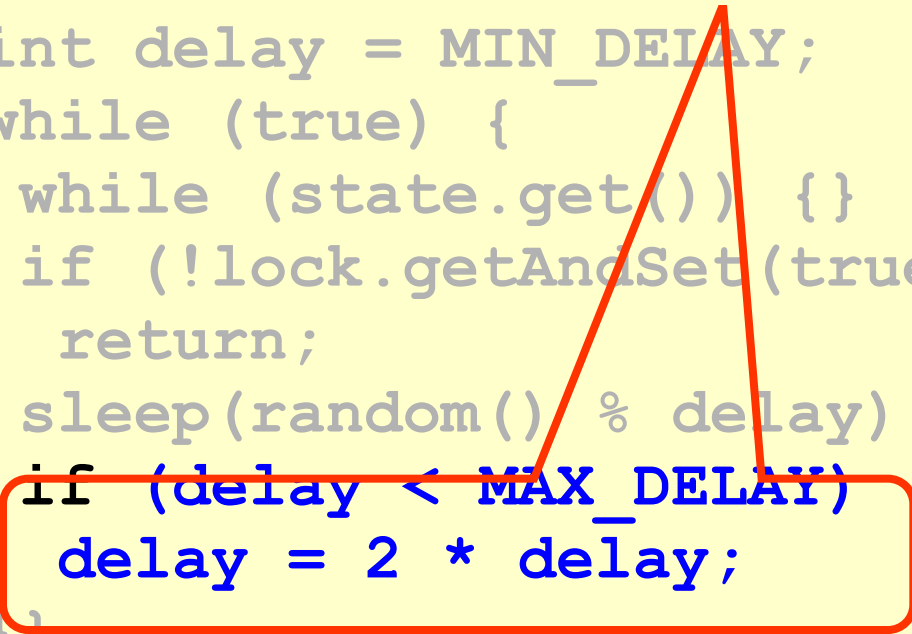
```
public class ExponentialBackoffLock {  
    public  
    int delay = MIN_DELAY;  
    while (true) {  
        while (state.get()) {}  
        if (!lock.getAndSet(true))  
            return;  
        sleep(random() % delay);  
        if (delay < MAX_DELAY)  
            delay = 2 * delay;  
    }  
}
```

**Back off for random duration**




# Exponential Backoff Lock

```
pub: ...
pul Double max delay, within reason
    int delay = MIN_DELAY;
    while (true) {
        while (state.get()) {}
        if (!lock.getAndSet(true))
            return;
        sleep(random() % delay);
        if (delay < MAX_DELAY)
            delay = 2 * delay;
    }
}
```

A red triangle is drawn with its base at the bottom, spanning the width of the 'if' and 'delay = 2 \* delay;' lines. Its apex points upwards towards the 'sleep' line. A red rounded rectangle is drawn around the 'if (delay < MAX\_DELAY)' and 'delay = 2 \* delay;' lines.


# Spin-Waiting Overhead



# Backoff: Other Issues

- Good
  - Easy to implement
  - Beats TTAS lock
- Bad
  - Must choose parameters carefully
  - Not portable across platforms


# Actual Data on 40-Core Machine




# Idea

- Avoid useless invalidations
  - By keeping a queue of threads
- Each thread
  - Notifies next in line
  - Without bothering the others

# Anderson Queue Lock




# Anderson Queue Lock







# Anderson Queue Lock




# Anderson Queue Lock




# Anderson Queue Lock




# Anderson Queue Lock




# Anderson Queue Lock




# Anderson Queue Lock



# Anderson Queue Lock



# Anderson Queue Lock





# Anderson Queue Lock

```
class ALock implements Lock {  
    boolean[] flags={true,false,...,false};  
    AtomicInteger next  
        = new AtomicInteger(0);  
    ThreadLocal<Integer> mySlot;
```

# Anderson Queue Lock

```
class ALock implements Lock {  
    boolean[] flags={true,false,...,false};  
    AtomicInteger next  
        = new AtomicInteger(0);  
    ThreadLocal<Integer> mySlot;
```

**One flag per thread**

# Anderson Queue Lock

```
class ALock implements Lock {  
    boolean[] flags={true,false,...,false};  
    AtomicInteger next  
    = new AtomicInteger(0);  
    ThreadLocal<Integer> mySlot;
```

**Next flag to use**

# Anderson Queue Lock

```
class ALock implements Lock {  
    boolean[] flags={true,false,...,false};  
    AtomicInteger next  
        = new AtomicInteger(0);  
    ThreadLocal<Integer> mySlot;
```

**Thread-local variable**

# Anderson Queue Lock

```
public lock() {  
    mySlot = next.getAndIncrement();  
    while (!flags[mySlot % n]) {};  
    flags[mySlot % n] = false;  
}  
  
public unlock() {  
    flags[(mySlot+1) % n] = true;  
}
```

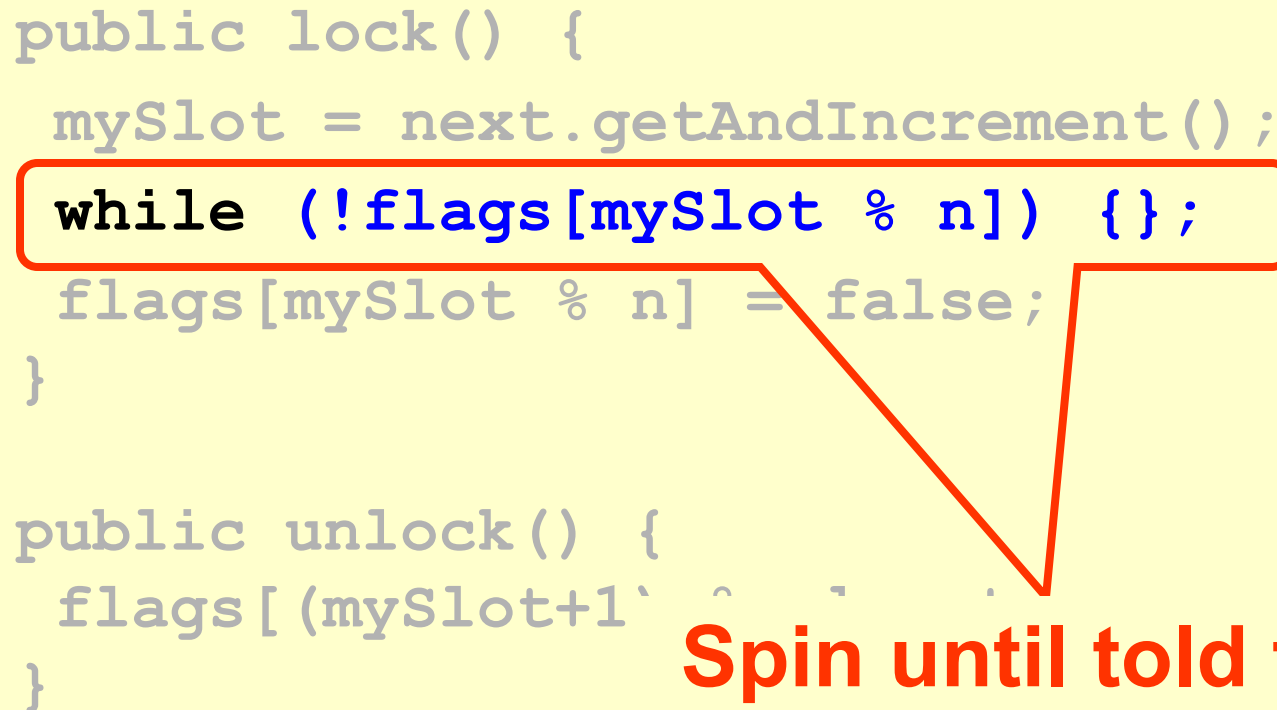
# Anderson Queue Lock

```
public lock() {  
    mySlot = next.getAndIncrement();  
    while (!flags[mySlot % n]) {};  
    flags[mySlot % n] = false;  
}  
  
public unlock() {  
    flags[(mySlot+1) % n]  
}
```

**Take next slot**

# Anderson Queue Lock

```
public lock() {  
    mySlot = next.getAndIncrement();  
    while (!flags[mySlot % n]) {};  
    flags[mySlot % n] = false;  
}  
  
public unlock() {  
    flags[(mySlot+1) % n] = true;  
}
```



**Spin until told to go**

# Anderson Queue Lock

```
public lock() {  
    myslot = next.getAndIncrement();  
    while (!flags[myslot % n]) {};  
    flags[myslot % n] = false;  
}
```

```
public unlock() {  
    flags[(myslot+1) % n] = true;  
}
```

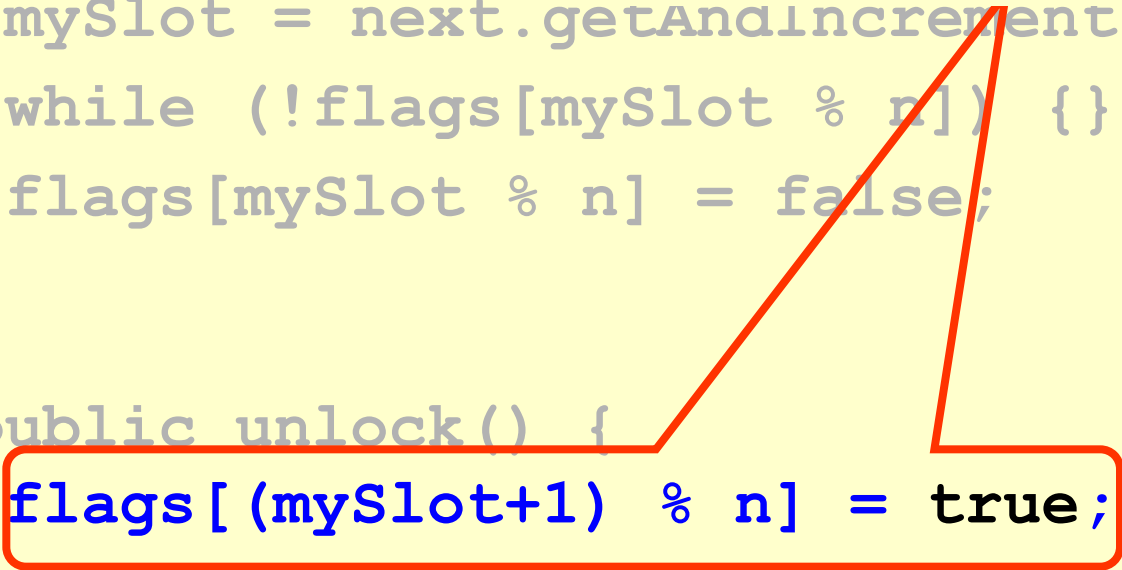
**Prepare slot for re-use**




# Anderson Queue Lock

```
public lock() { Tell next thread to go
    mySlot = next.getAndIncrement();
    while (!flags[mySlot % n]) {};
    flags[mySlot % n] = false;
}

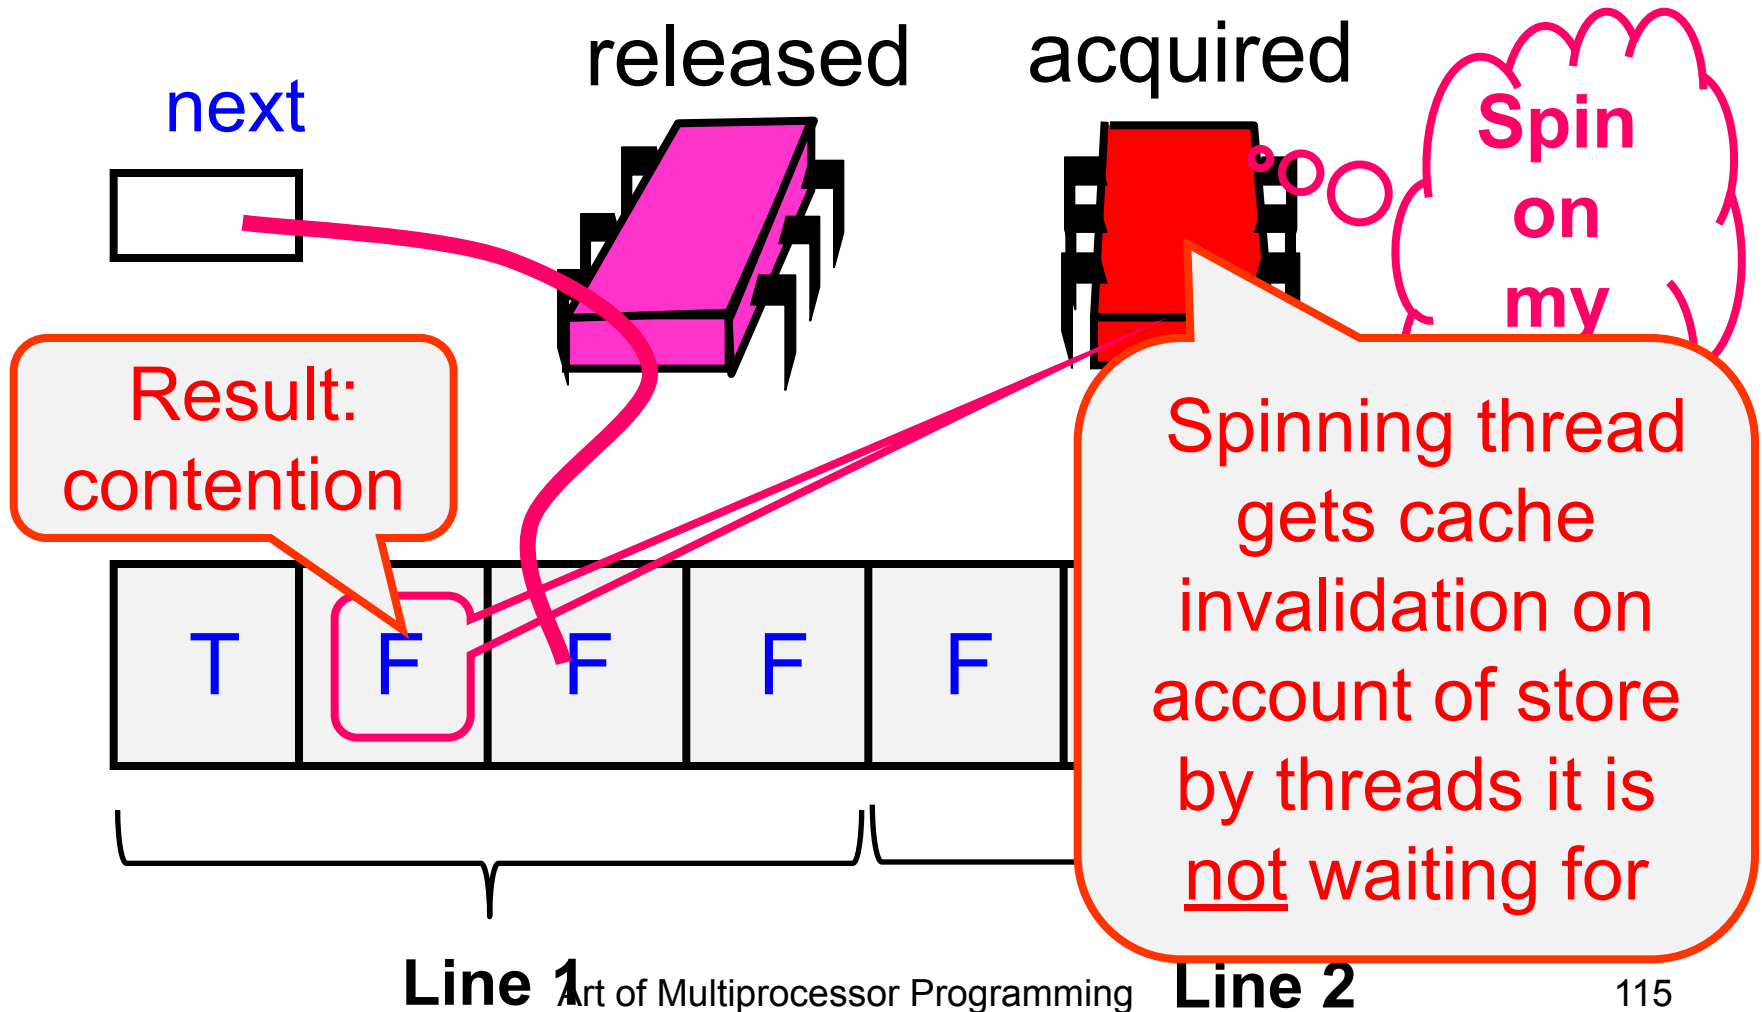
public unlock() {
    flags[(mySlot+1) % n] = true;
}
```




# Local Spinning




# False Sharing



# The Solution: Padding



# Performance



TTAS

queue

- Shorter handover than backoff
- Curve is practically flat
- Scalable performance

# Anderson Queue Lock

## Good

- First truly scalable lock
- Simple, easy to implement
- Back to FCFS order (like Bakery)

# Anderson Queue Lock

## Bad



- Space hog...
- One bit per thread ☐ one cache line per thread
  - What if unknown number of threads?
  - What if small number of actual contenders?

# CLH Lock



- FCFS order
- Small, constant-size overhead per thread





# Initially





# Initially



# Initially




# Initially




# Purple Wants the Lock


acquiring




tail



# Purple Wants the Lock




# Purple Wants the Lock



# Purple Has the Lock


acquired




tail



false

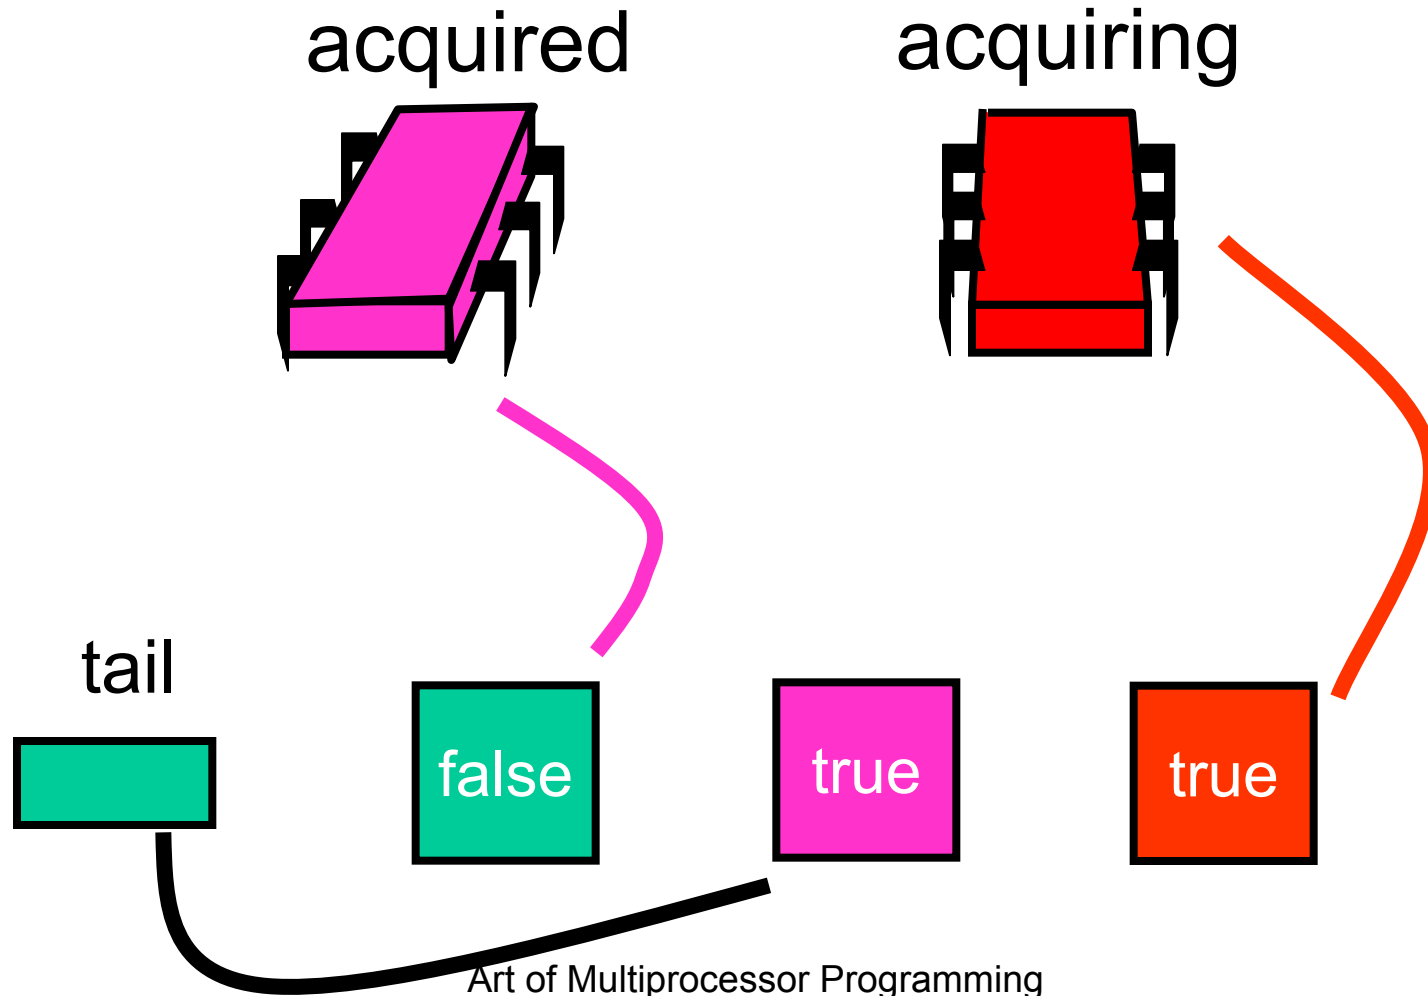


true

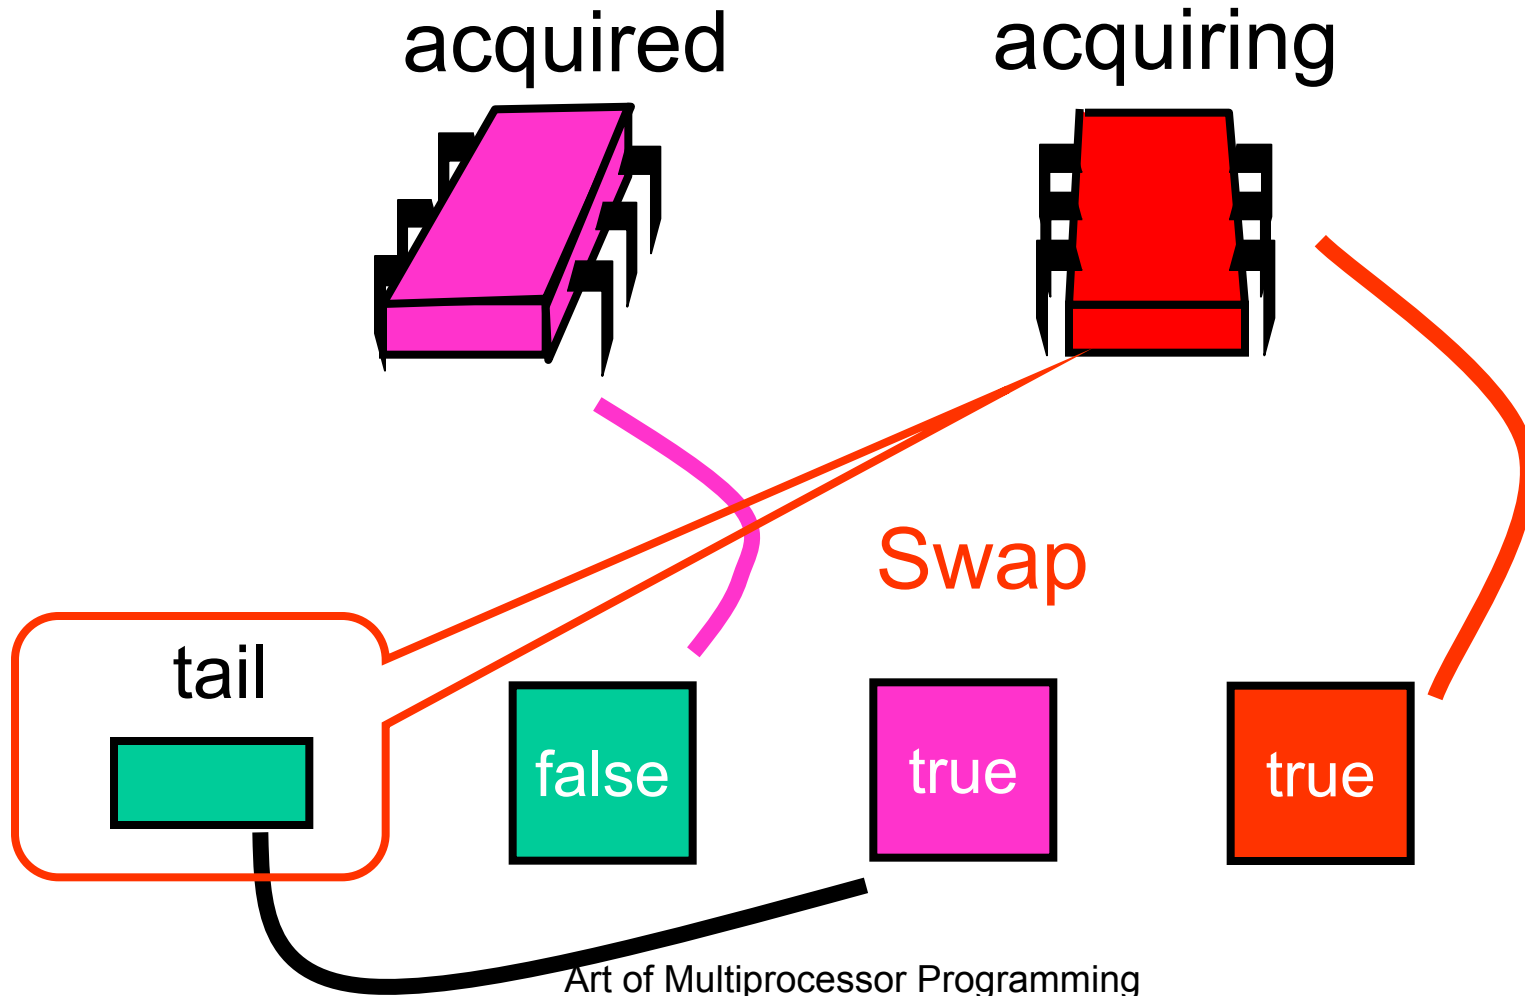




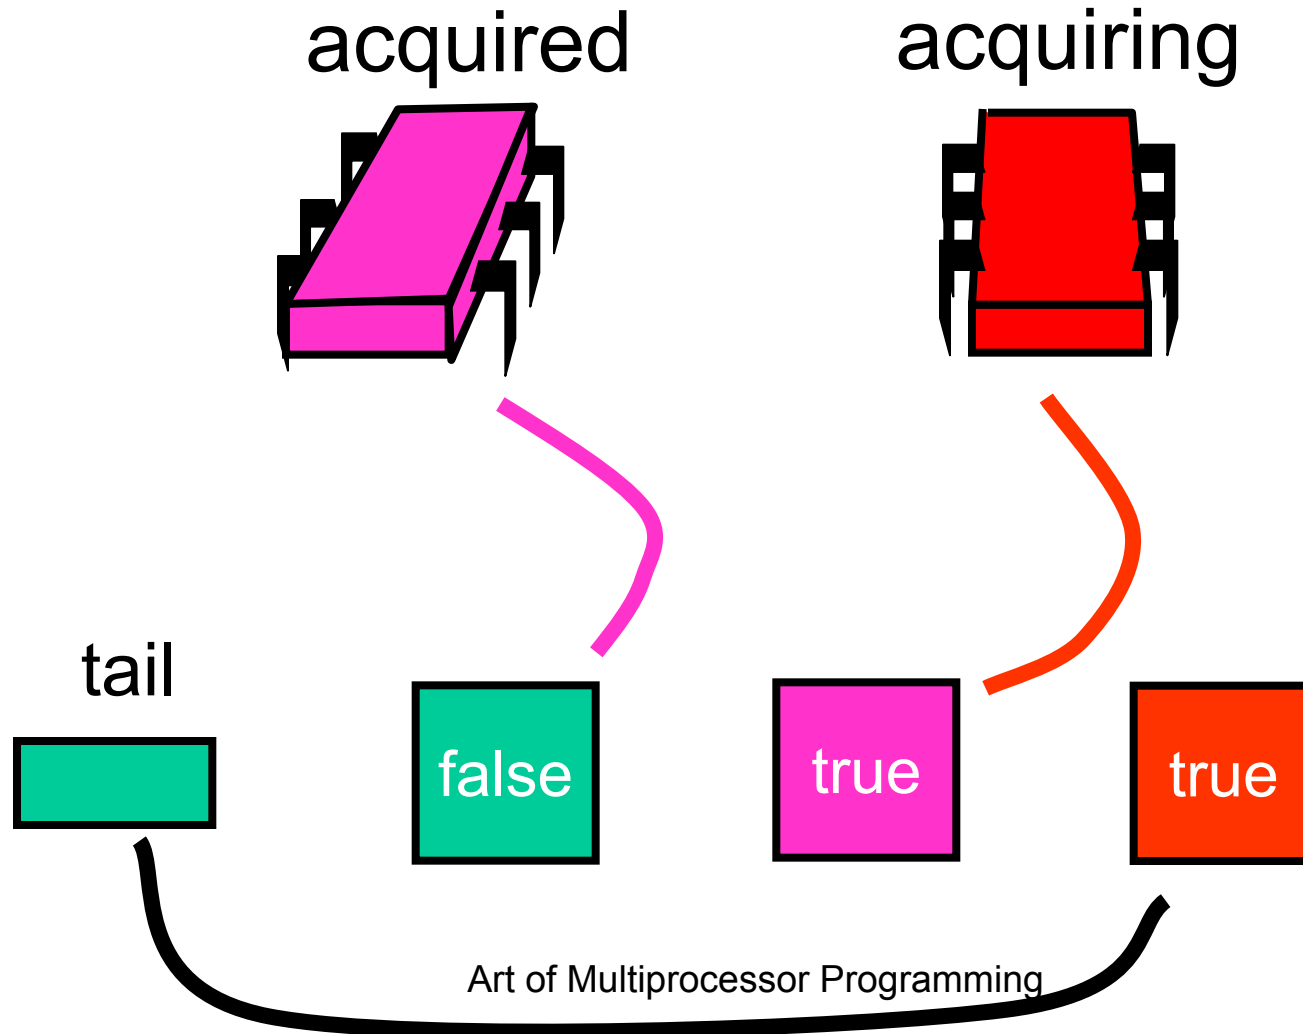
# Red Wants the Lock



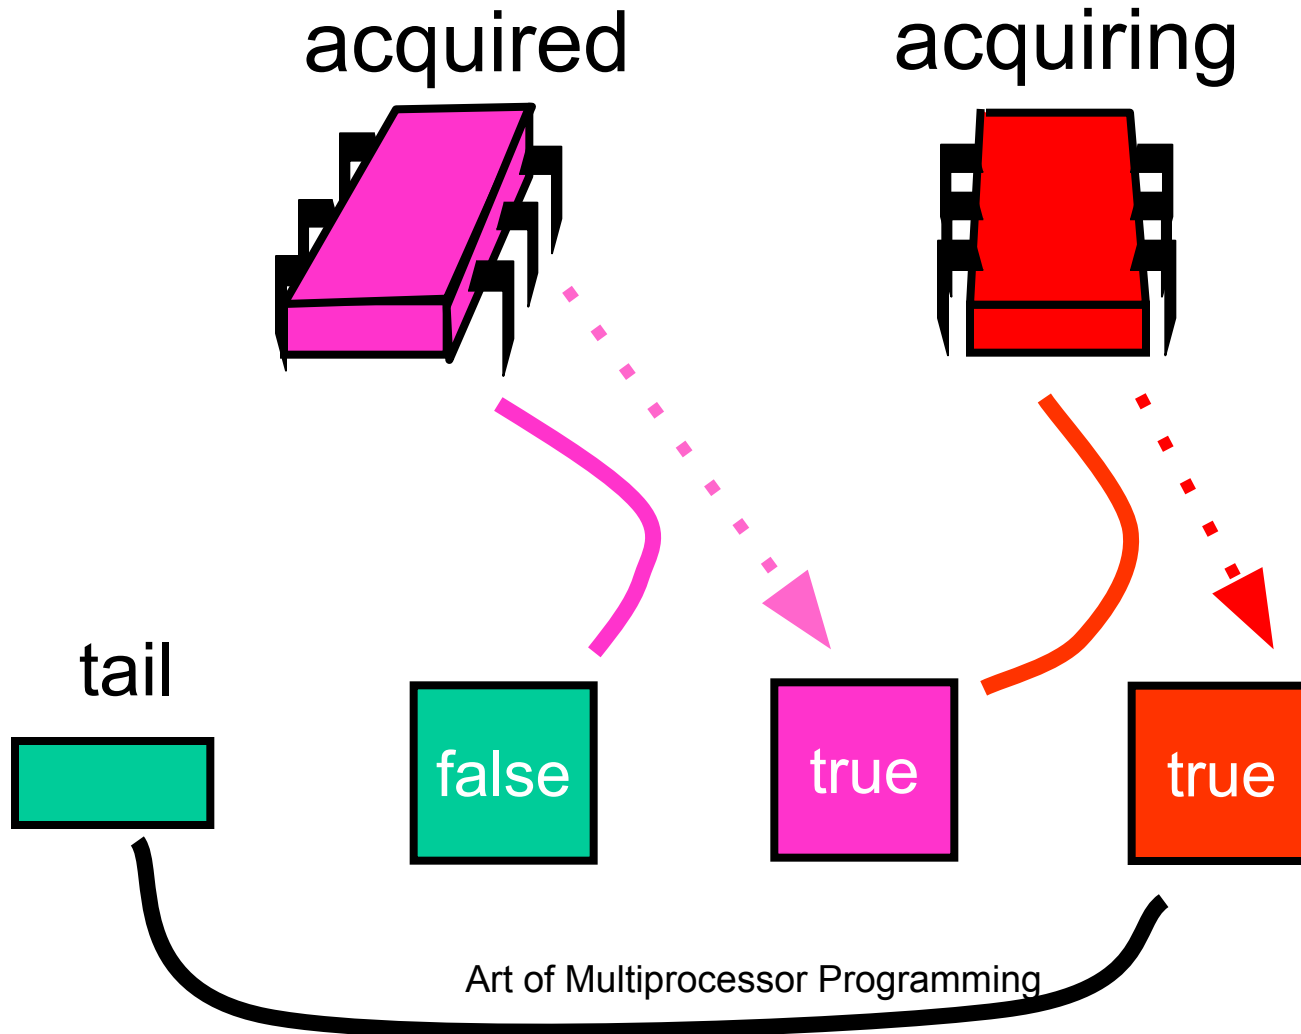
# Red Wants the Lock



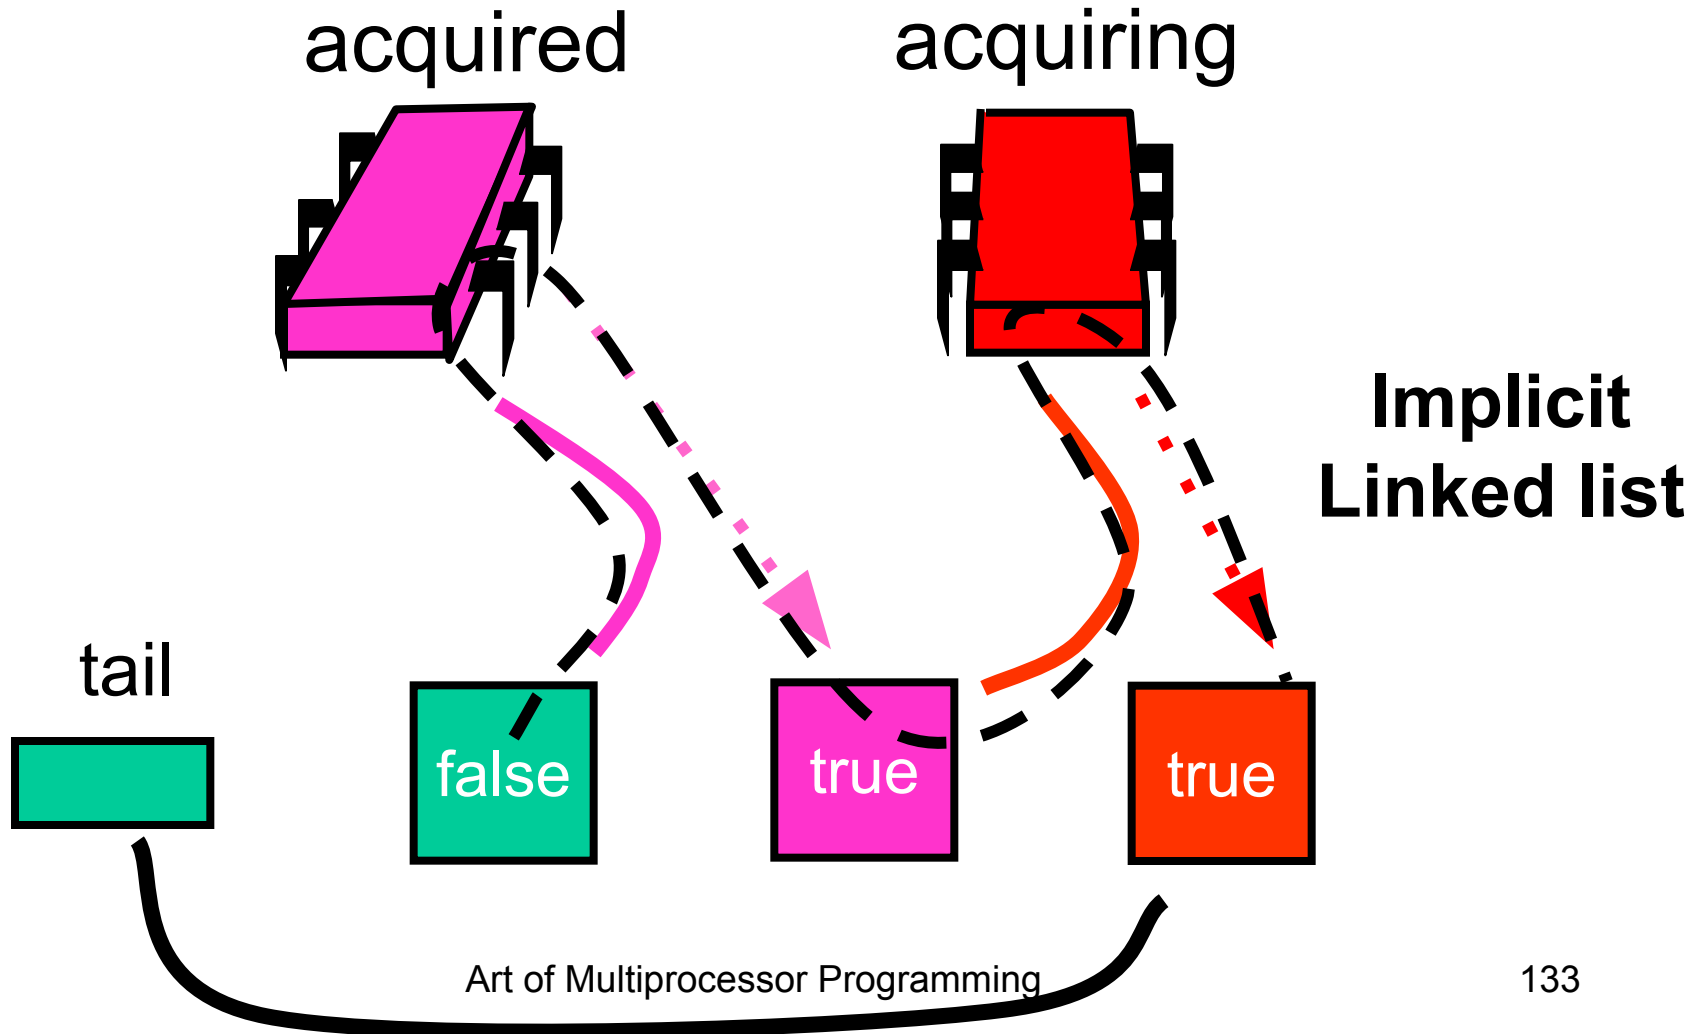
# Red Wants the Lock



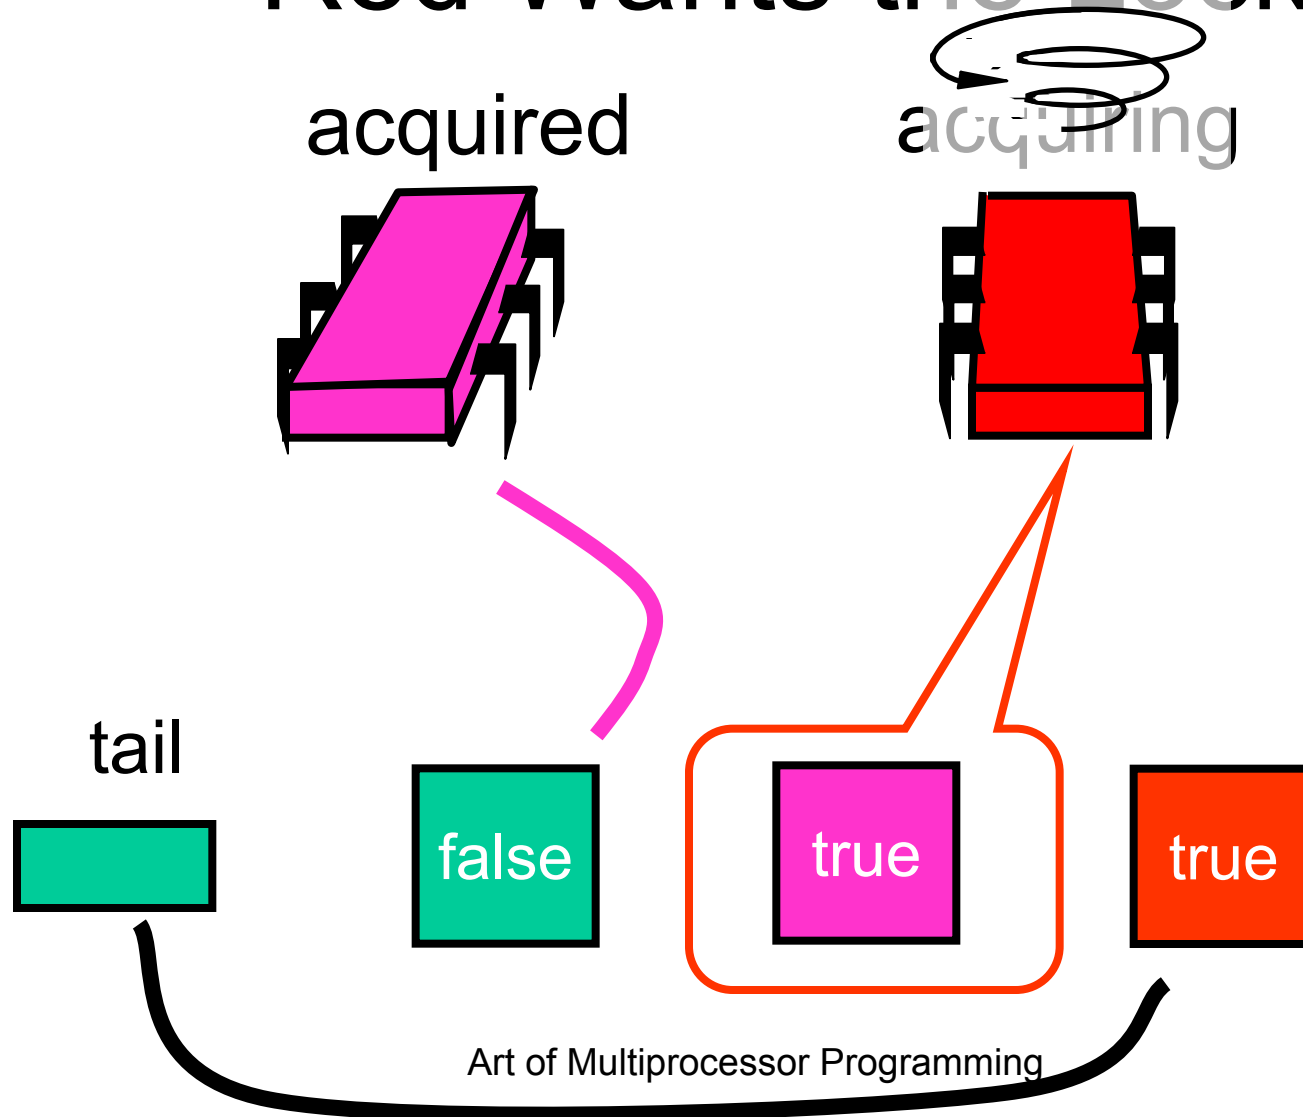
# Red Wants the Lock



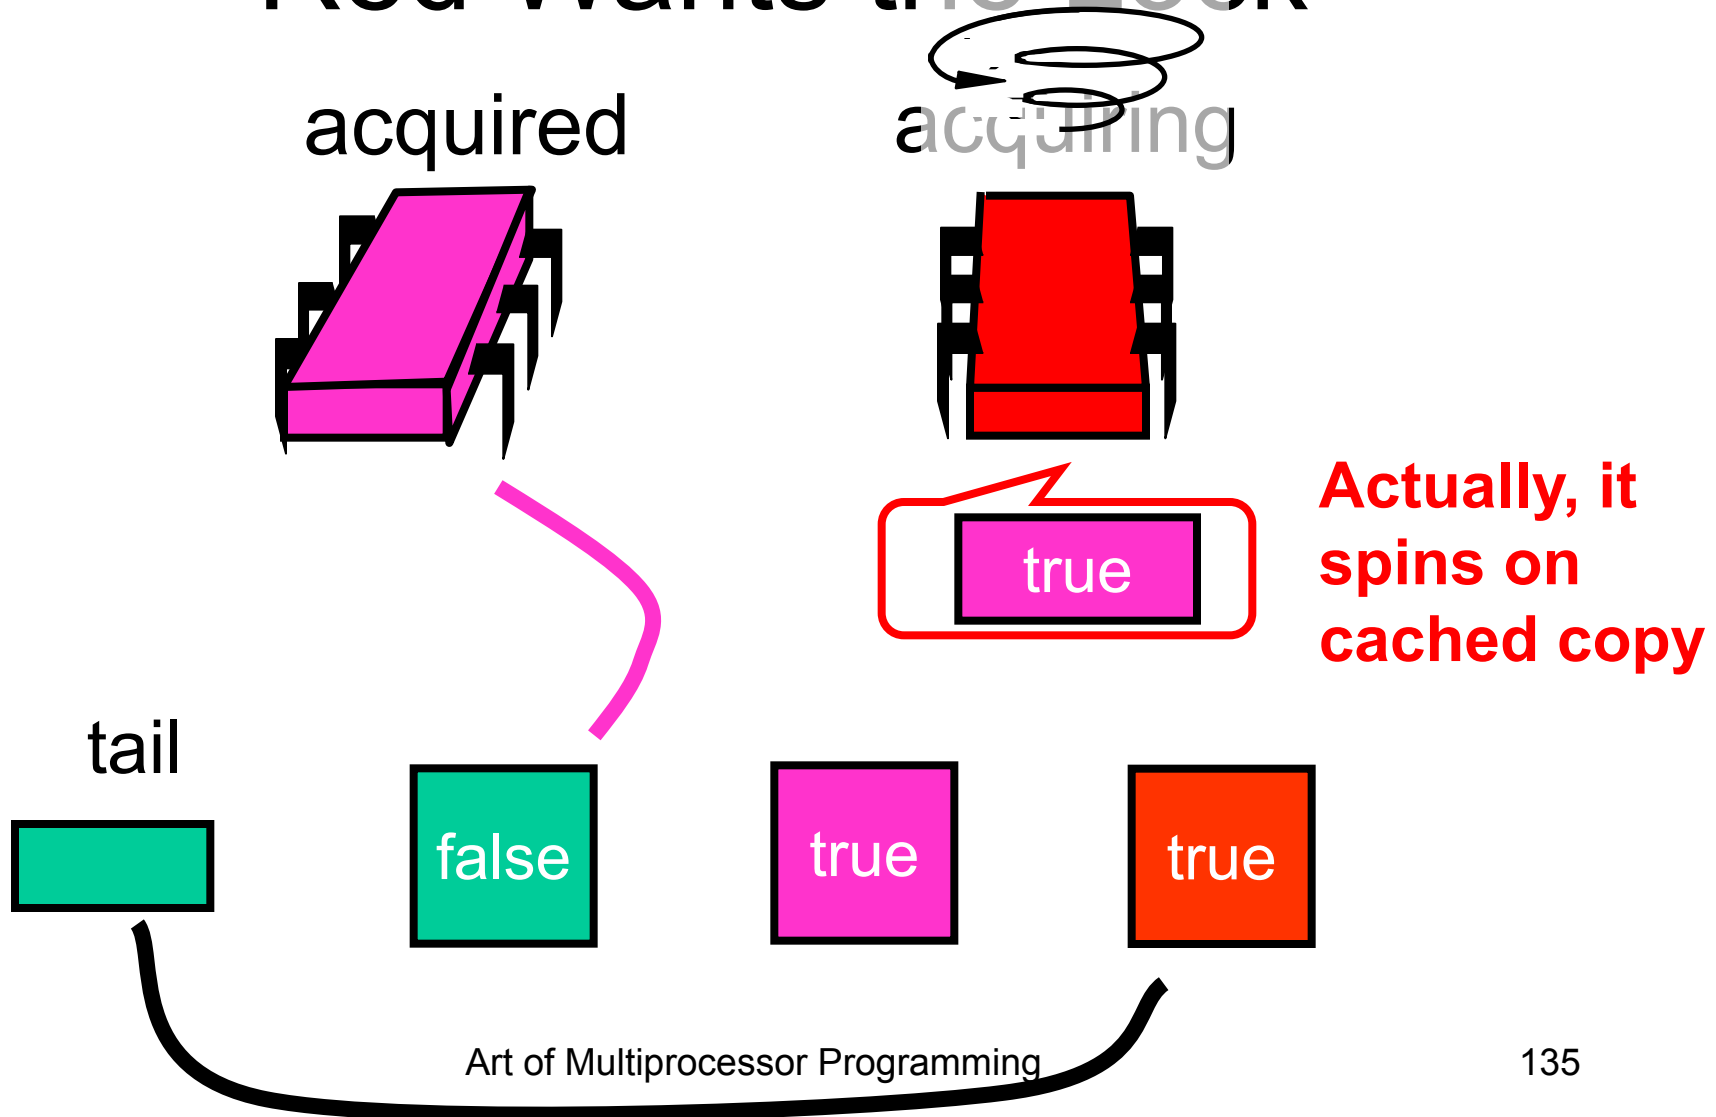
# Red Wants the Lock



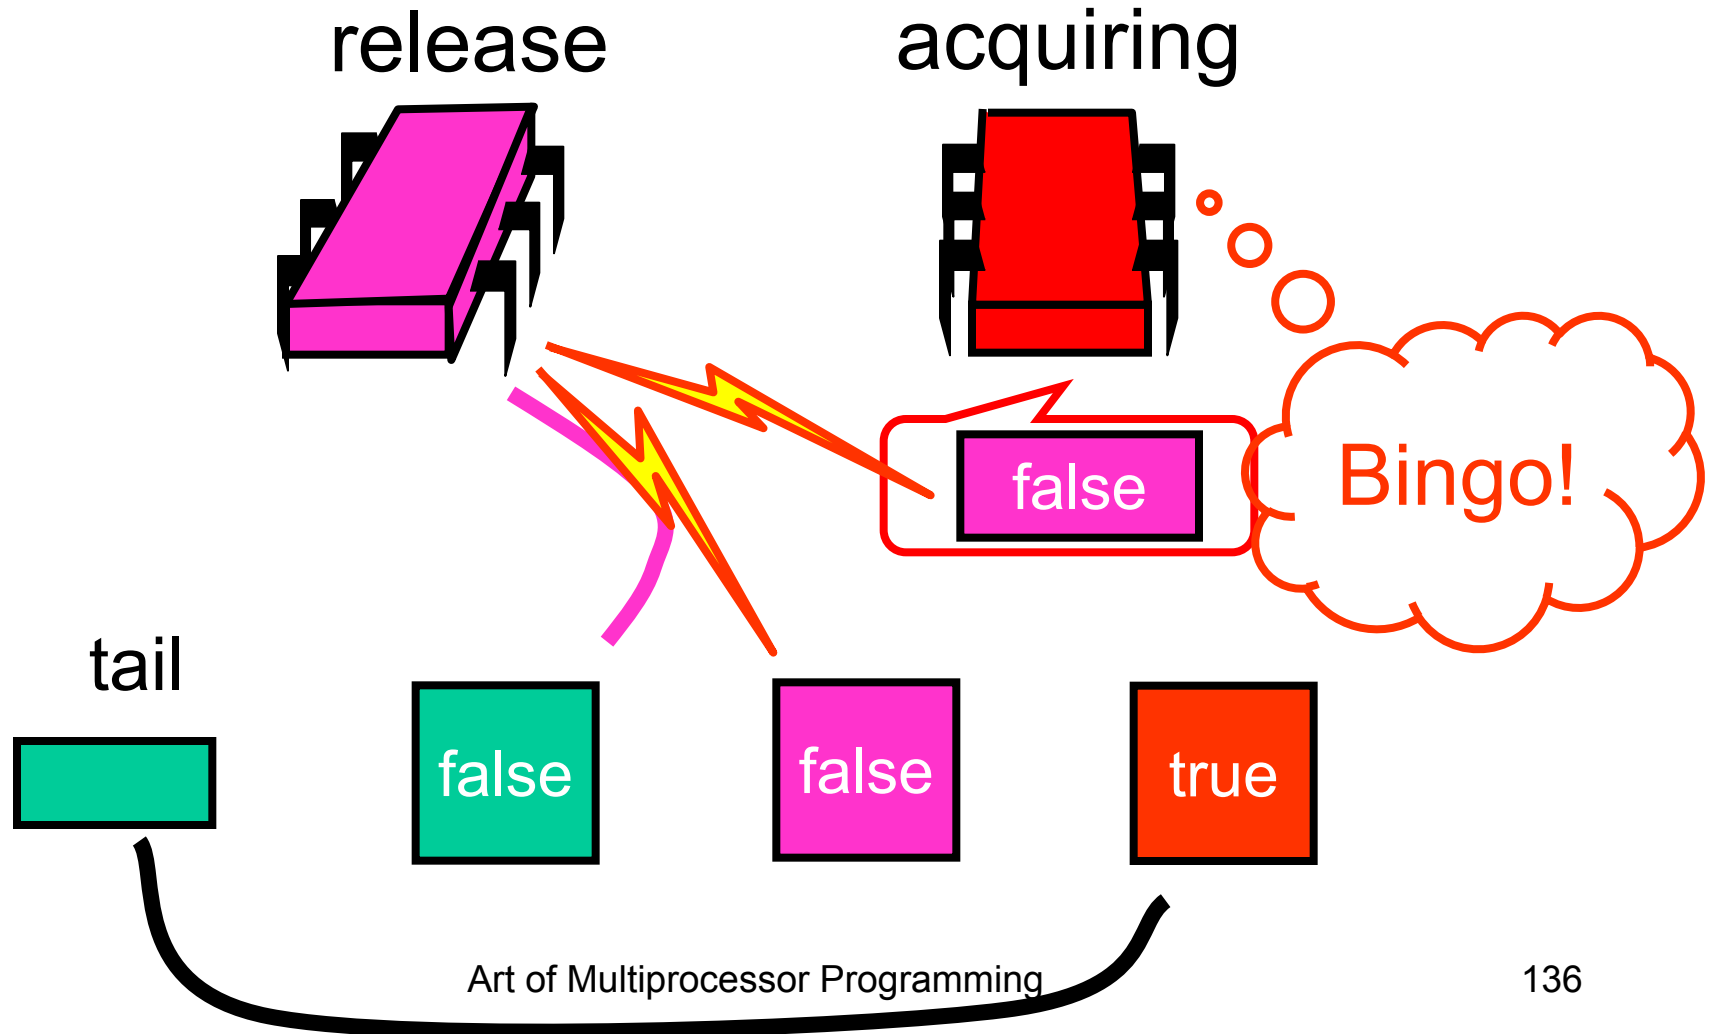
# Red Wants the Lock



# Red Wants the Lock




# Purple Releases






# Purple Releases

released



acquired



tail



true



# Space Usage

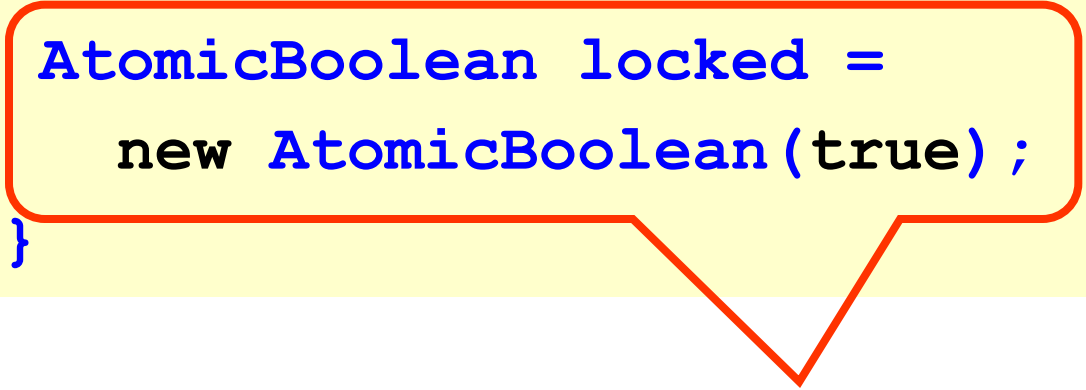
- Let
  - $L$  = number of locks
  - $N$  = number of threads
- ALock
  - $O(LN)$
- CLH lock
  - $O(L+N)$

# CLH Queue Lock

```
class QNode {  
    AtomicBoolean locked =  
        new AtomicBoolean(true) ;  
}
```

# CLH Queue Lock

```
class QNode {  
    AtomicBoolean locked =  
        new AtomicBoolean(true) ;  
}
```



**Not released yet**

# CLH Queue Lock

```
class CLHLock implements Lock {  
    AtomicReference<QNode> tail;  
    ThreadLocal<QNode> myNode  
        = new QNode();  
    public void lock() {  
        QNode pred  
            = tail.getAndSet(myNode);  
        while (pred.locked) {}  
    }  
}
```

# CLH Queue Lock

```
class CLHLock implements Lock {  
    AtomicReference<QNode> tail;  
    ThreadLocal<QNode> myNode  
        = new QNode();  
    public void lock() {  
        QNode pred  
            = tail.getAndSet(myNode);  
        while (pred.locked) {}  
    }  
}
```

**Queue tail**

# CLH Queue Lock

```
class CLHLock implements Lock {  
    AtomicReference<QNode> tail;  
    ThreadLocal<QNode> myNode  
        = new QNode();  
    public void lock() {  
        QNode pred  
            = tail.getAndSet(myNode);  
        while (pred.locked) {}  
    }  
}
```

Thread-local QNode

# CLH Queue Lock

```
class CLHLock implements Lock {  
    AtomicReference<QNode> tail;  
    ThreadLocal<QNode> myNode  
        = new QNode();  
    public void lock() {  
        QNode pred  
            = tail.getAndSet(myNode);  
        while (pred.locked) {}  
    }  
}
```

Swap in my node





# CLH Queue Lock

```
class CLHLock implements Lock {  
    AtomicReference<QNode> tail;  
    ThreadLocal<QNode> myNode  
        = new QNode();  
    public void lock() {  
        QNode pred  
            = tail.getAndSet(myNode);  
        while (pred.locked) {}  
    }  
}
```

Spin until predecessor  
releases lock



# CLH Queue Lock

```
Class CLHLock implements Lock {  
    ...  
    public void unlock() {  
        myNode.locked.set(false);  
        myNode = pred;  
    }  
}
```

# CLH Queue Lock

```
Class CLHLock implements Lock {  
    ...  
    public void unlock() {  
        myNode.locked.set(false);  
        myNode = pred;  
    }  
}
```

Notify successor

# CLH Queue Lock

```
Class CLHLock implements Lock {  
    ...  
    public void unlock() {  
        myNode.locked.set(false);  
        myNode = pred;  
    }  
}
```

Recycle  
predecessor's node

# CLH Queue Lock

```
Class CLHLock implements Lock {  
    ...  
    public void unlock() {  
        myNode.locked.set(false);  
        myNode = pred;  
    }  
}
```

**(we don't actually reuse myNode.  
Code in book shows how it's done.)**


# CLH Lock

- Good
  - Lock release affects predecessor only
  - Small, constant-sized space
- Bad
  - Doesn't work for uncached NUMA architectures

# NUMA and cc-NUMA Architectures

- Acronym:
  - **N**on-**U**niform **M**emory **A**rchitecture
  - ccNUMA = cache coherent NUMA
- Illusion:
  - Flat shared memory
- Truth:
  - No caches (sometimes)
  - Some memory regions faster than others


# NUMA Machines



**Spinning on local  
memory is fast**



# NUMA Machines



**Spinning on remote  
memory is slow**



# CLH Lock

- Each thread spins on predecessor's memory
- Could be far away ...

# MCS Lock


- FCFS order
- Spin on local memory only
- Small, Constant-size overhead

# Initially





# Acquiring


acquiring




(allocate QNode)




# Acquiring



# Acquiring




# Acquired







# Acquiring



# Acquiring







# Acquiring



# Acquiring

acquiring

acquired




tail




# Acquiring

acquiring

acquired



# Acquiring



# MCS Queue Lock

```
class QNode {  
    volatile boolean locked = false;  
    volatile qnode next = null;  
}
```

# MCS Queue Lock

```
class MCSLock implements Lock {  
    AtomicReference tail;  
    public void lock() {  
        QNode qnode = new QNode();  
        QNode pred = tail.getAndSet(qnode);  
        if (pred != null) {  
            qnode.locked = true;  
            pred.next = qnode;  
            while (qnode.locked) {}  
        }  
    }  
}
```



# MCS Queue Lock

```
class MCSLock implements Lock {  
    AtomicReference tail;  
    public void lock() {  
        QNode qnode = new QNode();  
        QNode pred = tail.getAndSet(qnode);  
        if (pred != null) {  
            qnode.locked = true;  
            pred.next = qnode;  
            while (qnode.locked) {}  
        }  
    }  
}
```

**Make a  
QNode**



# MCS Queue Lock

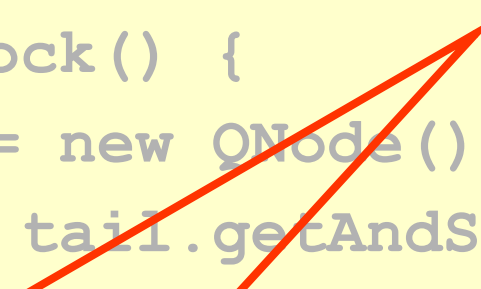
```
class MCSLock implements Lock {  
    AtomicReference tail;  
    public void lock() {  
        QNode qnode = new QNode();  
        QNode pred = tail.getAndSet(qnode);  
        if (pred != null) {  
            qnode.locked = true;  
            pred.next = qnode;  
            while (qnode.locked) {}  
        }  
    }  
}
```

**add my Node to  
the tail of queue**

# MCS Queue Lock

```
class MCSLock implements Lock {  
    AtomicReference tail;  
    public void lock() {  
        QNode qnode = new QNode();  
        QNode pred = tail.getAndSet(qnode);  
        if (pred != null) {  
            qnode.locked = true;  
            pred.next = qnode;  
        }  
        while (qnode.locked) {}  
    }  
}
```

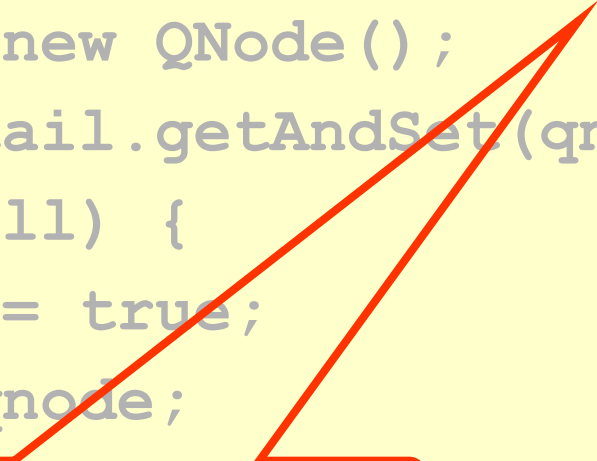
**Fix if queue was  
non-empty**




# MCS Queue Lock

```
class MCSLock implements Lock {  
    AtomicReference tail;  
    public void lock() {  
        QNode qnode = new QNode();  
        QNode pred = tail.getAndSet(qnode);  
        if (pred != null) {  
            qnode.locked = true;  
            pred.next = qnode;  
            while (qnode.locked) {}  
        }  
    }  
}
```

**Wait until  
unlocked**



# Purple Release




Pu

I don't see a successor. But by looking at the queue, I see another thread is active

releasing


swap






Pu

I don't see a successor. But by looking at the queue, I see another thread is active

releasing




swap





I have to release that thread so must wait for it to identify its node

# Purple Release

releasing




prepare to spin






# Purple Release




# Purple Release





# Purple Release

releasing



Acquired lock



# MCS Queue Unlock

```
class MCSLock implements Lock {  
    AtomicReference tail;  
    public void unlock() {  
        if (qnode.next == null) {  
            if (tail.CAS(qnode, null)  
                return;  
            while (qnode.next == null) {}  
        }  
        qnode.next.locked = false;  
    }  
}
```

# MCS Queue Lock

```
class MCSLock implements Lock {  
    AtomicReference tail;  
    public void unlock() {  
        if (qnode.next == null) {  
            if (tail.CAS(qnode, null)  
                return;  
            while (qnode.next == null) {}  
        }  
        qnode.next.locked = false;  
    }  
}
```

**Missing  
successor?**

# MCS Queue Lock

**If really no successor,  
return**

```
public void unlock() {  
    if (qnode.next == null) {  
        if (tail.CAS(qnode, null)  
            return;  
        while (qnode.next == null) {}  
    }  
    qnode.next.locked = false;  
}}
```

# MCS Queue Lock

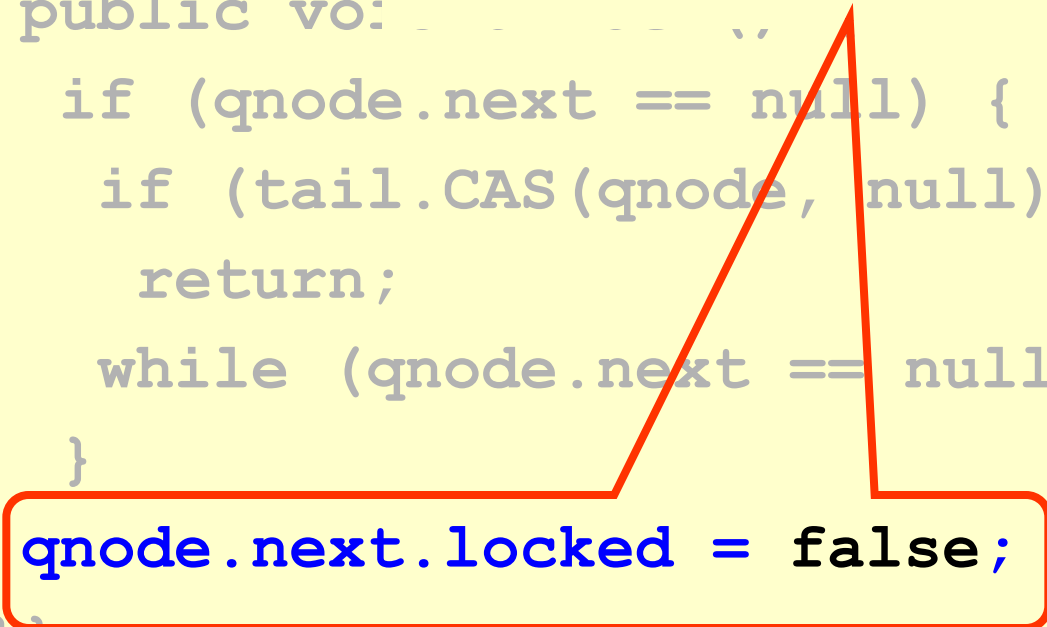
**Otherwise wait for  
successor to catch up**

```
public void unlock() {  
    if (qnode.next == null) {  
        if (tail.CAS(qnode, null)  
            return;  
        while (qnode.next == null) {}  
    }  
    qnode.next.locked = false;  
}}
```

# MCS Queue Lock

```
class MCSLock implements Lock {
    AtomicReference<Node> tail;
    public void lock() {
        if (qnode.next == null) {
            if (tail.CAS(qnode, null))
                return;
            while (qnode.next == null) {}
        }
        qnode.next.locked = false;
    }
}
```

**Pass lock to successor**





# Abortable Locks

- What if you want to give up waiting for a lock?
- For example
  - Timeout
  - Database transaction aborted by user


# Back-off Lock

- Aborting is trivial
  - Just return from lock() call
- Extra benefit:
  - No cleaning up
  - Wait-free
  - Immediate return


# Queue Locks

- Can't just quit
  - Thread in line behind will starve
- Need a graceful way out


# Queue Locks




# Queue Locks



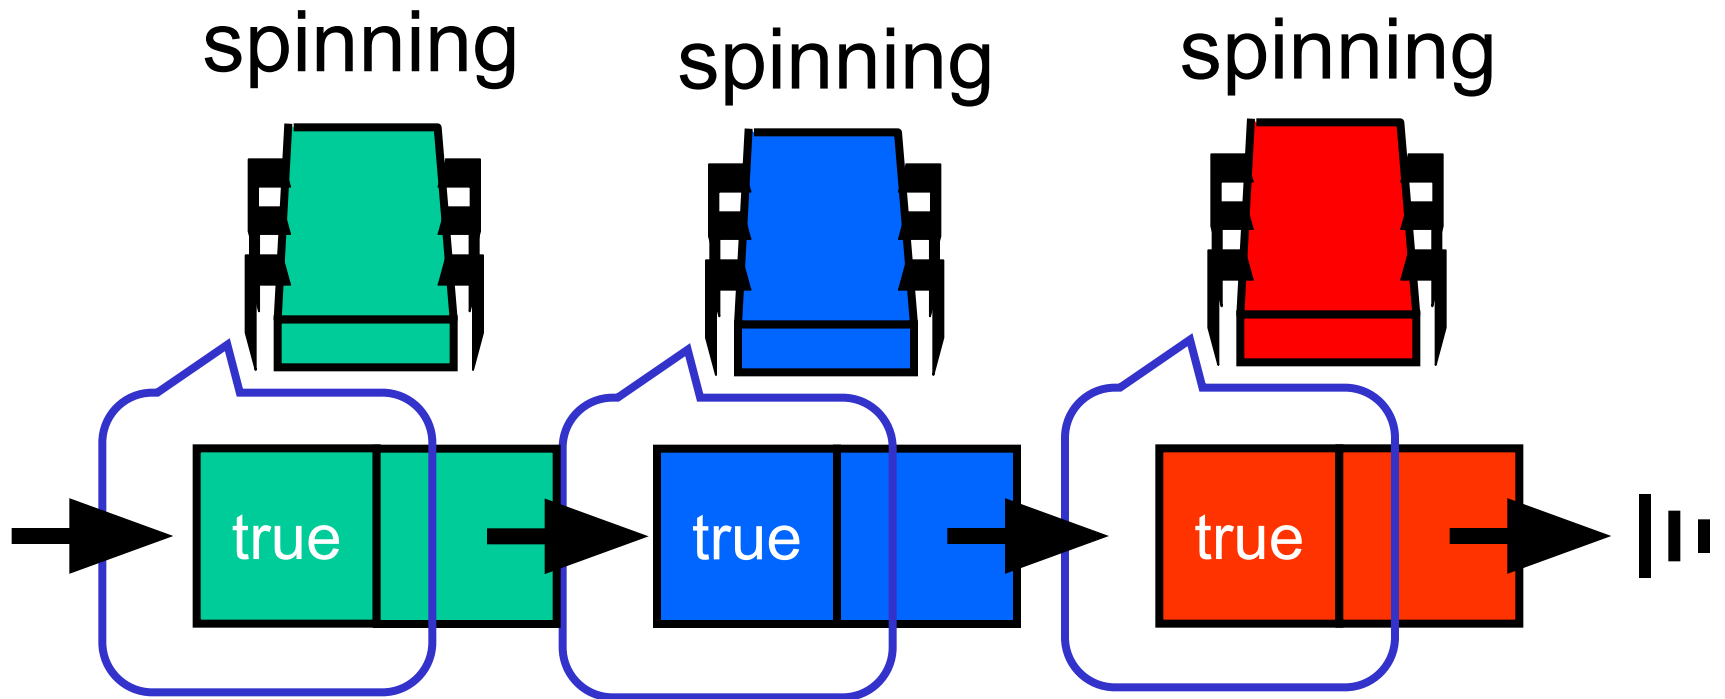
# Queue Locks



# Queue Locks




# Queue Locks







# Queue Locks




# Queue Locks



# Queue Locks




# Queue Locks




# Abortable CLH Lock

- When a thread gives up
  - Removing node in a wait-free way is hard
- Idea:
  - let successor deal with it.


# Initially



# Initially





Distinguished  
available node  
means lock is  
free



# Acquiring

acquiring






# Acquiring


Null predecessor  
means lock not  
released or  
aborted

acquiring




# Acquiring

acquiring

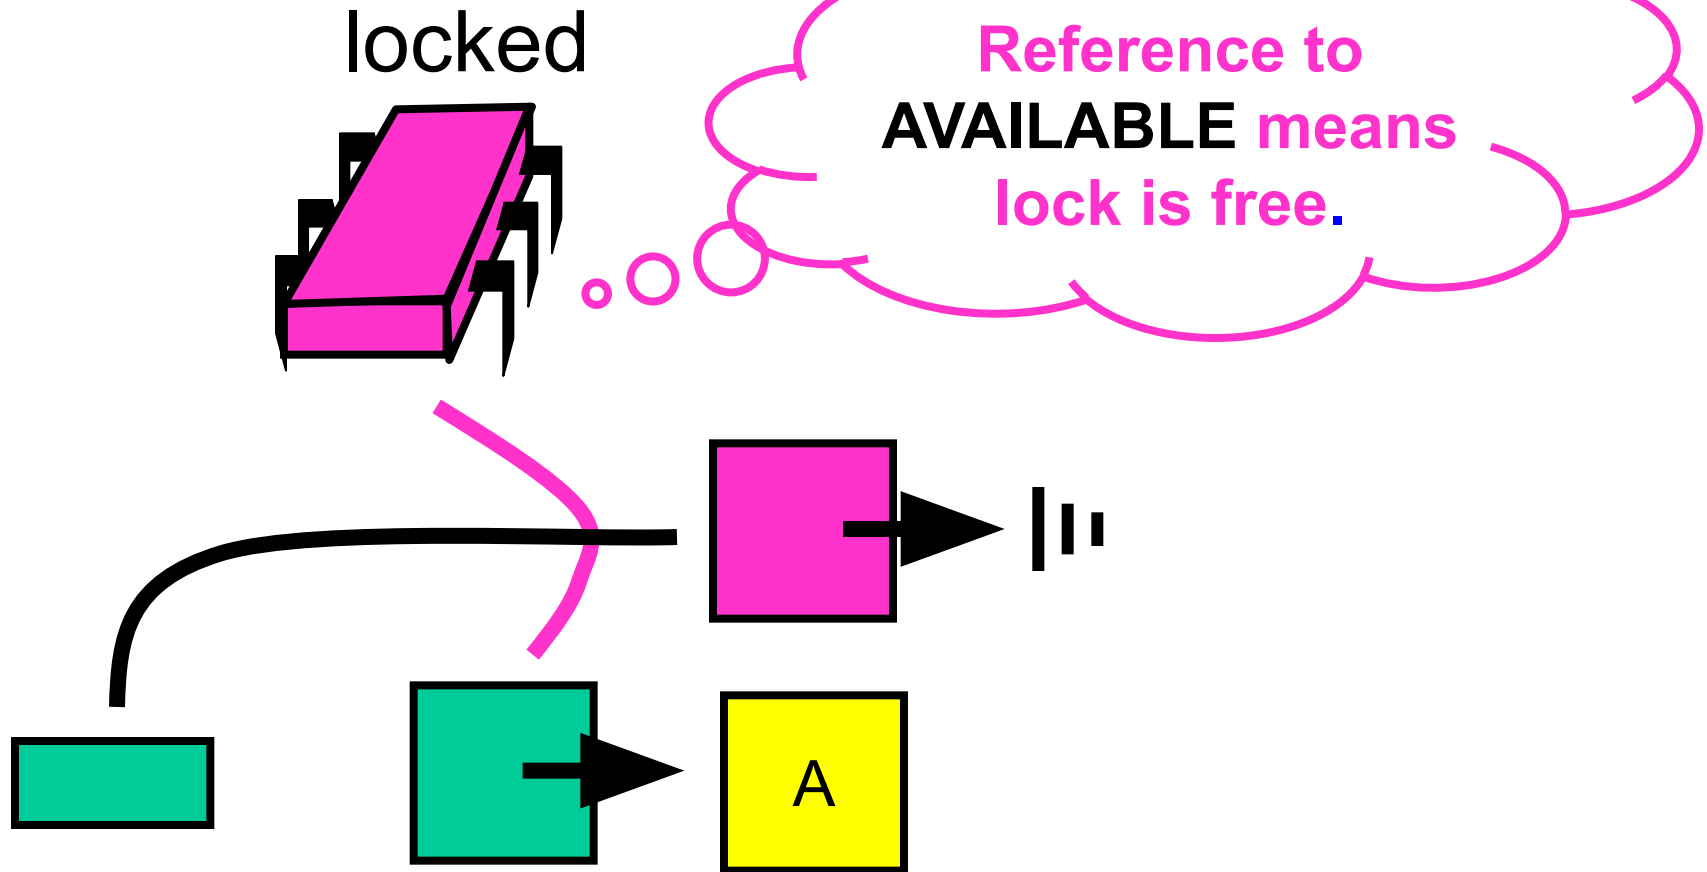


# Acquiring

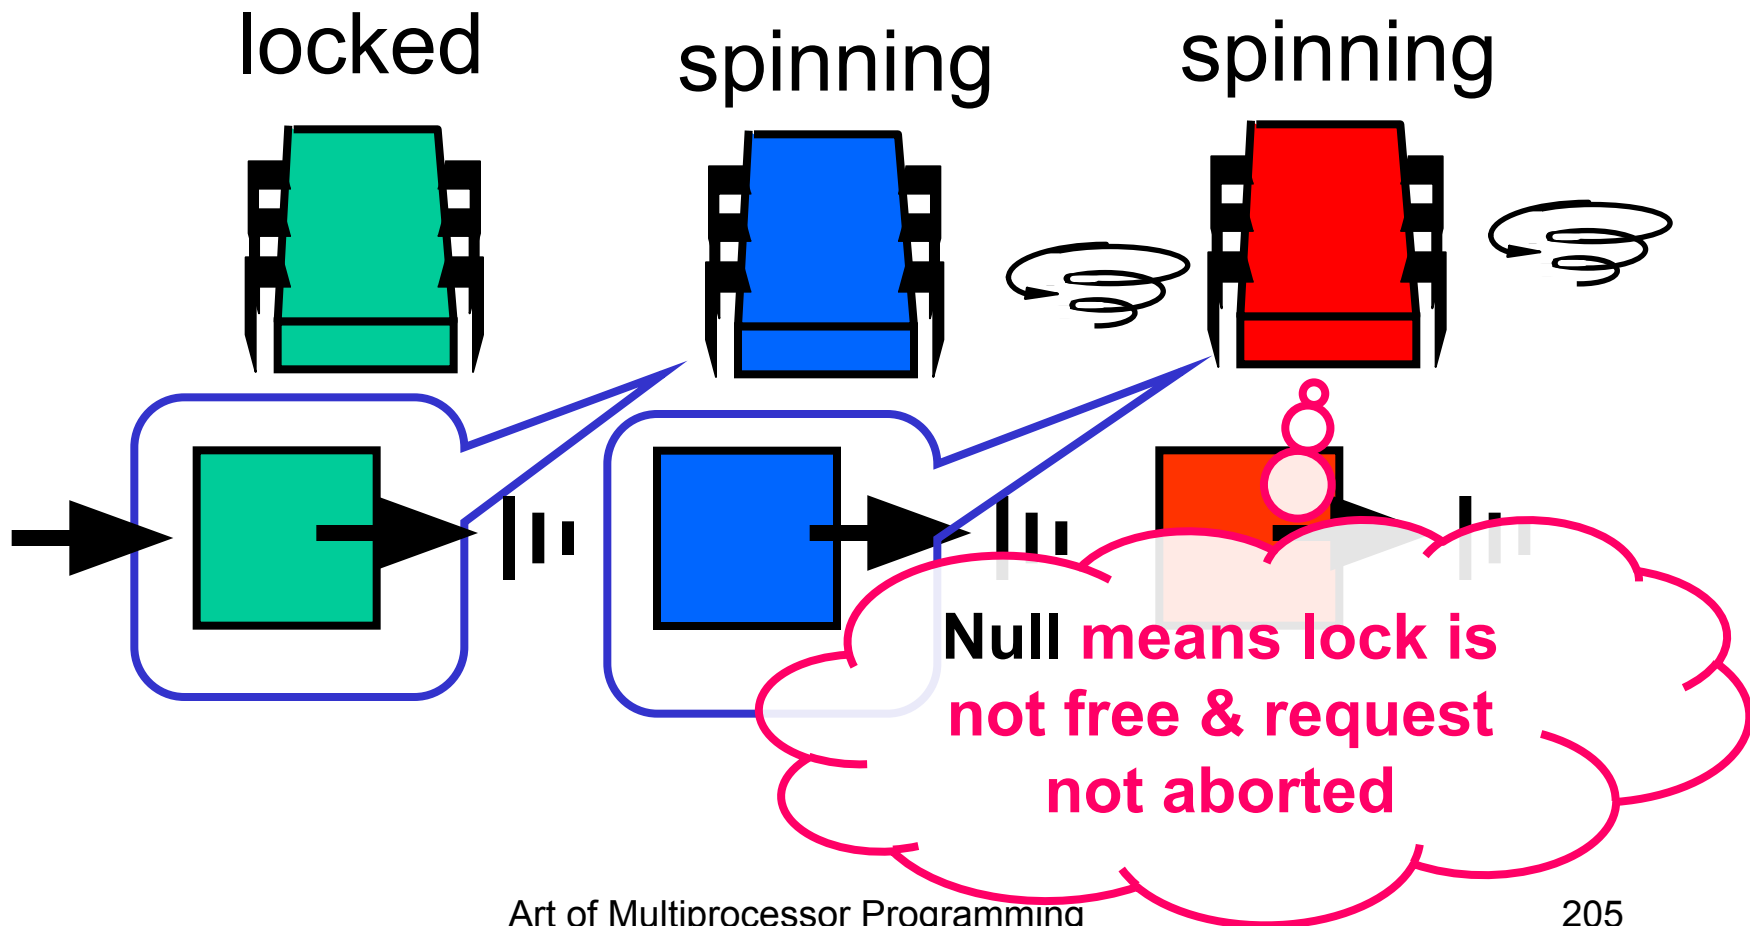
acquiring




# Acquired



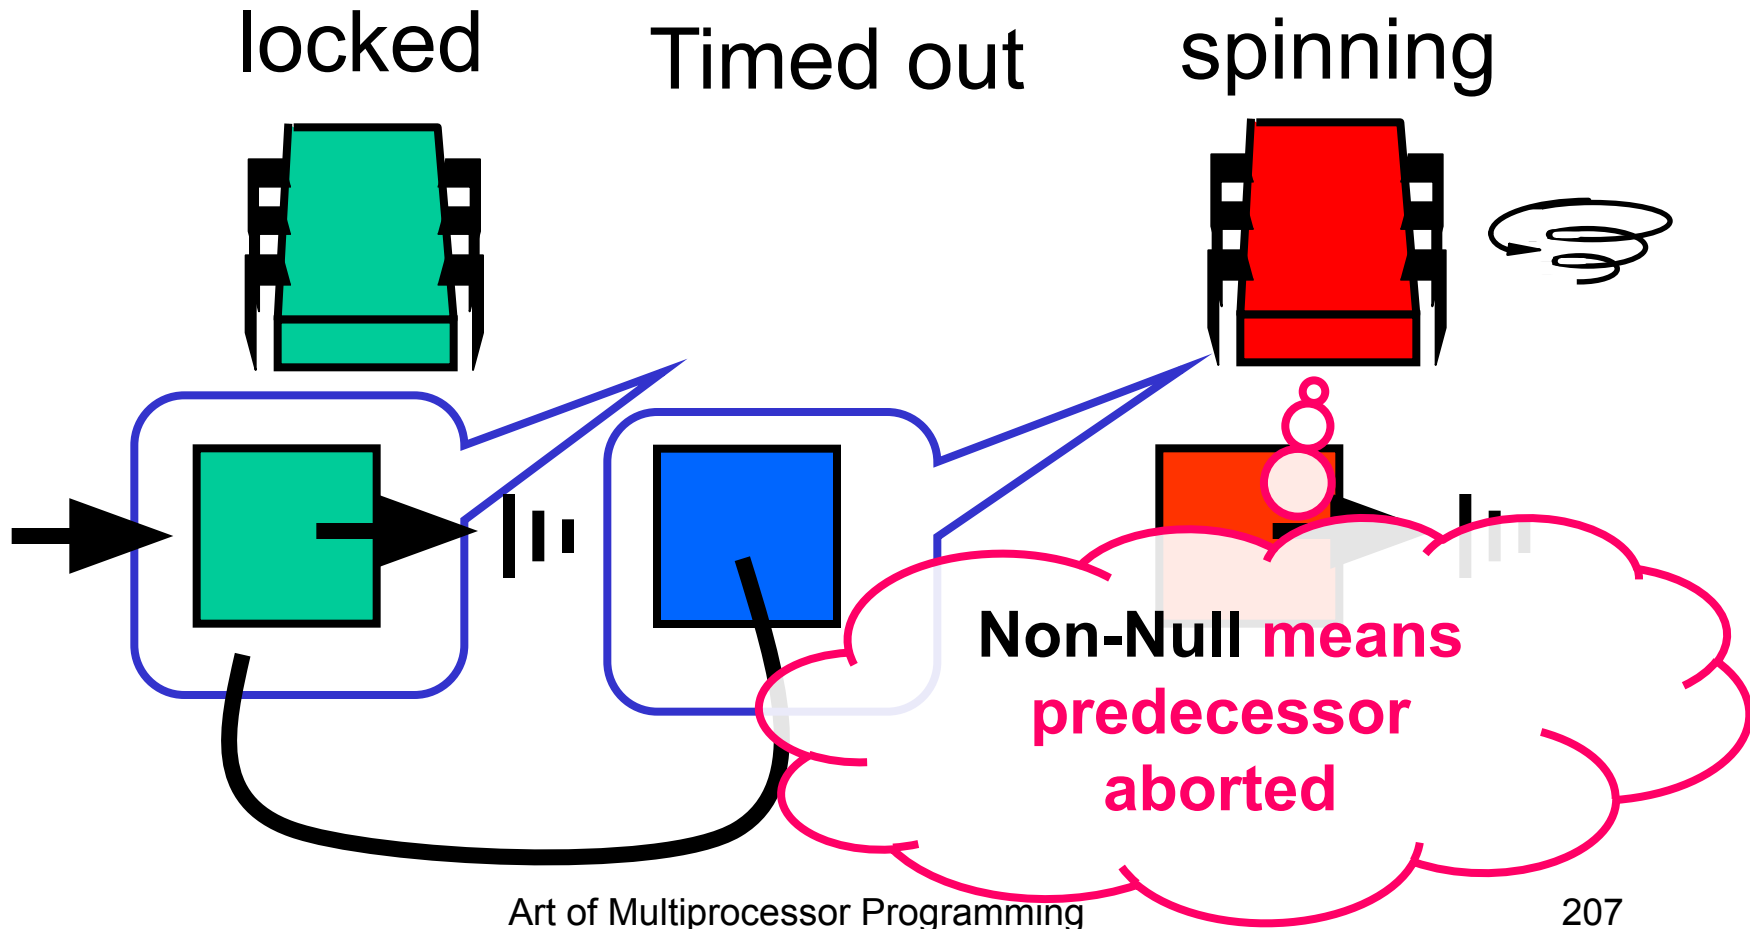
# Normal Case



# One Thread Aborts





# Successor Notices





# Recycle Predecessor's Node

locked




spinning

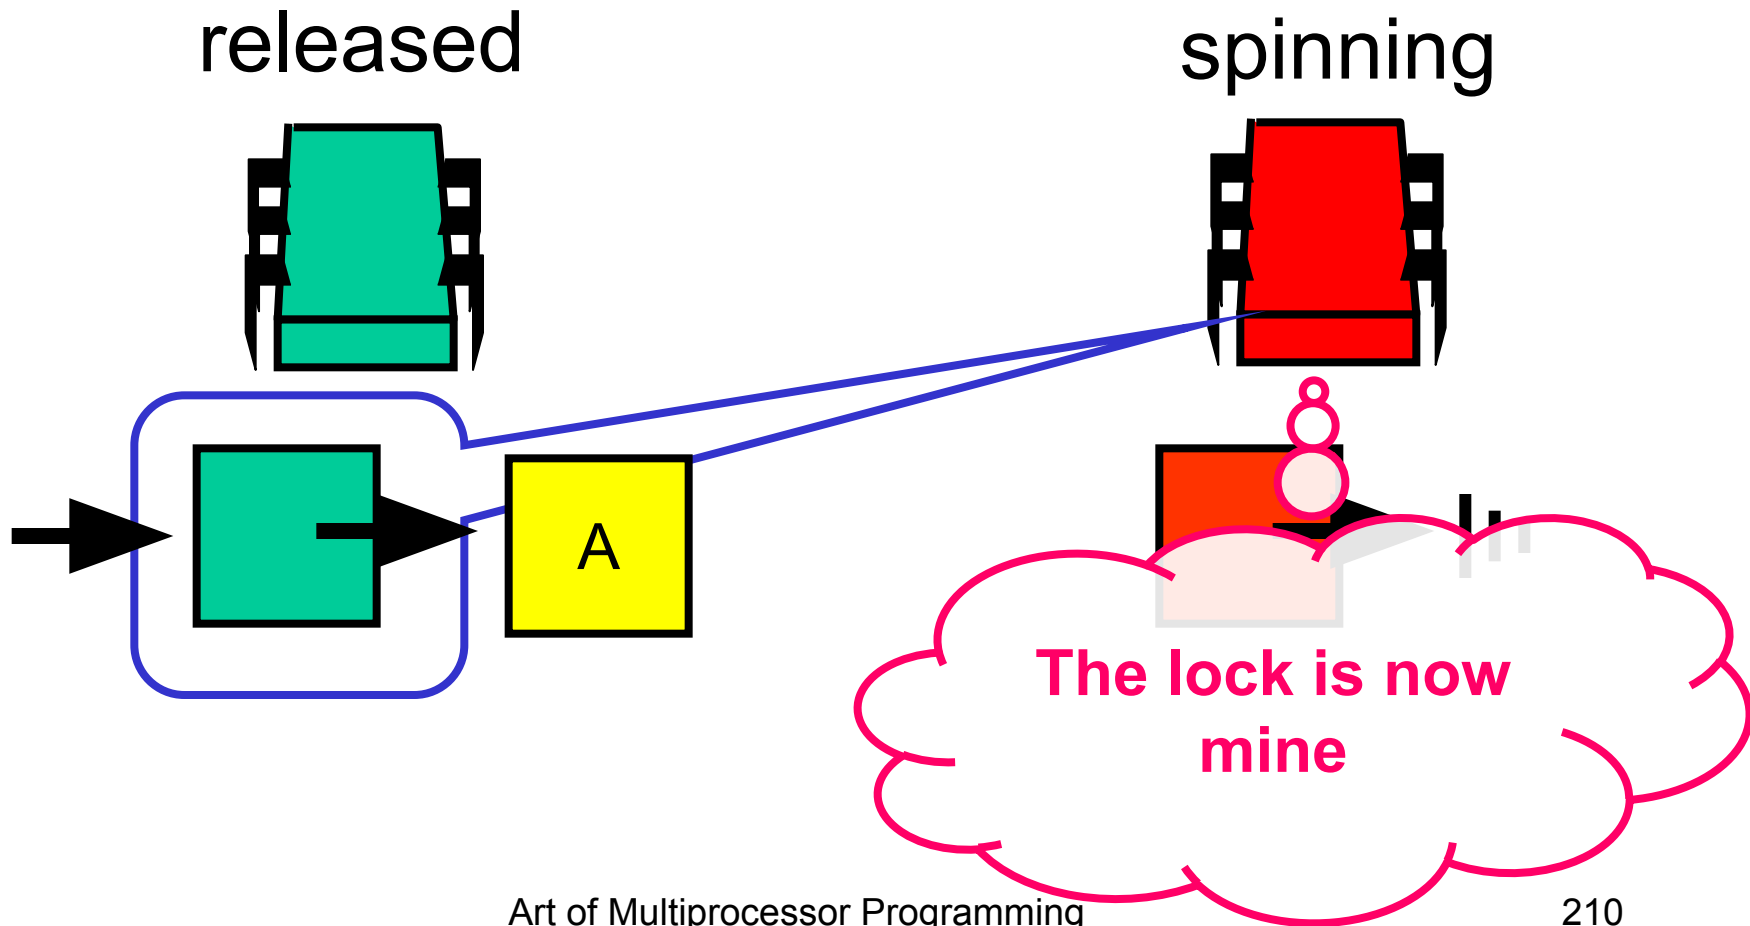




# Spin on Earlier Node



# Spin on Earlier Node



# Time-out Lock

```
public class TOLock implements Lock {  
    static QNode AVAILABLE  
        = new QNode();  
    AtomicReference<QNode> tail;  
    ThreadLocal<QNode> myNode;
```

# Time-out Lock

```
public class TOLock implements Lock {  
    static QNode AVAILABLE  
        = new QNode();  
    AtomicReference<QNode> tail;  
    ThreadLocal<QNode> myNode;
```

**AVAILABLE node  
signifies free lock**

# Time-out Lock

```
public class TOLock implements Lock {  
    static QNode AVAILABLE  
        = new QNode();  
    AtomicReference<QNode> tail;  
    ThreadLocal<QNode> myNode;
```

**Tail of the queue**

# Time-out Lock

```
public class TOLock implements Lock {  
    static QNode AVAILABLE  
        = new QNode();  
    AtomicReference<QNode> tail;  
    ThreadLocal<QNode> myNode;
```

**Remember my node ...**

# Time-out Lock

```
public boolean lock(long timeout) {  
    QNode qnode = new QNode();  
    myNode.set(qnode);  
    qnode.prev = null;  
    QNode myPred = tail.getAndSet(qnode);  
    if (myPred == null  
        || myPred.prev == AVAILABLE) {  
        return true;  
    }  
    ...  
}
```

# Time-out Lock

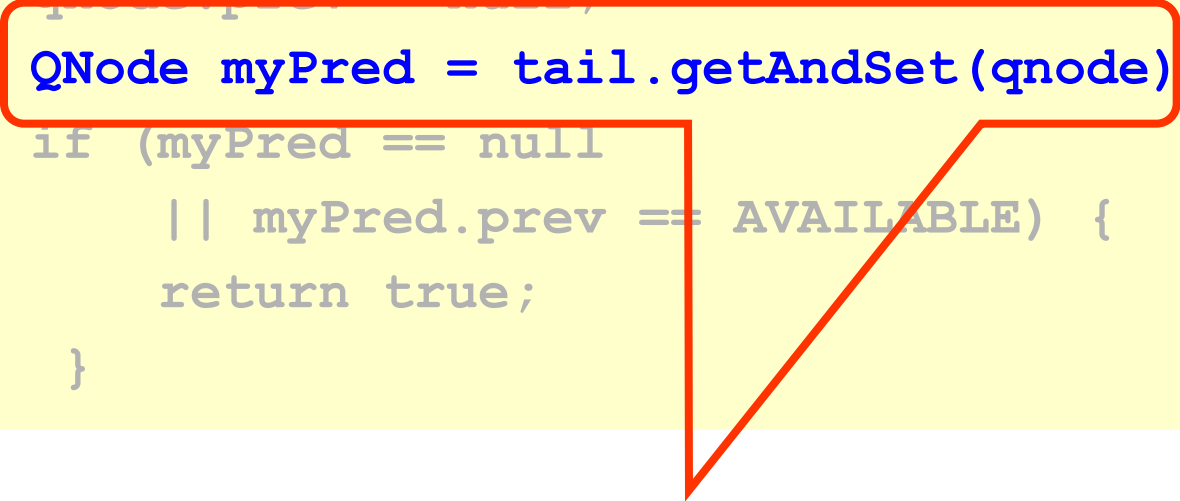
```
public boolean lock(long timeout) {  
    QNode qnode = new QNode();  
    myNode.set(qnode);  
    qnode.prev = null;  
    QNode myPred = tail.getAndSet(qnode);  
    if (myPred == null  
        || myPred.prev == AVAILABLE) {  
        return true;  
    }  
}
```

**Create & initialize node**



# Time-out Lock

```
public boolean lock(long timeout) {  
    QNode qnode = new QNode();  
    myNode.set(qnode);  
    qnode.prev = null;  
    QNode myPred = tail.getAndSet(qnode);  
    if (myPred == null  
        || myPred.prev == AVAILABLE) {  
        return true;  
    }  
}
```




**Swap with tail**

# Time-out Lock

```
public boolean lock(long timeout) {  
    QNode qnode = new QNode();  
    myNode.set(qnode);  
    qnode.prev = null;  
    QNode myPred = tail.getAndSet(qnode);  
    if (myPred == null  
        || myPred.prev == AVAILABLE) {  
        return true;  
    }  
    ...  
}
```

**If predecessor absent or released, we are done**

# Time-out Lock



...

```
long start = now();  
while (now() - start < timeout) {  
    QNode predPred = myPred.prev;  
    if (predPred == AVAILABLE) {  
        return true;  
    } else if (predPred != null) {  
        myPred = predPred;  
    }  
}
```

...

# Time-out Lock

...

```
long start = now();  
while (now() - start < timeout) {
```

```
    QNode predPred = myPred.prev;  
    if (predPred == AVAILABLE) {  
        return true;  
    } else if (predPred != null) {  
        myPred = predPred;  
    }  
}
```

```
}
```

...

...

**Keep trying for a while**

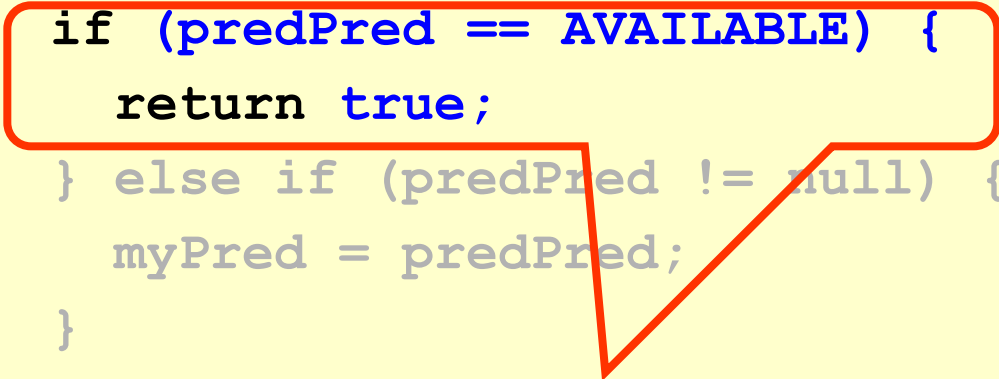
# Time-out Lock

```
...  
long start = now();  
while (now() - start < timeout) {  
    QNode predPred = myPred.prev;  
    if (predPred == AVAILABLE) {  
        return true;  
    } else if (predPred != null) {  
        myPred = predPred;  
    }  
}  
...
```

**Spin on predecessor's  
prev field**

# Time-out Lock

```
...  
long start = now();  
while (now() - start < timeout) {  
    QNode predPred = myPred.prev;  
    if (predPred == AVAILABLE) {  
        return true;  
    } else if (predPred != null) {  
        myPred = predPred;  
    }  
}  
...
```



**Predecessor released lock**

# Time-out Lock

...

```
long start = now();  
while (now() - start < timeout) {  
    QNode predPred = myPred.prev;  
    if (predPred == AVAILABLE) {  
        return true;  
    } else if (predPred != null) {  
        myPred = predPred;  
    }  
}
```

}

...

**Predecessor aborted,  
advance one**

# Time-out Lock

```
...  
if (!tail.compareAndSet(qnode, myPred))  
    qnode.prev = myPred;  
    return false;  
}  
}
```

**What do I do when I time out?**



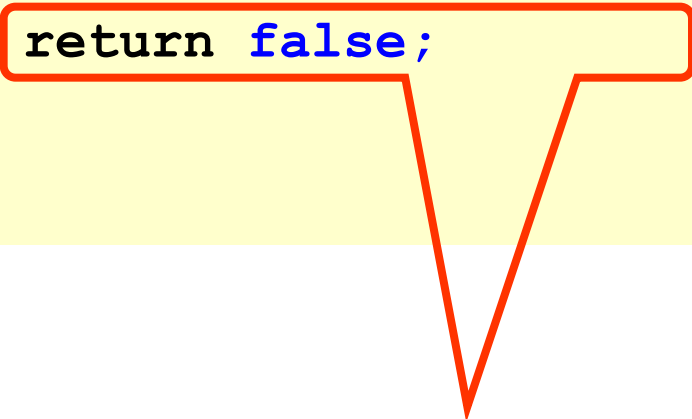
# Time-out Lock

```
...  
if (!tail.compareAndSet(qnode, myPred))  
    qnode.prev = myPred;  
return false;  
}  
}
```

**Do I have a successor?  
If CAS fails, I do.  
Tell it about myPred**

# Time-out Lock

```
...  
if (!tail.compareAndSet(qnode, myPred))  
    qnode.prev = myPred;  
    return false;  
}  
}
```



**If CAS succeeds: no  
successor, simply return false**

# Time-Out Unlock

```
public void unlock() {  
    QNode qnode = myNode.get();  
    if (!tail.compareAndSet(qnode, null))  
        qnode.prev = AVAILABLE;  
}
```

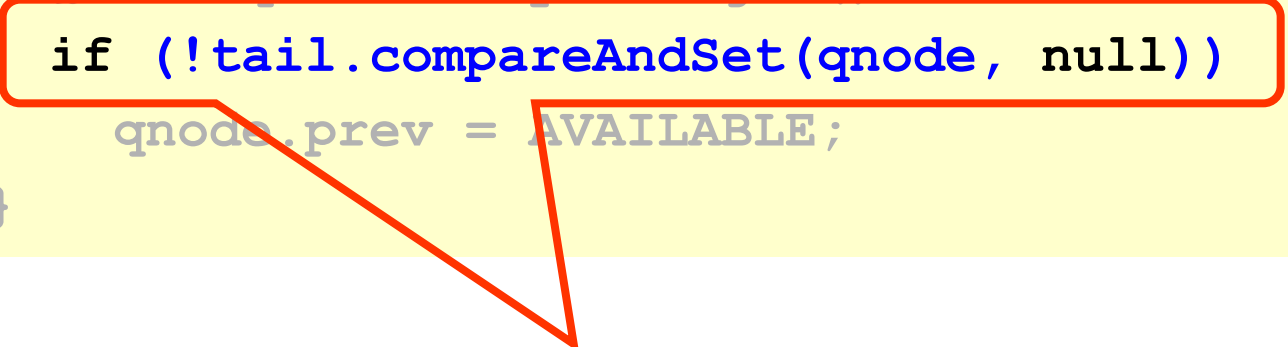
# Time-out Unlock

```
public void unlock() {  
    QNode qnode = myNode.get();  
    if (!tail.compareAndSet(qnode, null))  
        qnode.prev = AVAILABLE;  
}
```

**If CAS failed:  
successor exists,  
notify it can enter**

# Timing-out Lock

```
public void unlock() {  
    ONode qnode = myNode.get();  
    if (!tail.compareAndSet(qnode, null))  
        qnode.prev = AVAILABLE;  
}
```




**CAS successful: set tail to null, no clean up since no successor waiting**

# Fairness and NUMA Locks


- MCS lock mechanics are aware of NUMA
- Lock Fairness is FCFS
- Is this a good fit with NUMA and Cache-Coherent NUMA machines?

# Lock Data Access in NUMA Machine

Node 1



Node 2



CS



MCS  
lock

various  
memory  
locations

# “Who’s the Unfairest of Them All?”





- locality crucial to NUMA performance
- Big gains if threads from same node/cluster obtain lock consecutively
- Unfairness pays



# Hierarchical Backoff Lock (HBO)

Back off less  
for thread from  
same node

Unfairness is key to  
performance



CS

Global  
T&T&S  
lock


# Hierarchical Backoff Lock (HBO)

- Advantages:
  - Simple, improves locality
- Disadvantages:
  - Requires platform specific tuning
  - Unstable
  - Unfair
  - Continuous invalidations on shared global lock word


# Hybrid CLH Lock

Each thread spins on cached copy of predecessor's node

Thread at local head splices local queue into global queue




# Hierarchical CLH Lock (HCLH)



**Threads access 4 cache lines in CS**

# Hierarchical CLH Lock (HCLH)

- Advantages:
  - Improved locality
  - Local spinning
  - Fair
- Disadvantages:
  - Complex code implies long common path
  - Splicing into both local and global requires CAS
  - Hard to get long local sequences

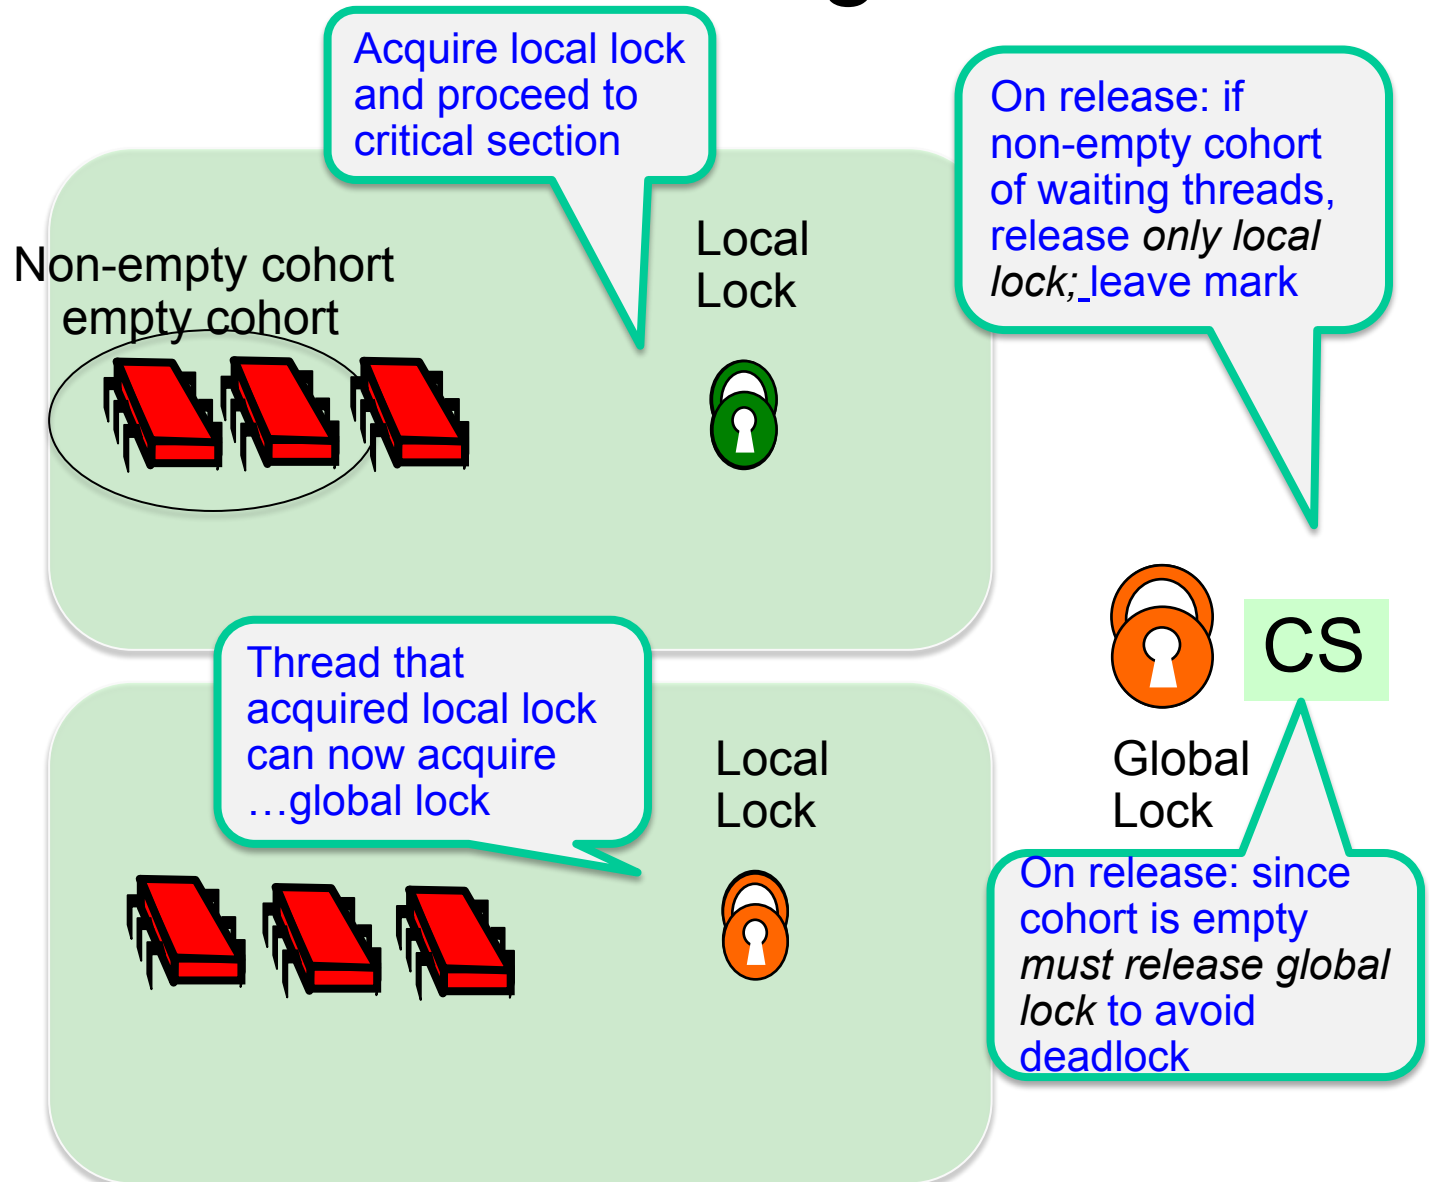


"Nothing yet. ... How about you, Newton?"

# Lock Cohorting

- General technique for converting almost any lock into a NUMA lock
- Allows combining different lock types
- But need these locks to have certain properties (will discuss shortly)

# Lock Cohorting





# Thread Obliviousness

- A lock is ***thread-oblivious*** if
  - After being acquired by one thread,
  - Can be released by another

# Cohort Detection

- A lock  $x$  provides ***cohort detection*** if
  - It can tell whether any thread is trying to acquire it

# Lock Cohorting


- Two levels of locking
- **Global lock:** thread oblivious
  - Thread acquiring the lock can be different than one releasing it
- **Local lock:** cohort detection
  - Thread releasing can detect if some thread is waiting to acquire it

Two new states: *acquire local* and *acquire global*.  
?Do we own global lock


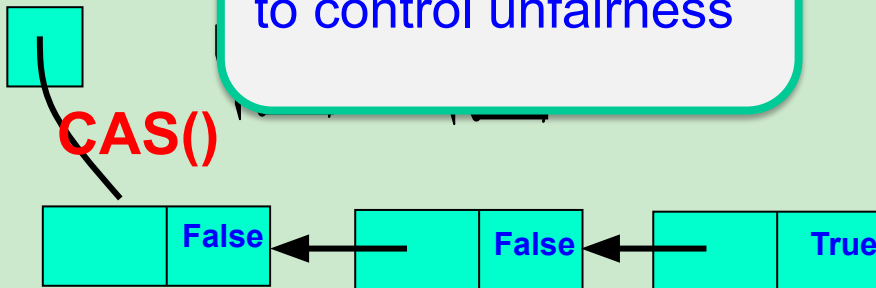
ing: C

In MCS Lock, *cohort*  
*detection* by checking  
successor pointer

### Local MCS lock tail



Bound number of  
consecutive acquires  
to control unfairness






BO Lock is *thread*  
*oblivious* by definition

# Lock Co

How to add *cohort detection*  
*?property* to BO lock

# BO Lock



As noted BO Lock is  
*thread oblivious*




# Lock Co

Add successorExists field before attempting to .acquire local lock

# -BO Lock

successorExists reset on .lock release


Release might overwrite another successor's write ... but we don't care...why



C-B

Aborting thread resets  
successorExists field  
before leaving local lock.  
Spinning threads set it to  
.true

e-Out




BO locks trivially abortable

Global  
backoff  
lock




If releasing thread finds  
successorExists false, it  
releases global lock





# Time-Out NUMA Lock

Local Time-Out locks  
have cohort detection  
?property ...why  
...Not enough

## Local time-out queue



## Local time-out queue







# Lock Cohorting

- Advantages:
  - Great locality
  - Low contention on shared lock
  - Practically no tuning
  - Has whatever properties you want:
    - Can be more or less fair, abortable...just choose the appropriate type of locks...
- Disadvantages:
  - Must tune fairness parameters

# Lock Cohorting



# Time-Out (Abortable) Lock Cohorting



# One Lock To Rule Them All?

- TTAS+Backoff, CLH, MCS, ToLock...
- Each better than others in some way
- There is no one solution
- Lock we pick really depends on:
  - the application
  - the hardware
  - which properties are important

This work is licensed under a [Creative Commons Attribution-ShareAlike 2.5 License](https://creativecommons.org/licenses/by-sa/3.0/).

- **You are free:**
  - **to Share** — to copy, distribute and transmit the work
  - **to Remix** — to adapt the work
- **Under the following conditions:**
  - **Attribution.** You must attribute the work to “The Art of Multiprocessor Programming” (but not in any way that suggests that the authors endorse you or your use of the work).
  - **Share Alike.** If you alter, transform, or build upon this work, you may distribute the resulting work only under the same, similar or a compatible license.
- For any reuse or distribution, you must make clear to others the license terms of this work. The best way to do this is with a link to
  - <http://creativecommons.org/licenses/by-sa/3.0/>.
- Any of the above conditions can be waived if you get permission from the copyright holder.
- Nothing in this license impairs or restricts the author's moral rights.