

# Concurrent programming

## Concurrent Queues and Stacks

Companion slides for

The Art of Multiprocessor Programming  
by Maurice Herlihy, Nir Shavit, Victor Luchangco,  
and Michael Spear

Modified by Piotr Witkowski

# The Five-Fold Path


- Coarse-grained locking
- Fine-grained locking
- Optimistic synchronization
- Lazy synchronization
- Lock-free synchronization

# Another Fundamental Problem


- We told you about
  - Sets implemented by linked lists
  - ~~– Hash Tables~~
- Next: queues
- Next: stacks

# Queues & Stacks


- pool of items



# Queues




# Stacks



**Total order**  
**Last in**  
**First out**


# Bounded

- Fixed capacity
- Good when resources an issue




# Unbounded

- Unlimited capacity
- Often more convenient






# Blocking



# Blocking




# Non-Blocking




# This Lecture

- Queue
  - Bounded, blocking, lock-based
  - Unbounded, non-blocking, lock-free
- Stack
  - Unbounded, non-blocking lock-free
  - Elimination-backoff algorithm


# Queue: Concurrency



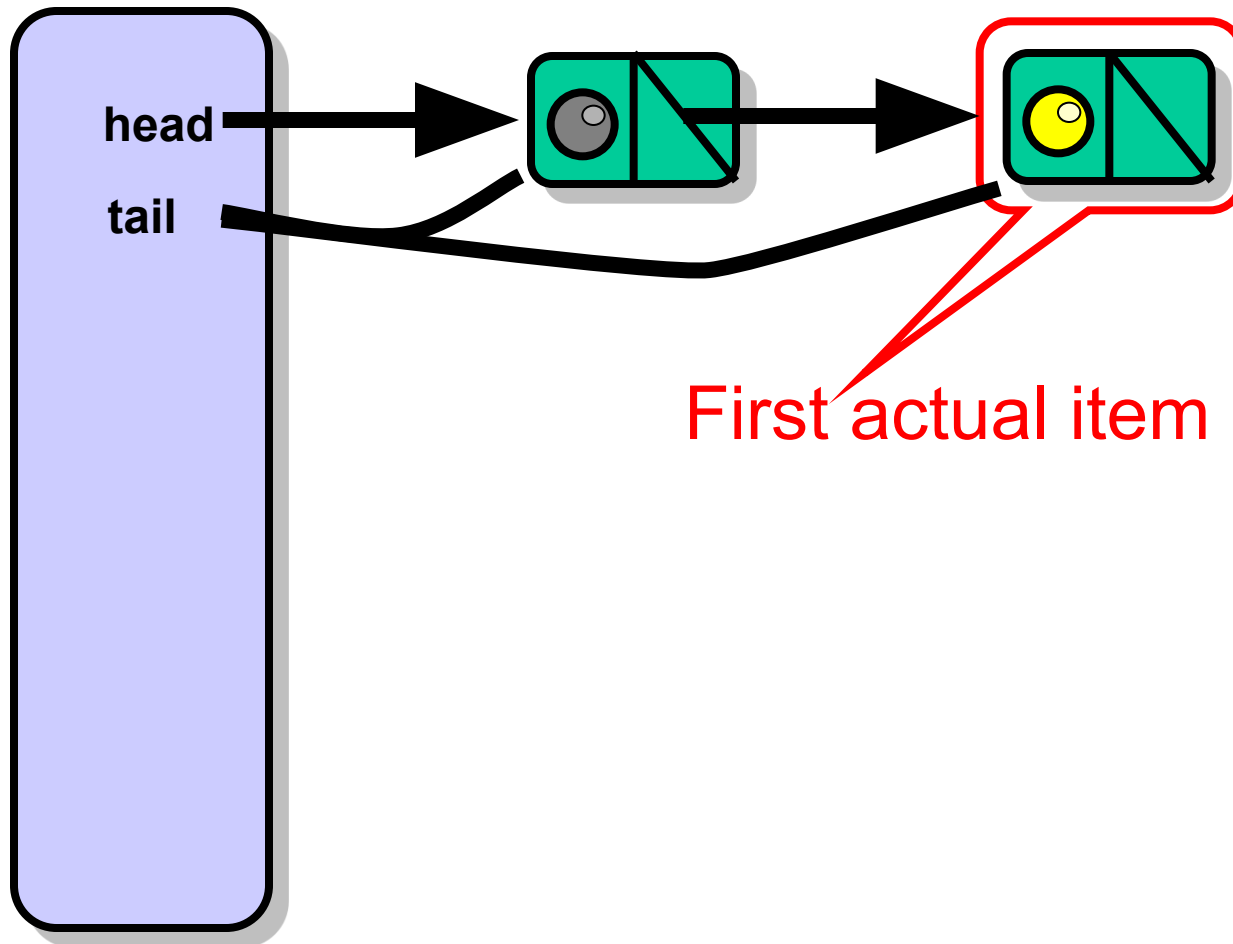
# Concurrency



# Bounded Queue




# Bounded Queue







# Bounded Queue



# Bounded Queue




# Not Done Yet



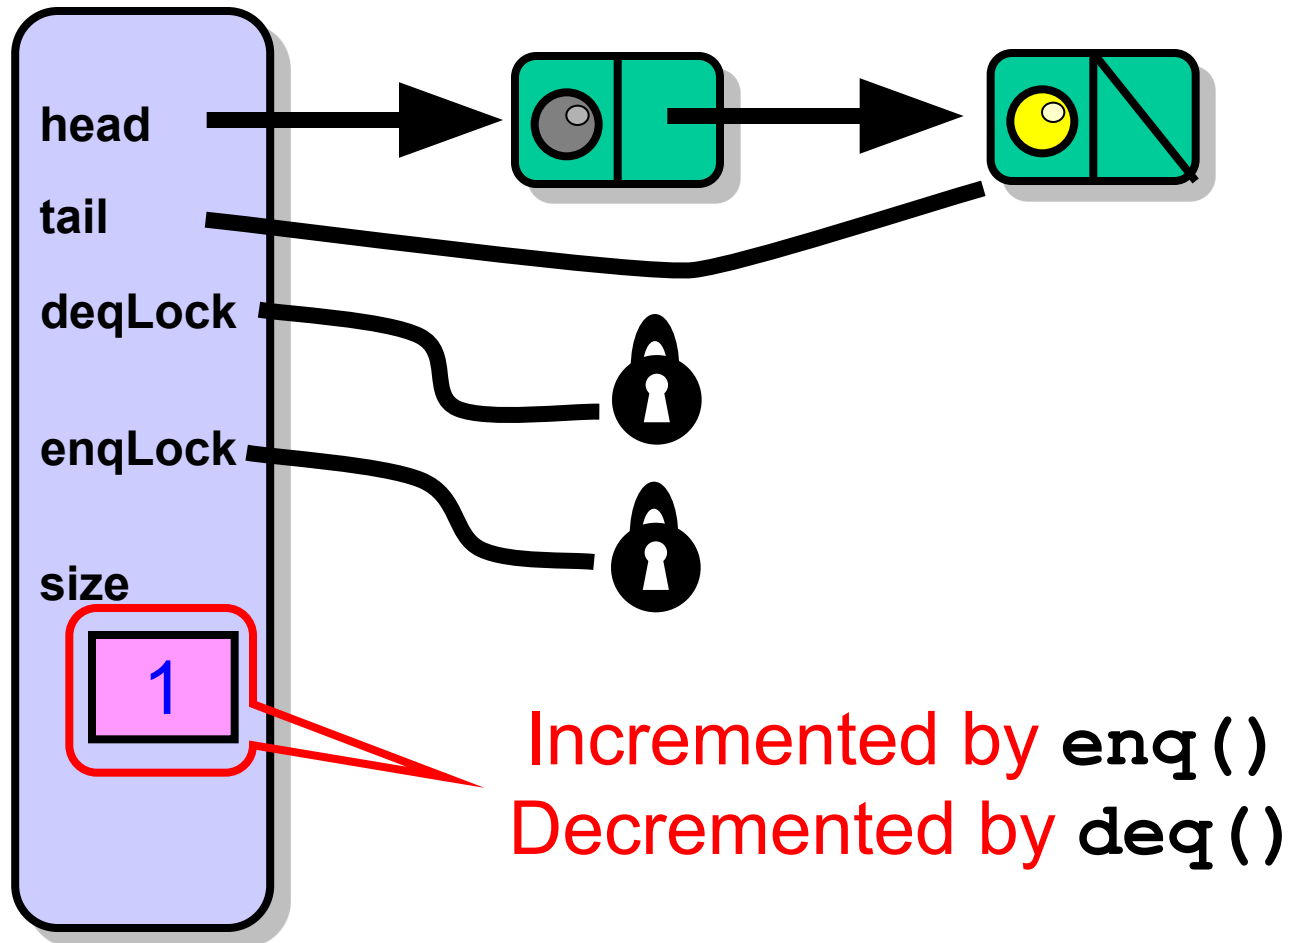
Need to tell whether  
queue is full or  
empty

# Not Done Yet




Max size is 8 items


# Not Done Yet




# Enqueuer



# Enqueuer




# Enqueuer







# Enqueuer




# Enqueuer




# Enqueuer




# Enqueuer




# Unsuccessful Enqueuer




# Dequeuer



# Dequeuer




# Dequeuer






Make first Node  
new sentinel


# Dequeuer




# Dequeuer



# Dequeuer



# Unsuccessful Dequeueer

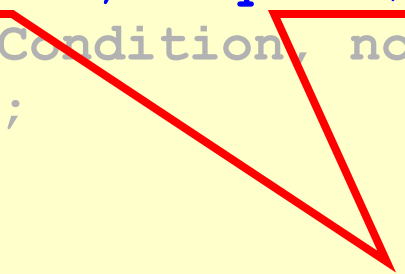


# Bounded Queue

```
public class BoundedQueue<T> {  
    ReentrantLock enqLock, deqLock;  
    Condition notEmptyCondition, notFullCondition;  
    AtomicInteger size;  
    Node head;  
    Node tail;  
    int capacity;  
    enqLock = new ReentrantLock();  
    notFullCondition = enqLock.newCondition();  
    deqLock = new ReentrantLock();  
    notEmptyCondition = deqLock.newCondition();  
}
```

# Bounded Queue

```
public class BoundedQueue<T> {  
    ReentrantLock enqLock, deqLock;  
    Condition notEmptyCondition, notFullCondition;  
    AtomicInteger size;  
    Node head;  
    Node tail;  
    int capacity;  
    enqLock = new ReentrantLock();  
    notFullCondition = enqLock.newCondition();  
    deqLock = new ReentrantLock();  
    notEmptyCondition = deqLock.newCondition();  
}
```



**enq & deq locks**

(push)

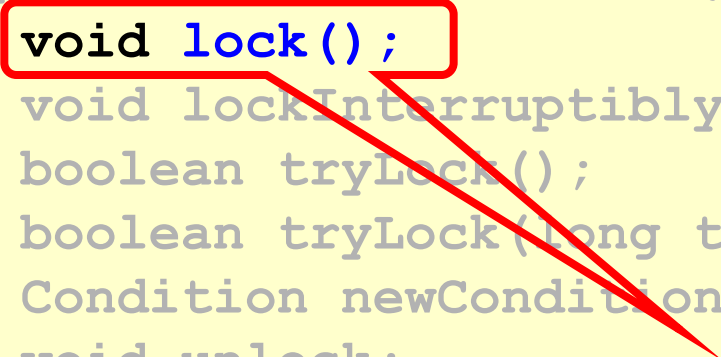
# Digression: Monitor Locks

- `Java` `synchronized` `objects` and `ReentrantLocks` `are` monitors
- Allow blocking on a condition rather than spinning
- Threads:
  - acquire `and` release `lock`
  - wait `on a condition`



# The Java Lock Interface


```
public interface Lock {  
    void lock();  
    void lockInterruptibly() throws InterruptedException;  
    boolean tryLock();  
    boolean tryLock(long time, TimeUnit unit);  
    Condition newCondition();  
    void unlock;  
}
```



**Acquire lock**

# The Java Lock Interface

```
public interface Lock {  
    void lock();  
    void lockInterruptibly() throws InterruptedException;  
    boolean tryLock();  
    boolean tryLock(long time, TimeUnit unit);  
    Condition newCondition();  
    void unlock;  
}
```



**Release lock**

# The Java Lock Interface

```
public interface Lock {  
    void lock();  
    void lockInterruptibly() throws InterruptedException;  
    boolean tryLock();  
    boolean tryLock(long time, TimeUnit unit);  
    Condition newCondition();  
    void unlock();  
}
```

**Try for lock, but not too hard**

# The Java Lock Interface

```
public interface Lock {  
    void lock();  
    void lockInterruptibly() throws InterruptedException;  
    boolean tryLock();  
    boolean tryLock(long time, TimeUnit unit);  
    Condition newCondition();  
    void unlock;  
}
```

**Create condition to wait on**

# The Java Lock Interface

```
public interface Lock {  
    void lock();  
    void lockInterruptibly() throws InterruptedException;  
    boolean tryLock();  
    boolean tryLock(long time, TimeUnit unit);  
    Condition newCondition();  
    void unlock;  
}
```

**Never mind what this method does**

# Lock Conditions

```
public interface Condition {  
    void await();  
    boolean await(long time, TimeUnit unit);  
    ...  
    void signal();  
    void signalAll();  
}
```

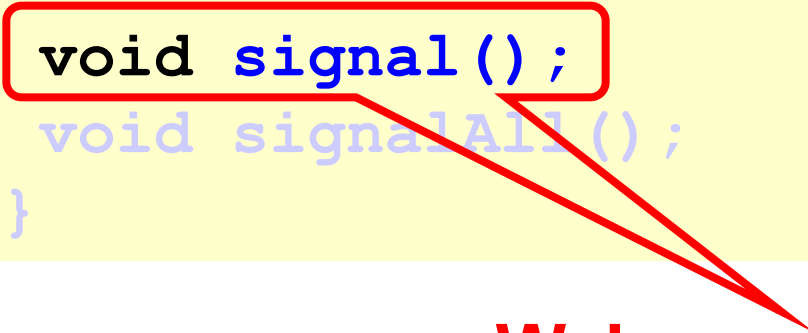
# Lock Conditions

```
public interface Condition {  
    void await();  
    boolean await(long time, TimeUnit unit);  
    ...  
    void signal();  
    void signalAll();  
}
```

**Release lock and  
wait on condition**

# Lock Conditions

```
public interface Condition {  
    void await();  
    boolean await(long time, TimeUnit unit);  
    ...  
    void signal();  
    void signalAll();  
}
```



**Wake up one waiting thread**



# Lock Conditions

```
public interface Condition {  
    void await();  
    boolean await(long time, TimeUnit unit);  
    ...  
    void signal();  
    void signalAll();  
}
```

**Wake up all waiting threads**

# Await

```
q.await()
```

- Releases lock associated with q
- Sleeps (gives up processor)
- Awakens (resumes running)
- Reacquires lock & returns

# Signal

```
q.signal();
```


- Awakens one waiting thread
  - Which will reacquire lock

# Signal All


```
q.signalAll();
```

- Awakens all waiting threads
  - Which will each reacquire lock

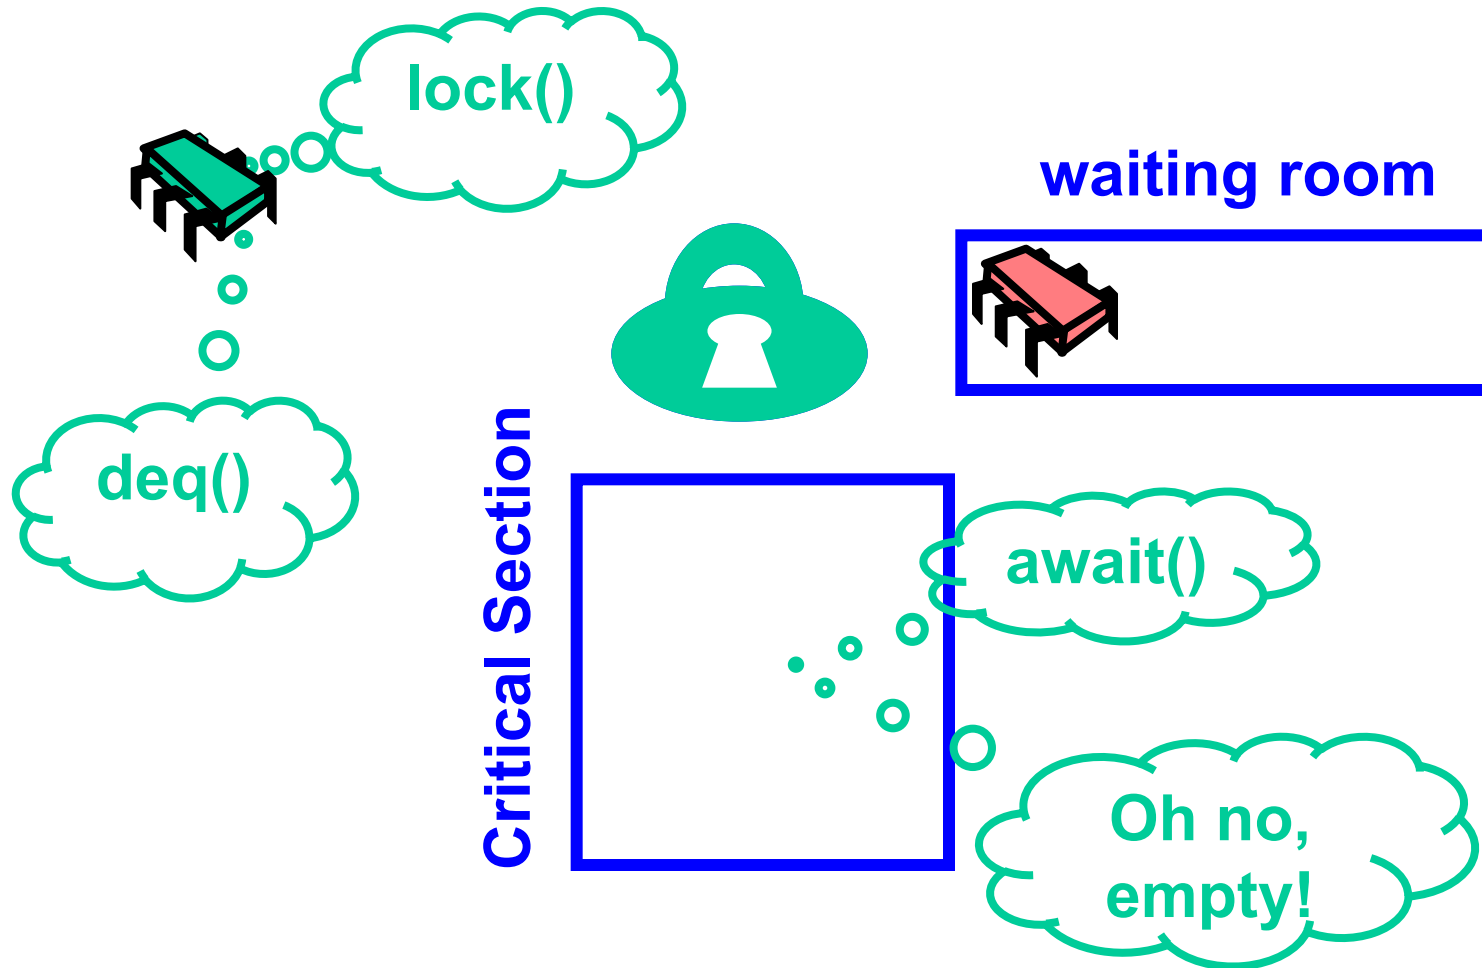
# A Monitor Lock




# Unsuccessful Deq



# Another One



# Enqueuer to the Rescue







# Monitor Signalling




# Dequeuers Signalled



# Dequeuers Signaled



# Dollar Short + Day Late



# Java Synchronized Methods

```
public class Queue<T> {  
  
    int head = 0, tail = 0;  
    T[QSIZE] items;  
  
    public synchronized T deq() {  
        while (tail - head == 0)  
            wait();  
        T result = items[head % QSIZE]; head++;  
        notifyAll();  
        return result;  
    }  
    ...  
}
```

# Java Synchronized Methods

```
public class Queue<T> {  
  
    int head = 0, tail = 0;  
    T[QSIZE] items;  
  
    public synchronized T deq() {  
        while (tail - head == 0)  
            wait();  
        T result = items[head % QSIZE]; head++;  
        notifyAll();  
        return result;  
    }  
    ...  
}
```

**Each object has an implicit lock with an implicit condition**

# Java Synchronized Methods

```
public class Queue<T> {  
  
    int head = 0, tail = 0;  
    T[QSIZE] items;  
  
    public synchronized T deq() {  
        while (tail - head == 0)  
            wait();  
        T result = items[head % QSIZE]; head++;  
        notifyAll();  
        return result;  
    }  
    ...  
}
```

**Lock on entry,  
unlock on return**

# Java Synchronized Methods

```
public class Queue<T> {  
  
    int head = 0, tail = 0;  
    T[QSIZE] items;  
  
    public synchronized T deq() {  
        while (tail - head == 0)  
            wait();  
        T result = items[head % QSIZE]; head++;  
        this.notifyAll();  
        return result;  
    }  
    ...  
}
```

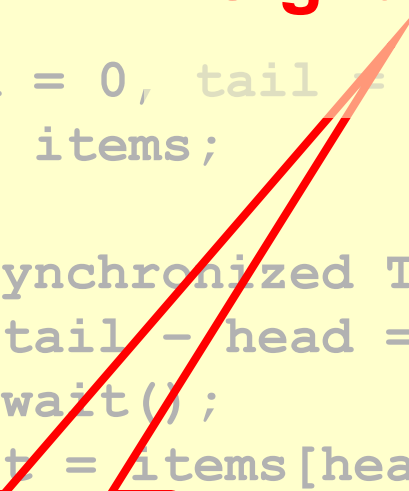
**Wait on implicit condition**



# Java Synchronized Methods

```
public class Queue<T> {  
    int head = 0, tail = 0;  
    T[QSIZE] items;  
  
    public synchronized T deq() {  
        while (tail - head == 0)  
            this.wait();  
        T result = items[head % QSIZE]; head++;  
        notifyAll();  
        return result;  
    }  
    ...  
}}
```

**Signal all threads waiting on condition**

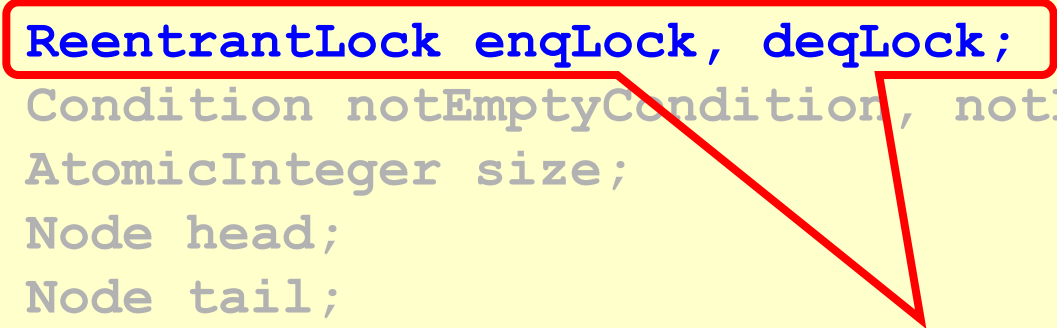


# (Pop!) The Bounded Queue

```
public class BoundedQueue<T> {  
    ReentrantLock enqLock, deqLock;  
    Condition notEmptyCondition, notFullCondition;  
    AtomicInteger size;  
    Node head;  
    Node tail;  
    int capacity;  
    enqLock = new ReentrantLock();  
    notFullCondition = enqLock.newCondition();  
    deqLock = new ReentrantLock();  
    notEmptyCondition = deqLock.newCondition();  
}
```

# Bounded Queue Fields

```
public class BoundedQueue<T> {  
    ReentrantLock enqLock, deqLock;  
    Condition notEmptyCondition, notFullCondition;  
    AtomicInteger size;  
    Node head;  
    Node tail;  
    int capacity;  
    enqLock = new ReentrantLock();  
    notFullCondition = enqLock.newCondition();  
    deqLock = new ReentrantLock();  
    notEmptyCondition = deqLock.newCondition();  
}
```

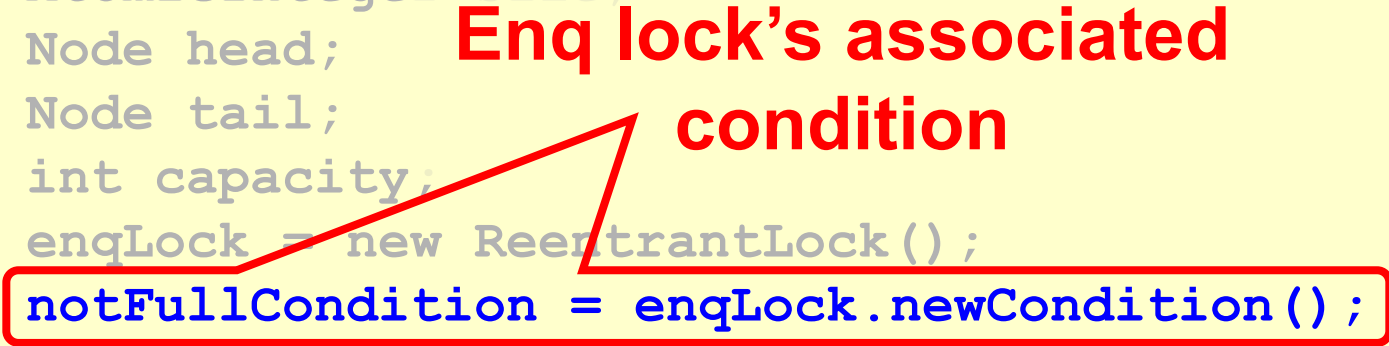


**Enq & deq locks**

# Bounded Queue Fields

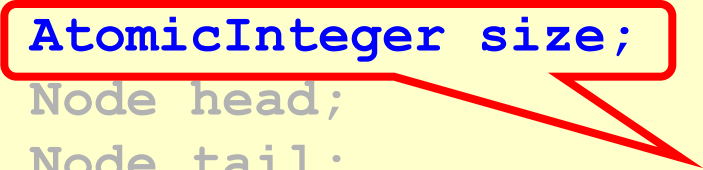
```
public class BoundedQueue<T> {  
    ReentrantLock enqLock, deqLock;  
    Condition notEmptyCondition, notFullCondition;  
    AtomicInteger size;  
    Node head;  
    Node tail;  
    int capacity;  
    enqLock = new ReentrantLock();  
    notFullCondition = enqLock.newCondition();  
    deqLock = new ReentrantLock();  
    notEmptyCondition = deqLock.newCondition();  
}
```

**Enq lock's associated condition**



# Bounded Queue Fields

```
public class BoundedQueue<T> {  
    ReentrantLock enqLock, deqLock;  
    Condition notEmptyCondition, notFullCondition;  
    AtomicInteger size;  
    Node head;  
    Node tail;  
    int capacity;  
    enqLock = new ReentrantLock();  
    notFullCondition = enqLock.newCondition();  
    deqLock = new ReentrantLock();  
    notEmptyCondition = deqLock.newCondition();  
}
```



**size: 0 to capacity**

# Bounded Queue Fields

```
public class BoundedQueue<T> {  
    ReentrantLock enqLock, deqLock;  
    Condition notEmptyCondition, notFullCondition;  
    AtomicInteger size;  
    Node head;  
    Node tail;  
    int capacity;  
    enqLock = new ReentrantLock();  
    notFullCondition = enqLock.newCondition();  
    deqLock = new ReentrantLock();  
    notEmptyCondition = deqLock.newCondition();  
}
```

**Head and Tail**

# Enq Method Part One

```
public void enq(T x) {
    boolean mustWakeDequeuers = false;
    enqLock.lock();
    try {
        while (size.get() == Capacity)
            notFullCondition.await();
        Node e = new Node(x);
        tail.next = e;
        tail = tail.next;
        if (size.getAndIncrement() == 0)
            mustWakeDequeuers = true;
    } finally {
        enqLock.unlock();
    }
    ...
}
```

# Enq Method Part One

```
public void enq(T x) {  
    boolean mustWakeDequeueurs = false;  
    enqLock.lock();
```

**Lock and unlock  
enq lock**

```
    try {  
        while (size.get() == capacity)  
            notFullCondition.await();  
        Node e = new Node(x);  
        tail.next = e;  
        tail = tail.next;  
        if (size.getAndIncrement() == 0)  
            mustWakeDequeueurs = true;  
    } finally {  
        enqLock.unlock();  
    }
```

```
    ...  
}
```



# Enq Method Part One

```
public void enq(T x) {  
    boolean mustWakeDequeuers = false;  
    enqLock.lock();  
    try {  
        while (size.get() == capacity)  
            notFullCondition.await();  
        Node e = new Node(x);  
        tail.next = e;  
        tail = tail.next;  
        if (size.getAndIncrement() == 0)  
            mustWakeDequeuers = true;  
    } finally {  
        enqLock.unlock();  
    }  
    ...  
}
```

**Wait while queue is full ...**

# Enq Method Part One

```
public void enq(T x) {  
    boolean mustWakeDequeueurs = false;  
    enqLock.lock();  
    try {  
        while (size.get() == capacity)  
            notFullCondition.await();  
        Node e = new Node(x);  
        tail.next = e;  
        tail = tail.next;  
        if (size.getAndIncrement() == 0)  
            mustWakeDequeueurs = true;  
    } finally {  
        enqLock.unlock();  
    }  
    ...  
}
```

**when await() returns, you  
might still fail the test !**

# Be Afraid

```
public void enq(T x) {  
    boolean mustWakeDequeueurs = false;  
    enqLock.lock();  
    try {  
        while (size.get() == capacity)  
            notFullCondition.await();  
        Node e = new Node(x);  
        tail.next = e;  
        tail = tail.next;  
        if (size.getAndIncrement() == 0)  
            mustWakeDequeueurs = true;  
    } finally {  
        enqLock.unlock();  
    }  
    ...  
}
```

**After the loop: how do we know the queue won't become full again?**

# Enq Method Part One

```
public void enq(T x) {  
    boolean mustWakeDequeueurs = false;  
    enqLock.lock();  
    try {  
        while (size.get() == capacity)  
            notFullCondition.await();  
        Node e = new Node(x);  
        tail.next = e;  
        tail = tail.next;  
        if (size.getAndIncrement() == 0)  
            mustWakeDequeueurs = true;  
    } finally {  
        enqLock.unlock();  
    }  
    ...  
}
```

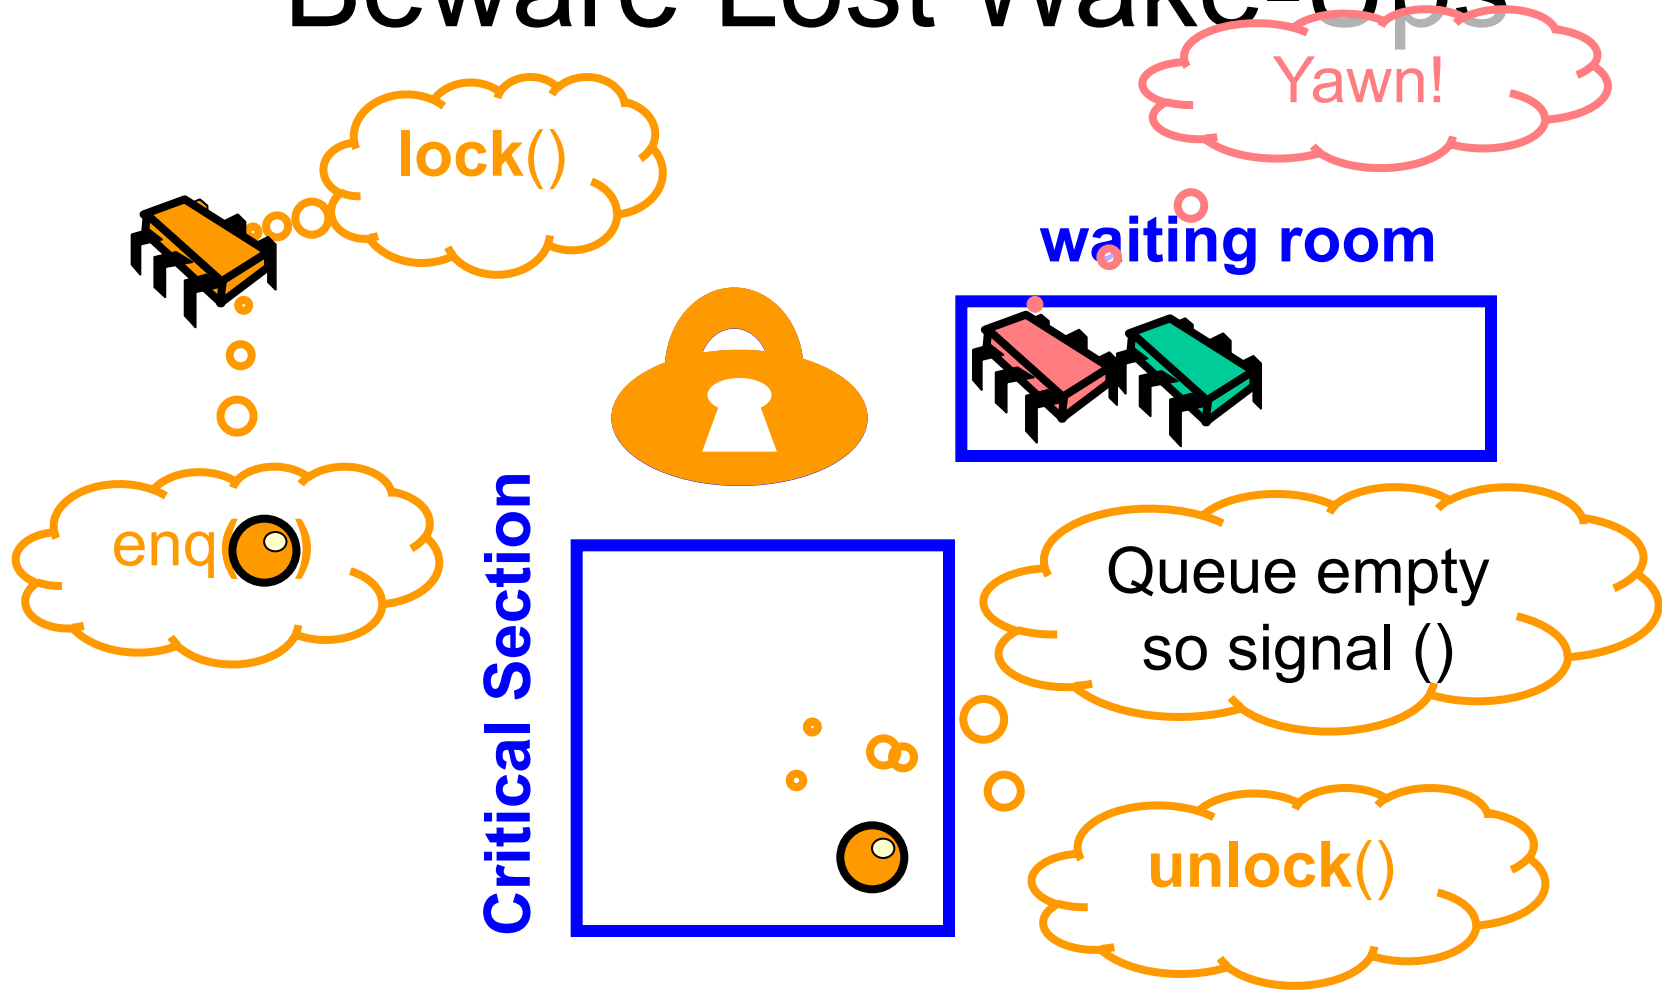
**Add new node**

# Enq Method Part One

```
public void enq(T x) {  
    boolean mustWakeDequeuers = false;  
    enqLock.lock();  
    try {  
        while (size.get() == capacity)  
            notFullCondition.await();  
        Node e = new Node(x);  
        tail.next = e;  
        tail = tail.next;  
        if (size.getAndIncrement() == 0)  
            mustWakeDequeuers = true;  
    } finally {  
        enqLock.unlock();  
    }  
    ...  
}
```

**If queue was empty, wake  
frustrated dequeuers**


# Beware Lost Wake-Ups



# Lost Wake-Up




# Lost Wake-Up

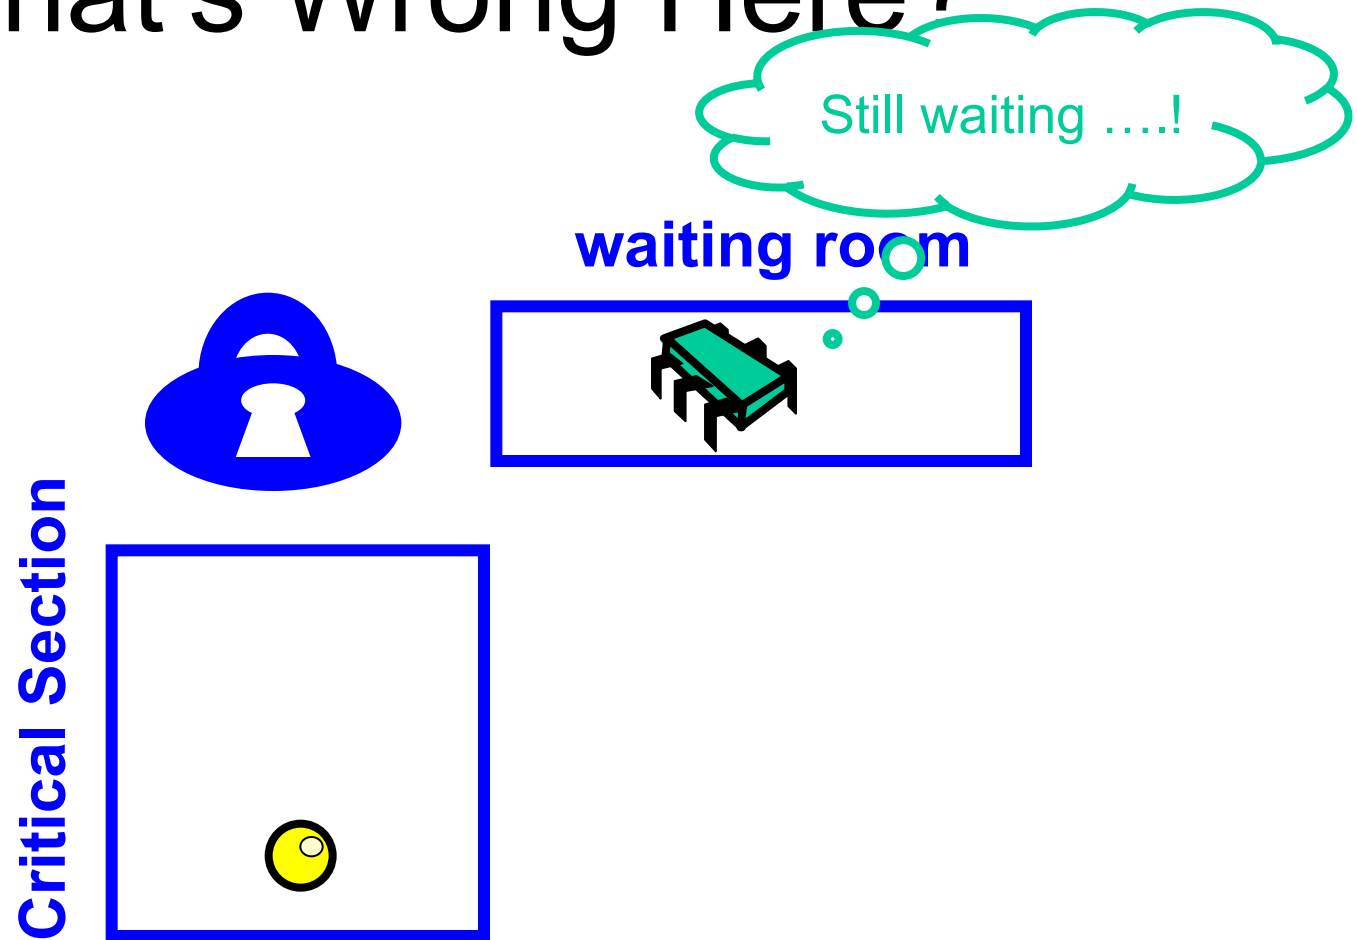




# Lost Wake-Up



# What's Wrong Here?



# Solution to Lost Wakeup

- Always use
  - `signalAll()` and `notifyAll()`
- Not
  - `signal()` and `notify()`

(pop)

# Enq Method Part Deux

```
public void enq(T x) {  
    ...  
    if (mustWakeDequeuers) {  
        deqLock.lock();  
        try {  
            notEmptyCondition.signalAll();  
        } finally {  
            deqLock.unlock();  
        }  
    }  
}
```

# Enq Method Part Deux

```
public void enq(T x) {  
    ...  
    if (mustWakeDequeuers) {  
        deqLock.lock();  
        try {  
            notEmptyCondition.signalAll();  
        } finally {  
            deqLock.unlock();  
        }  
    }  
}
```

**Are there dequeuers to be signaled?**

# Enq Method Part Deux

```
public void enq(T x) {
```

```
...
```

```
if (mustWakeDequeueers) {
```

```
    deqLock.lock();
```

```
    try {
```

```
        notEmptyCondition.signalAll();
```

```
    } finally {
```

```
        deqLock.unlock();
```

```
    }
```

```
}
```

```
}
```

**Lock and  
unlock deq lock**

# Enq Method Part Deux

**Signal dequeuers that  
queue is no longer empty**

```
    deqLock.lock();  
    try {  
        notEmptyCondition.signalAll();  
    } finally {  
        deqLock.unlock();  
    }  
}
```



# The `enq()` & `deq()` Methods

- Share no locks
  - That's good
- But do share an atomic counter
  - Accessed on every method call
  - That's not so good
- Can we alleviate this bottleneck?


# Split the Counter

- The **enq()** method
  - Increments only
  - Cares only if value is **capacity**
- The **deq()** method
  - Decrements only
  - Cares only if value is **zero**


# Split Counter

- Enqueueer increments **enqSize**
- Dequeueer increments **deqSize**
- When enqueueer hits capacity
  - Locks **deqLock**
  - Sets **size = enqSize - DeqSize**
- Intermittent synchronization
  - Not with each method call
  - Need both locks! (careful ...)


# A Lock-Free Queue




# Compare and Set




# Enqueue



# Enqueue




# Logical Enqueue





# Physical Enqueue



# Enqueue

- These two steps are not atomic
- The tail field refers to either
  - Actual last Node (good)
  - Penultimate Node (not so good)
- Be prepared!


# Enqueue

- What do you do if you find
  - A trailing **tail**?
- Stop and help fix it
  - If **tail** node has non-*null* next field
  - CAS the queue's **tail** field to **tail.next**
- As in the universal construction

# When CASs Fail


- During logical enqueue
  - Abandon hope, restart
  - Still lock-free (why?)
- During physical enqueue
  - Ignore it (why?)

# Dequeuer



Make first Node  
new sentinel


# Dequeuer



# Memory Reuse?

- What do we do with nodes after we dequeue them?
- Java: let garbage collector deal?
- Suppose there is no GC, or we prefer not to use it?

# Dequeuer







# Simple Solution

- Each thread has a free list of unused queue nodes
- Allocate node: pop from list
- Free node: push onto list
- Deal with underflow somehow ...


# Why Recycling is Hard




# Both Nodes Reclaimed




# One Node Recycled




# Why Recycling is Hard





# Recycle FAIL



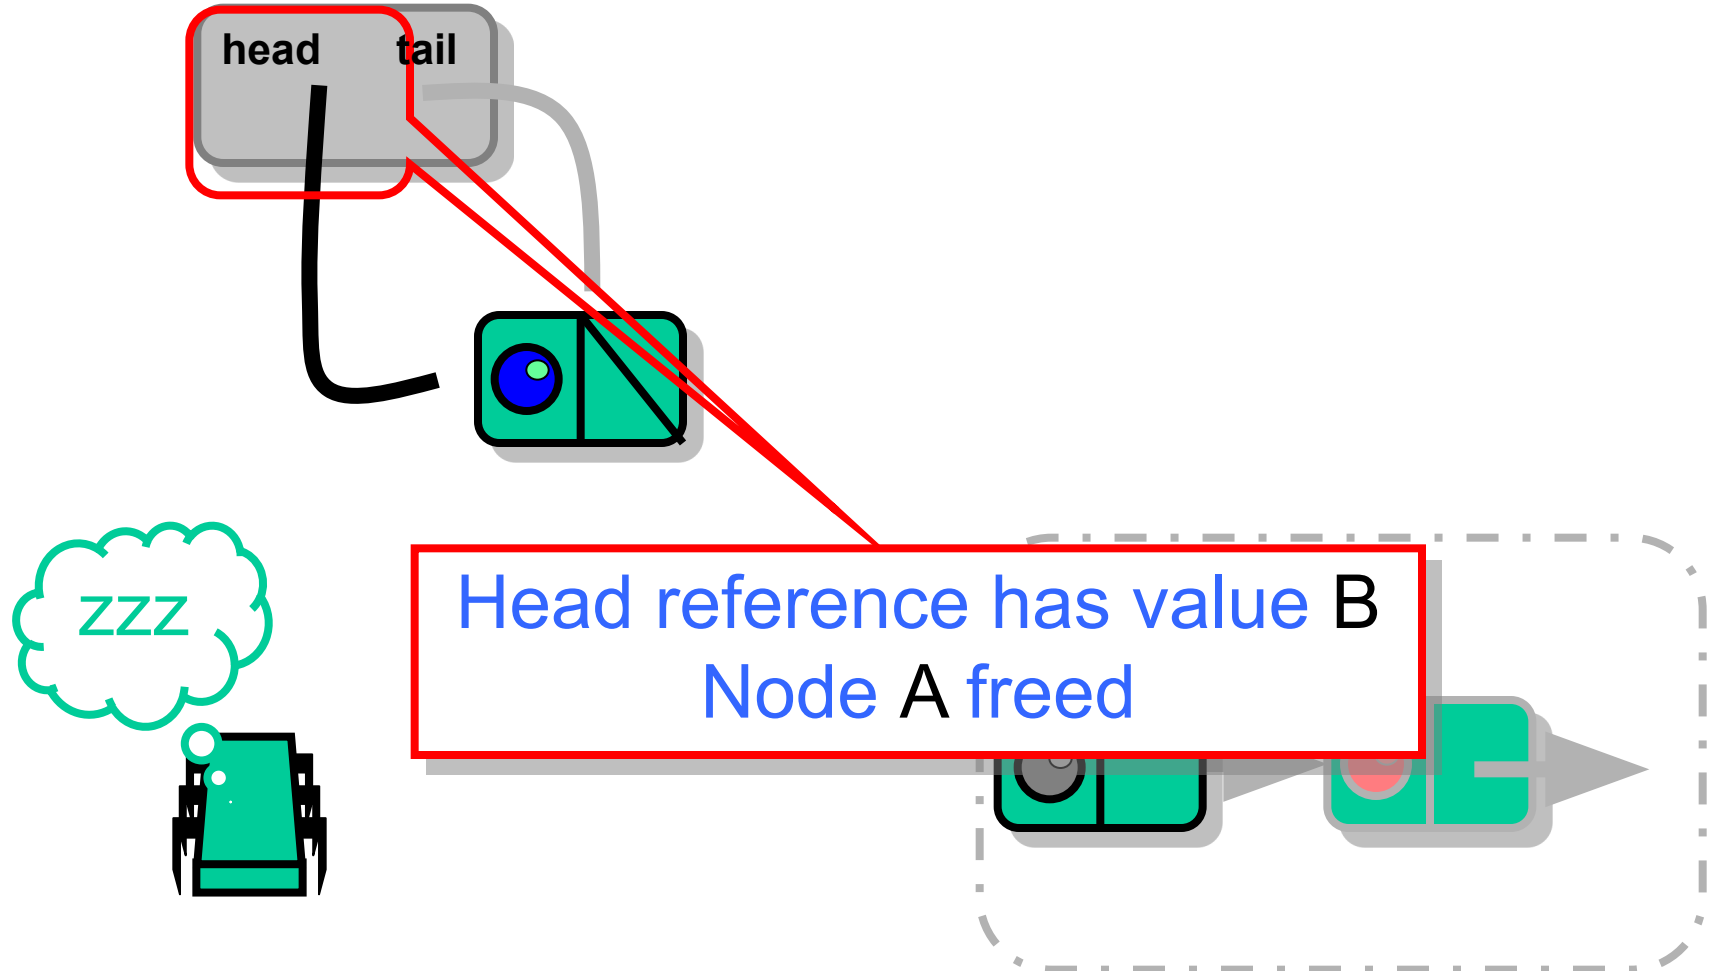
zOMG what went wrong?



# The Dreaded ABA Problem

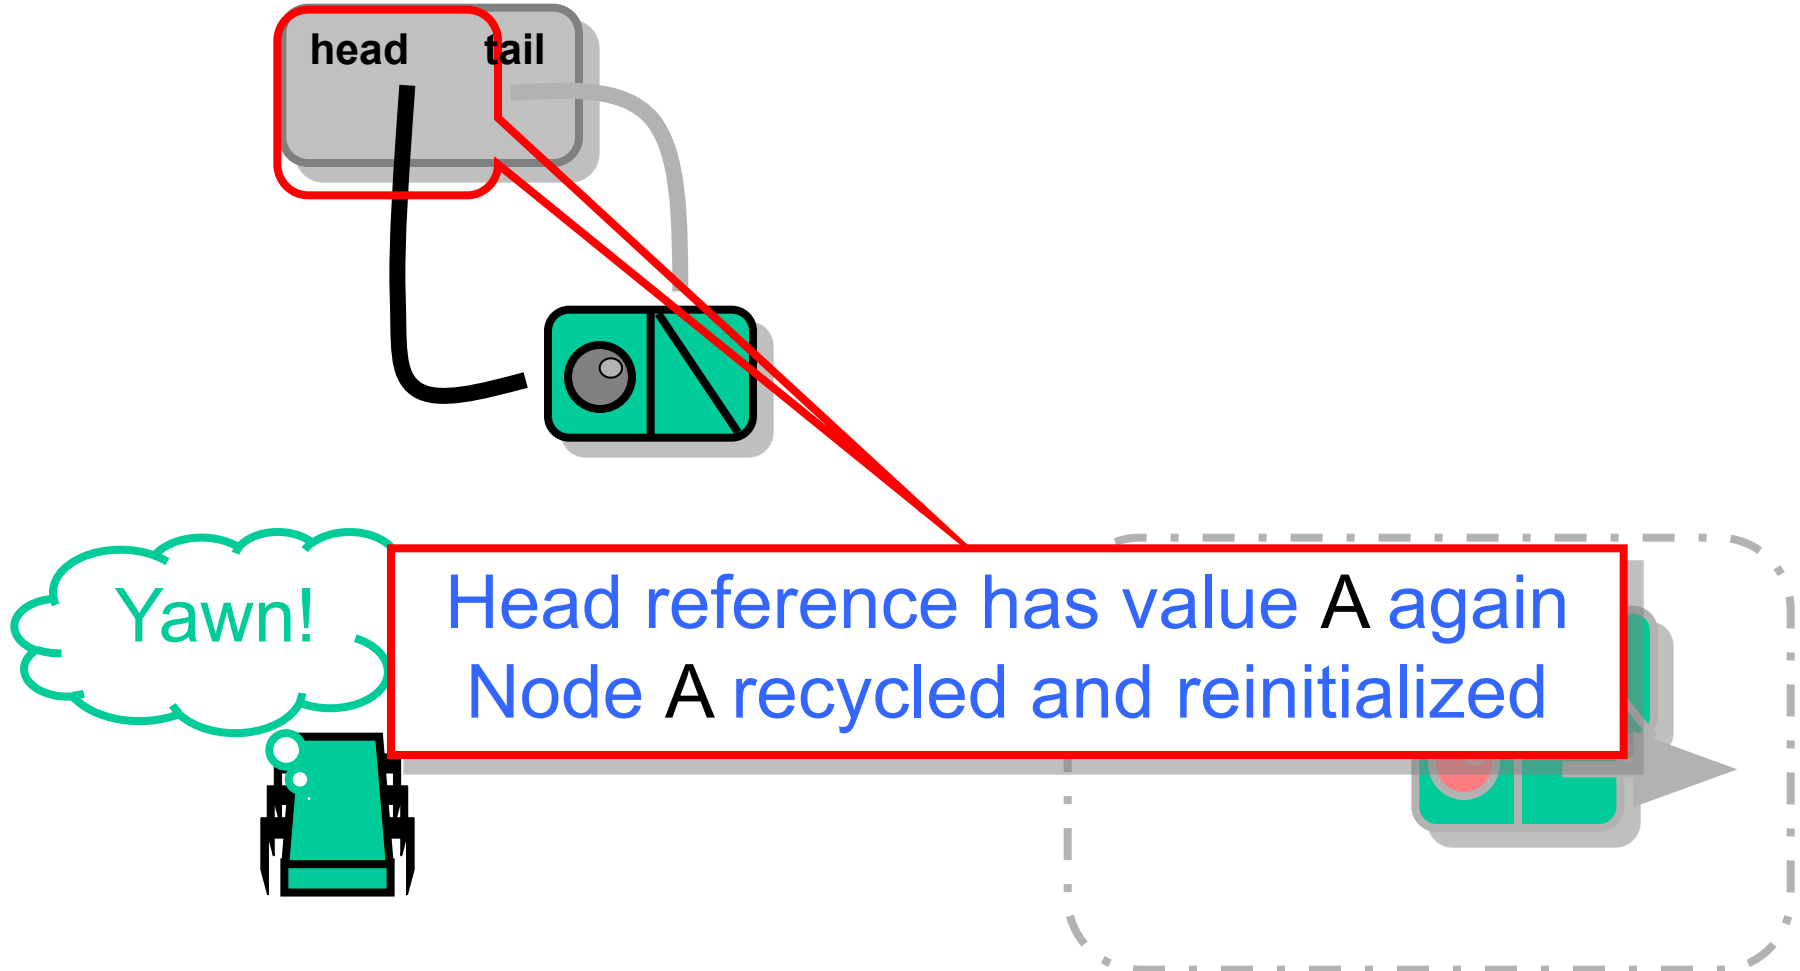


# Dreaded ABA continued






# Dreaded ABA continued



# Dreaded ABA continued



CAS succeeds because references match, even though reference's meaning has changed

# The Dreaded ABA FAIL

- Is a result of CAS() semantics
  - Oracle, Intel, AMD, ...
- Not with Load-Locked/Store-Conditional
  - IBM ...


# Dreaded ABA – A Solution

- Tag each pointer with a counter
- Unique over lifetime of node
- Pointer size vs word size issues
- Overflow?
  - Don't worry be happy?
  - Bounded tags?
- AtomicStampedReference **class**

# Atomic Stamped Reference

- **AtomicStampedReference** class
  - Java.util.concurrent.atomic package


Can get reference & stamp atomically




# Concurrent Stack

- **Methods**
  - **push(x)**
  - **pop()**
- **Last-in, First-out (LIFO) order**
- **Lock-Free!**

# Empty Stack




# Push








# Push




# Push




# Push




# Push




# Push




# Push



# Pop




# Pop








# Pop



# Pop



# Pop



# Lock-free Stack

```
public class LockFreeStack {
    private AtomicReference top =
        new AtomicReference(null);
    public boolean tryPush(Node node) {
        Node oldTop = top.get();
        node.next = oldTop;
        return (top.compareAndSet(oldTop, node))
    }
    public void push(T value) {
        Node node = new Node(value);
        while (true) {
            if (tryPush(node)) {
                return;
            } else backoff.backoff();
        }
    }
}
```

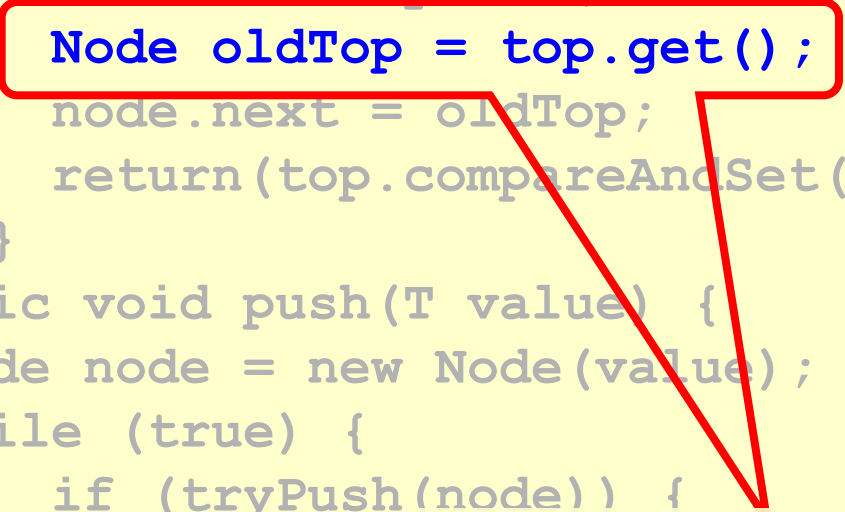
# Lock-free Stack

```
public class LockFreeStack {  
    private AtomicReference top = new  
AtomicReference(null);  
public Boolean tryPush(Node node) {  
    Node oldTop = top.get();  
    node.next = oldTop;  
    return(top.compareAndSet(oldTop, node))  
    }  
    public void push(T value) {  
        Node node = new Node(value);  
        while (true) {  
            if (tryPush(node))  
                return;  
        }  
    }  
}
```

**tryPush attempts to push a node**

# Lock-free Stack

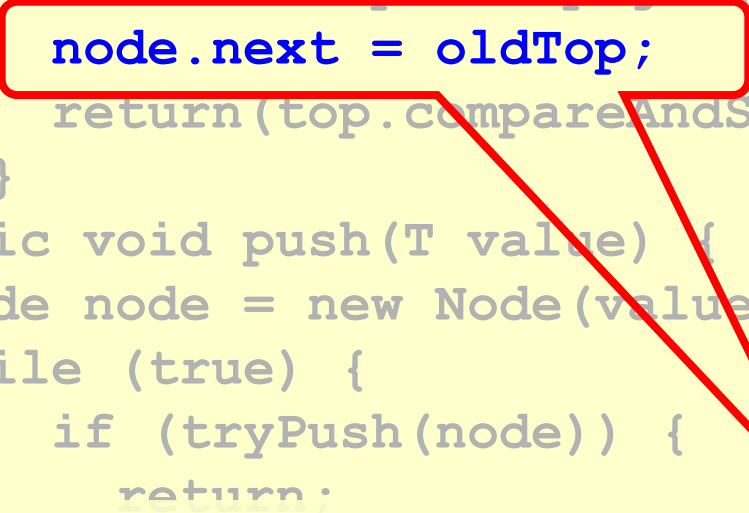
```
public class LockFreeStack {
    private AtomicReference top = new
AtomicReference(null);
    public boolean tryPush(Node node) {
        Node oldTop = top.get();
        node.next = oldTop;
        return (top.compareAndSet(oldTop, node))
    }
    public void push(T value) {
        Node node = new Node(value);
        while (true) {
            if (tryPush(node)) {
                return;
            } else backoff.backoff()
        }
    }
}
```



**Read top value**

# Lock-free Stack

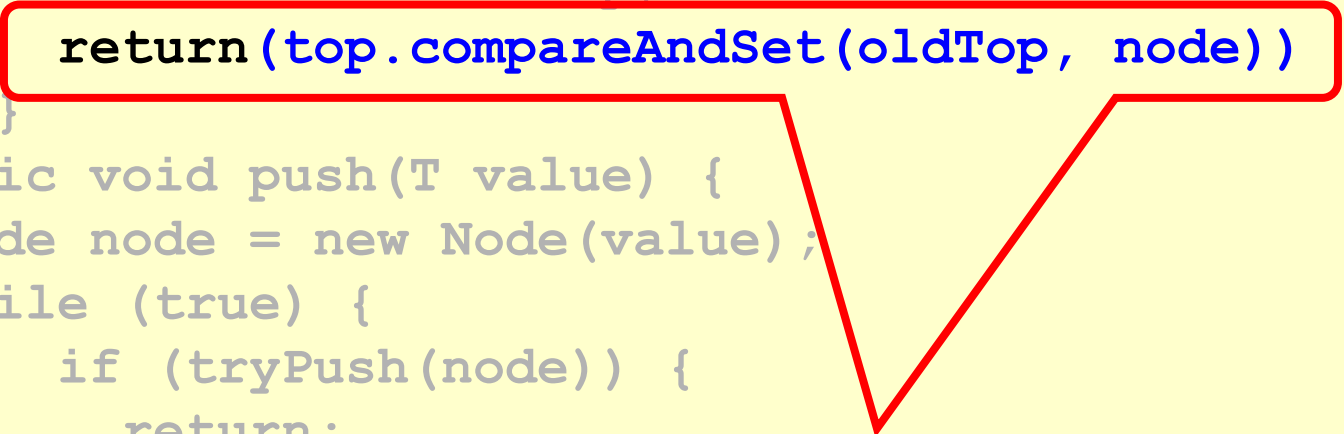
```
public class LockFreeStack {
    private AtomicReference top = new
AtomicReference(null);
    public boolean tryPush(Node node) {
        Node oldTop = top.get();
        node.next = oldTop;
        return(top.compareAndSet(oldTop, node))
    }
    public void push(T value) {
        Node node = new Node(value);
        while (true) {
            if (tryPush(node)) {
                return;
            }
        }
    }
}
```



**current top will be new node's successor**

# Lock-free Stack

```
public class LockFreeStack {
    private AtomicReference top = new
AtomicReference(null);
    public boolean tryPush(Node node){
        Node oldTop = top.get();
        node.next = oldTop;
        return (top.compareAndSet(oldTop, node))
    }
    public void push(T value) {
        Node node = new Node(value);
        while (true) {
            if (tryPush(node)) {
                return;
            }
        }
    }
}
```

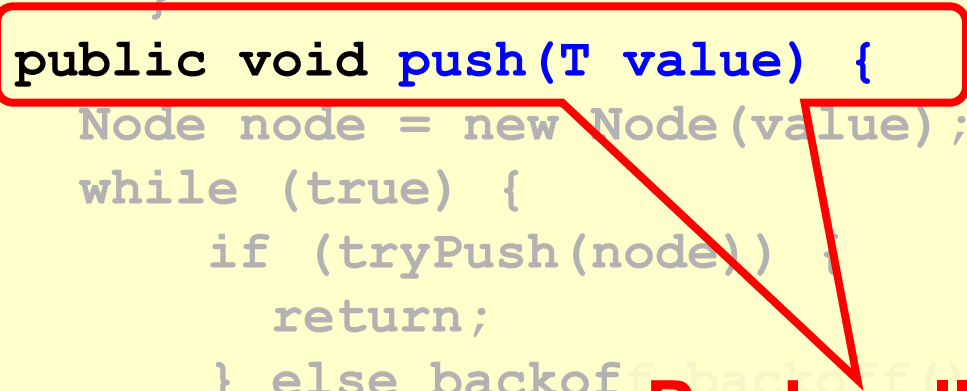


**Try to swing top, return success or failure**



# Lock-free Stack

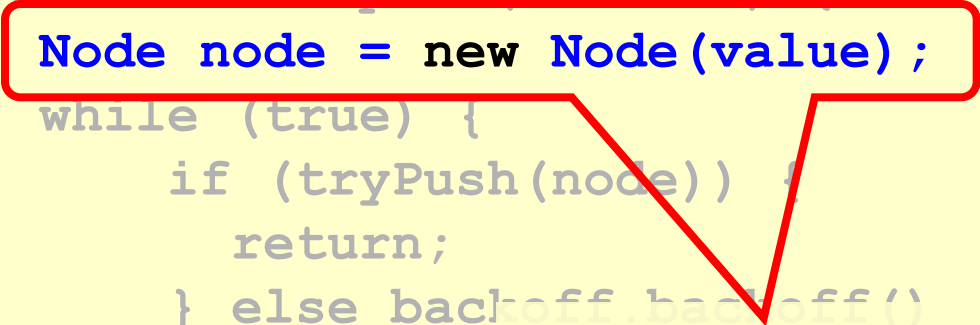
```
public class LockFreeStack {
    private AtomicReference top = new
AtomicReference(null);
    public boolean tryPush(Node node) {
        Node oldTop = top.get();
        node.next = oldTop;
        return (top.compareAndSet(oldTop, node))
    }
    public void push(T value) {
        Node node = new Node(value);
        while (true) {
            if (tryPush(node))
                return;
            else backoff();
        }
    }
}
```



**Push calls tryPush**

# Lock-free Stack

```
public class LockFreeStack {
    private AtomicReference top = new
AtomicReference(null);
    public boolean tryPush(Node node) {
        Node oldTop = top.get();
        node.next = oldTop;
        return (top.compareAndSet(oldTop, node))
    }
    public void push(T value) {
        Node node = new Node(value);
        while (true) {
            if (tryPush(node)) {
                return;
            } else backoff.backoff()
        }
    }
}
```

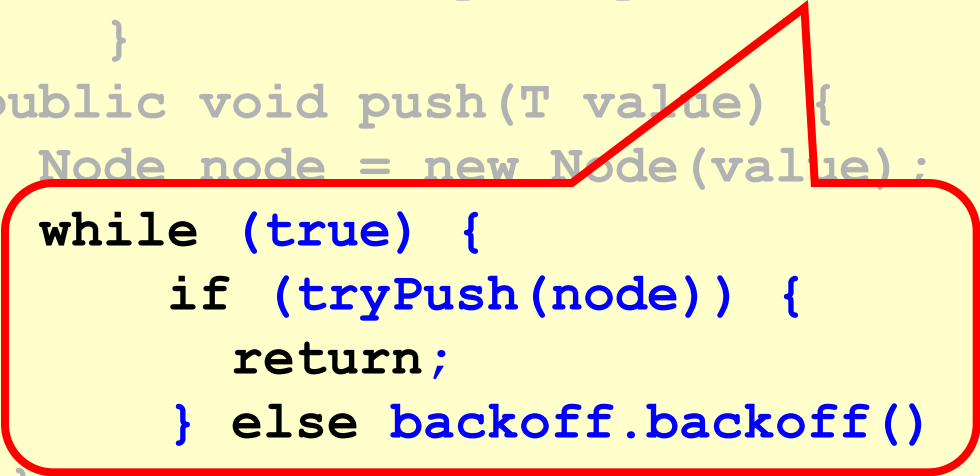


**Create new node**

# Lock-free Stack

```
public class LockFreeStack {
    private AtomicReference top = new
AtomicReference(null);
    public boolean tryPush(Node node){
        Node oldTop = top.get();
        node.next = oldTop;
        return (top.compareAndSet(oldTop, node));
    }
    public void push(T value) {
        Node node = new Node(value);
        while (true) {
            if (tryPush(node)) {
                return;
            } else backoff.backoff()
        }
    }
}
```

**If tryPush() fails,  
back off before retrying**



# Lock-free Stack

- Good
  - No locking
- Bad
  - Without GC, fear ABA
  - Without backoff, huge contention at top
  - In any case, no parallelism


# Big Question

- Are stacks *inherently* sequential?
- Reasons why
  - Every **pop()** call fights for top item
- Reasons why not
  - Stay tuned ...

# Elimination-Backoff Stack


- How to
  - “turn contention into parallelism”
- Replace familiar
  - **exponential backoff**
- With alternative
  - **elimination-backoff**

# Observation




**Push()**

**linearizable stack**




**Pop()**



Yes!


After an equal number  
of pushes and pops,  
stack stays the same

# Idea: Elimination Array






# Push Collides With Pop




# No Collision




# Elimination-Backoff Stack

- Lock-free stack + elimination array
- Access Lock-free stack,
  - If **uncontended**, apply operation
  - if **contended**, back off to elimination array and attempt elimination


# Elimination-Backoff Stack




# Dynamic Range and Delay




# 50-50, Random Slots



# Asymmetric Rendezvous




## 50% Push/50% Pop/Random ids







## 50% Push/50% Pop/Random ids




## (a) Instructions/op



## (c) CAS/op




## (b) Cache misses/op



# Effect of Backoff and Slot Choice



## 50% Push/50% Pop



# Effect of Slot Choice


Random choice

Sequential choice




Darker shades mean more exchanges


# Effect of Slot Choice




## Rendezvous array probes: sequential vs. random ids



## Rendezvous instruction count: sequential vs. random ids



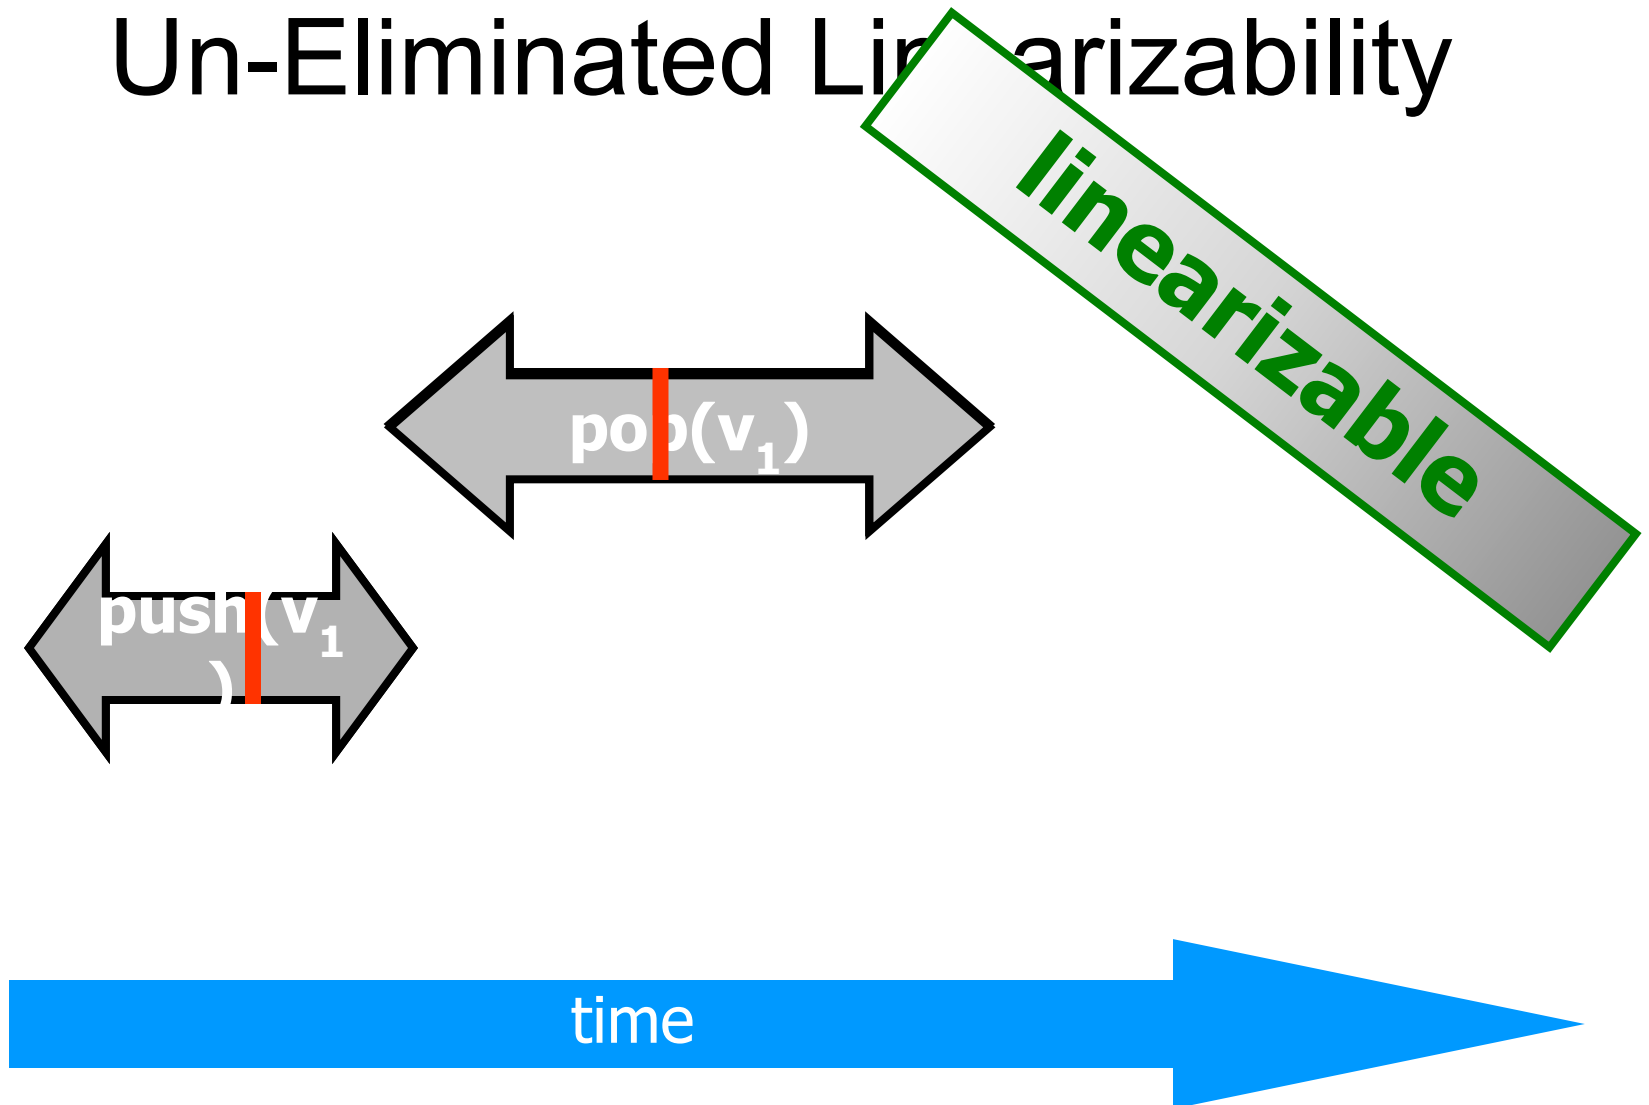
## Rendezvous Pop spin iterations: sequential vs. random ids



# Linearizability

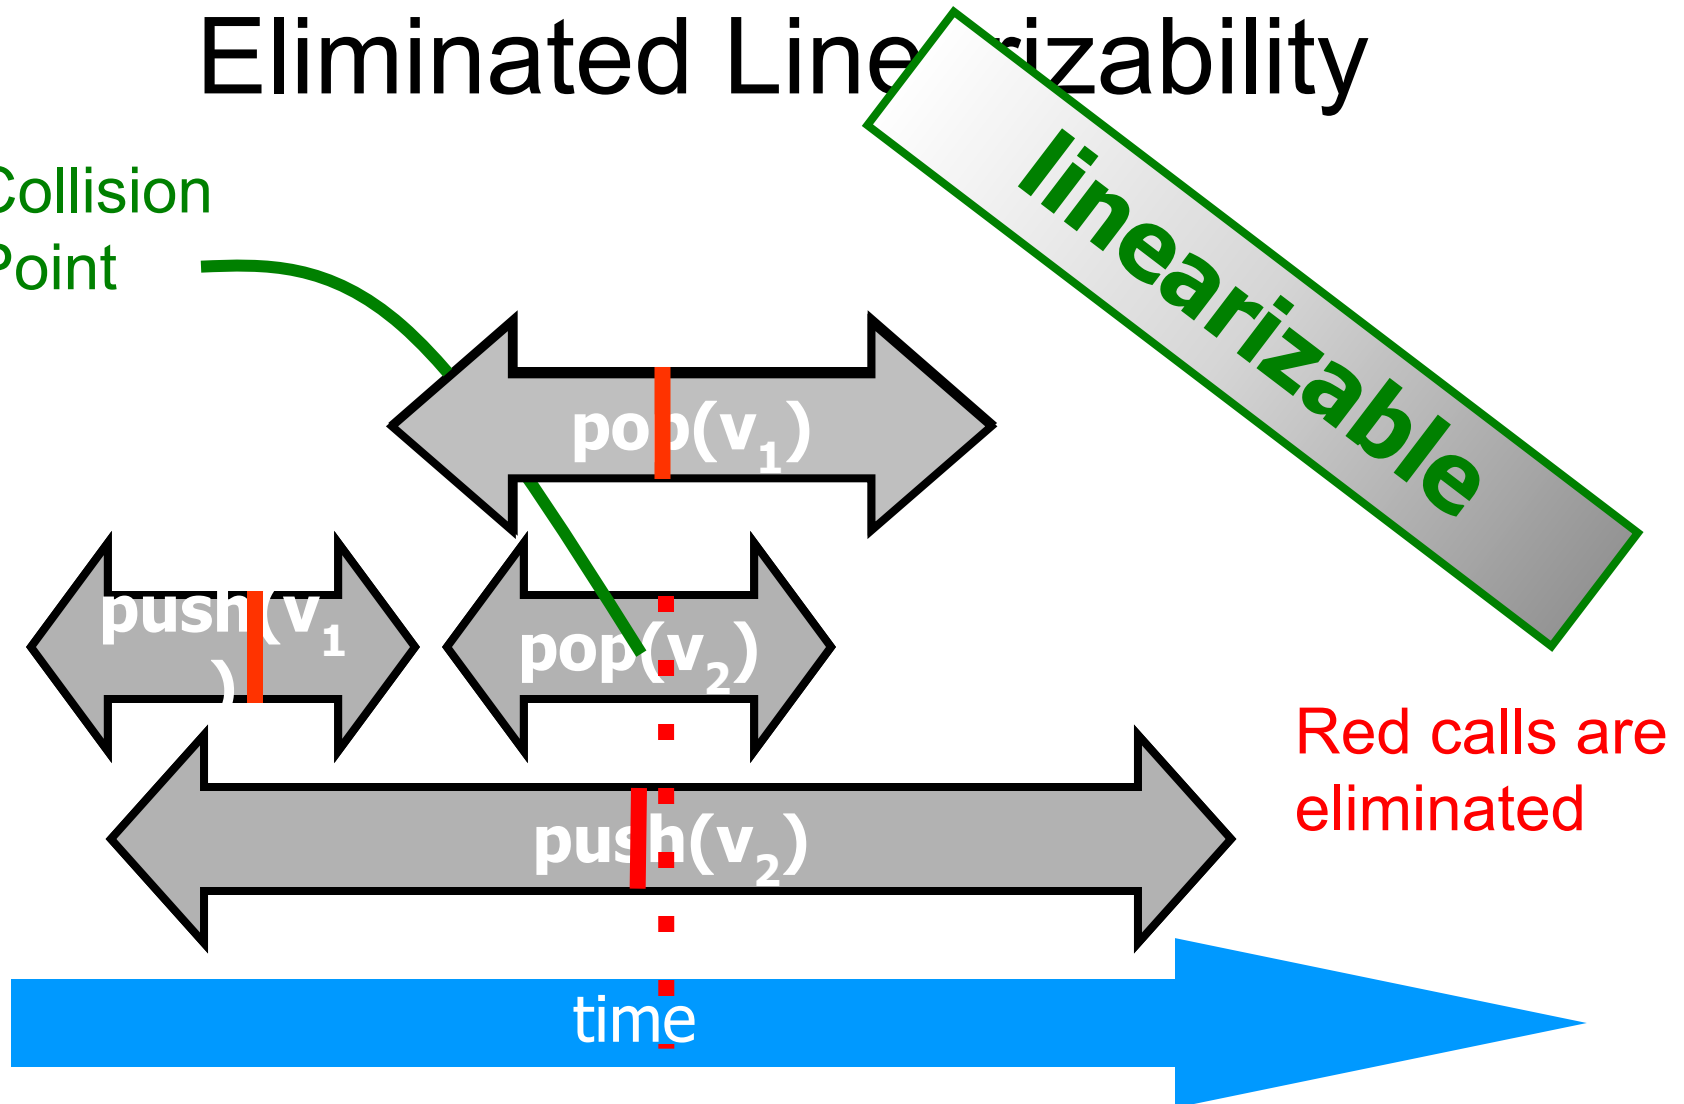
- **Un-eliminated calls**
  - **linearized as before**
- **Eliminated calls:**
  - **linearize pop() immediately after matching push()**
- **Combination is a linearizable stack**

# Un-Eliminated Linearizability



# Eliminated Linearizability

Collision  
Point



# Backoff Has Dual Effect

- Elimination introduces parallelism
- Backoff to array cuts contention on lock-free stack
- Elimination in array cuts down number of threads accessing lock-free stack



# Elimination Array

```
public class EliminationArray {  
    private static final int duration = ...;  
    private static final int timeUnit = ...;  
    Exchanger<T>[] exchanger;  
    public EliminationArray(int capacity) {  
        exchanger = new Exchanger[capacity];  
        for (int i = 0; i < capacity; i++)  
            exchanger[i] = new Exchanger<T>();  
        ...  
    }  
    ...  
}
```

# Elimination Array

```
public class EliminationArray {  
    private static final int duration = ...;  
    private static final int timeUnit = ...;  
    Exchanger<T>[] exchanger;  
    public EliminationArray(int capacity) {  
        exchanger = new Exchanger[capacity];  
        for (int i = 0; i < capacity; i++)  
            exchanger[i] = new Exchanger<T>();  
        ...  
    }  
    ...  
}
```

An array of *Exchangers*

# Digression: A Lock-Free Exchanger

```
public class Exchanger<T> {  
    AtomicStampedReference<T> slot  
    = new AtomicStampedReference<T>(null, 0);  
}
```


# A Lock-Free Exchanger

```
public class Exchanger<T> {  
    AtomicStampedReference<T> slot  
    = new AtomicStampedReference<T>(null, 0);  
}
```

Atomically modifiable  
reference + status

# Atomic Stamped Reference

- **AtomicStampedReference** class
  - Java.util.concurrent.atomic package
- In C or C++:



# Extracting Reference & Stamp

```
public T get(int[] stampHolder) ;
```

# Extracting Reference & Stamp

```
public T get(int[] stampHolder) ;
```

Returns reference to  
object of type T

Returns stamp at  
array index 0!

# Exchanger Status

```
enum Status {EMPTY, WAITING, BUSY};
```



# Exchanger Status

```
enum Status {EMPTY, WAITING, BUSY};
```



Nothing yet

# Exchange Status

```
enum Status { EMPTY, WAITING, BUSY};
```

Nothing yet

One thread is waiting  
for rendez-vous

# Exchange Status

```
enum Status { EMPTY, WAITING, BUSY };
```

Nothing yet

One thread is waiting  
for rendez-vous

Other threads busy  
with rendez-vous

# The Exchange

```
public T Exchange(T myItem, long nanos)
    throws TimeoutException {
    long timeBound = System.nanoTime() + nanos;
    int[] stampHolder = {EMPTY};
    while (true) {
        if (System.nanoTime() > timeBound)
            throw new TimeoutException();
        T herItem = slot.get(stampHolder);
        int stamp = stampHolder[0];
        switch(stamp) {
            case EMPTY: ...      // slot is free
            case WAITING: ...    // someone waiting for me
            case BUSY: ...       // others exchanging
        }
    }
}
```

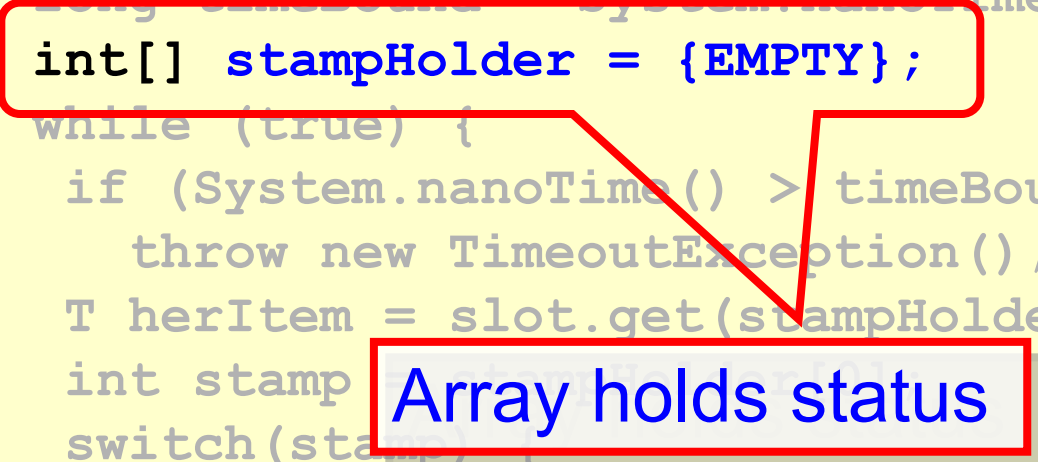
# The Exchange

```
public T Exchange(T myItem, long nanos)
    throws TimeoutException {
    long timeBound = System.nanoTime() + nanos;
    int[] stampHolder = {EMPTY};
    while (true) {
        if (System.nanoTime() > timeBound)
            throw new TimeoutException();
        T herItem = slot.get(stampHolder);
        int stamp = stampHolder[0];
        switch(stamp) {
            case EMPTY: ... // slot is free
            case WAITING: ... // someone waiting for me
            case BUSY: ... // others exchanging
        }
    }
}
```

Item and timeout

# The Exchange

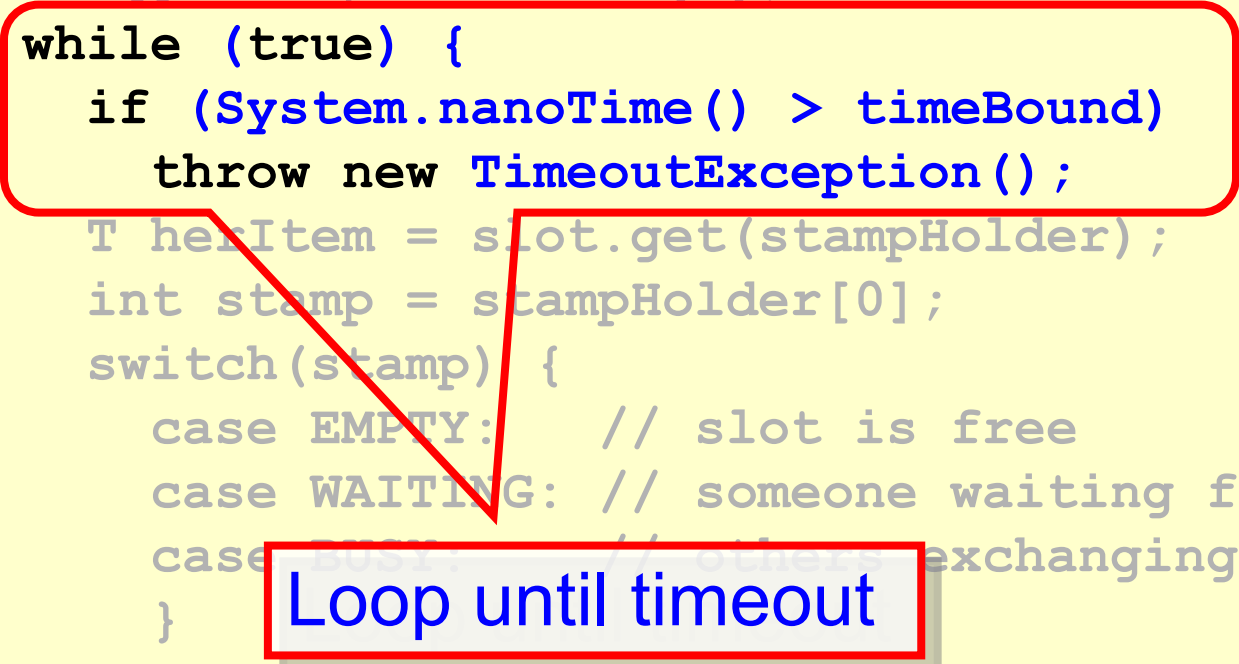
```
public T Exchange(T myItem, long nanos)
    throws TimeoutException {
    long timeBound = System.nanoTime() + nanos;
    int[] stampHolder = {EMPTY};
    while (true) {
        if (System.nanoTime() > timeBound)
            throw new TimeoutException();
        T herItem = slot.get(stampHolder);
        int stamp
        switch(stamp) {
            case EMPTY: ...    // slot is free
            case WAITING: ...  // someone waiting for me
            case BUSY: ...     // others exchanging
        }
    }
}
```



Array holds status

# The Exchange

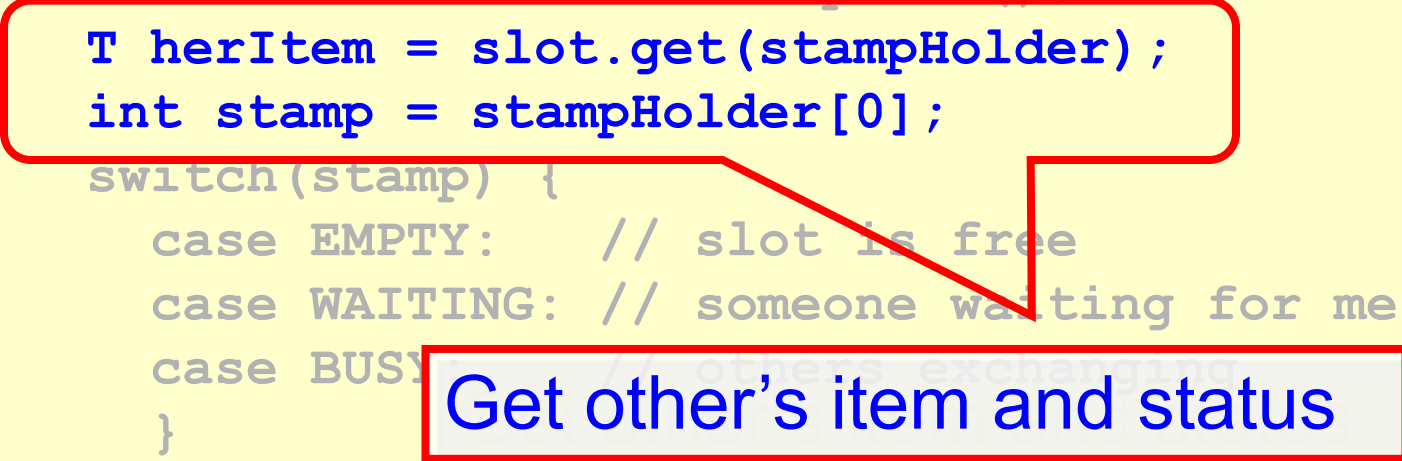
```
public T Exchange(T myItem, long nanos) throws
TimeoutException {
    long timeBound = System.nanoTime() + nanos;
    int[] stampHolder = {0};
    while (true) {
        if (System.nanoTime() > timeBound)
            throw new TimeoutException();
        T herItem = slot.get(stampHolder);
        int stamp = stampHolder[0];
        switch(stamp) {
            case EMPTY: // slot is free
            case WAITING: // someone waiting for me
            case BUSY: // others exchanging
        }
    }
}
```



Loop until timeout

# The Exchange

```
public T Exchange(T myItem, long nanos) throws
TimeoutException {
    long timeBound = System.nanoTime() + nanos;
    int[] stampHolder = {0};
    while (true) {
        if (System.nanoTime() > timeBound)
            throw new TimeoutException();
        T herItem = slot.get(stampHolder);
        int stamp = stampHolder[0];
        switch(stamp) {
            case EMPTY:    // slot is free
            case WAITING:  // someone waiting for me
            case BUSY:     // others exchanging
        }
    }
}
```



**T herItem = slot.get(stampHolder);**  
**int stamp = stampHolder[0];**

**Get other's item and status**





# The Exchange

```
public T Exchange(T myItem, long nanos) throws  
TimeoutException {
```


An *Exchanger* has three possible states

```
    long timeBound = System.nanoTime() + nanos;  
    while (true) {  
        if (System.nanoTime() > timeBound)  
            throw new TimeoutException();  
        T herItem = slot.get(stampHolder);  
        int stamp = stampHolder[0];  
        switch(stamp) {  
            case EMPTY: ... // slot is free  
            case WAITING: ... // someone waiting for me  
            case BUSY: ... // others exchanging  
        }  
    }  
}
```


# Lock-free Exchanger




# Lock-free Exchanger



# Lock-free Exchanger






# Lock-free Exchanger



# Lock-free Ex




Still waiting ...

Try to exchange  
item and set  
status to **BUSY**



# Look-free Exchanger

Partner showed  
up, take item and  
reset to EMPTY




**item**

**status**

# Look-free Exchanger


Partner showed  
up, take item and  
reset to EMPTY

Slot



item

status





# Exchanger State EMPTY

```
case EMPTY: // slot is free
    if (slot.CAS(herItem, myItem, EMPTY, WAITING)) {
        while (System.nanoTime() < timeBound){
            herItem = slot.get(stampHolder);
            if (stampHolder[0] == BUSY) {
                slot.set(null, EMPTY);
                return herItem;
            }
        }
        if (slot.CAS(myItem, null, WAITING, EMPTY)) {
            throw new TimeoutException();
        } else {
            herItem = slot.get(stampHolder);
            slot.set(null, EMPTY);
            return herItem;
        }
    }
    break;
```

# Exchanger State EMPTY

```
case EMPTY: // slot is free
    if (slot.CAS(herItem, myItem, EMPTY, WAITING)) {
        while (System.nanoTime() < timeBound) {
            herItem = slot.get(stampHolder);
            if (stampHolder[0] == BUSY) {
                slot.set(null, EMPTY);
                return herItem;
            }
        }
        if (slot.CAS(myItem, null, WAITING, EMPTY)) {
            throw new TimeoutException();
        } else {
            herItem = slot.get(stampHolder);
            slot.set(null, EMPTY);
            return herItem;
        }
    }
    break;
```

Try to insert *myItem* and  
change state to *WAITING*

# Exchanger State EMPTY

```
case EMPTY: // slot is free
    if (slot.CAS(herItem, myItem, EMPTY, WAITING)) {
        while (System.nanoTime() < timeBound) {
            herItem = slot.get(stampHolder);
            if (stampHolder[0] == BUSY) {
                slot.set(null, EMPTY);
                return herItem;
            }
        }
        if (slot.CAS(myItem, null, WAITING, EMPTY)) {
            throw new TimeoutException();
        } else {
            herItem = slot.get(stampHolder);
            slot.set(myItem, herItem);
            return herItem;
        }
    }
} break;
```

Spin until either  
*myItem* is taken or timeout

# Exchanger State EMPTY

```
case EMPTY: // slot is free
    if (slot.CAS(herItem, myItem, EMPTY, WAITING)) {
        while (System.nanoTime() < timeBound) {
            herItem = slot.get(stampHolder);
            if (stampHolder[0] == BUSY) {
                slot.set(null, EMPTY);
                return herItem;
            }
        }
        if (slot.CAS(myItem, null, WAITING, EMPTY)) {
            throw new TimeoutException();
        } else {
            herItem = slot.get(stampHolder);
            slot.set(myItem, herItem);
            return herItem;
        }
    }
} break;
```

*myItem* was taken,  
so return *herItem*  
that was put in its place

# Exchanger State EMPTY

```
case EMPTY: // slot is free
```

Otherwise we ran out of time,  
try to reset status to EMPTY  
and time out

```
slot.set(null, EMPTY);
```

```
return herItem;
```

```
}}
```

```
if (slot.CAS(myItem, null, WAITING, EMPTY)) {  
    throw new TimeoutException();
```

```
} else {
```

```
herItem = slot.get(stampHolder);
```

```
slot.set(null, EMPTY);
```

```
return herItem;
```

```
}
```

```
} break;
```

# Exchanger State EMPTY

```
case EMPTY: // slot is free
    if (slot.compareAndSet(herItem, myItem, WAITING,
BUSY)) {
        while (System.currentTimeMillis() < timeBound)
            herItem = slot.get(stampHolder);
        if (stampHolder == 0) {
            slot.set(herItem, EMPTY);
            return herItem;
        }
        if (slot.compareAndSet(myItem, null, WAITING,
EMPTY)) {
            throw new TimeoutException();
        }
        else {
            herItem = slot.get(stampHolder);
            slot.set(null, EMPTY);
            return herItem;
        }
    }
} break;
```

If reset failed,  
someone showed up after all,  
so take that item

# Exchanger State EMPTY

```
case EMPTY: // slot is free
    if (slot.CAS(herItem, myItem, EMPTY, WAITING)) {
        while (System.nanoTime() < timeBound) {
            herItem = slot.get(stampHolder);
            if (stampHolder[0] == BUSY) {
                slot.set(null, EMPTY);
                return herItem;
            }
        }
        if (slot.CAS(myItem, null, WAITING, EMPTY)) {
            throw new TimeoutException();
        } else {
            herItem = slot.get(stampHolder);
            slot.set(null, EMPTY);
            return herItem;
        }
    }
    break;
```

Clear slot and take that item

**slot.set(null, EMPTY);**  
**return herItem;**

# Exchanger State EMPTY

```
case EMPTY: // slot is free
    if (slot.CAS(herItem, myItem, EMPTY, WAITING)) {
        while (System.nanoTime() < timeBound) {
            herItem = slot.get(stampHolder);
            if (stampHolder[0] == BUSY) {
                // If initial CAS failed,
                // then someone else changed status
                // from EMPTY to WAITING,
                // so retry from start
            }
        }
        herItem = slot.get(stampHolder);
        slot.set(null, EMPTY);
        return herItem;
    }
} break;
```



# States WAITING and BUSY

```
case WAITING:    // someone waiting for me
    if (slot.CAS(herItem, myItem, WAITING, BUSY))
        return herItem;
    break;
case BUSY:       // others in middle of exchanging
    break;
default:        // impossible
    break;
    }
    }
    }
}
```

# States WAITING and BUSY

```
case WAITING: // someone waiting for me
    if (slot.CAS(herItem, myItem, WAITING, BUSY))
        return herItem;
    break;
case BUSY: // others in middle of exchanging
    break;
default: // impossible
    break;
}
}
```

someone is waiting to exchange,  
so try to CAS my item in  
and change state to BUSY

# States WAITING and BUSY

```
case WAITING:    // someone waiting for me
    if (slot.CAS(herItem, myItem, WAITING, BUSY))
        return herItem;
    break;
case BUSY:       // others in middle of exchanging
    break;
default:        // impossible
    break;
}
}
}
}
```

If successful, return other's item,  
otherwise someone else took it,  
so try again from start

# States WAITING and BUSY

```
case WAITING:    // someone waiting for me
    if (slot.CAS(herItem, myItem, WAITING, BUSY))
        return herItem;
    break;
case BUSY:      // others in middle of exchanging
    break;
default:         // impossible
    break;
    }
    }
    }
}
```

If **BUSY**,  
other threads exchanging,  
so start again

# The Exchanger Slot

- Exchanger is lock-free
- Because the only way an exchange can fail is if others repeatedly succeeded or no-one showed up
- The slot we need does not require symmetric exchange

# Back to the Stack: the Elimination Array

```
public class EliminationArray {  
    ...  
    public T visit(T value, int range)  
        throws TimeoutException {  
        int slot = random.nextInt(range);  
        int nanodur = convertToNanos(duration, timeUnit);  
        return (exchanger[slot].exchange(value, nanodur)  
    }  
}
```

# Elimination Array

```
public class EliminationArray {  
    public T visit(T value, int range)  
        throws TimeoutException {  
        int slot = random.nextInt(range);  
        int nanodur = convertToNanos(duration, timeUnit);  
        return (exchanger[slot].exchange(value, nanodur));  
    }  
}
```

visit the elimination array  
with fixed value and range

# Elimination Array

```
public class EliminationArray {  
    ...  
    public T visit(T value, int range)  
        throws TimeoutException {  
        int slot = random.nextInt(range);  
        int nanodur = convertToNanos(duration, timeUnit);  
        return (exchanger[slot].exchange(value, nanodur)  
    }  
}
```

**Pick a random array entry**



# Elimination Array

```
public class EliminationArray {  
    ...  
    public T visit(T value, long range)  
        throws TimeoutException {  
        int slot = random.nextInt(range);  
        int nanodur = convertToNanos(duration, timeUnit);  
        return (exchanger[slot].exchange(value, nanodur));  
    }  
}
```

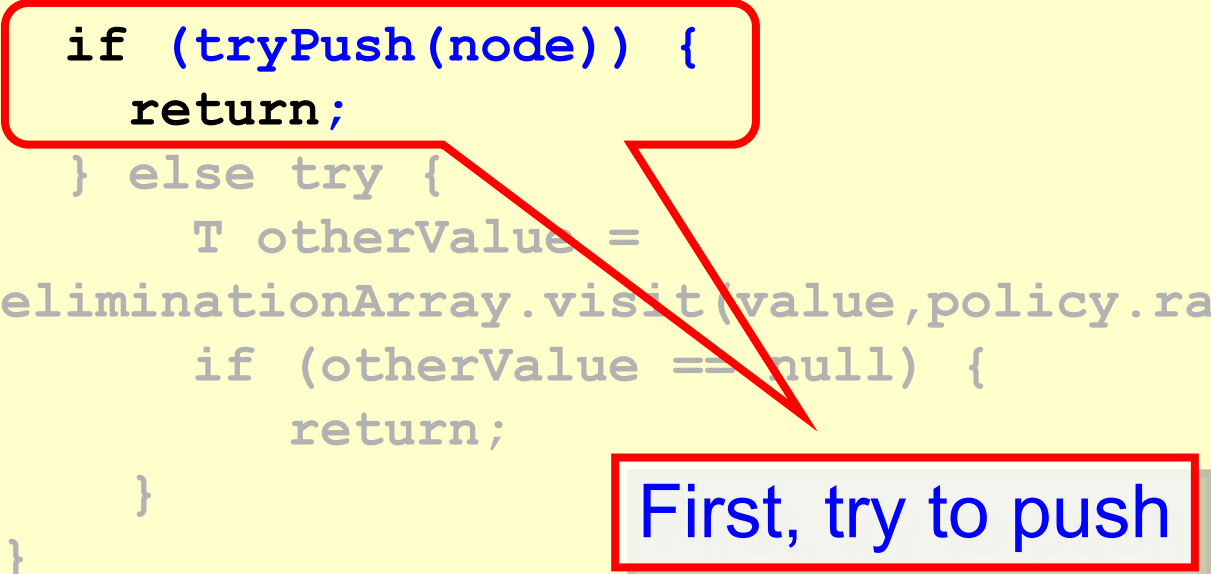
Exchange value or time out

# Elimination Stack Push

```
public void push(T value) {  
    ...  
    while (true) {  
        if (tryPush(node)) {  
            return;  
        } else try {  
            T otherValue =  
eliminationArray.visit(value,policy.range) ;  
            if (otherValue == null) {  
                return;  
            }  
        }  
    }  
}
```

# Elimination Stack Push

```
public void push(T value) {  
    ...  
    while (true) {  
        if (tryPush(node)) {  
            return;  
        } else try {  
            T otherValue =  
eliminationArray.visit(value, policy.range) ;  
            if (otherValue == null) {  
                return;  
            }  
        }  
    }  
}
```



First, try to push

# Elimination Stack Push

```
public void push(T value) {  
    ...  
    while (true) {  
        if (tryPush(node)) {  
            return;  
        } else try {  
            T otherValue =  
eliminationArray.visit(value, policy.range) ;  
            if (otherValue == null) {  
                return;  
            }  
        }  
    }  
}
```

If I failed, backoff & try to eliminate

# Elimination Stack Push

```
public void push(T value) {  
    ...  
    while (true) {  
        if (tryPush(node)) {  
            return;  
        } else try {  
            T otherValue =  
eliminationArray.visit(value, policy.range) ;  
            if (otherValue == null) {  
                return;  
            }  
        }  
    }  
}
```

Value pushed and range to try

# Elimination Stack Push

```
public void push(T value) {  
    ...  
    while (true) {  
        if (tryPush(value)) {  
            return;  
        } else try {  
            T otherValue =  
eliminationArray.visit(value, policy.range);  
            if (otherValue == null) {  
                return;  
            }  
        }  
    }  
}
```

Only `pop()` leaves null,  
so elimination was successful

# Elimination Stack Push

```
public void push(T value) {
```

Otherwise, retry **push()** on lock-free stack

```
... while (true) {  
    if (tryPush(node)) {  
        return;  
    } else try {  
        T otherValue =  
eliminationArray.visit(value, policy.range);  
        if (otherValue == null) {  
            return;  
        }  
    }  
}
```

```
}
```

# Elimination Stack Pop

```
public T pop() {  
    ...  
    while (true) {  
        if (tryPop()) {  
            return returnNode.value;  
        } else  
            try {  
                T otherValue =  
                    eliminationArray.visit(null, policy.range;  
                if (otherValue != null) {  
                    return otherValue;  
                }  
            }  
    }  
}
```



# Elimination Stack Pop

```
public T pop() {
```

```
...
```

```
while (true) {
```

If value not null, other thread is a **push()**,  
so elimination succeeded

```
return returnNode.value;  
} else
```

```
try {
```

```
    T otherValue =
```

```
        eliminationArray.visit(null, policy.range;
```

```
    if ( otherValue != null) {  
        return otherValue;
```

```
    }
```

```
}
```

```
}}
```

# Summary

- We saw both lock-based and lock-free implementations of
- queues and stacks
- Don't be quick to declare a data structure inherently sequential
  - Linearizable stack is not inherently sequential (though it is in the worst case)
- ABA is a real problem, pay attention

This work is licensed under a [Creative Commons Attribution-ShareAlike 2.5 License](https://creativecommons.org/licenses/by-sa/2.5/).

- **You are free:**
  - **to Share** — to copy, distribute and transmit the work
  - **to Remix** — to adapt the work
- **Under the following conditions:**
  - **Attribution.** You must attribute the work to “The Art of Multiprocessor Programming” (but not in any way that suggests that the authors endorse you or your use of the work).
  - **Share Alike.** If you alter, transform, or build upon this work, you may distribute the resulting work only under the same, similar or a compatible license.
- For any reuse or distribution, you must make clear to others the license terms of this work. The best way to do this is with a link to
  - <http://creativecommons.org/licenses/by-sa/3.0/>.
- Any of the above conditions can be waived if you get permission from the copyright holder.
- Nothing in this license impairs or restricts the author's moral rights.